

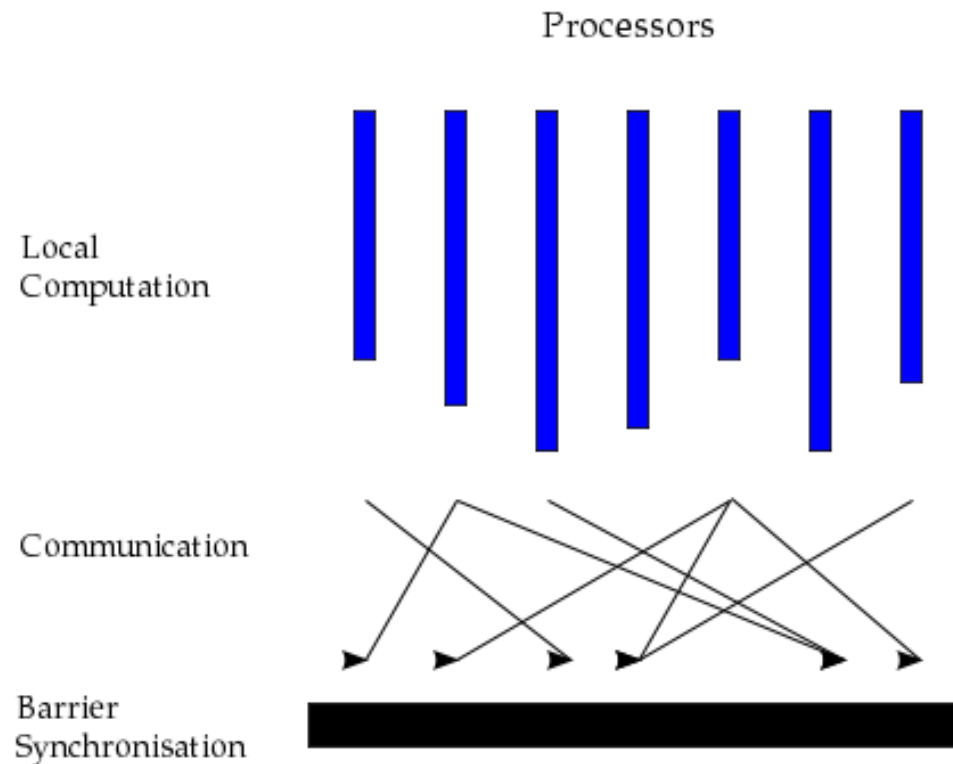


Whatever is left

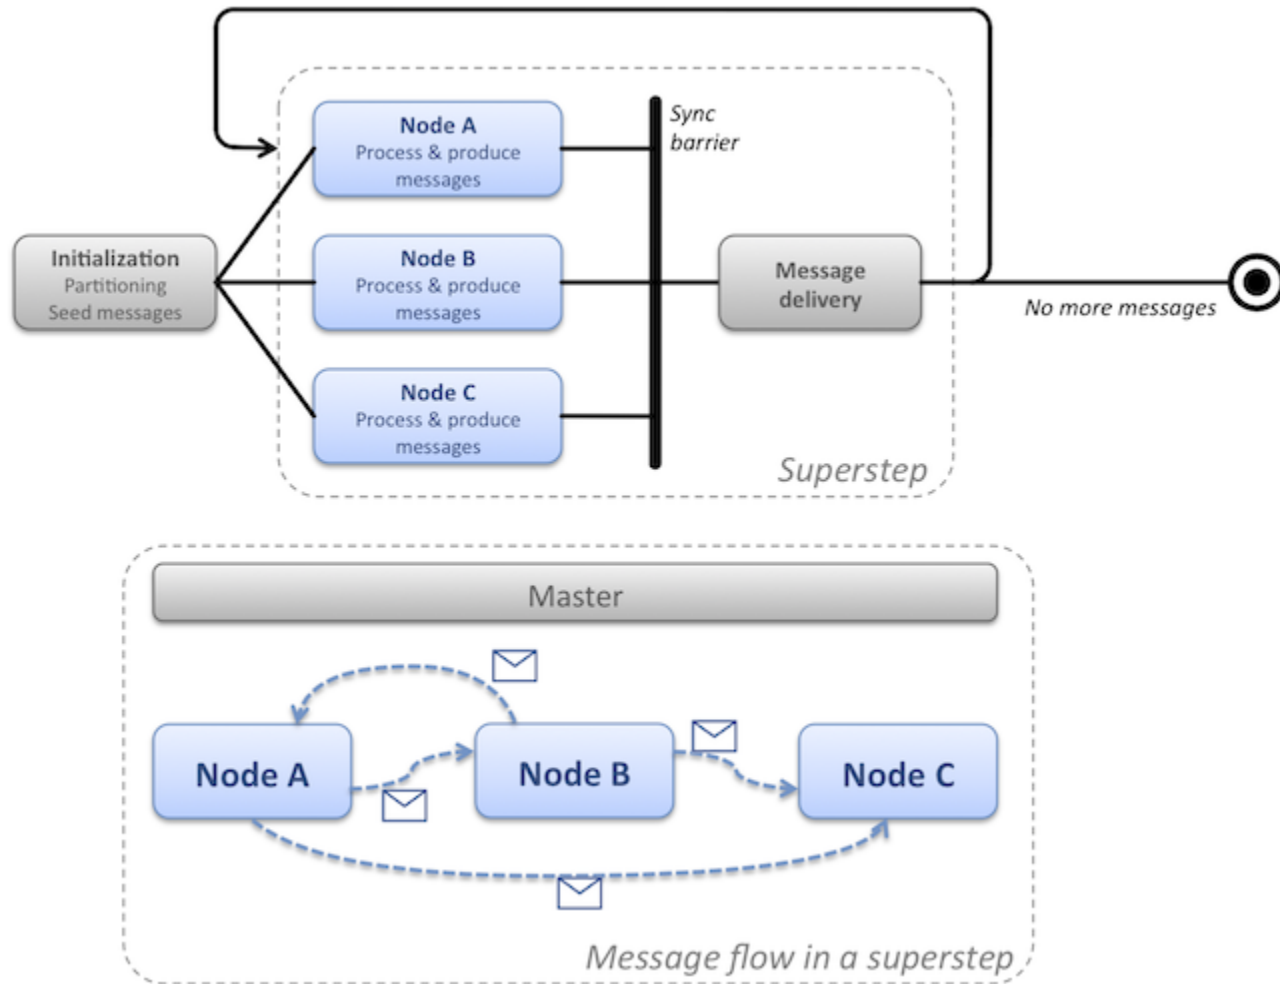


Alexander Lazovik
a.lazovik@rug.nl


Bulk Synchronous Parallel (1990)



BSP



BSP on distributed graphs





Map/Reduce implementation

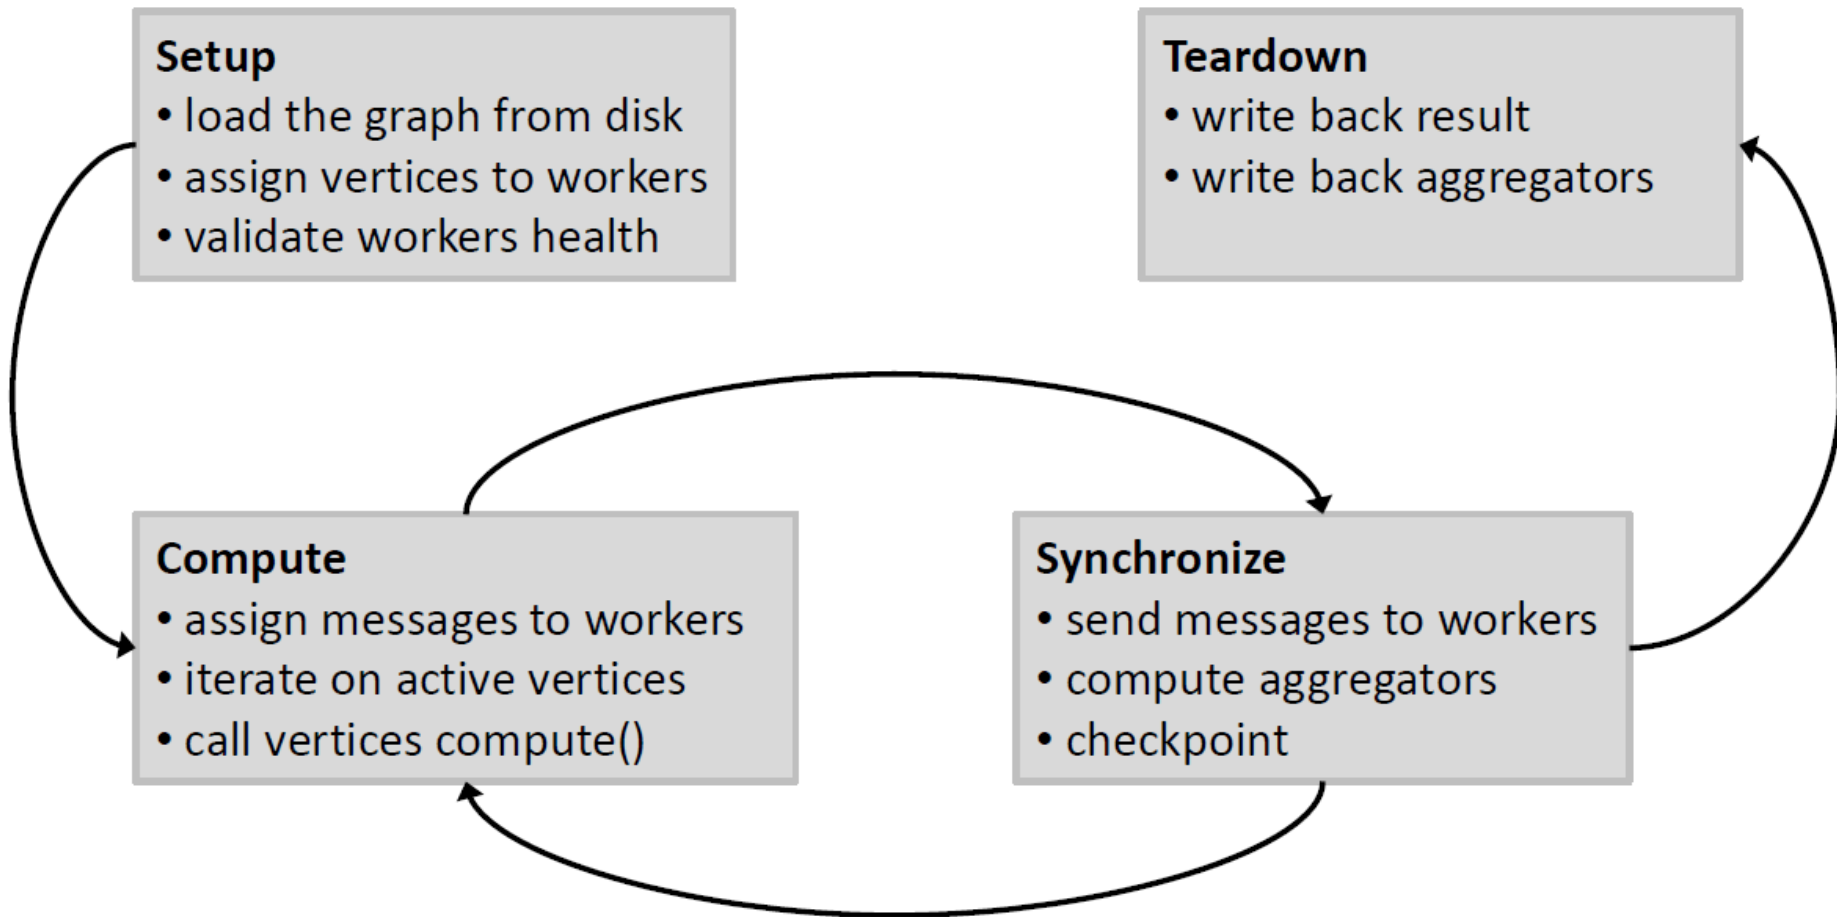
- ▶ *<https://highlyscalable.wordpress.com/2012/02/01/mapreduce-patterns/>*

Pregel: BSP applied to distributed graphs


- ▶ Apache Giraph: an open source implementation of Pregel
- ▶ graph is distributed across several machines
- ▶ computation “node” is a vertex



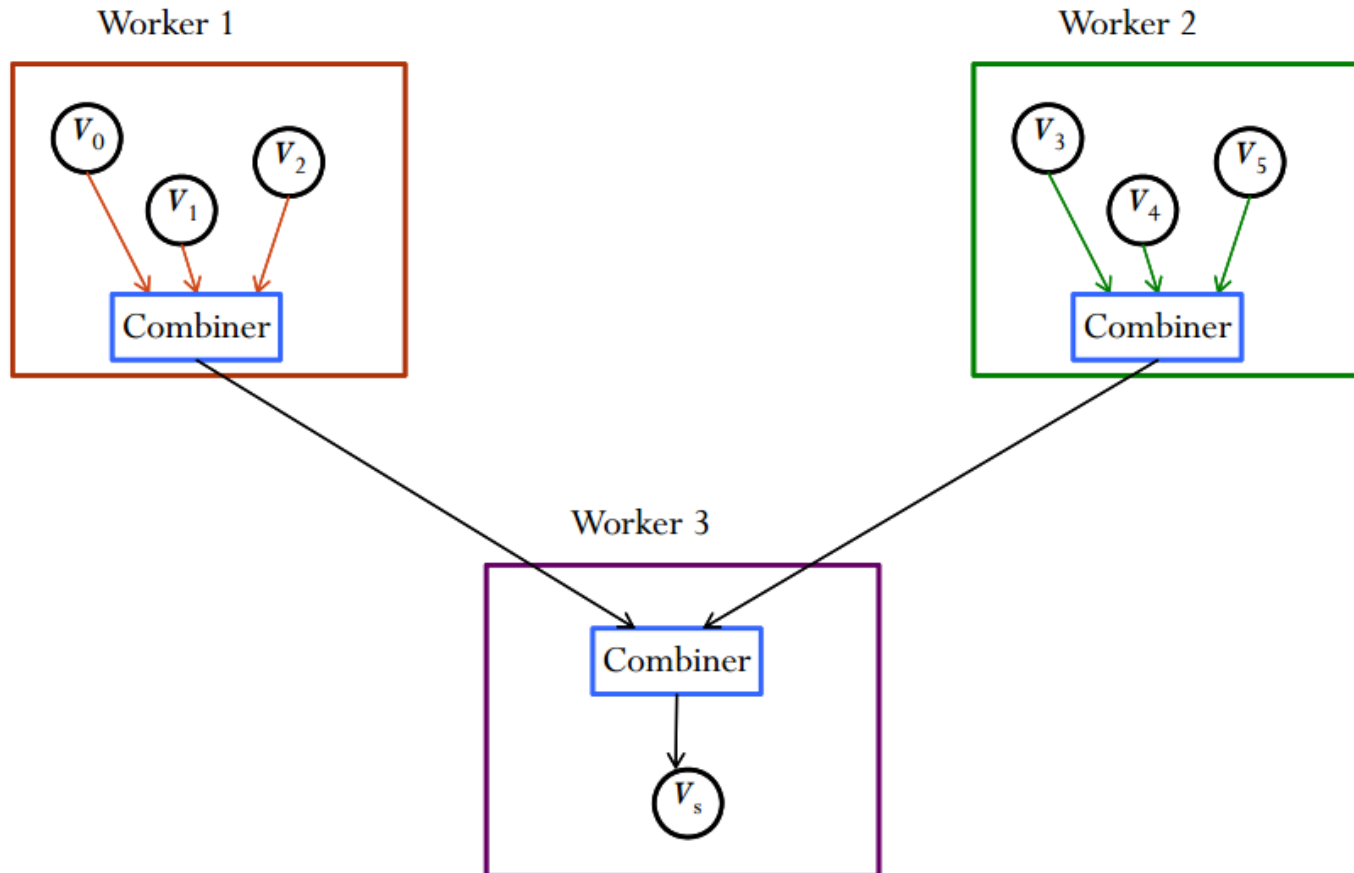
Apache Giraph



Combiners (user-defined)



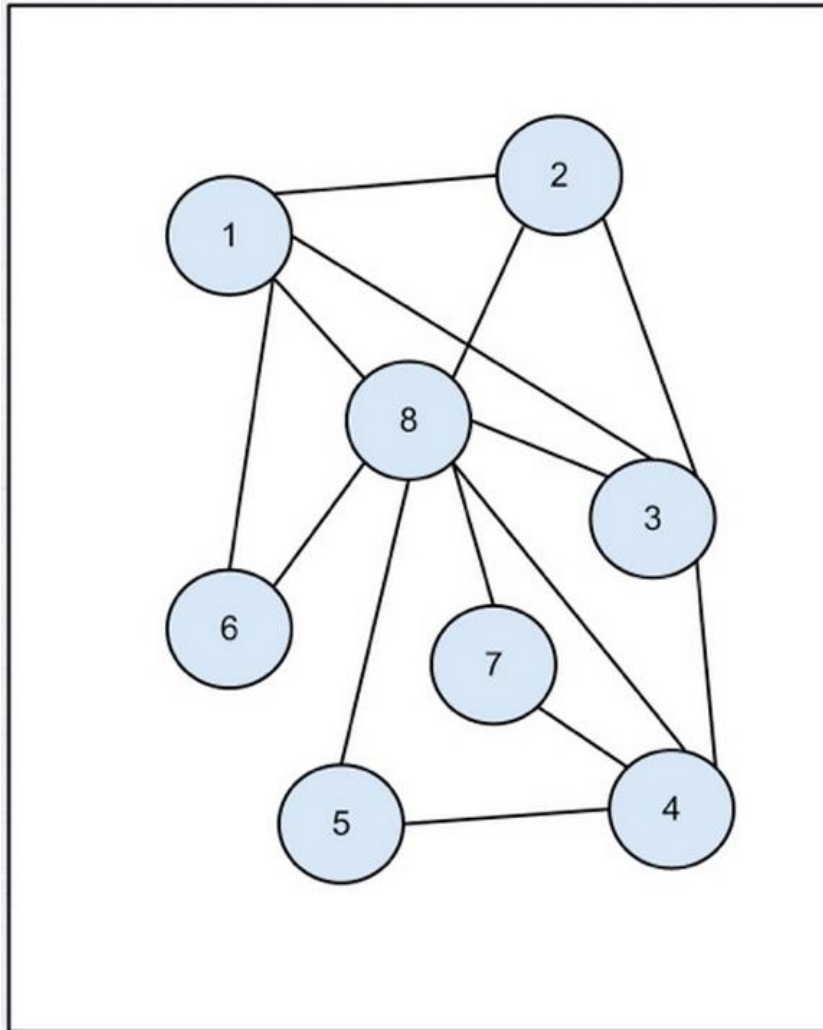
Combiners



main goal: reduce network bandwidth / number of supersteps



Aggregators



4, 3, 4, 4, 2, 2, 2, 7



Aggregator

8 Nodes


Average

Node

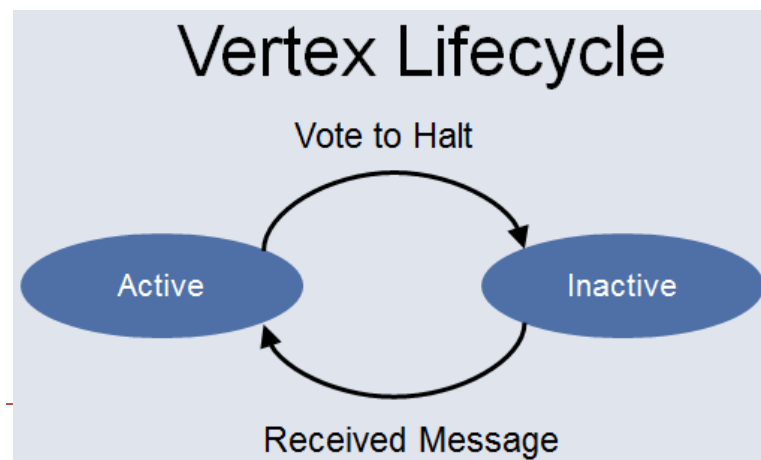
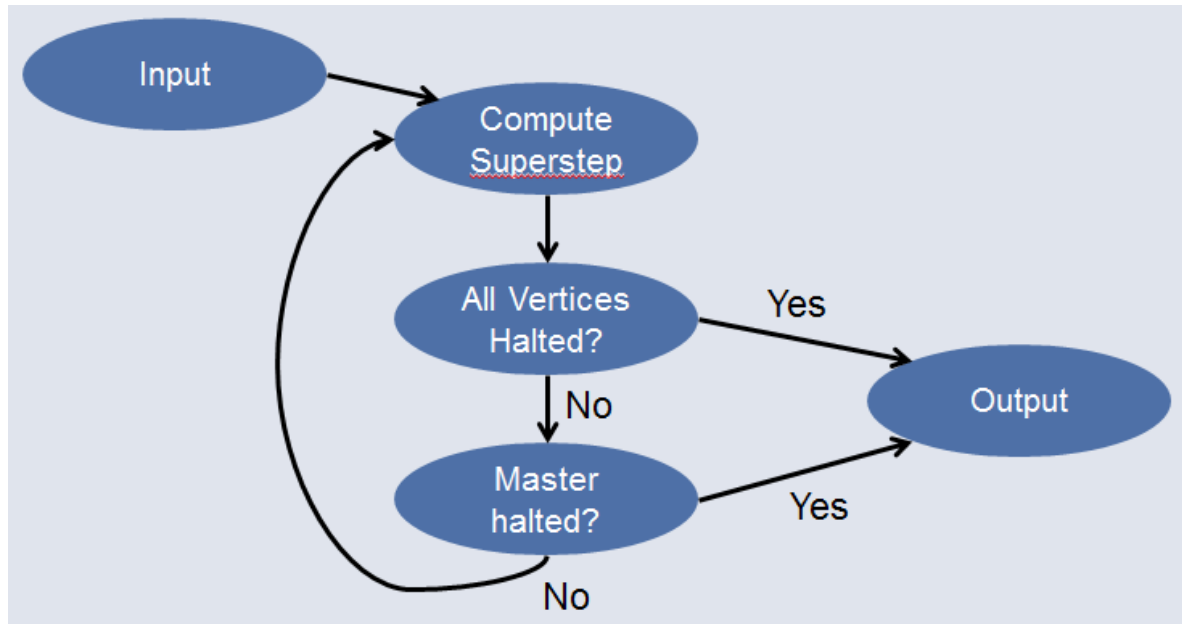
Degree: 3,5

Aggregators

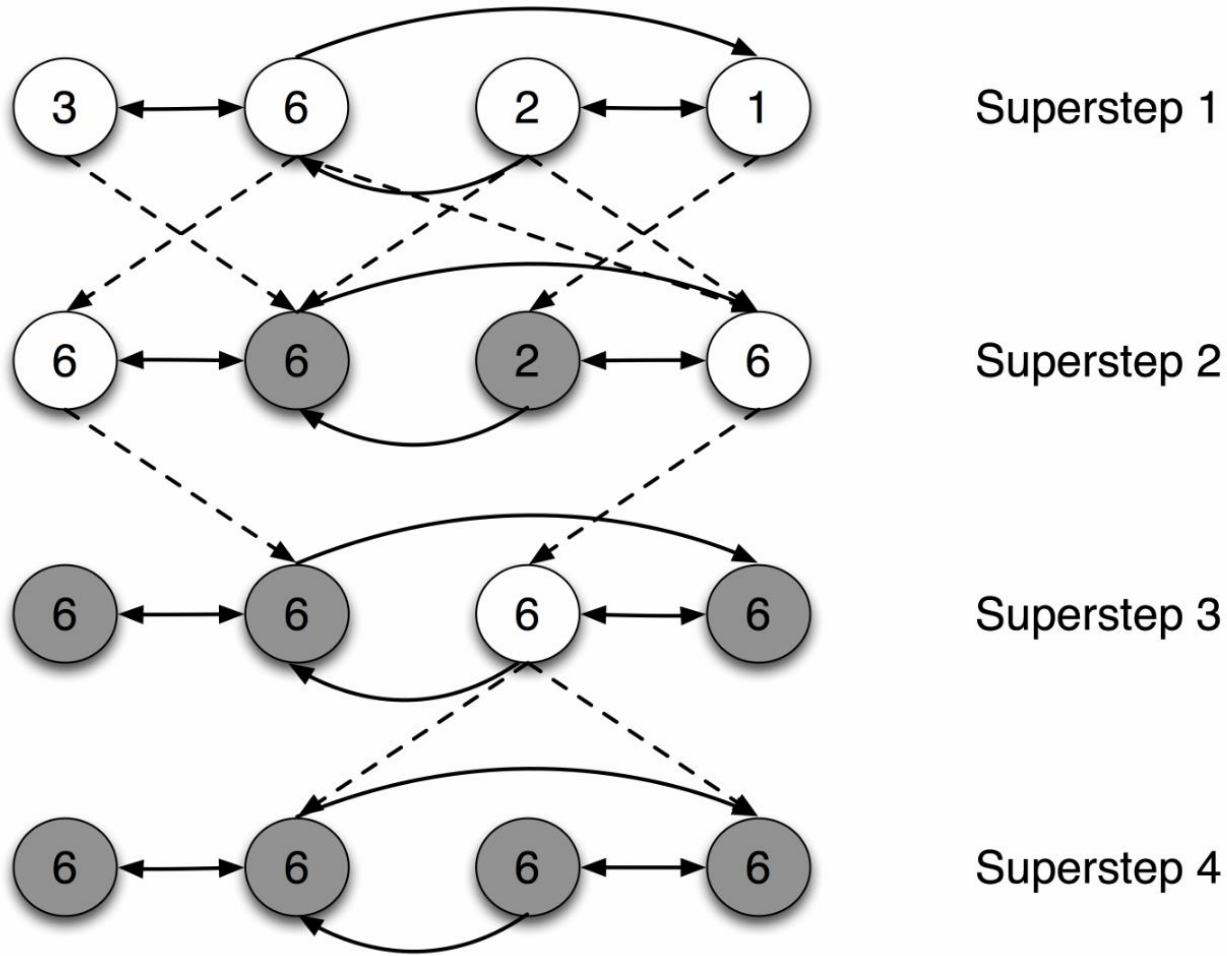
- each worker aggregates values from its vertices
- partially aggregated values are aggregated in a tree structure




Apache Giraph Lifecycle




Finding max



Spark GraphX: Resilient Distributed Graphs




Edge Cut



Vertex Cut

vertex vs. edge cutting



Spark GraphX: Resilient Distributed Graphs

```
// Vertex collection
class VertexRDD[VD] extends RDD[(VertexId, VD)]

// Edge collection
class EdgeRDD[ED] extends RDD[Edge[ED]]
case class Edge[ED](srcId: VertexId = 0, dstId: VertexId = 0,
                    attr: ED = null.asInstanceOf[ED])


// Edge Triple
class EdgeTriplet[VD, ED] extends Edge[ED]
```

```
// VD: the type of the vertex attribute
// ED: the type of the edge attribute
class Graph[VD, ED] {
  val vertices: VertexRDD[VD]
  val edges: EdgeRDD[ED]
}
```

```
class Graph[VD, ED] {
  def mapVertices[VD2](map: (VertexId, VD) => VD2): Graph[VD2, ED]

  def mapEdges[ED2](map: Edge[ED] => ED2): Graph[VD, ED2]

  def mapTriplets[ED2](map: EdgeTriplet[VD, ED] => ED2): Graph[VD, ED2]
}
```



Streaming (unbounded)

- ▶ You cannot get answers to some questions
 - ▶ e.g., what is the average of all elements?
 - ▶ you can answer them for a given subset though
 - ▶ last 20 elements, last 20 seconds
 - ▶ effectively, generating a new stream
 - ▶ not necessarily after each 20 elements
- ▶ Two approaches:
 - ▶ Microbatching (combine elements and then use existing tools)
 - ▶ e.g., Spark
 - ▶ additional buffering may help better distribute load
 - do not have a convincing example (though, you can think of some)
 - ▶ Process-each-element
 - ▶ e.g., Twitter Storm, Apache Flink
 - ▶ response time/availability of the first processed element is better



Streaming frameworks



Apache Storm

- True streaming, low latency - lower throughput
- Low level API (Bolts, Spouts) + Trident



Spark Streaming

- Stream processing on top of batch system, high throughput - higher latency
- Functional API (DStreams), restricted by batch runtime



Apache Samza


- True streaming built on top of Apache Kafka, state is first class citizen
- Slightly different stream notion, low level API



Apache Flink


- True streaming with adjustable latency-throughput trade-off
- Rich functional API exploiting streaming runtime; e.g. rich windowing semantics

Enterprise Integration



point-to-point communication

Enterprise Integration

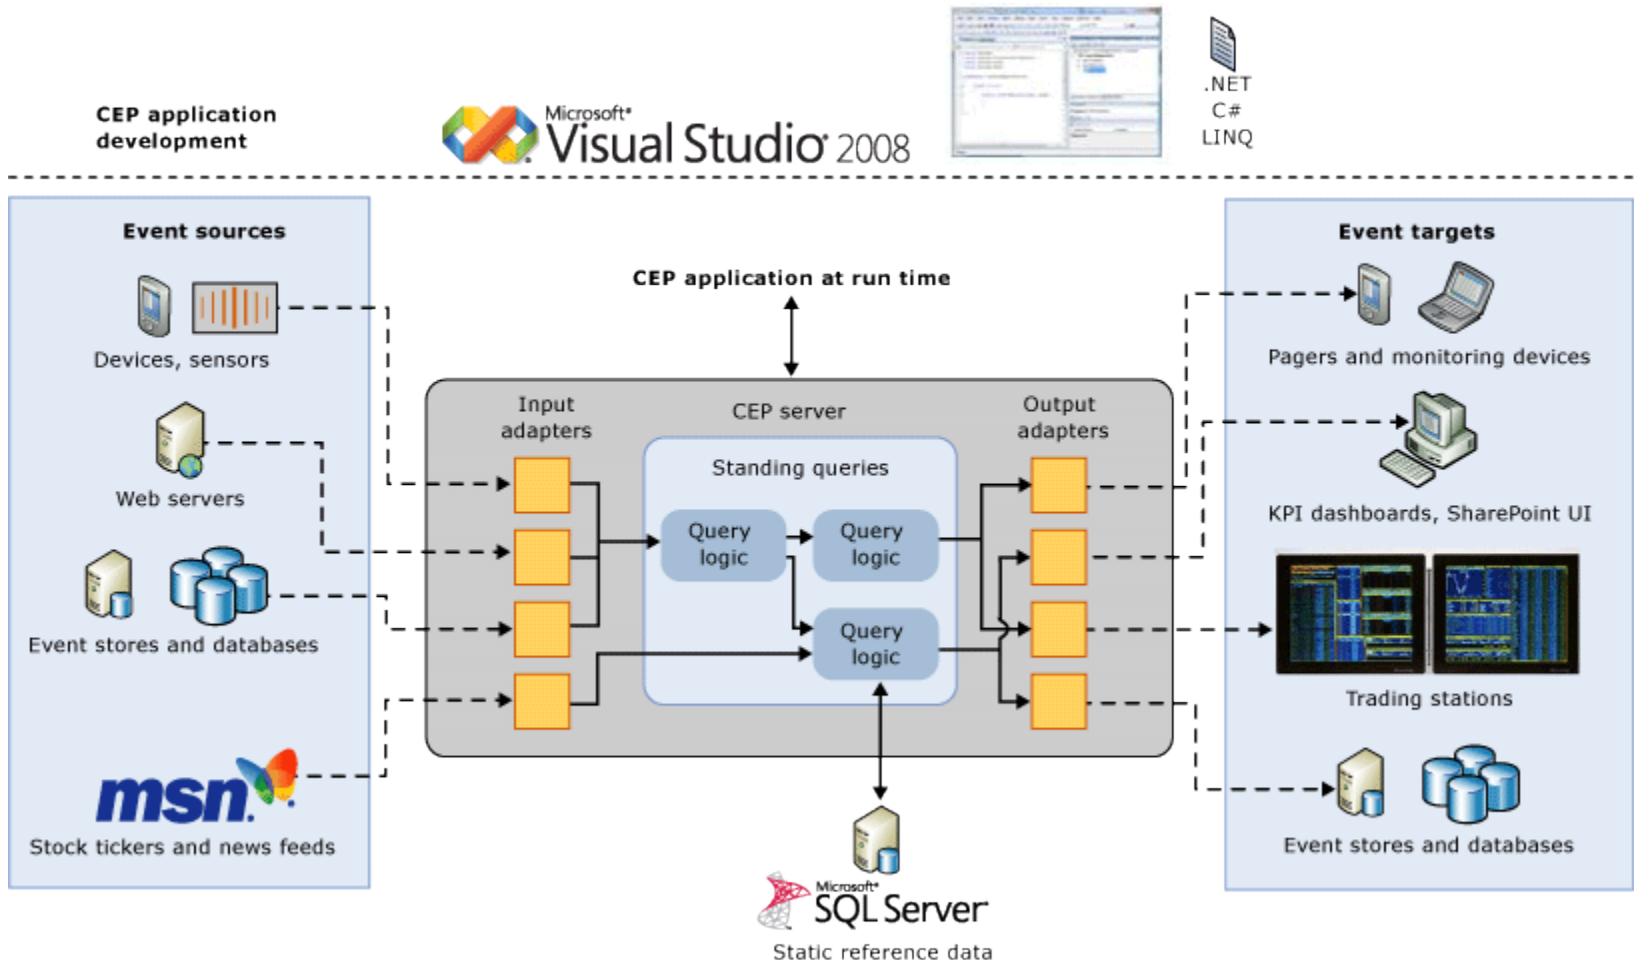


enterprise service bus

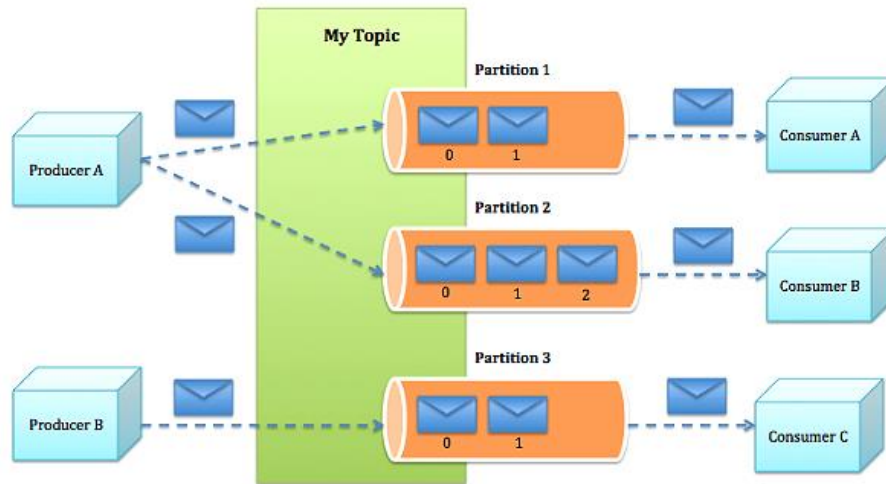
consider integration across organizations

Complex Event Processing (CEP)


StreamInsight platform



Kafka

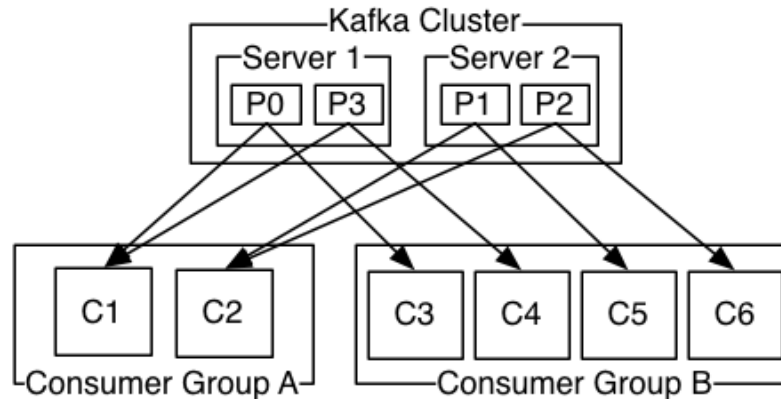


Anatomy of a Topic





- partitions are immutable sequences
 - and stored (on disk) for a configurable retention period
- producer is responsible to choose topic/partition
- within a partition, order is preserved
- partitions are replicated over several servers (one acting as a leader)

Kafka



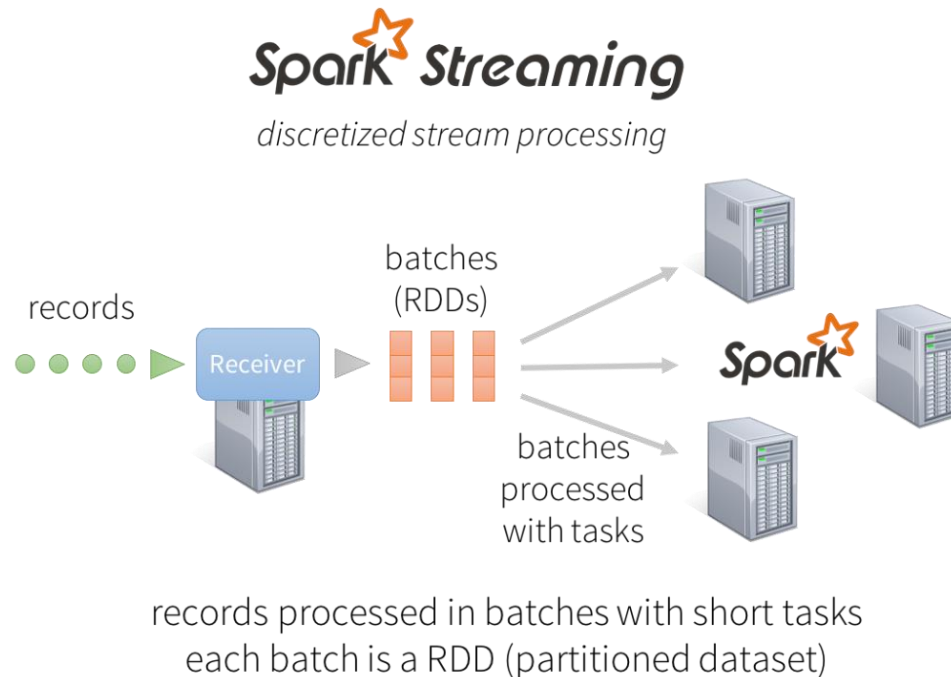
- each record is delivered to one consumer from each consumer group
 - one consumer group => load balancing
 - consumer group per consumer => broadcasting
- each consumer is assigned a partition exclusively
 - one partition => max one consumer per consumer group

Kafka Streaming API



map, filter, reduceByKey, ...

Spark Microbatching approach



- ▶ DStream is a stream of RDDs
 - ▶ each element (RDD) representing one batch

Windows and Sliding



Sensor → , 9, 6, 8, 4, 7, 3, 8, 4, 2, 1, 3, 2, → out

Sensor → , 9, 6, 8, 4, 7, 3, 8, 4, 2, 1, 3, 2, →

rolling
sum → , 57, 48, 42, 34, 30, 23, 20, 12, 8, 6, 5, 2, → out



Windows and Sliding



Tumbling vs hopping vs sliding windows

- ▶ Tumbling
 - ▶ non-overlapping windows
- ▶ Hopping
 - ▶ overlapping windows with a fixed jump-ahead period
- ▶ Sliding
 - ▶ overlapping “continuous” window
 - ▶ all possible windows of a given size

Reactive Streams

- ▶ (Akka) Reactive Streams
 - ▶ good for streaming-like APIs

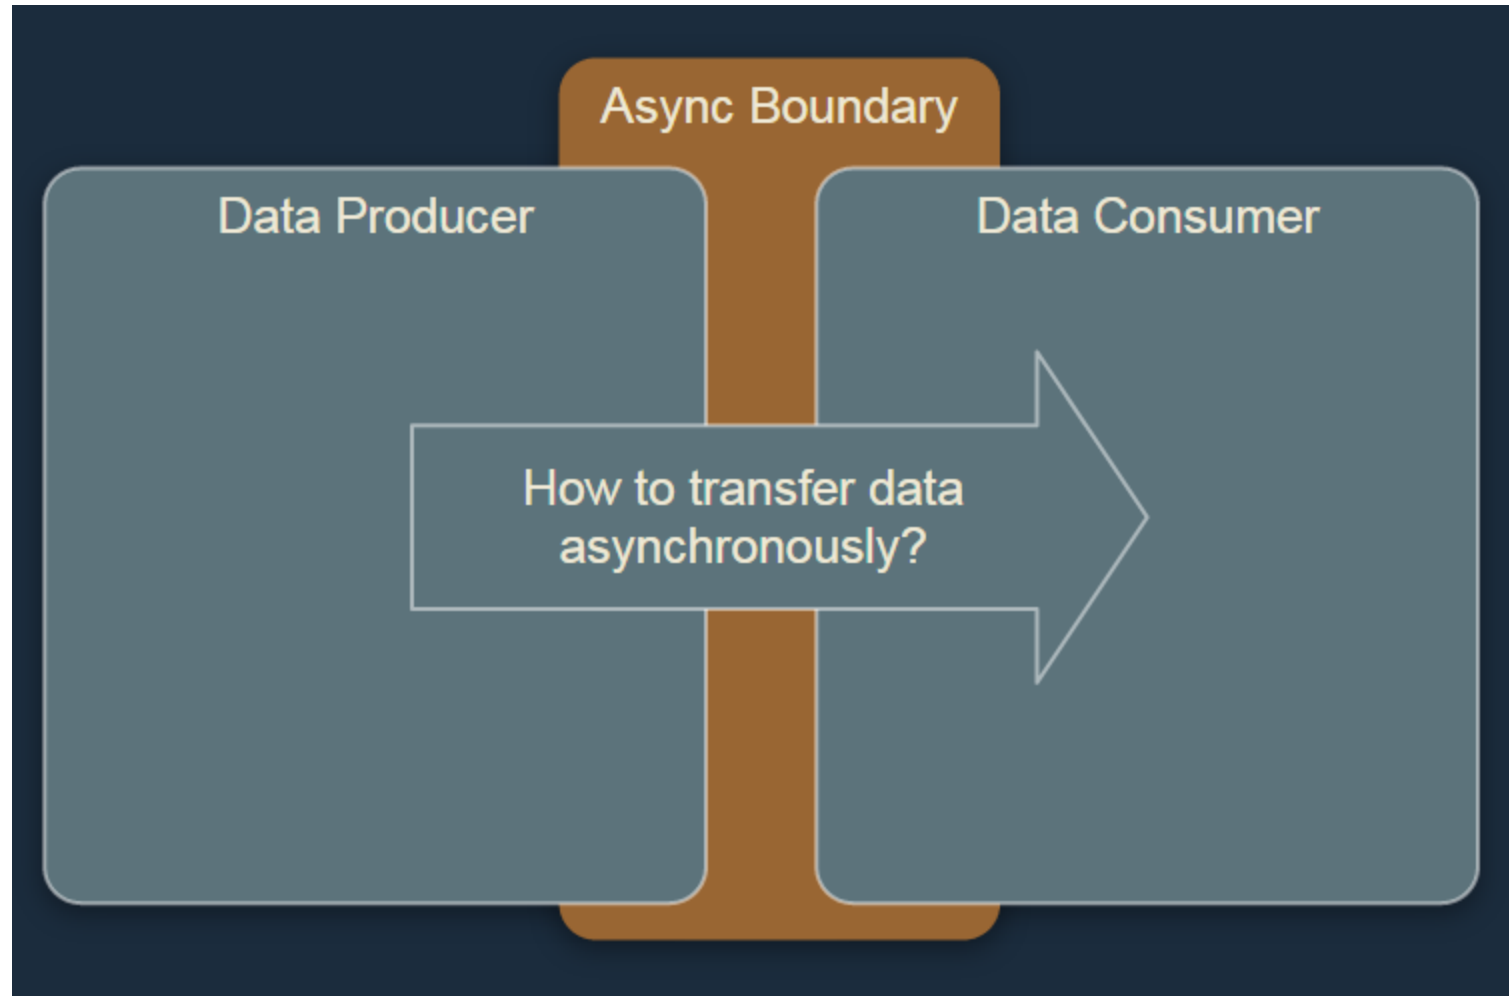
```
public interface Subscriber<T> {  
    public void onSubscribe(Subscription s);  
    public void onNext(T t);  
    public void onError(Throwable t);  
    public void onComplete();  
}
```

- ▶ Goal:
 - let the consumer explicitly tell the producer
“I am ready to accept new data”



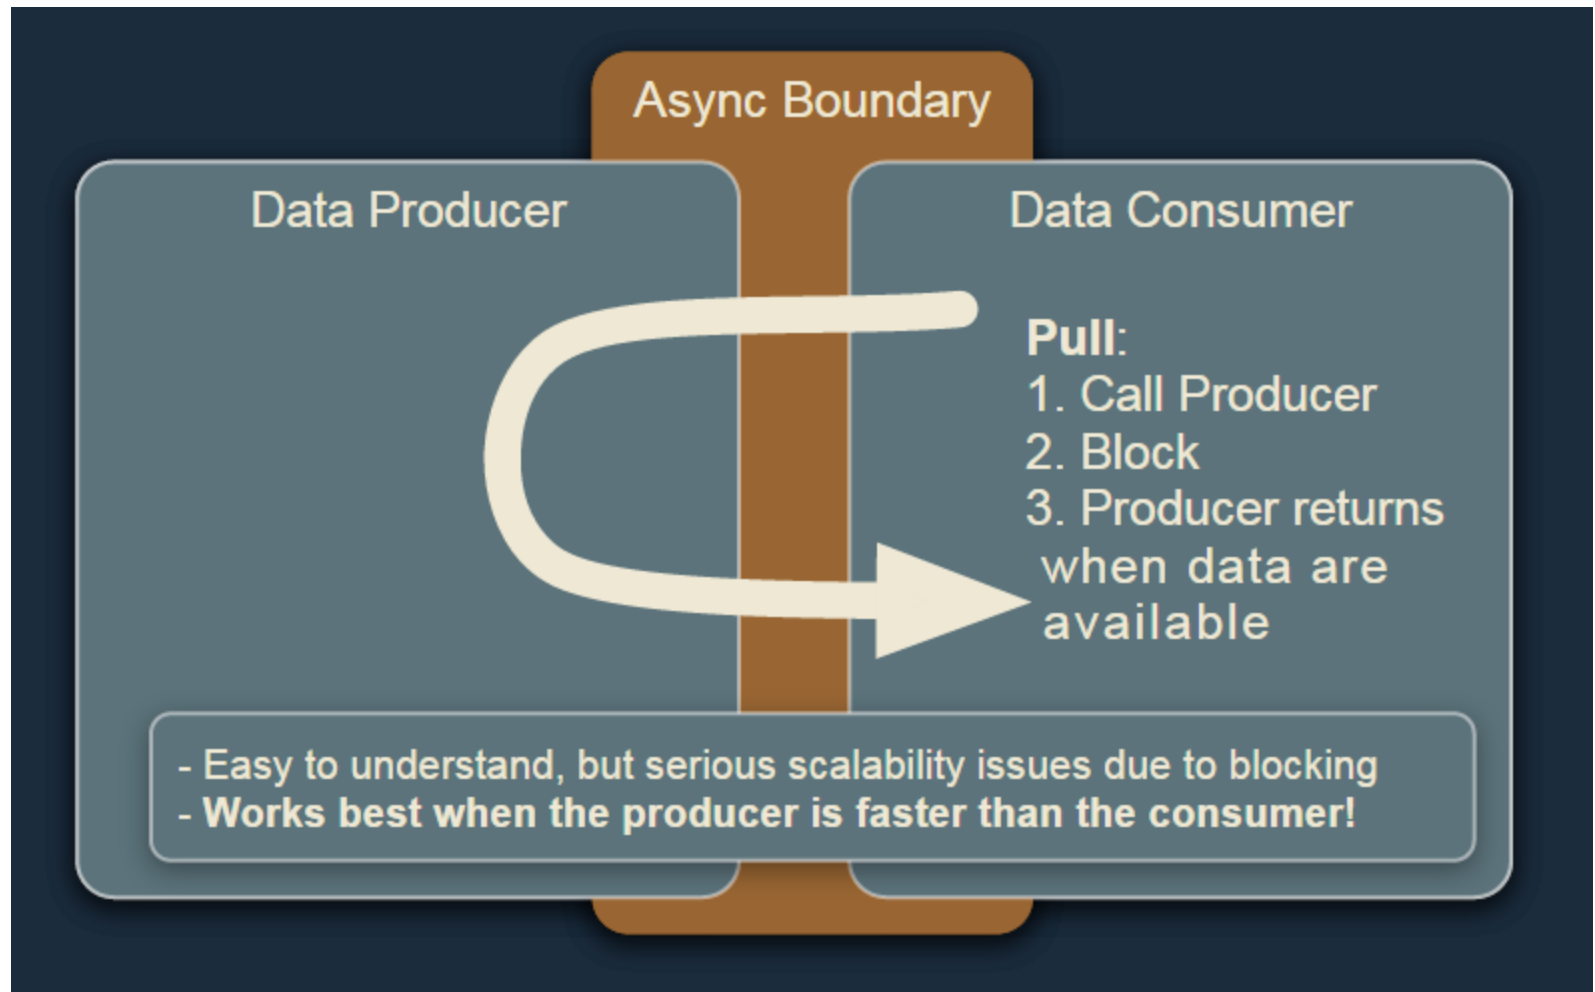
Slides from Mathias Doenitz

(spray.io, akka-http)



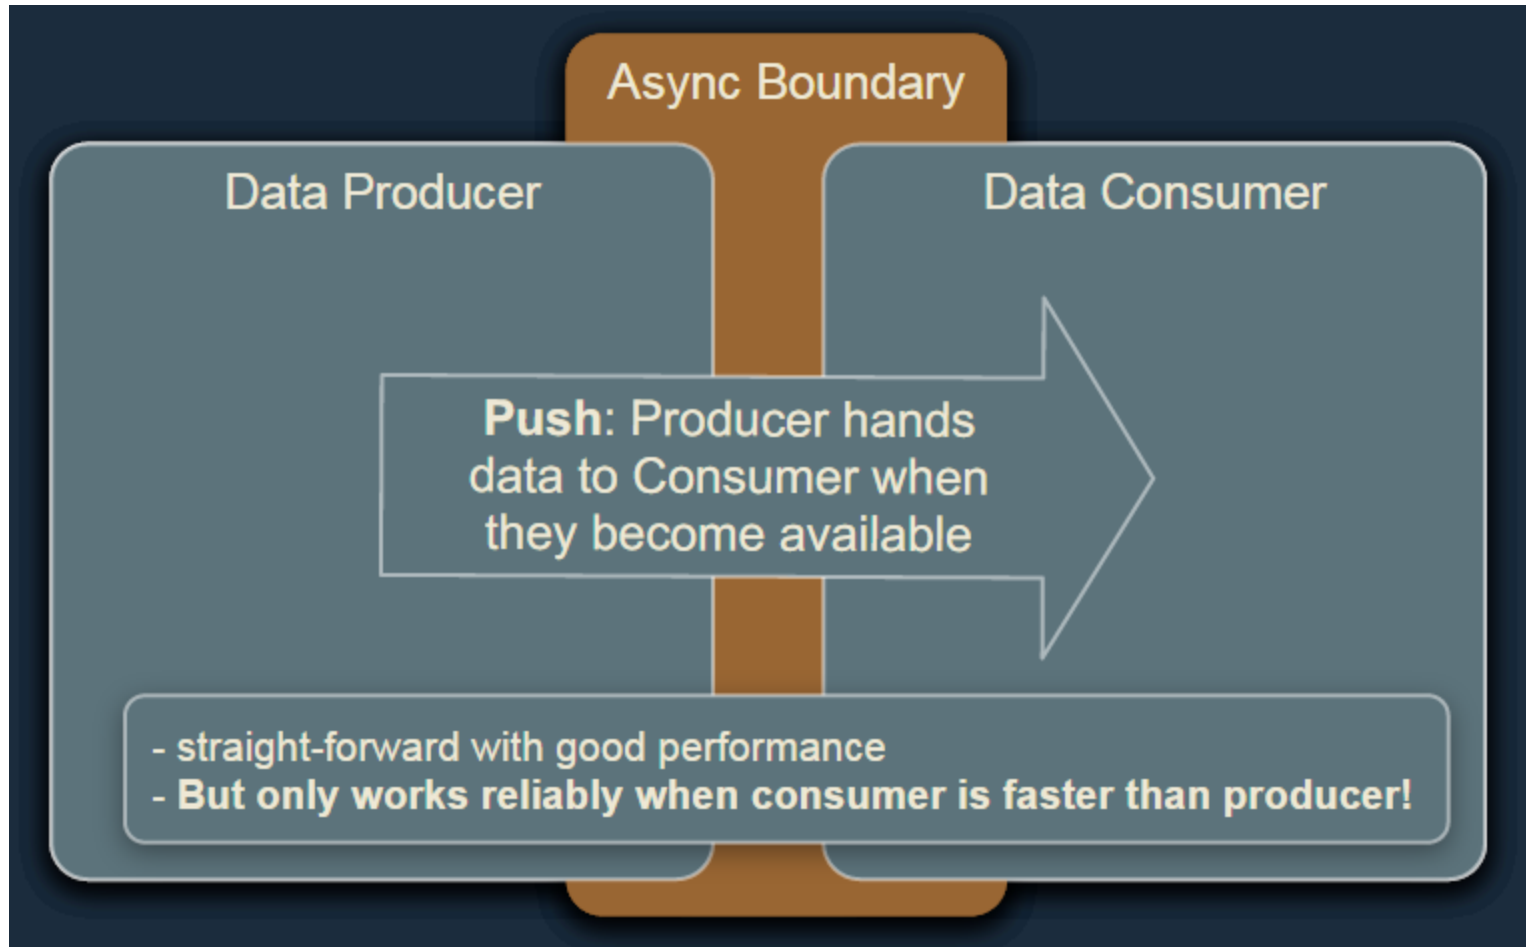
Slides from Mathias Doenitz

(spray.io, akka-http)



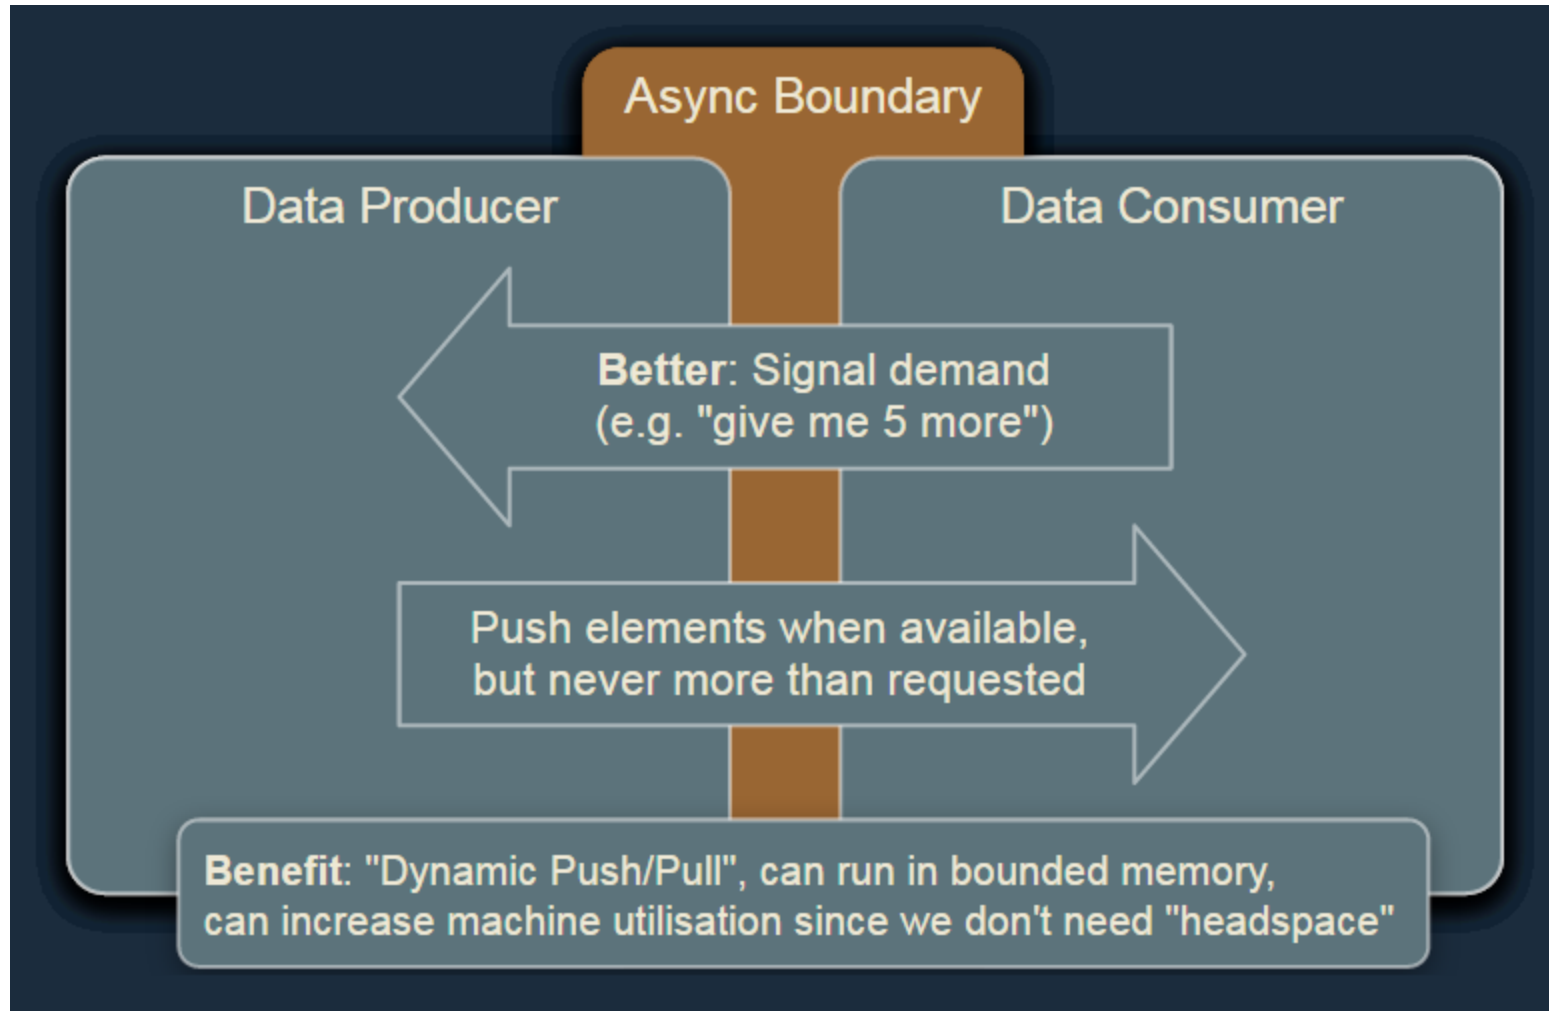
Slides from Mathias Doenitz

(spray.io, akka-http)

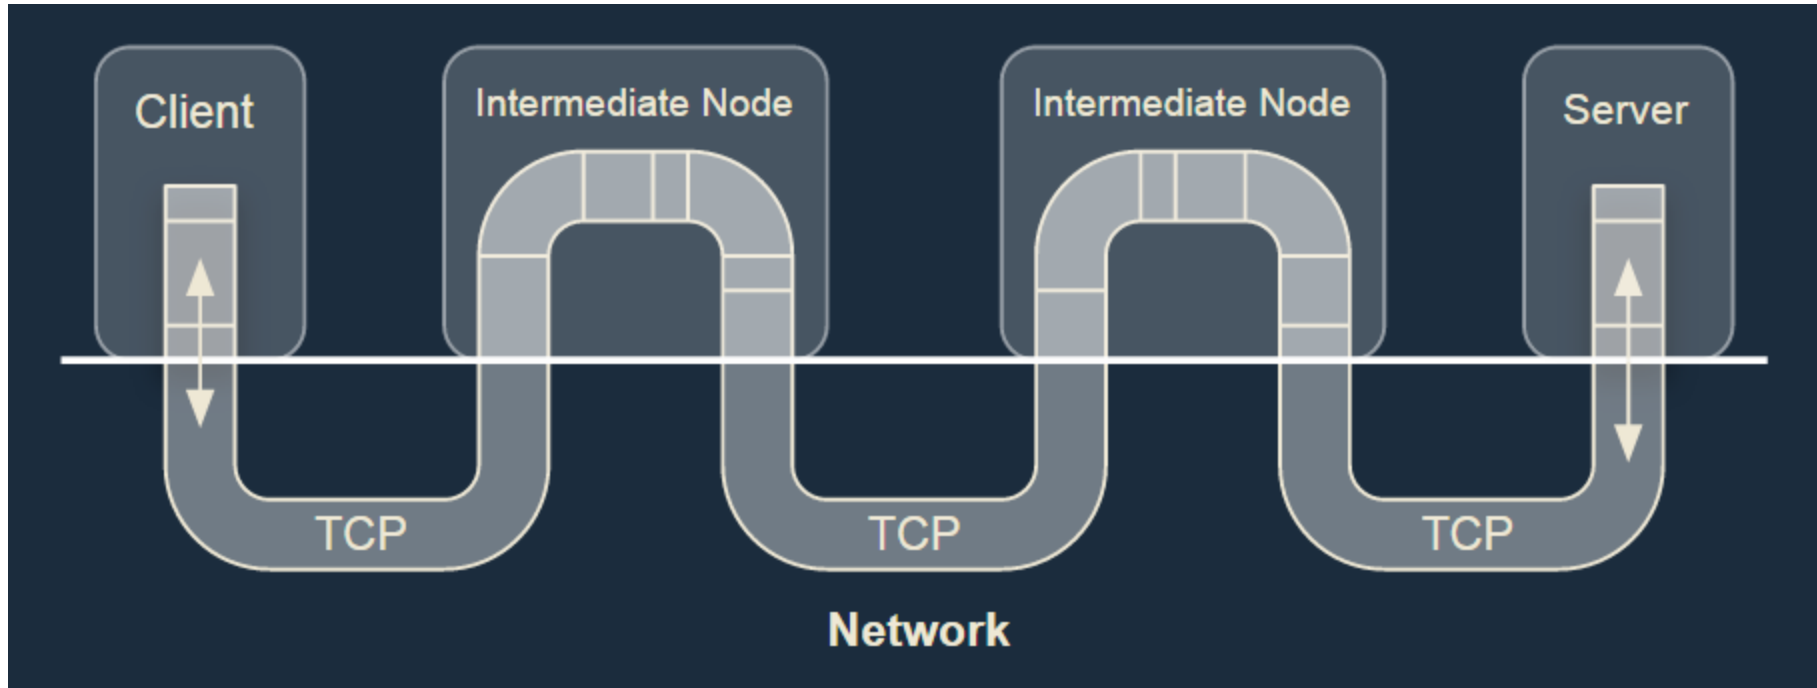


Slides from Mathias Doenitz

(spray.io, akka-http)

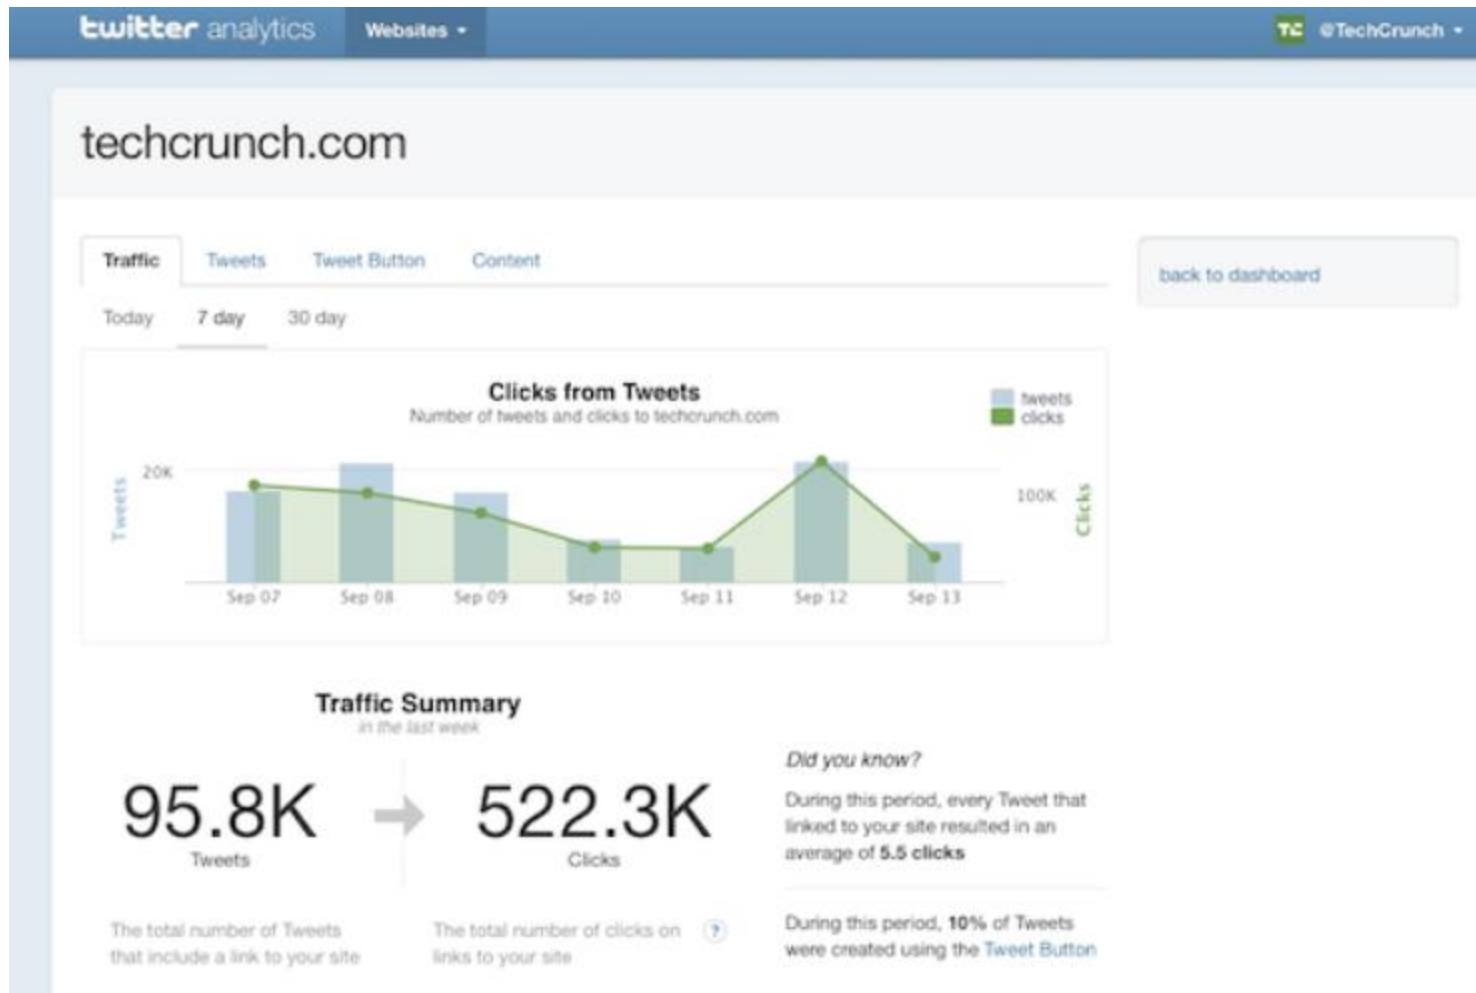


Slides from Mathias Doenitz (spray.io, akka-http)




full resource utilization if used across the whole pipeline!


Storm @ Twitter




Use cases



Stream
processing



Distributed
RPC



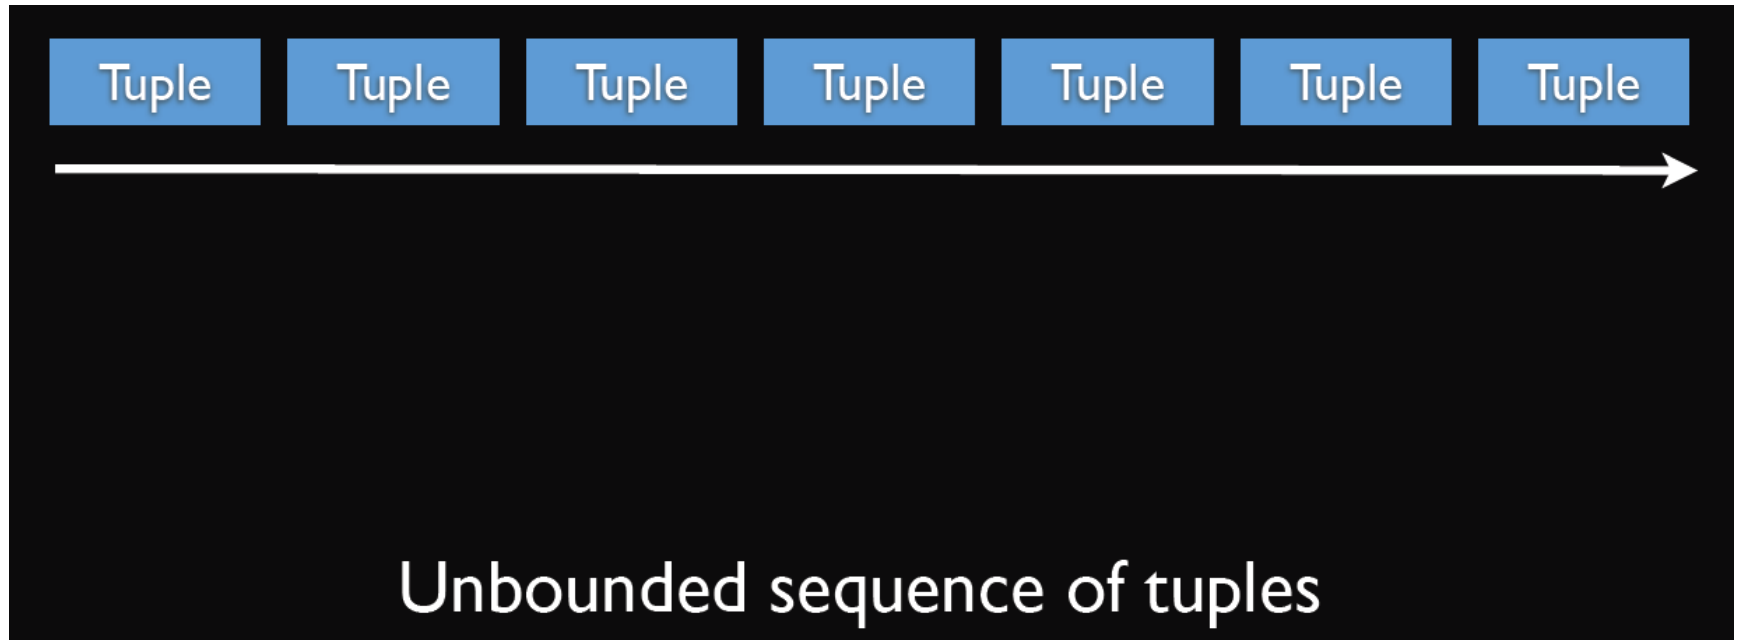
Continuous
computation

Concepts


- Streams
- Spouts
- Bolts
- Topologies




Streams



Spouts



Bolts




Processes input streams and produces new streams

Bolts


- Functions
- Filters
- Aggregation
- Joins
- Talk to databases



Topology



Tasks

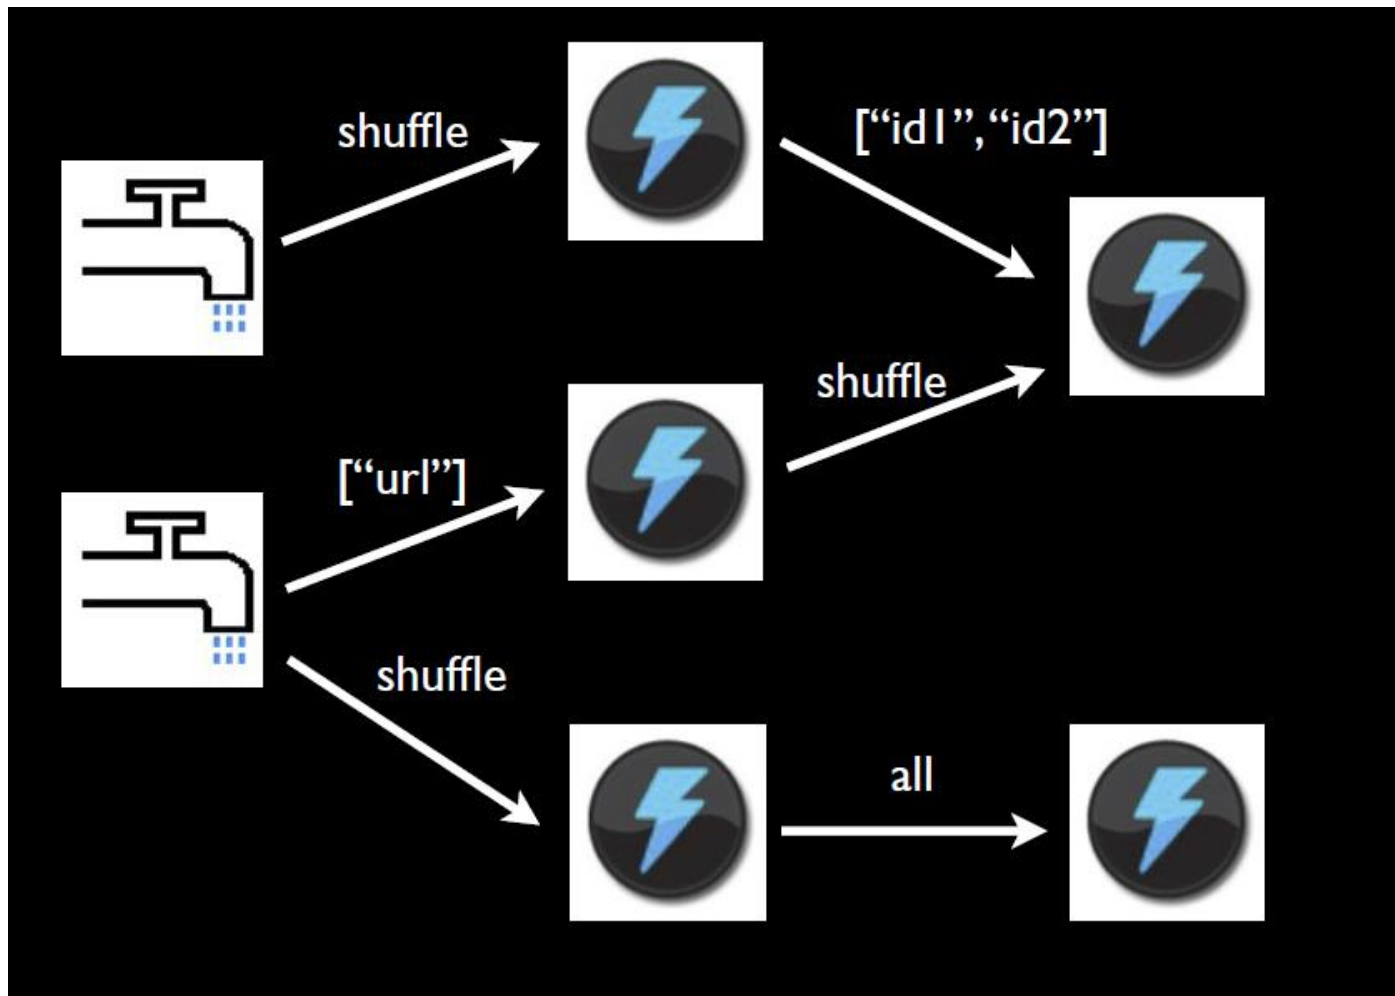


Spouts and bolts execute as many tasks across the cluster

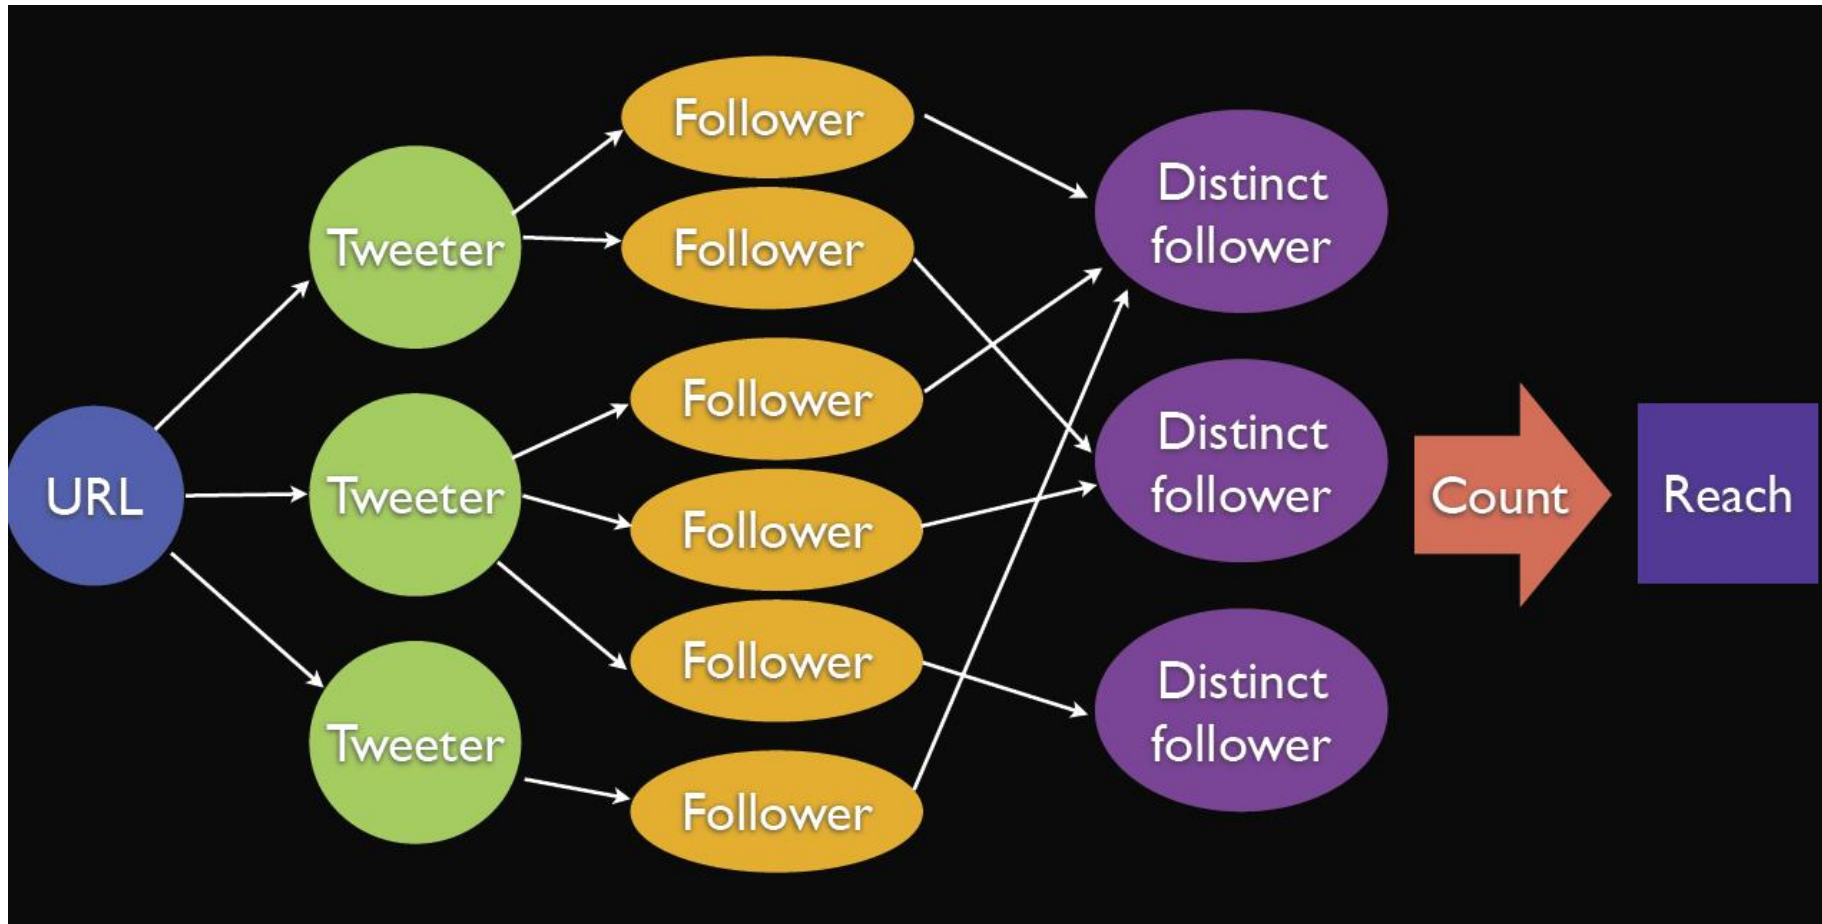
Stream grouping

- **Shuffle grouping:** pick a random task
- **Fields grouping:** consistent hashing on a subset of tuple fields
- **All grouping:** send to all tasks
- **Global grouping:** pick task with lowest id

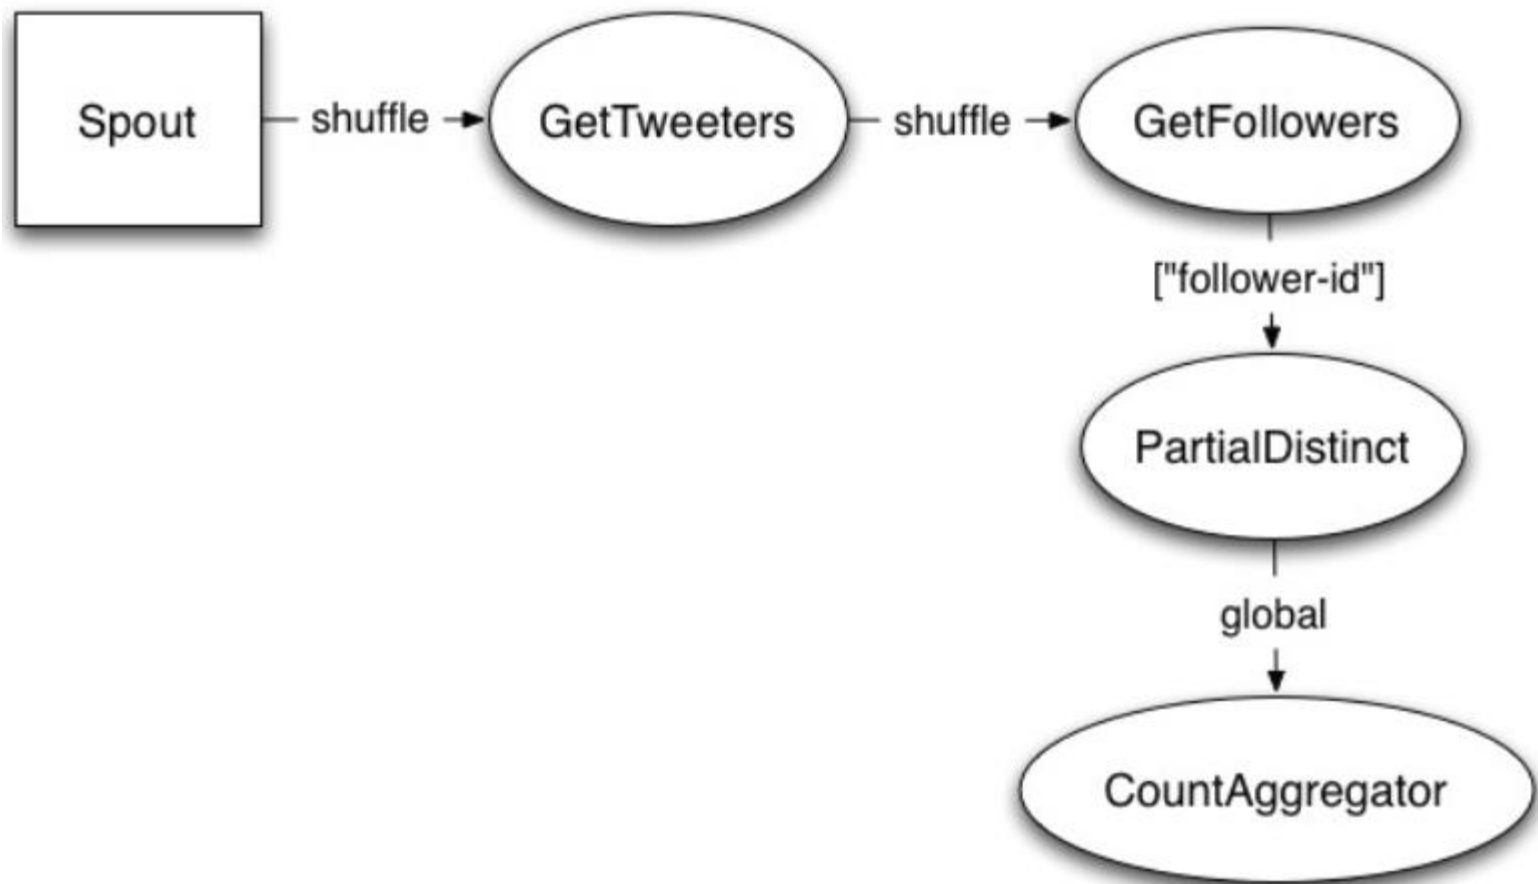
Filtering and grouping



Sample Application



Filtering and grouping



Cluster Coordination

