

*Article*

ISPRS International Journal of

# An Efficient Parallel Algorithm for Multi-Scale Analysis of Connected Components in Gigapixel Images

Michael H.F. Wilkinson<sup>1,\*</sup>, Martino Pesaresi<sup>2,†</sup> and Georgios K. Ouzounis<sup>3,†</sup>

<sup>1</sup> Johann Bernoulli Institute for Mathematics and Computer Science, University of Groningen, P.O. Box 407, 9700 AK Groningen, The Netherlands

<sup>2</sup> Global Security And Crisis Management Unit, Institute for the Protection and Security of the Citizen, Joint Research Centre, European Commission, via Enrico Fermi 2749, I-21027 Ispra (VA), Italy

<sup>3</sup> DigitalGlobe, Inc., 1300 W 120th Ave, Westminster, CO 80234, USA

<sup>†</sup> These authors contributed equally to this work.

\* Author to whom correspondence should be addressed; e-mail: m.h.f.wilkinson@rug.nl; tel: +31 50 363 8140; and fax: +31 50 363 3800.

Version March 28, 2017 submitted to *Ijgi*. Typeset by *LATEX* using class file *mdpi.cls*

<sup>1</sup> **Abstract:** Differential Morphological Profiles (DMPs) and their generalized Differential Attribute Profiles (DAPs) are spatial signatures used in the classification of earth observation data. The Characteristic-Salience-Leveling or CSL is a model allowing to compress and store the multi-scale information contained in the DMP and DAP into raster data layers, used for further analytic purposes. Computing DMPs or DAPs is often constrained by the size of the input data and scene complexity. Addressing very high resolution remote sensing gigascale images, this paper presents a new concurrent algorithm based on the Max-Tree structure that allows the efficient computation of CSL. The algorithm extends the “one-pass” method for computation of DAPs, and delivers an attribute zone segmentation of the underlying trees. The DAP vector field and the set of multi-scale characteristics are computed separately and in a similar fashion to concurrent attribute filters. Experiments on test images of 3.48 to 3.96 Gpixel showed an average computational speed of 59.85 Mpixel per second, or 3.59 Gpixel per minute on a single 2U rack server with 64 opteron cores. The new algorithms could be extended to morphological keypoint detectors capable of handling gigascale images.

<sup>15</sup> **Keywords:** Differential attribute profile; connected filters; spatial signature; CSL model;  
<sup>16</sup> image decomposition; giga-pixel images

## 17 1. Introduction

18 Multi-scale analysis is pivotal in both human and computer vision [1–6]. This is in part due to the fact  
19 that features of importance might be present at a range of scales, depending on distance to the observer or  
20 camera. Conceptually, multi-scale analysis maps the image into a higher-dimensional space or stack of  
21 images. The problem then is to identify the most important or salient features in this scale space [5,7,8].

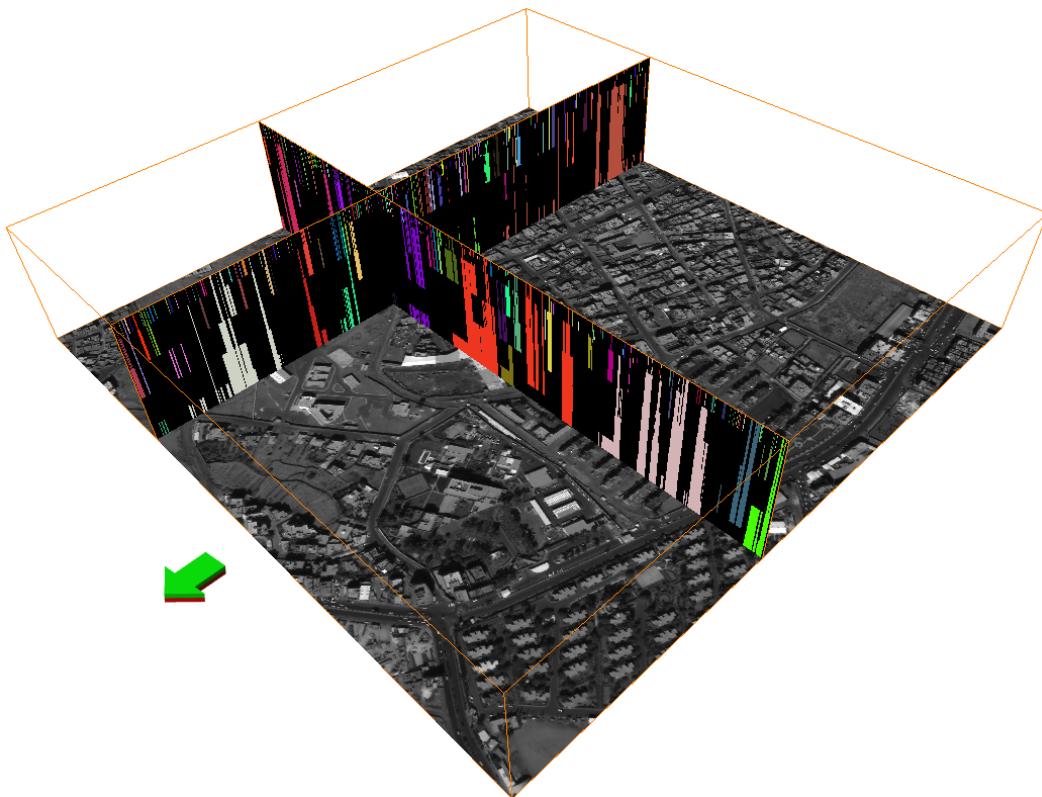
22 Multi-scale methods based on connected morphological filters [1,2,9–11] of various types have  
23 been useful in remote sensing applications, both in multi-band processing and in gray-scale (typically  
24 panchromatic) images. In the former case alpha-trees [12] or other partition hierarchies are used  
25 most [13,14], whereas in the latter methods based on either reconstruction or attribute filters are most  
26 common. The most popular of these gray-scale methods are Differential Morphological Profiles (DMP)  
27 [7,8], Differential Area Profiles [15,16], and their generalisation Differential Attribute Profiles (DAP)  
28 [17,18].

29 In this paper we will focus on these latter gray-scale methods and review recent developments in  
30 this field, in particular adaptation of algorithms for parallel computation of multiscale representations.  
31 Besides, we will introduce a number of improvements to existing algorithms, and compare the  
32 computational performance and memory load of these methods. A common factor limiting the usability  
33 of both DMPs and DAPs is the lack of efficient means for handling the massive input-data size. The  
34 memory load and computation-time of standard iterative methods increase linearly with respect to both  
35 input image size and the number of scales used [8,16].

36 DMPs [7,8] are sets of top-hat scale-spaces based on openings by reconstruction. Traditionally, they  
37 can be computed by first performing a series of erosions with an increasing set of structuring elements  
38 (SEs), computing a geodesic reconstruction from each of these marker images, and finally taking the  
39 differences between successive reconstructions, where the reconstruction of the smallest SE is subtracted  
40 from the original. Thus a stack of images is obtained, representing a top-hat scale space, each of which  
41 stores bright details of a certain width range, dictated by the widths of the SEs used. To obtain the  
42 dark details, a similar process is performed after first inverting the input image. The DMPs have large  
43 spectrum of applications in the field of Earth and planetary science, computer science, engineering,  
44 physics and astronomy, and mathematics. In the field of remote sensing they are mostly applied for  
45 image feature extraction phase and filtering or knowledge-driven image information selection purposes.  
46 The output of DMPs are typically used as input of supervised or unsupervised classification processes.  
47 The computation of DMPs in the classical way involves multiple reconstruction operations, which is  
48 costly. In [8] it was shown that the process could be speeded up by (i) using a single Max-tree [19] to  
49 compute all reconstructions in a single pass, and (ii) using parallel processing based on the algorithm in  
50 [20].

51 Differential attribute profiles [15,17,21] are generalizations of DMPs based on connected attribute  
52 filters [22]. Instead of a series of openings-by-reconstruction, a series of attribute filters are used which  
53 form either a size or shape granulometry [23,24]. Again differences between consecutive filtered images  
54 are taken, to create a stack of images containing different size or shape classes of details, and again, this  
55 stack of images is usually summarized in a single multi-band image. An example of sets of DAPs in the

56 cross-section of an area opening top hat scale space and an area closing bottom hat scale space is shown  
 57 in Fig. 1.



**Figure 1.** Response vectors from the cross-sections of a DAP vector field. The cross-section facing the direction of the arrow corresponds to positive, and the other to negative response vectors, i.e. describing bright and dark information respectively.

58 Applications of DAPs vary and typical examples are scene classification [18,25], image  
 59 segmentation [26,27], urban pattern characterisation [28–32] and human settlement visualisation [33].  
 60 Note that both the DAP and DMP can be considered alternatively as a set of images at different scales, or  
 61 as a vector image, in which each pixel is represented by a vector of values corresponding to the response  
 62 in each morphological filter band.

63 The CSL is a general model suitable for compact representation of multi-scale image decomposition  
 64 schema, including DMPs and DAPs. It consists of three raster layers derived from the image multi-scale  
 65 decomposition; the Characteristic scale (C), the Saliency (S) and the Leveling (L). The basic components  
 66 of the CSL (C,S) were firstly introduced in [34] supporting a classification schema made with a neural  
 67 supervised classifier of image connected components (also called regions or objects) and their attributes.  
 68 Subsequently, they were presented in [7,35] for general image segmentation purposes. In [36] the DMP  
 69 concept was reintroduced with the name of "ultimate opening" and the same C,S elements of the DMP  
 70 were presented with the name of parameters of maximal residue  $v$  and the size  $q$  of the opening leading  
 71 to this residue. With the latter naming they have been applied for segmentation of text regions in digital  
 72 images in [37] and shape recognition tasks in [38]. In [33,39] the CSL schema was completed with  
 73 the explicit introduction of the leveling (L) descriptor and applied for reduction of the dimensionality  
 74 and the memory storage/allocation cost of the image multi-scale decomposition. In particular, the need

75 of performing massive unsupervised image processing tasks in a statistical-model-free approach, i.e.  
76 avoiding clustering based on the statistical distribution of the DMP/DAP features, was addressed. The  
77 image data scale-decomposed and described by the CSL schema can be used for knowledge-driven  
78 image information mining and selection [39,40], supervised and unsupervised classification [41–44] and  
79 visualization purposes [39,45].

80 In the case of DAPs as with DMPs the problems with memory load and compute time can be addressed  
81 using Max-Trees, and their duals Min-Trees [19]. One key problem is the need to store the entire  
82 scale space, or vector field, consisting of as many images as there are scales. In the DMP case this  
83 is unavoidable, because the marker images are needed for each scale. These markers can be reused to  
84 store the DMP vector field itself, as in [8]. In the case of the DAP this can be avoided *if* all we need  
85 is a summary such as the CSL representation. Extending the work on fast computation of connected  
86 granulometries and pattern spectra in [24], in [16] a scheme was proposed that is referred to as the  
87 *attribute zone decomposition*. Tree nodes are assigned a unique zone label depending on the attribute  
88 value of the image component they associate with. Organising the tree nodes into zones requires a  
89 single pass through the tree structure and has a minimal dependency on the length of the signature. This  
90 scheme is called the *one-pass* method. A performance comparison against other iterative methods is  
91 given in [16].

92 Apart from zone labeling, the one pass method can be applied to compute the set of features that  
93 constitute the CSL model [33]. The DAP vector field, i.e. the set of DAPs accounting for the entire input  
94 image, is extracted by mapping each image component to the respective DAP plane. There are twice as  
95 many planes as the number of attribute thresholds defining the decomposition. The model is computed  
96 directly from the two tree-based data structures and with no need of exporting the DMP/DAP vector  
97 fields.

98 Building upon a parallel method for computing the Max-Tree structure on shared memory platforms,  
99 in this paper a new concurrent version of the one-pass method is presented. The algorithm is based on  
100 the concurrent algorithm for Differential Morphological Profiles [8], which in turn was based on the  
101 shared-memory parallel algorithm for component trees from [20]. We adapt the existing algorithm for  
102 DMP to the DAP, and develop this further to produce an algorithm which does not compute the entire  
103 vector field, but just the data needed for the CSL model computation. The latter is both faster and much  
104 more memory efficient. The code can handle images up to 4 Gpixel, the limit imposed by GeoTIFF,  
105 and processing on a 64-core machine shows speeds of over 3.5 Gpixel per minute. Extension to larger  
106 volumes requires only using 64-bit rather than 32-bit indexing of the arrays used, and of course a different  
107 input format.

108 The structure of this paper is as follows. Section 2 gives a brief overview of connected attribute  
109 filters, differential attribute profiles and the CSL model. In Section 3, the Max-Tree structure is  
110 revisited followed by a brief description of the one-pass method. The proposed concurrent algorithm  
111 for computing the DAPs and the CSL bands is presented in Section 4. A set of timing experiments on  
112 very high resolution satellite images follows in Section 5. A discussion and analysis of the experimental  
113 results is given in Section 6, and a summary of the work presented together with conclusions, in  
114 Section 7.

115 **2. Background**116 **2.1. Attribute Filters**

117 Attribute filters [22] are morphological filters designed to preserve or remove connected image  
 118 structures based on their properties or *attributes*. These attributes can be anything, ranging from simple  
 119 size measures such as area [46], through shape numbers [23,24] to feature vectors. In the binary case  
 120 they simply remove all connected foreground or background components based on a selection criterion,  
 121 based on the attribute values and some decision boundary.

Let  $E$  be the definition domain of a gray-scale image  $f : E \rightarrow \mathbb{R}$ , and let the information content of  $f$  be organised into a set of *flat zones* [11], the union of which, equals  $E$ . A flat zone  $F$  is a connected set of maximal extent. It consists of isotone pixels that are path-wise connected. Given a point  $x \in E$ ,  $F$  is formally defined as:

$$F_x(f) = \{x\} \cup \{y \in E \mid \exists \pi(x \rightsquigarrow y) \wedge f(y) = f(x)\}, \quad (1)$$

in which  $\pi(x \rightsquigarrow y)$  denotes a path from  $x$  to  $y$ . The set of image flat zones constitutes a partition of  $E$ . To introduce a notion of order with respect to image intensity a more general type of component is defined. A peak component marked by a point  $x \in E$  at level  $h$  is given by:

$$P_x^h(f) = \{x\} \cup \{y \in E \mid \exists \pi(x \rightsquigarrow y) \wedge f(y) \geq f(x)\}. \quad (2)$$

122 A peak component that coincides with a single flat-zone is called a regional maximum.

123 Connected filters consider connected operators that access peak components instead of individual  
 124 pixels. In particular, gray-scale attribute filters [22] are image transforms that reduce the information  
 125 content by removing peak components that fail an attribute criterion. Attribute filters are idempotent, i.e.,  
 126 re-iterating the operator does not modify the result further. Attribute filters are edge preserving operators  
 127 and can be categorised based on the intensity order they are configured with, and the properties of the  
 128 attribute they employ. A filter that treats bright components as foreground information against a dark  
 129 background is denoted as  $\gamma_\lambda^A$  and is called an attribute opening if the criterion is increasing, or thinning  
 130 otherwise. By contrast, filter that treats dark components as foreground information against a bright  
 131 background is denoted as  $\varphi_\lambda^A$  and is called an attribute closing if the criterion is increasing, or thickening  
 132 otherwise. The term  $A$  specifies the attribute type, and  $\lambda$  the attribute threshold.

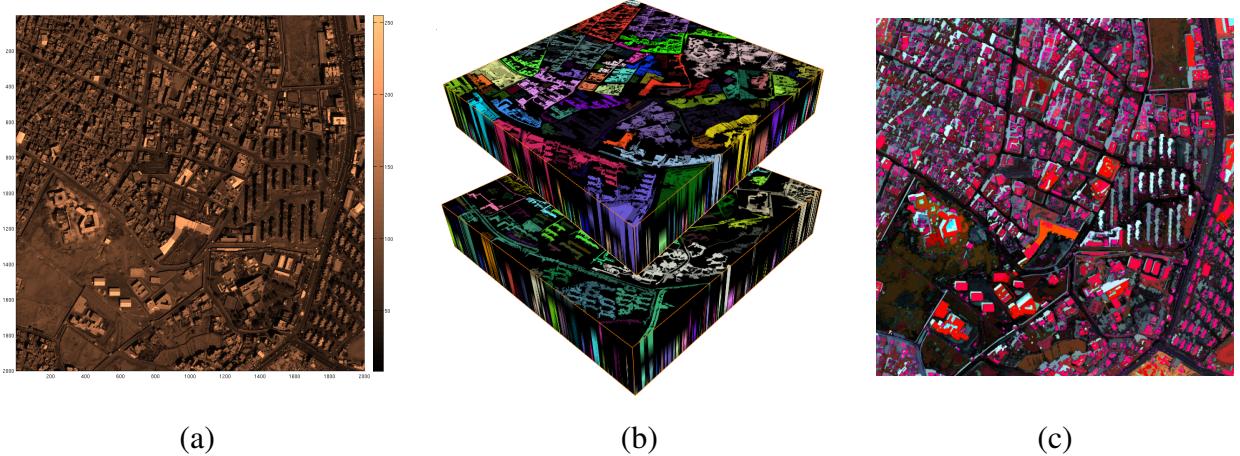
133 **2.2. Differential Attribute Profiles and the CSL Model**

Consider a dual pair of attribute filters  $(\gamma_\lambda^A, \varphi_\lambda^A)$  and an attribute threshold vector  $\vec{\lambda} = [\lambda_i]$  with  $i \in [0, 1, 2, \dots, I - 1]$  and  $\lambda_0 = 0$ . For each filter, let  $\Delta^\Pi$  be a differential profile that given a point  $x \in E$  is defined as follows:

$$\Delta^\Pi(\gamma_\lambda^A(f))(x) = \left( (\gamma_{\lambda_{i-1}}^A(f) - \gamma_{\lambda_i}^A(f))(x) \mid \lambda_i > \lambda_{i-1}, \forall i \in [1, \dots, I - 1] \right) \quad (3)$$

for the filter  $\gamma_\lambda^A$ , and

$$\Delta^\Pi(\varphi_\lambda^A(f))(x) = \left( -(\varphi_{\lambda_{i-1}}^A(f) - \varphi_{\lambda_i}^A(f))(x) \mid \lambda_i > \lambda_{i-1}, \forall i \in [1, \dots, I - 1] \right) \quad (4)$$



**Figure 2.** Example of two instances of a DAP vector field. The original image in (a); the DAP vector field with the first and last attribute zone at the bottom and top of each instance in (b) respectively. The top volume in each image corresponds to the opening, and the bottom to the closing instance. (c) The blended CSL triplet.

for  $\varphi_\lambda^A$ . Each profile is a response vector associated to  $x$ .  $\Delta^\Pi(\gamma_\lambda^A)$  is called the positive, and  $\Delta^\Pi(\varphi_\lambda^A)$  the negative response vector respectively. The concatenation of the two, denoted with  $\sqcup$ , is a  $2 \times (I - 1)$  long vector, called the differential attribute profile [17] of  $x$ :

$$\text{DAP}(x) = (\Delta^\Pi(\gamma_\lambda^A(f)) \sqcup \Delta^\Pi(\varphi_\lambda^A(f)))(x). \quad (5)$$

The set of DAPs computed from the full extent of image definition domain is called the DAP vector field. An instance  $i$  of a DAP vector field is a difference image with respect to the given operator and is called a DAP plane. Figure 2 shows an example. Image (a) shows a residential section of the city of Sana'a, Yemen. It is a 2009, Quickbird panchromatic image ©DigitalGlobe, Inc., at 0.6m spatial resolution in colormap projection. Image (b) shows the DAP vector field computed from (a). The visualisation in (b) employs 3D connected component color labels. The top volume set in (b) shows the area opening top-hat scale-space and the lower volume set shows the area closing bottom-hat scale space.

The extreme values of an entry in any of the two response vectors are 0 and  $h_{\max}$  with the latter being the maximal image intensity value. Consider a point  $x \in E$  and let  $\check{d}h_{\gamma_\lambda^A}$  be the highest value in the positive response vector of  $\text{DAP}(x)$ , i.e.:

$$\check{d}h_{\gamma_\lambda^A}(x) = \vee \Delta^\Pi(\gamma_\lambda^A(f))(x), \quad (6)$$

which is given at a scale  $i$ . Note that  $i$  may not be necessarily unique, i.e. the same response can be observed at other scales too. However, since the most relevant information of the decomposition are contained within planes associated with smaller scales, the parameter:

$$\hat{i}_{\gamma_\lambda^A}(x) = \wedge i : (\gamma_{\lambda_{i-1}}^A - \gamma_{\lambda_i}^A)(x) = \check{d}h_{\gamma_\lambda^A}(x), \quad (7)$$

is identified. That is,  $\hat{i}_{\gamma_\lambda^A}$  is the smallest scale at which the positive response vector of the  $\text{DAP}(x)$  registers the maximal response. The equivalent scale parameter for the negative response vector is defined analogously:

$$\hat{i}_{\varphi_\lambda^A}(x) = \wedge i : (\varphi_{\lambda_i}^A - \varphi_{\lambda_{i-1}}^A)(x) = \check{d}h_{\varphi_\lambda^A}(x). \quad (8)$$

141 They are referred to as the *multi-scale opening* and *closing characteristic* respectively; **MOC** and **MCC**.

142 Consider the following labeling scheme based on the maxima of the two response vectors:

143 1. **convex**, if  $\check{dh}_{\gamma_\lambda^A}(x) > \check{dh}_{\varphi_\lambda^A}(x)$ ,

144 2. **concave**, if  $\check{dh}_{\gamma_\lambda^A}(x) < \check{dh}_{\varphi_\lambda^A}(x)$ ,

145 3. **flat**, if  $\check{dh}_{\gamma_\lambda^A}(x) = \check{dh}_{\varphi_\lambda^A}(x)$ .

146 This is the equivalent taxonomy to the ones proposed based on the differential morphological and area  
 147 profile decompositions in [7] and [16] respectively. The set of labels is referred to as the local curvature  
 148 of the gray level function surface, where the attribute value determines the local spatial domain.

An intuitive multi-band segmentation scheme is then:

$$\bar{i}(x) = \begin{cases} \hat{i}_{\gamma_\lambda^A}(x) & \text{if } \mathbf{convex}, \\ \hat{i}_{\varphi_\lambda^A}(x) & \text{if } \mathbf{concave}, \\ 0 & \text{if } \mathbf{flat}. \end{cases} \quad (9a)$$

$$(9b)$$

$$(9c)$$

149 This multi-level segmentation scheme can be further simplified to a 3-band segmentation by replacing  
 150 the scale parameter with a fixed gray-level.

151 The winning scale from the comparison in (9) is referred to as the *characteristic scale* or simply  
 152 *characteristic C*. The response associated to the winning scale is called the *saliency S* and is denoted  
 153 by  $\bar{dh}$ . The two, complemented by the level  $L$  of the pixel  $x$  after the iteration of the respective attribute  
 154 filter at the reference scale  $i$ , denoted as  $\hbar$ , constitute the three bands of the CSL model [33]. The latter  
 155 is a non-linear mixture model used for compact representation of multi-scale image information to be  
 156 used for classification and other analytic purpose. An example of CSL used for visualisation purposes is  
 157 shown in Fig. 2(c); computed from (a).

### 158 2.3. Attribute Zone Decomposition

159 Attribute filters acting on a gray-scale image produce a dichotomy, i.e. they separate the image  
 160 contents in two sets of components; those that satisfy the attribute criterion and those that fail it.

161 Extending this for multiple attribute thresholds gives raise to the concept of *attribute zone decomposition*.

162 Each zone is a grouping of peak components with attribute values being within some prespecified range.

Let  $\Gamma_\lambda^A$  and  $\Phi_\lambda^A$  be the binary counterparts of  $\gamma_\lambda^A$  and  $\varphi_\lambda^A$  respectively, and consider a threshold decomposition of  $f$  [47]. Each threshold set  $T_h(f)$ , with  $h \in \{h_{min}, \dots, h_{max}\}$ , contains  $k \in K_h^f$  peak components  $P_k^h$ , and for each pixel  $x \in E$  the characteristic function  $\chi$  of  $T_h(f)$  is given by:

$$(\chi(T_h(f)))(x) = \begin{cases} 1 & \text{if } x \in T_h(f), \\ 0 & \text{otherwise,} \end{cases} \quad (10)$$

**Definition 1.** The attribute zone  $\zeta_{\lambda_a, \lambda_b}^A$  of a mapping  $f : E \rightarrow \mathbb{Z}$ , bounded by two attribute thresholds  $\lambda_a$  and  $\lambda_b$  such that  $\lambda_a < \lambda_b$ , is given by:

$$\zeta_{\lambda_a, \lambda_b}^A(f) = \sum_{h=h_{min}+1}^{h_{max}} \left( \sum_{k \in K_h^f} \chi\left(\Gamma_{\lambda_a}^A(P_k^h) \setminus \Gamma_{\lambda_b}^A(P_k^h)\right) \right). \quad (11)$$

163 In brief, the intensity of any point  $x$  of the image domain, that marks a component in a zone  $\zeta_{\lambda_a, \lambda_b}^A$   
 164 can be obtained by initializing it to zero and updating it by adding the value of 1 for each level at  
 165 which,  $x \in P_k^h$  is non-zero in the difference between the two attribute filters  $\Gamma_{\lambda_a}^A, \Gamma_{\lambda_b}^A$ . The same  
 166 decomposition is defined for the operator  $\Phi_{\lambda}^A$ . The attribute zone operator is a connected operator and it  
 167 generates an attribute-based decomposition as described in [16], i.e. any two zones do not overlap and  
 168 the decomposition is unique.

169 **2.4. The DMP vs. the DAP**

170 The DMP is a special case of a DAP, just as reconstruction filters are a special case of attribute  
 171 filters. By replacing the general attribute filters in the DAP by reconstructions from a set of markers we  
 172 obtain the DMP. The set of markers replaces the set of attribute thresholds, and this is where there is an  
 173 important difference in terms of computation, as will be discussed in the following sections. The key  
 174 point is that the markers themselves need to be stored in memory as a stack of images. As has been  
 175 shown [8], we can reuse this space to store first the stack of reconstructed images, and finally, the output  
 176 DMP. The latter can then be processed into an MOC or MCC. Though the algorithm in [8] is highly  
 177 efficient, we will explore a one-pass method and further economies to increase the performance further.

178 **3. The Max-Tree Algorithm and the One-Pass Method**

179 An efficient method for computing the DAP vector field and the CSL model has been presented in  
 180 [16]. It relies on a combination of the Max-Tree and Min-Tree image representation structures [19] for  
 181 computing the attribute zone decomposition discussed in the previous section.

The Max-Tree of an image  $f$  is a tree for which the node hierarchy corresponds to the nesting of the peak components of  $f$ . Each node  $N_k^h$  is addressed by its level  $h$  and its index  $k$ , and corresponds to a set of flat-zones [11] for which there exists a unique mapping to a single peak component:

$$N_k^h = \{x \in P_k^h \mid f(x) = h\}. \quad (12)$$

182 The “leaves” of the tree correspond to the regional maxima of  $f$  and its root is defined at the lowest gray  
 183 level of the structure, representing the set of background pixels.

184 In the current implementation, as in [20], each pixel is represented by a Max-Tree node, stored in  
 185 an array of the same size as the image. All nodes contain a *parent* reference, represented as a 32-bit  
 186 unsigned integer. A boolean field *valid* indicates whether a node has been processed or not. In the  
 187 general case of attribute filters, each node contains a pointer to a set of auxiliary attribute variables that  
 188 are updated from its member pixels during the construction of the tree [48]. In the case of area filters,  
 189 memory is saved by using an *area* field, also represented as a 32-bit unsigned integer instead. A peak  
 190 component’s area is computed by summing the area of its descendants with the total area of its flat-zones.  
 191 Within each peak component a single pixel is chosen as the representative, or canonical element. We call  
 192 this pixel the *level root*. All pixels of a peak component, except the level root point to the level root. The  
 193 level root itself points to a pixel in the parent peak component. The root of the tree points to a global  
 194 root  $\perp$ .

195 The Max-Tree is computed using a hierarchical, depth-first, flood-filling algorithm introduced by  
 196 Salembier et al. [19]. Alternative methods were presented in [49–52]. The recursive method discussed in  
 197 [19] separates the construction of the tree from the computation of attributes and the actual filtering. This  
 198 architecture is particularly appreciated in applications requiring interactive filtering or multiple filtering  
 199 instances. Note that the Min-Tree is an equivalent representation to the Max-Tree of the inverted input  
 200 image.

The one-pass method for computing the attribute zones on each of the two trees was presented in [16]. A pseudo-code listing can be found in the same article. The methods follows a similar approach to the filtering functions described in [19,48]. Attribute zones are independent of the various root-paths and they are identified by a unique integer corresponding to the zone's  $\lambda_b$  address in the  $\vec{\lambda}$  vector:

$$\text{ZoneID}(P_k^h) = \wedge \left\{ i \in \{1, \dots, I - 1\} \mid \text{Attr}(P_k^h) < \lambda_i \right\}. \quad (13)$$

201 In words, the zone ID member of each node is set to the address of the lowest attribute threshold from  
 202  $\vec{\lambda}$  for which the peak component under study fails the attribute criterion. Once the image is decomposed  
 203 into its attribute zones, the two instances of the DAP field each of which is an  $(I - 1) \times \text{ImageSize}$   
 204 volume set, can be extracted by visiting each pixel  $I - 1$  times, i.e. once for each area zone.

## 205 4. A Concurrent One-Pass Method

206 In this section we study how the above approach can be adapted for the concurrent computation of the  
 207 Differential Attribute Profile. We first focus on the case of increasing attributes, and in particular the case  
 208 that the attribute is the area of each component. Two cases can be separated: (i) explicit computation of  
 209 the DAP vector field, and (ii) direct computation of the MOC, and by duality the MCC.

### 210 4.1. Parallel DAP computation

211 Unlike the DMP, we do not have to compute any marker images. We simply need a Max-Tree with  
 212 suitable attributes computed for all nodes. We do need an array *out* of *numscales* images to store the  
 213 output. An array *lambda* containing the scale-class boundaries is required as input. Rather than filtering  
 214 the image at each level explicitly, and computing the difference at consecutive scales, we approach the  
 215 problem in a similar way to that in [8]. A single filtering stage computes the entire DAP. The algorithm  
 216 is shown in Alg. 1.

217 Each process or thread  $p$  scans the section  $V^p$  assigned to it. Any pixel for which the corresponding  
 218 node has not yet been marked as valid is processed. Starting at the smallest scale, the algorithm follows  
 219 the root path until it either finds the global root, or a valid node, or it finds a node with an area larger  
 220 than the current scale. It does this for all scales, storing the node location in an array *ws*, and the  
 221 corresponding gray level in an array *val*. After all scales are processed, a second pass along the root path  
 222 is made, setting all values in the DAP array *out* to the correct values for all pixels in its section  $V_p$ .

223 After this pass, each slice of the array *out* contains the area opening for the appropriate scale  $\lambda_i$   
 224 represented in the pseudo code by *lambda*[*i*]. What remains is to compute the differences between  
 225 successive slices to obtain the final DAP.

---

**Algorithm 1** The filtering stage of the parallel DAP algorithm

---

```

procedure MaxTreeMakeDAP (Vp : Section;
    var node : Max-Tree; var out : DAP;
    lambda : integer[numscales])

for all  $v \in V_p$  do
    if not  $\text{node}[v].valid$  then
         $w := v;$ 
        for all scales  $i$  increasing order do
            while  $\text{Par}(w) \neq \perp \wedge \text{not } \text{node}[w].valid \wedge \text{node}[w].area < \lambda[i]$  do
                 $w := \text{Par}(w);$ 
            end;
             $ws[i] := w;$  (* temporary storage of for each scale *)
            if  $\text{node}[w].valid$  then
                for all scales  $j \geq i$  do (* filtered node found *)
                     $val[j] := out[j][w];$ 
                end;
            else if  $\text{node}[w].area \geq \lambda[i]$  then (*  $w$  is filtered *)
                 $val[i] := out[i][w];$ 
            else (*  $\lambda[i]$  too large *)
                 $val[j] := 0;$ 
            end;
            end;
        end;
    end;
     $u := v;$ 
    for all scales  $i$  increasing order do
        repeat
            if  $u \in V_p$  then
                for all scales  $j < i$  do  $out[j][u] := f[u];$  end;
                for all scales  $j \geq i$  do  $out[j][u] := val[j];$  end;
                 $\text{node}[u].valid := \text{true};$ 
            end;
             $u := \text{node}[u].parent;$ 
        until  $u = ws[i];$ 
    end;
    if  $u \in V_p$  then (* Process  $ws[numscales - 1]$  *)
        for all scales  $j < i$  do  $out[j][u] := f[u];$  end;
         $\text{node}[u].valid := \text{true};$ 
    end;
    end;
end.

```

---

226 The disadvantage of this approach is that both memory use and computational time scale linearly with  
227 the number of scales. This becomes prohibitive for large numbers of scales.

228 *4.2. Direct MOC computation*

229 Like the previous algorithm, we assume that the Max-tree has been computed with all attribute  
230 information in place. Computing the MOC directly can be done without the need for storing a DAP  
231 vector field. Instead, we only compute the supremum of the DAP vector for each scale, stored in image  
232 *outDH*, along with a relevant scale, stored in image *outScale*, and the orginal gray level at that scale for  
233 each pixel in the image, stored in image *outOrig*. To compute this we equip each Max-Tree node with  
234 fields *maxDH* and *scale* both of gray-level type (unsigned char in our case). As before, array *lambda*  
235 stores the scale class boundaries.

236 In a first version of the algorithm, we reused the *area* field to store the scale class (as an integer). This  
237 algorithm worked in two stages (after building the tree). In the first stage, we simply compute the scale  
238 class for each level root in the tree. This can be done completely in parallel, because the decision is local.  
239 After this the tree was traversed in much the same way as the original one-pass method.

240 Some initial experiments showed that the first phase of this two pass approach was surprisingly  
241 inefficient in terms of speed-up, achieving only 16-18 times faster computation on 64 nodes. The reason  
242 for this is most likely that it is a very short phase with barriers on either side. This means that load  
243 imbalance and synchronization costs have a large impact. We therefore implemented a one-pass version.  
244 Here we do need to introduce a *scale* field in the tree, and though this is of course more costly in terms  
245 of memory, it has an important advantage, in that the area data is not overwritten. This means that we  
246 can recompute MOCs with different parameter settings from the same tree, should we need to.

247 The one-pass algorithm for the MOC computation is given in Alg. 2, and Alg. 3. The latter simply  
248 scans the section of the image assigned to the process, and calls procedure *NodeSetMOC* from Alg. 2  
249 for each node that is not yet valid. This checks if the current node is a level root, and if so, computes the  
250 scale and gray level contrast with its parent. Next it checks whether the scale is maximal, indicating that  
251 the current node has area outside the range of interest, and its MOC parameters are set to “out-of-bounds”  
252 conditions. No further ascent is needed as all ancestors must have larger areas.

253 Otherwise we look up the parent, and check if it is valid or not. If not, *NodeSetMOC* is called  
254 recursively for the parent. Otherwise, the parent data are copied (as they are valid). We now have the  
255 correct MOC values of the root path up to the parent of the current node. If the current node is a level  
256 root we check whether its scale is the same as the parent’s value stored in *curScale*. If so, the contrast  
257 value *curDH* must be incremented with the contrast of the current node, stored in *DH*. If the scales are  
258 not the same, the current scale and contrast are updated to those of the current node.

259 After this, we check whether the current contrast *curDH* is larger than or equal to the maximum  
260 contrast *maxDHm*, and if so, *maxDH*, *maxScale*, and *outOrig* are updated (corresponding to the C, S,  
261 and L parts of the MOC, respectively).

262 Finally, if the current node is in the processor private section *Vp* of the image, the parameters are  
263 written to the fields in the node, and to the output MOC. This approach does mean scale values may be  
264 computed multiple times by different threads, but this is outweighed by the fact that we lose a barrier.

---

**Algorithm 2** The filtering stages of the parallel MOC

---

```

procedure NodeSetMOC ( Vp : Section; var node : Max-Tree; current : integer;
    var maxDH, curScale : grayval; var outDH, outScale,
    outOrig : array[0..N-1] of pixel, lambda : integer[numscales] );

    var scale, DH : grayval;
    if IsLevelRoot(node[current]) then                                (* compute scale and DH for current nodes *)
        scale := FindScale(node[current], lambda); DH := getDH(node[current]);
    end;
    if IsLevelRoot(node[current]) and scale = numscales then          (* Initialize to out of scale range *)
        maxScale := numscales; maxDH := 0; curDH := 0;
        maxOrig := 0; curScale := numscales;
    else
        parent = node[current].parent;
        if not node[parent].valid then                                     (* go into recursion to set parent values correctly *)
            NodeSetMOC(Vp, node, parent, maxDH,
                        curScale, outDH, outScale, outOrig, lambda);
        else                                                 (* if the parent is valid, copy relevant values *)
            maxScale := outScale[parent]; maxDH := outDH[parent]; maxOrig := outOrig[parent];
            curScale := node[parent].scale; curDH := node[parent].curDH;
        end;
        if IsLevelRoot(node[current]) then          (* if I have a level root, some things might change *)
            if scale = curScale then                  (* same scale class: add current pixel's curDH *)
                curDH = curDH + curDH;
            else                                         (* scale class change, update current scale and DH *)
                curDH = DH; curScale = scale;
            end;
            if curDH  $\geq$  maxDH then                  (* If updated curDH is higher than or equal to the maximum
                                                       DH found update maxDH, maxScale, and outOrig *)
                maxDH := curDH; maxScale := scale;
                outOrig = gval[current];
            end
        end;
    if current  $\in$  Vp then                                     (* Store the information *)
        outScale[current] := maxScale; outDH[current] := maxDH;
        outOrig[current] := maxOrig; node[current].scale := curScale;
        node[current].curDH := curDH; node[current].valid := true;
    end;
end.

```

---

**Algorithm 3** The final filtering stage of the parallel MOC

---

```

procedure MaxTreeComputeMOC (Vp : Section;
    var outDH, outScale,
    outOrig : array[0..N-1] of pixel,
    lambda : integer[numscales] );

var curScale, maxScale, maxDH, curDH, maxOrig : grayval;

for all v ∈ Vp do
    if not node[v].valid then
        NodeSetMOC(Vp, v,
            curDH, maxDH,
            maxScale, maxOrig,
            curScale, outDH,
            outScale, outOrig, lambda);
    end;
end;
end.;

```

---

**Table 1.** Images used, with size and mean and standard deviation of the gray levels.

Name	Size ( $\times 10^6$ pixels)	mean	std. dev.
Sana'a	3482	121.78	75.57
Airport	3600	96.58	47.55
Haiti1	3888	107.71	62.99
Haiti2	3963	81.34	52.39

265 **5. Experiments**

266 The algorithm was implemented in C, and tested on a Dell R815 compute server with 4 Opteron  
 267 processors with 16 cores each. A total of 512 GB RAM was installed. The source code is available  
 268 at [www.cs.rug.nl/~michael/ParMaxTree/](http://www.cs.rug.nl/~michael/ParMaxTree/). The data sets consisted of a panchromatic Quickbird image  
 269 of Sana'a, Yemen, courtesy of DigitalGlobe, Inc., two panchromatic post-earthquake images of Haiti,  
 270 courtesy of Google, and one panchromatic satellite image of an airport. All images were duplicated and  
 271 tiled to achieve a total image size of between 3.4 and 3.9 Gpixel. For the image properties see Table 1.

272 *5.1. Concurrent CSL computation*

273 In all cases, timings were performed at 1, 2, 4, 8, 12, 16, 24, 32, 48, 64, 128 and 256 threads. In  
 274 previous work, we had observed that the Max-Tree building phase could actually speed up further if  
 275 more threads than cores were used, due to reduced cache thrashing [20]. As the CSL computation was  
 276 expected to be dominated by the tree-building phase, we wanted to know the impact of this effect.

277 In the CSL case the final program has the following structure:

- 278 1. Read image, and create inverted copy of image

- 279     2. Build Max-Tree from original image
- 280     3. Compute MOC
- 281     4. Build Min-Tree from inverted image
- 282     5. Build MCC
- 283     6. Combine MOC and MCC to final CSL
- 284     7. Write output

285     The parallel CSL-model segmentation program was run 24 times per timing, and the shortest time  
286     was considered indicative of the performance of the algorithm in the absence of interference by other  
287     programs. In the timing results, we added the Min-Tree and Max-Tree times into one tree-building  
288     timing, and likewise for the MOC and MCC times.

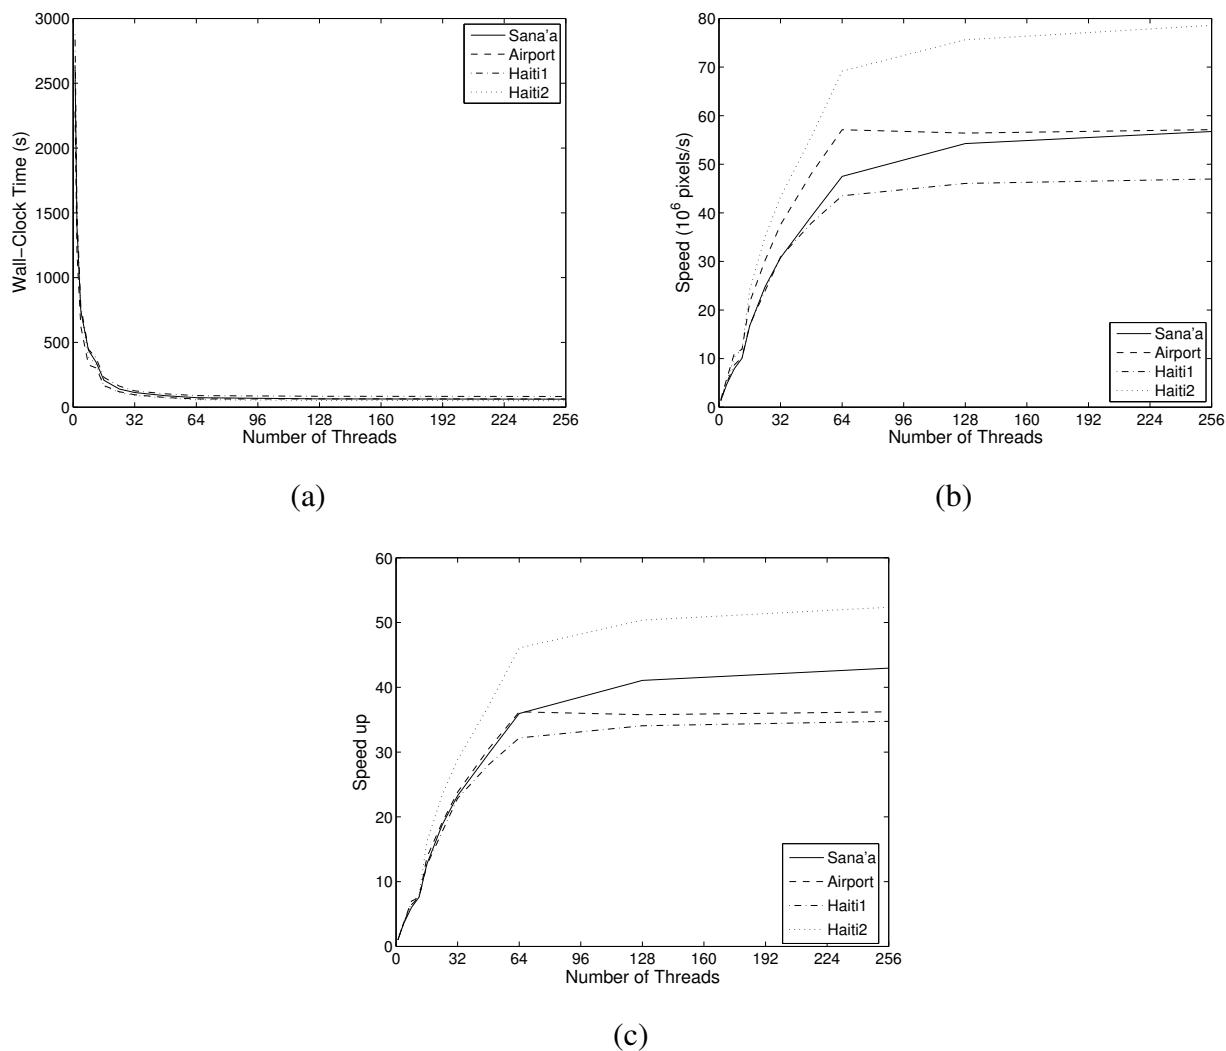
289     Experiments were run in three parameter settings used in practical cases. Two settings with a  
290     maximum scale of  $16.7 \cdot 10^6 \text{ m}^2$  were used, with 12 and 32 scales respectively. The last setting had  
291     64 scales and a maximum area scale of 16,384 m<sup>2</sup>. It turned out that these settings have no significant  
292     impact on the performance of the algorithms. In the following, we discuss the latter timings, with the  
293     maximum number of scales.

294     Figure 3 shows the performance of the algorithm in terms of wall-clock time, speed in millions of  
295     pixels per second, and speed up. Total computing times varied from an average of 2609 s at 1 thread,  
296     down to 70.8 s at 64 threads, and a further reduction to 64.4 s at 256. The latter is due to reduced cache  
297     thrashing during the Max-Tree construction phase [20]. The mean speed up for the complete process is  
298     36.85 at 64 threads, and 40.51 at 256 threads.

299     Figure 3(b) shows the computational speed in millions of pixels per second. This calibrates for the  
300     differences in processing time due to differences in image size. Despite this correction, quite significant  
301     differences are seen. At maximum performance, the processing speed varies from 46.96 Mpixel/s to  
302     78.58 Mpixel/s, with an average of 59.85 Mpixels/s, whereas speed-up varies from 34.74 to 52.36.  
303     These differences stem mainly from different complexities of the image. Haiti1 contains only built-up  
304     areas, which means that the Max-Tree contains many more nodes. Haiti2 contains large sections of rural  
305     area, and some extensive black regions where no data were available. This simplifies the Max-Tree. The  
306     Sana'a and Airport images are more mixed.

307     Construction of the Max-Tree took the most time, using the algorithm from [20], from a mean of  
308     1912 s at a single thread, down to 46.4 s at 64 threads, and further down to 38.8 s at 256 threads. The  
309     MOC-MCC phase took 576.75 s on a single thread on average, down to 18.7 s at 64 threads, rising to  
310     21.0 s at 256 threads. Combining the MOC and MCC information into the final CSL-model segmentation  
311     is the cheapest by far, costing 105.7 s on one core on average, dropping to 2.25 s on average for 64 cores,  
312     with no significant changes beyond.

313     There are striking differences in speedup between the phases, as seen in Figure 4. The building phase  
314     has the highest speed-up on average, but also the largest variation, ranging from 33.05 for Haiti1, to  
315     53.67 for Haiti2 at 64 threads, with an average of 42.22. Beyond 64 threads, there is a further increase in

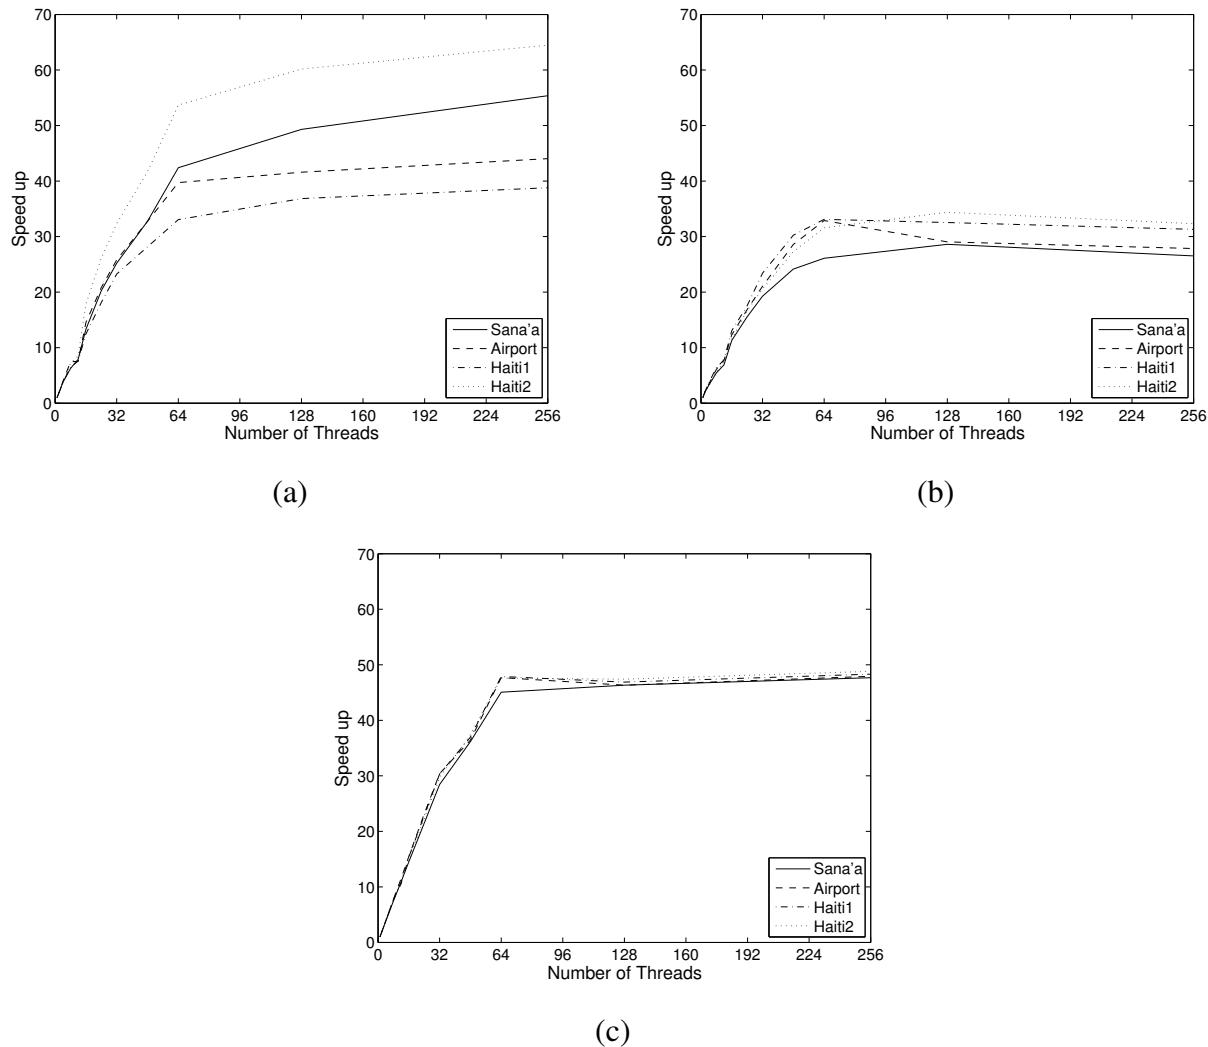


**Figure 3.** Performance of the algorithm for four images ranging from 3.48 to 3.96 Gpixel, on a 64 core machine, using a 64-scale CSL-model segmentation, as a function of number of threads: (a) wall-clock times; (b) speed in  $10^6$  pixels/s; (c) speed-up. The continued rise of speed and speed up for more threads than cores (64) is due to reduced cache thrashing during the Max-Tree building phase [20].

316 speed up, and 256 threads the speed-up varies from 38.80 for Haiti1 to 64.49 for Haiti2, with an average  
 317 of 50.67.

318 The mean speed-up of the MOC-MCC phase is a lot lower, achieving 30.92 on average at 64 threads,  
 319 with no significant increase beyond that. The final phase has a very good speed up, averaging 47.07 at  
 320 64 threads.

321 By comparison, running the parallel algorithm for differential morphological profiles on the 3.8  
 322 Gpixel Sana'a image, using just 7 scales, on the same machine, and on a single core, the wall clock  
 323 time was 5 h 9 min 6.7 s, dropping to 8 min 33.7 s, or a speed up of 36.1, on 64 cores. About half of the  
 324 wall-clock time is taken up by computation of markers, which is done by computing erosions by exact  
 325 Euclidean discs by the algorithm of [53]. When using square SE with the same algorithm, the memory  
 326 use goes down dramatically, and the wall clock time for 7 scales drops to 4 h 10 min 46.8 s and 6 min  
 327 19.5 s for 64 cores, or a speed up of 39.6.



**Figure 4.** Breakdown of the speed-up: (a) Tree building; (b) MOC + MCC computation; (c) Combining results into CSL model

328     The CSL-model segmentation algorithm has a memory complexity of  $O(N + N_s)$ , with  $N$  the number  
 329     of pixels, and  $N_s$  the number of scales. Because  $N_s \ll N$  this becomes  $O(N)$  in practice. What we  
 330     observe is that maximum memory allocation is 24 times the size of the image data, at least for the area  
 331     attribute. Other attributes would require more memory, as would the use of 64-bit integers for area  
 332     computation and pixel addressing, as required when going beyond 4 Gpixel images. In that case we  
 333     would require 32 times the data size.

334     By contrast, computing the CSL-model segmentation from the DAP explicitly has a memory  
 335     complexity of  $O(NN_s)$ , which becomes prohibitive for large  $N_s$ . In the 64 scale case above the memory  
 336     usage would rise to 164 times the size of the image (because two arrays of size of  $64N$  need to be stored).  
 337     We ran a few experiments on small images, which indicated that computation of the DAP vector field  
 338     at 64 scales is more than 10 times slower than CSL-model segmentation. On the 870 Mpixel original  
 339     Sana'a image, the DAP computation took 177.25 s on 64 threads, as opposed to 15.42 s using direct  
 340     CSL-model segmentation. Even for 12 scales the difference is factor of two to three. In the case of  
 341     the 870 Mpixel Sana'a image the DAP computation took 39.61 s, whereas the CSL-model segmentation  
 342     took 15.43 s.

343 For the DMP the situation is similar to that of explicit DAP computation in terms of memory use.  
344 A stack of  $N_s$  images is first used to store the markers. These markers are then transformed into  
345 reconstructed images, after which pixel-wise differencing is used to obtain the DMP vector field (stored  
346 in the same location). This DMP field must then be transformed into a final CSL segmentation. The  
347 computational load is typically higher, due to the need for  $N_s$  erosions, each of at least  $O(N)$  yielding  
348  $O(NN_s)$  time complexity. If rotation invariance is needed, this readily increases to  $O(NN_sR_{\max})$  time  
349 complexity, with  $R_{\max}$  the maximum radius of the SEs used. This initial phase is much more costly than  
350 using area as attribute in the DAP, which has constant time complexity per pixel. For more complicated  
351 attributes than area, memory load does increase. When using higher order moments the storage per  
352 maxtree node goes up roughly as  $O(M^2)$ , with  $M$  the order of the moment used, and the compute time  
353 increases similarly. As the number of nodes can equal the number of pixels the worst-case memory  
354 complexity of the explicit DAP computation becomes  $O(N(N_s + M^2))$ . For direct CSL computation it  
355 becomes  $O(NM^2)$ , so again, the dependency on the number of scales is removed.

356 To indicate the severity of the problem it should be noted that even though our server was fitted with  
357 the maximum of 512 GB of RAM, the large data sets could not be run at 64 scales due to memory  
358 limitations.

## 359 6. Discussion

360 The results show that parallel computation of CSL-model segmentation can be done very efficiently  
361 using the proposed algorithm. We obtain an overall average speed of nearly 60 Mpixel per second, or  
362 3.59 Gpixel per minute, even for 64 scales in the CSL-model segmentation. This means that routine  
363 processing of these very large images becomes possible on fairly standard compute servers as well as  
364 high-end workstations. Furthermore, if the Max-Tree is built once, it could be stored, and different  
365 CSL-model segmentations could be made with different scale parameters. For images of the sizes used  
366 here, this would take between 14.09 and 26.62 s. Whilst this is not exactly a real-time response, it  
367 is tolerably fast. These later phases of the algorithm have linear time complexity, and for the original  
368 Sana'a image (870 Mpixel) the MOC-MCC and CSL-model segmentation phases can take as little as  
369 3.90 s (out of a total compute time of 15.49 s).

370 It is curious that the MOC-MCC phase of the algorithm has the least speed up, and seems to be limited  
371 to about 32. We considered whether this might be due to the fact that some floating point math was used  
372 in this phase. The R815 server may have 64 cores, but these share only 32 floating point units between  
373 them. We replaced the floating point code with integer math, but obtained only a slight speed increase.  
374 Therefore, we surmised that the fact that area scales of a given node might be computed multiple times  
375 is the problem. However, an earlier version computed the area scale for each nodes in parallel, and then  
376 used the results in a slightly simpler MOC-MCC phase. In this case the MOC-MCC-phase was no faster:  
377 20.37 s for the 3.48 Gpixel Sana'a image, vs 20.22 s, when area scales are included in the MOC-MCC  
378 phase. This difference is not significant, suggesting that this is not the solution. By putting the scale  
379 computation inside the MOC-MCC phase, a barrier was removed, and the overall performance increased  
380 a great deal. For the 3.48 Gpixel Sana'a image, the wall-clock time at 256 threads dropped from 73.11 s  
381 to 61.37 s.

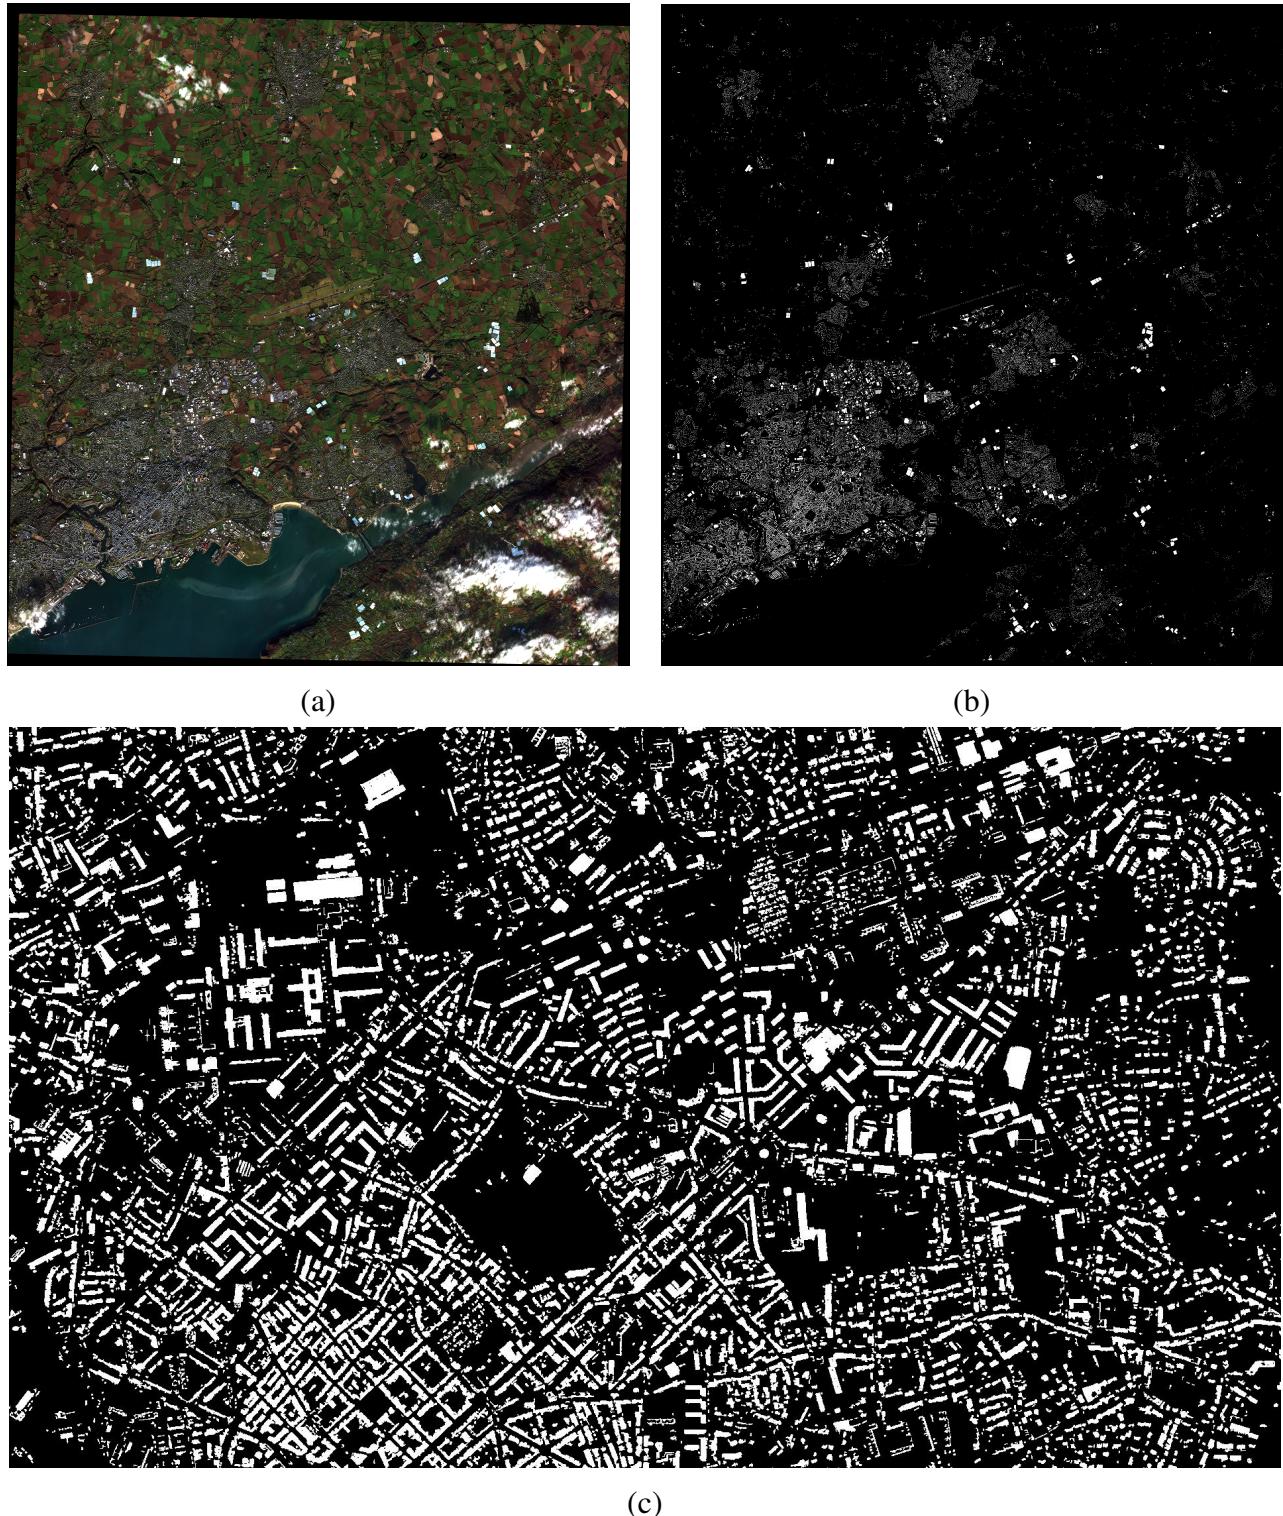
382 It therefore seems reasonable that the main cost lies in having to traverse the Max-Tree in parallel,  
383 with multiple threads traversing the same root-paths several times. Adding extra computations per node  
384 visited has only a very minor impact on either wall-clock time or speed-up. Thus we find that the fairly  
385 complex MOC-MCC phase costs only little more than a single filtering stage of an area opening. This  
386 has large implications for the speed gained by the proposed algorithm, quite apart from the parallelism  
387 gained.

388 By comparison, performing a single area filter step on a precomputed Max-Tree of these large images  
389 takes roughly 3-5 s, depending mainly on image content. For 64 scales, this would mean 128 filter steps,  
390 which comes to a value between 384 and 640 s, assuming no slow-down due to synchronization issues.  
391 With the new algorithm this range is roughly 14 to 27 s (some 25 times faster) regardless of the number  
392 of scales.

393 A variant of this parallelization strategy for the Max-Tree algorithm and the CSL computation is  
394 hosted in DigitalGlobe's Geo Spatial Big Data platform (GBDX) (<http://www.digitalglobe.com/>) for  
395 computing the raw CSL layers of each queued image prior to radiometric classification. This is a  
396 building foot-print extraction service operated on the cloud which can be optimized to process all of  
397 DigitalGlobe's WorldView 2 and 3 multispectral imagery in near real time. The respective workflow  
398 computes the CSL triplet that is configured to contain the most relevant/descriptive image structures  
399 compared to the full DAP vector field. These in turn, can be classified efficiently using radiometric  
400 features at a meta-process stage and have possible noise elements removed. Efficiency in this case  
401 comes as a consequence of reduced number of test samples (dominant CSL structures). An example  
402 is shown in Fig. 5. Image (a) shows an RGB view of an 8-band multi-spectral data-set covering the  
403 city of Brest, France and its surroundings. This is a WorldView 2 image © DigitalGlobe, Inc, at 1.6  
404 m. spatial resolution following ortho-rectification and atmospheric compensation (14 bpp data-depth).  
405 The image was acquired on Oct. 15<sup>th</sup>, 2014 and covers a total area of 313 km<sup>2</sup>. Image (b) shows  
406 the respective CSL-segmentation following natural element refinement procedures. The segmentation  
407 targets all built-up structures independent of size, structural or radiometric characteristics. Image (c)  
408 shows a zoom-in view of the city's center north of its port. This workflow makes use of MOCs only  
409 (i.e. a Max-Tree structure alone) thanks to a novel band blending method reducing WorldView 2 or  
410 3, 8-band atmospherically compensated data-sets to gray-scale layers that show a strong response to  
411 building materials exclusively.

## 412 7. Conclusions

413 In this paper we have presented a new algorithm for efficient, shared memory, parallel computation  
414 of the CSL-model segmentation and CSL profile of images in the gigapixel range, besides several  
415 improvements on existing DMP, MSLS, and DAP algorithms. The CSL algorithm can be adapted to any  
416 increasing attribute, by incorporating the auxiliary data handling described in [20]. One such extension  
417 to 2D pattern spectra has already been provided, using area and non-compactness as attributes [54].  
418 More complicated attributes will impact on Max-Tree building times and memory use, but not on the  
419 MOC-MCC and CSL-model segmentation times. Non-increasing attributes require several changes,



**Figure 5.** (a) Brest, France; a WorldView 2 RGB image of the city and its surroundings at 1.6m spatial resolution. Approximate image coverage is  $313 \text{ km}^2$ . (b) The CSL-segmentation (a) following radiometric refinement. The  $\vec{\lambda}$  vector was configured to fine-tune the CSL-model for detecting built-up in the input scene. (c) A zoom-in view of the city's center.

420 which we intend to address in future work. For pattern spectra, non-increasing attributes do not pose  
421 problems.

422 In comparison to both the DMP and the DAP, the CSL-model segmentation approach has a clear  
423 speed advantage, especially at a large number of scales. Besides, the memory use of the new method  
424 is a great deal lower than the DMP method, allowing much larger images to be processed in the same  
425 memory. Most importantly, the memory requirement is independent of the number of scales used.

426 The current algorithm is limited to 4 Gpixel only by the fact that 32-bit integers are used. Changing  
427 a single declaration to 64-bit integers makes it suitable for any image size that will fit into memory. We  
428 are currently extending our code to distributed memory, and aim at processing terapixel images in the  
429 near future. A further limitation is that of gray-level resolution. The current algorithm is suitable up to  
430 roughly 16 bits per pixel. Beyond that, the parallel algorithm of [20] on which this is based no longer  
431 works efficiently, due to the fact that the computational cost of the connect algorithm scales linearly with  
432 the number of gray levels. Work is in progress on parallel algorithms to handle any number of gray levels  
433 efficiently. Finally, the current algorithm is a shared memory algorithm, which poses upper bounds to  
434 the maximum image size, as memory sizes on shared memory machines are typically limited. Work on  
435 distributed memory algorithms is in progress.

436 The results of this paper already allow fast processing of vast data sets under tight time constraint, such  
437 as might be required in applications providing support to crisis management, to disaster preparedness  
438 and relief efforts, or when rapid changes are occurring on the ground. Our single rack server could  
439 already process about 5 terapixel of data in a single day, provided care is taken that I/O does not form a  
440 bottleneck. This can be achieved by careful set-up of the processing pipeline. The new algorithm also  
441 allows seamless processing of areas on the scales of entire cities. Moving to distributed computation will  
442 allow us do seamless processing at the scale of countries, or even the world.

#### 443 Acknowledgments

444 This work was part of the High Performance Algorithms for Remote Sensing (HIPARS) project  
445 and was carried out between 2011 and 2013. The HIPARS project was supported by the institutional  
446 research program of the European Commission's Joint Research Centre, Global Security and Crisis  
447 Management Unit, and the Johann Bernoulli Institute, University of Groningen, Netherlands. The Dell  
448 R815 server was funded by the Netherlands Organisation for Scientific Research (NWO), under project  
449 number 612.001.110. All overhead imagery illustrated in this article is courtesy of DigitalGlobe Inc. and  
450 was kindly made available for demonstrating the capabilities of the proposed methodology.

#### 451 Author Contributions

452 All authors contributed equally to algorithm development and analysis. Images and test settings  
453 were provided by Georgios Ouzounis and Martino Pesaresi respectively. Timings were run by Michael  
454 Wilkinson.

#### 455 Conflicts of Interest

456 The authors declare no conflict of interest.

457 **References**

- 458 1. Bangham, J.A.; Ling, P.D.; Harvey, R. Scale-space from nonlinear filters. *IEEE Transactions on*  
459 *Pattern Analysis and Machine Intelligence* **1996**, *18*, 520–528.
- 460 2. Bangham, J.A.; Chardaire, P.; Pye, C.J.; Ling, P.D. Multiscale nonlinear decomposition: the  
461 sieve decomposition theorem. *IEEE Transactions on Pattern Analysis and Machine Intelligence*  
462 **1996**, *18*, 529–538.
- 463 3. Jackway, P.T.; Deriche, M. Scale-Space Properties of the Multiscale Morphological  
464 Dilation-Erosion. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **1996**,  
465 *18*, 38–51.
- 466 4. Koenderink, J.J. The Structure of Images. *Biol. Cybernet.* **1984**, *50*, 363–370.
- 467 5. Lowe, D.G. Distinctive Image Features from Scale-Invariant Keypoints. *International Journal*  
468 *of Computer Vision* **2004**, *60*, 91–110.
- 469 6. Wandell, B.A. *Foundations of Vision*; Sinauer Associates: Sunderland, Mass., 1995.
- 470 7. Pesaresi, M.; Benediktsson, J. A new approach for the morphological segmentation of  
471 high-resolution satellite imagery. *IEEE Trans. Geosci. Remote Sensing* **2001**, *39*, 309–320.
- 472 8. Wilkinson, M.H.F.; Soille, P.; Pesaresi, M.; Ouzounis, G.K. Concurrent computation  
473 of differential morphological profiles on giga-pixel images. In *Proc. 10th Int. Symp.*  
474 *Morphology, ISMM 2011*; Soille, P.; Pesaresi, M.; Ouzounis, G.K., Eds.; Springer  
475 Berlin/Heidelberg, 2011; Vol. 6671, *Lecture Notes in Computer Science*, pp. 331–342.
- 476 9. Gimenez, D.; Evans, A.N. An Evaluation of Area Morphology Scale-Spaces for Colour Images.  
477 *Computer Vision and Image Understanding* **2008**, *110*, 32–42.
- 478 10. Salembier, P.; Wilkinson, M.H.F. Connected Operators: A review of region-based morphological  
479 image processing techniques. *IEEE Signal Processing Magazine* **2009**, *26*, 136–157.
- 480 11. Salembier, P.; Serra, J. Flat Zones Filtering, Connected Operators, and Filters by Reconstruction.  
481 *IEEE Transactions on Image Processing* **1995**, *4*, 1153–1160.
- 482 12. Ouzounis, G.K.; Soille, P. The Alpha-Tree algorithm. Technical Report JRC74511, Publications  
483 Office of the European Union, 2012.
- 484 13. Ehrlich, D.; Kemper, T.; Blaes, X.; Soille, P. Extracting building stock information from optical  
485 satellite imagery for mapping earthquake exposure and its vulnerability. *Natural Hazards* **2013**,  
486 *68*, 79–95.
- 487 14. Ouzounis, G.K. Automatic extraction of built-up footprints from high resolution overhead  
488 imagery through manipulation of Alpha-Tree data structures. US 8682079 B1, 2014.
- 489 15. Ouzounis, G.K.; Soille, P. Differential area profiles. Proc. 20th Int. Conf. on Pattern  
490 Recognition, ICPR 2010; IEEE Computer Society: Istanbul, Turkey, 2010; pp. 4085–4088.
- 491 16. Ouzounis, G.K.; Pesaresi, M.; Soille, P. Differential area profiles: decomposition properties and  
492 efficient computation. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **2012**,  
493 *34*, 1533–1548.
- 494 17. Murra, M.D.; an B. Waske, J.B.; Bruzzone, L. Morphological attribute profiles for the analysis  
495 of very high resolution images. *IEEE Transactions on Geoscience and Remote Sensing* **2010**,  
496 *48*, 3747–3762.

- 497 18. Murra, M.D.; Villa, A.; Benediktsson, J.; Bruzzone, L. Classification of  
498 hyperspectral images by using extended morphological attribute profiles and independent  
499 component analysis. *IEEE Geoscience and Remote Sensing Letters* **2011**, *8*, 541–545.
- 500 19. Salembier, P.; Oliveras, A.; Garrido, L. Anti-extensive Connected Operators for Image and  
501 Sequence Processing. *IEEE Transactions on Image Processing* **1998**, *7*, 555–570.
- 502 20. Wilkinson, M.H.F.; Gao, H.; Hesselink, W.H.; Jonker, J.E.; Meijster, A. Concurrent Computation  
503 of Attribute Filters using Shared Memory Parallel Machines. *IEEE Transactions on Pattern  
504 Analysis and Machine Intelligence* **2008**, *30*, 1800–1813.
- 505 21. Mura, M.D.; Benediktsson, J.A.; Bruzzone, L. Modelling structural information for building  
506 extraction with morphological attribute filters. *Image and Signal Processing for Remote Sensing  
507 XV*; Bruzzone, L.; Notarnicola, C.; Posa, F., Eds. SPIE, 2009, Vol. 7477.
- 508 22. Breen, E.J.; Jones, R. Attribute openings, thinnings and granulometries. *Computer Vision and  
509 Image Understanding* **1996**, *64*, 377–389.
- 510 23. Urbach, E.R.; Wilkinson, M.H.F. Shape-Only Granulometries and Grey-Scale Shape Filters.  
511 Proc. Int. Symp. Math. Morphology (ISMM) 2002, 2002, pp. 305–314.
- 512 24. Urbach, E.R.; Roerdink, J.B.T.M.; Wilkinson, M.H.F. Connected Shape-Size Pattern Spectra for  
513 Rotation and Scale-Invariant Classification of Gray-Scale Images. *IEEE Transactions on Pattern  
514 Analysis and Machine Intelligence* **2007**, *29*, 272–285.
- 515 25. Gueguen, L.; Soille, P.; Pesaresi, M. Differential morphological decomposition segmentation: a  
516 multi-scale object based image representation. Proc. 20th Int. Conf. on Pattern Recognition,  
517 ICPR 2010; IEEE Computer Society: Istanbul, Turkey, 2010; pp. 938–941.
- 518 26. Gueguen, L.; Pesaresi, M.; Soille, P. An Interactive Image Mining Tool Handling Gigapixel  
519 Images. IEEE IGARSS'11; IEEE Computer Society: Vancouver, Canada, 2011; pp. 1581–1584.
- 520 27. Akçay, H.G.; Aksoy, S. Automatic Detection of Geospatial Objects Using Multiple Hierarchical  
521 Segmentations. *IEEE Transactions on Geoscience and Remote Sensing* **2008**, *46*, 2097–2111.
- 522 28. Murra, M.D.; Benediktsson, J.; Waske, B.; Bruzzone, L. Extended profiles with morphological  
523 attribute filters for the analysis of hyperspectral data. *Int. Journal of Remote Sensing* **2010**,  
524 *31*, 5975–5991.
- 525 29. Ouzounis, G.K.; Soille, P.; Pesaresi, M. Rubble detection from VHR aerial imagery data using  
526 differential morphological profiles. 34th Int. Symp. Remote Sensing of the Environment; , 2011.
- 527 30. Gueguen, L.; Soille, P.; Pesaresi, M. Structure extraction and characterization from differential  
528 morphological profile. Proc. 7th Conf. Image Information Mining; , 2011; pp. 53–57.
- 529 31. Jin, X.; Davis, C.H. Automated building extraction from high-resolution satellite imagery in  
530 urban areas using structural, contextual, and spectral information. *EURASIP Journal on Applied  
531 Signal Processing* **2005**, *14*, 20196–2206.
- 532 32. Shyu, C.R.; Scott, G.; Klaric, M.; Davis, C.H.; Palaniappan, K. Automatic Object Extraction  
533 From Full Differential Morphological Profile in Urban Imagery for Efficient Object Indexing and  
534 Retrievals. Proc. 3rd Int. Symp. Remote Sensing and Data Fusion Over Urban Areas (URBAN  
535 2005); , 2005.
- 536 33. Pesaresi, M.; Ouzounis, G.K.; Gueguen, L. A new compact representation of morphological  
537 profiles: report on first massive VHR image processing at the JRC. Proc. SPIE 8390, Algorithms

- 538 and Technologies for Multispectral, Hyperspectral, and Ultraspectral Imagery XVIII; , 2012; pp.  
539 839025–839025–6.
- 540 34. Pesaresi, M.; Kanellopoulos, I. Detection of urban features using morphological based  
541 segmentation and very high resolution remotely sensed data. In *Machine Vision and Advanced*  
542 *Image Processing in Remote Sensing*; Springer, 1999; pp. 271–284.
- 543 35. Pesaresi, M.; Benediktsson, J.A. Image segmentation based on the derivative of the  
544 morphological profile. In *Mathematical morphology and its applications to image and signal*  
545 *processing*; Springer, 2000; pp. 179–188.
- 546 36. Beucher, S. Numerical residues. *Image and Vision Computing* **2007**, *25*, 405–415.
- 547 37. Retornaz, T.; Marcotegui, B. Scene text localization based on the ultimate opening. International  
548 Symposium on Mathematical Morphology, 2007, Vol. 1, pp. 177–188.
- 549 38. Hernández, J.; Marcotegui, B. Shape ultimate attribute opening. *Image and Vision Computing*  
550 **2011**, *29*, 533–545.
- 551 39. Pesaresi, M.; Huadong, G.; Blaes, X.; Ehrlich, D.; Ferri, S.; Gueguen, L.; Halkia, M.; Kauffmann,  
552 M.; Kemper, T.; Lu, L.; Marin-Herrera, M.; Ouzounis, G.; Scavazzon, M.; Soille, P.; Syrris, V.;  
553 Zanchetta, L. A global human settlement layer from optical HR/VHR RS data: Concept and first  
554 results. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*  
555 **2013**, *6*, 2102–2131.
- 556 40. Florczyk, A.J.; Ferri, S.; Syrris, V.; Kemper, T.; Halkia, M.; Soille, P.; Pesaresi, M. A New  
557 European Settlement Map From Optical Remotely Sensed Data. *IEEE Journal of Selected Topics*  
558 *in Applied Earth Observations and Remote Sensing* **2015**. Article in Press.
- 559 41. Ehrlich, D.; Tenerelli, P. Optical satellite imagery for quantifying spatio-temporal dimension of  
560 physical exposure in disaster risk assessments. *Natural Hazards* **2013**, *68*, 1271–1289.
- 561 42. Lu, L.; Guo, H.; Pesaresi, M.; Soille, P.; Ferri, S. Automatic recognition of built-up areas in  
562 China using CBERS-2B HR data. 2013, pp. 65–68.
- 563 43. Song, B.; Li, J.; Dalla Mura, M.; Li, P.; Plaza, A.; Bioucas-Dias, J.; Benediktsson, J.; Chanussot,  
564 J. Remotely sensed image classification using sparse representations of morphological attribute  
565 profiles. *IEEE Transactions on Geoscience and Remote Sensing* **2014**, *52*, 5122–5136.
- 566 44. Kemper, T.; Mudau, N.; Mangara, P.; Pesaresi, M. Towards an automated monitoring of human  
567 settlements in South Africa using high resolution SPOT satellite imagery. 2015, Vol. 40, pp.  
568 1389–1394.
- 569 45. Horvat, D.; Aalik, B.; Rupnik, M.; Mongus, D. Visualising the attributes of biological cells,  
570 based on human perception. *Lecture Notes in Computer Science (including subseries Lecture*  
571 *Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* **2013**, *7947 LNCS*, 386–399.
- 572 46. Vincent, L. Greyscale area openings and closings, their efficient implementation and applications.  
573 Proc. EURASIP Workshop on Mathematical Morphology and its Applications to Signal  
574 Processing; , 1993; pp. 22–27.
- 575 47. Maragos, P.; Ziff, R.D. Threshold Superposition in Morphological Image Analysis Systems.  
576 *IEEE Transactions on Pattern Analysis and Machine Intelligence* **1990**, *12*, 498–504.
- 577 48. Meijster, A.; Wilkinson, M.H.F. A comparison of algorithms for connected set openings and  
578 closings. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **2002**, *24*, 484–494.

- 579 49. Najman, L.; Couprie, M. Building the component tree in quasi-linear time. *IEEE Transactions  
580 on Image Processing* **2006**, *15*, 3531–3539.
- 581 50. Menotti-Gomes, D.; Najman, L.; de Albuquerque Araújo, A. 1D Component tree in linear time  
582 and space and its application to gray-level image multithresholding. Proc. Int. Symp. Math.  
583 Morphology, ISMM 20071; Banon, G.J.F.; Barrera, J.; Braga-Neto, U.; Hirata, N.S.T., Eds.  
584 INPE, 2007, Vol. 1, pp. 437–448.
- 585 51. Berger, C.; Geraud, T.; Levillain, R.; Widynski, N.; Baillard, A.; Bertin, E. Effective Component  
586 Tree Computation with Application to Pattern Recognition in Astronomical Imaging. *Image  
587 Processing*, 2007. ICIP 2007. IEEE International Conference on, 2007, Vol. 4, pp. 41–44.
- 588 52. Wilkinson, M.H.F. A Fast Component-Tree Algorithm for High Dynamic-Range Images and  
589 Second Generation Connectivity. Proc. Int. Conf. on Image Processing, ICIP 2011, 2011, pp.  
590 1041–1044.
- 591 53. Urbach, E.R.; Wilkinson, M.H.F. Efficient 2-D gray-scale morphological transformations with  
592 arbitrary flat structuring elements. *IEEE Transactions on Image Processing* **2008**, *17*, 1–8.
- 593 54. Wilkinson, M.H.F.; Moschini, U.; Ouzounis, G.K.; Pesaresi, M. Concurrent computation of  
594 connected pattern spectra for very large image information mining. Proc. of ESA-EUSC-JRC  
595 8th Conference on Image Information Mining; , 2012; pp. 21–25.

596 © March 28, 2017 by the authors; submitted to *ISPRS Int. J. Geo-Inf.* for possible open  
597 access publication under the terms and conditions of the Creative Commons Attribution license  
598 <http://creativecommons.org/licenses/by/4.0/>.