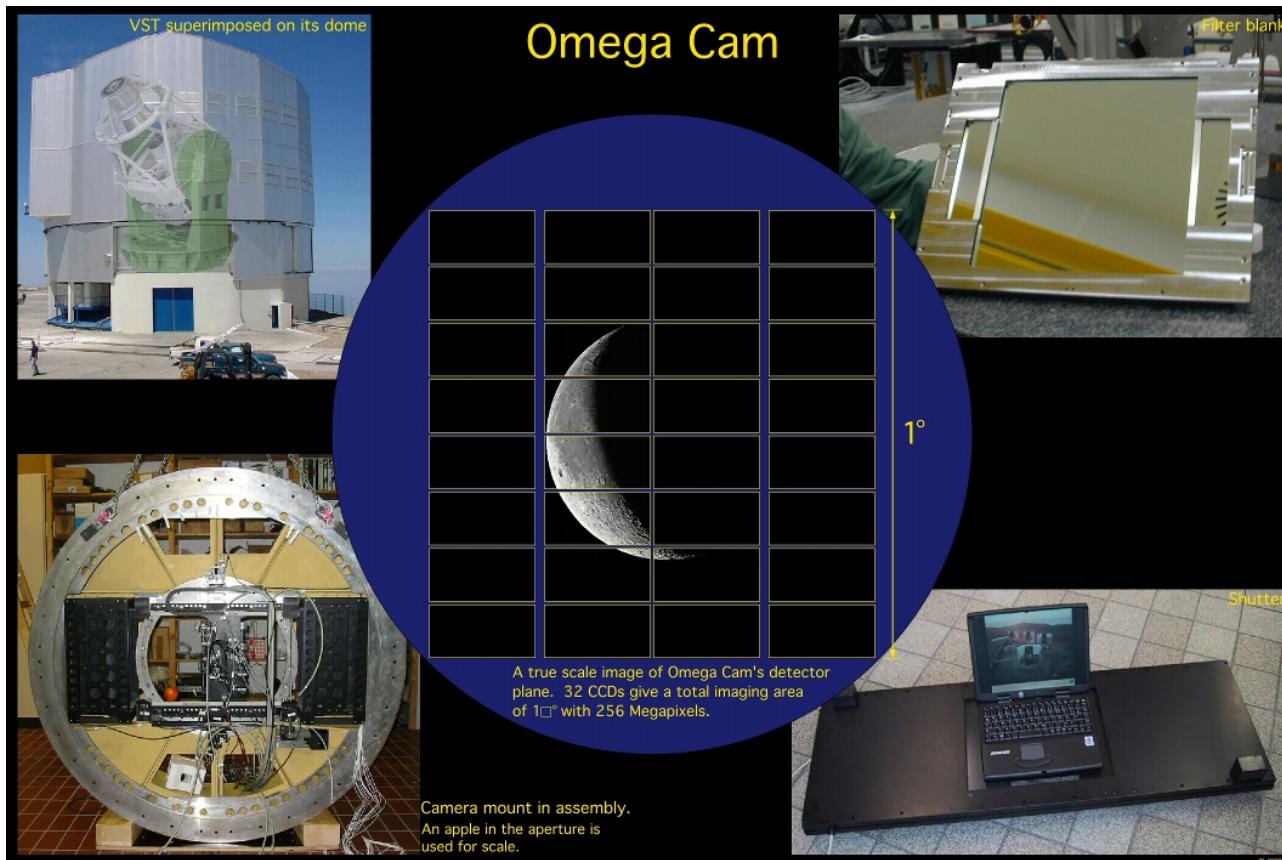


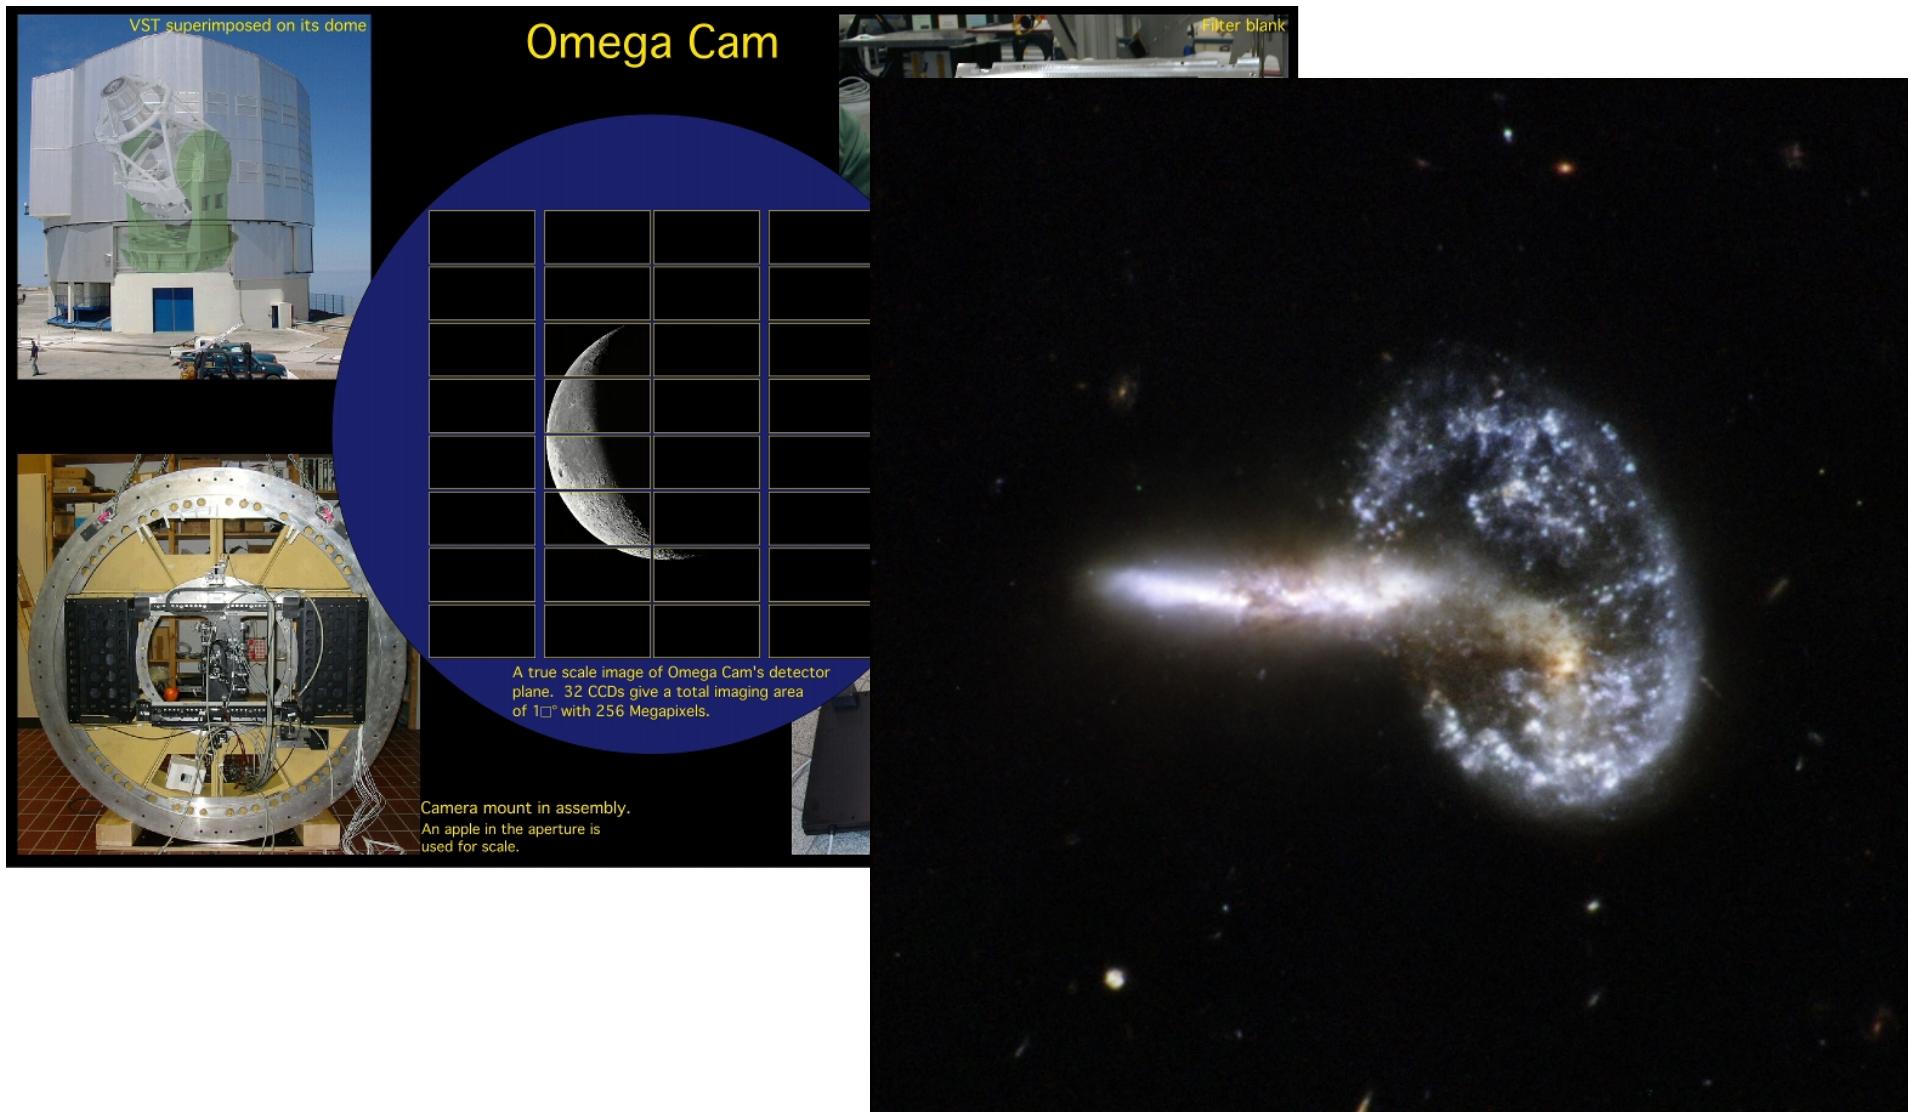


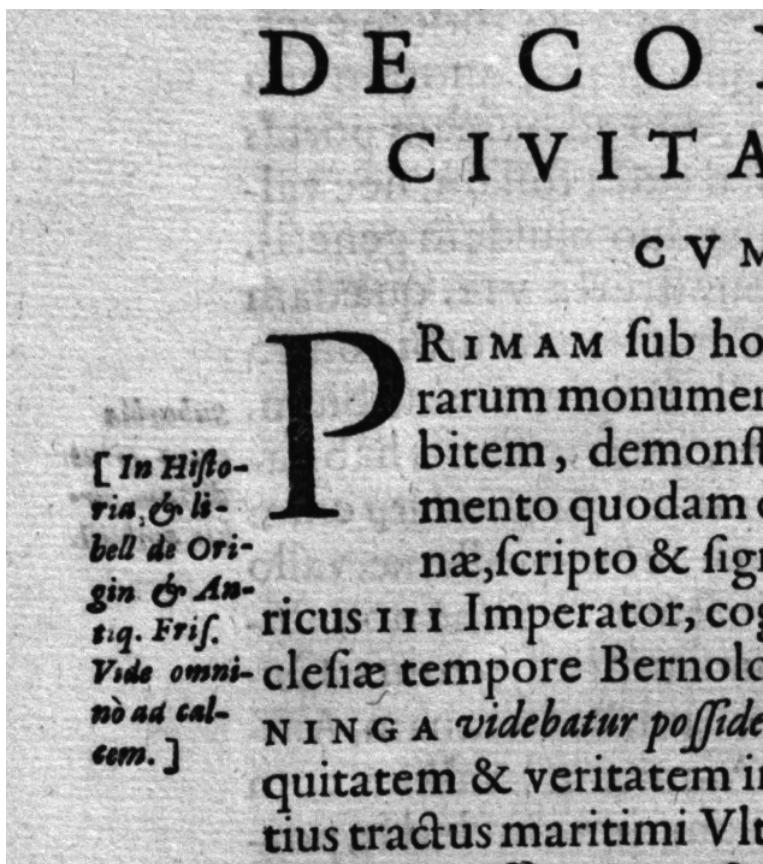
Supercomputing for Giga- and Tera-Pixel Image Scales

Michael H. F. Wilkinson

*Johann Bernoulli Institute
University of Groningen*

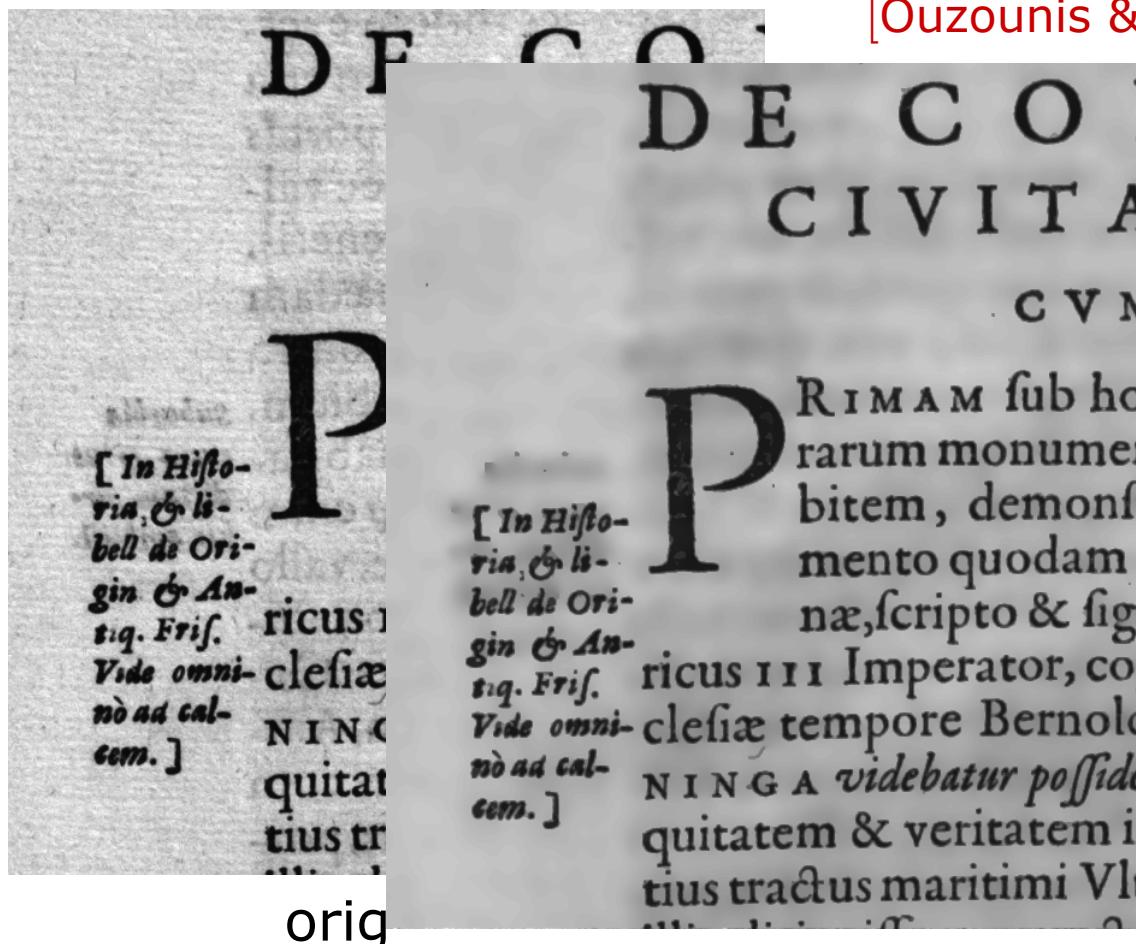






original

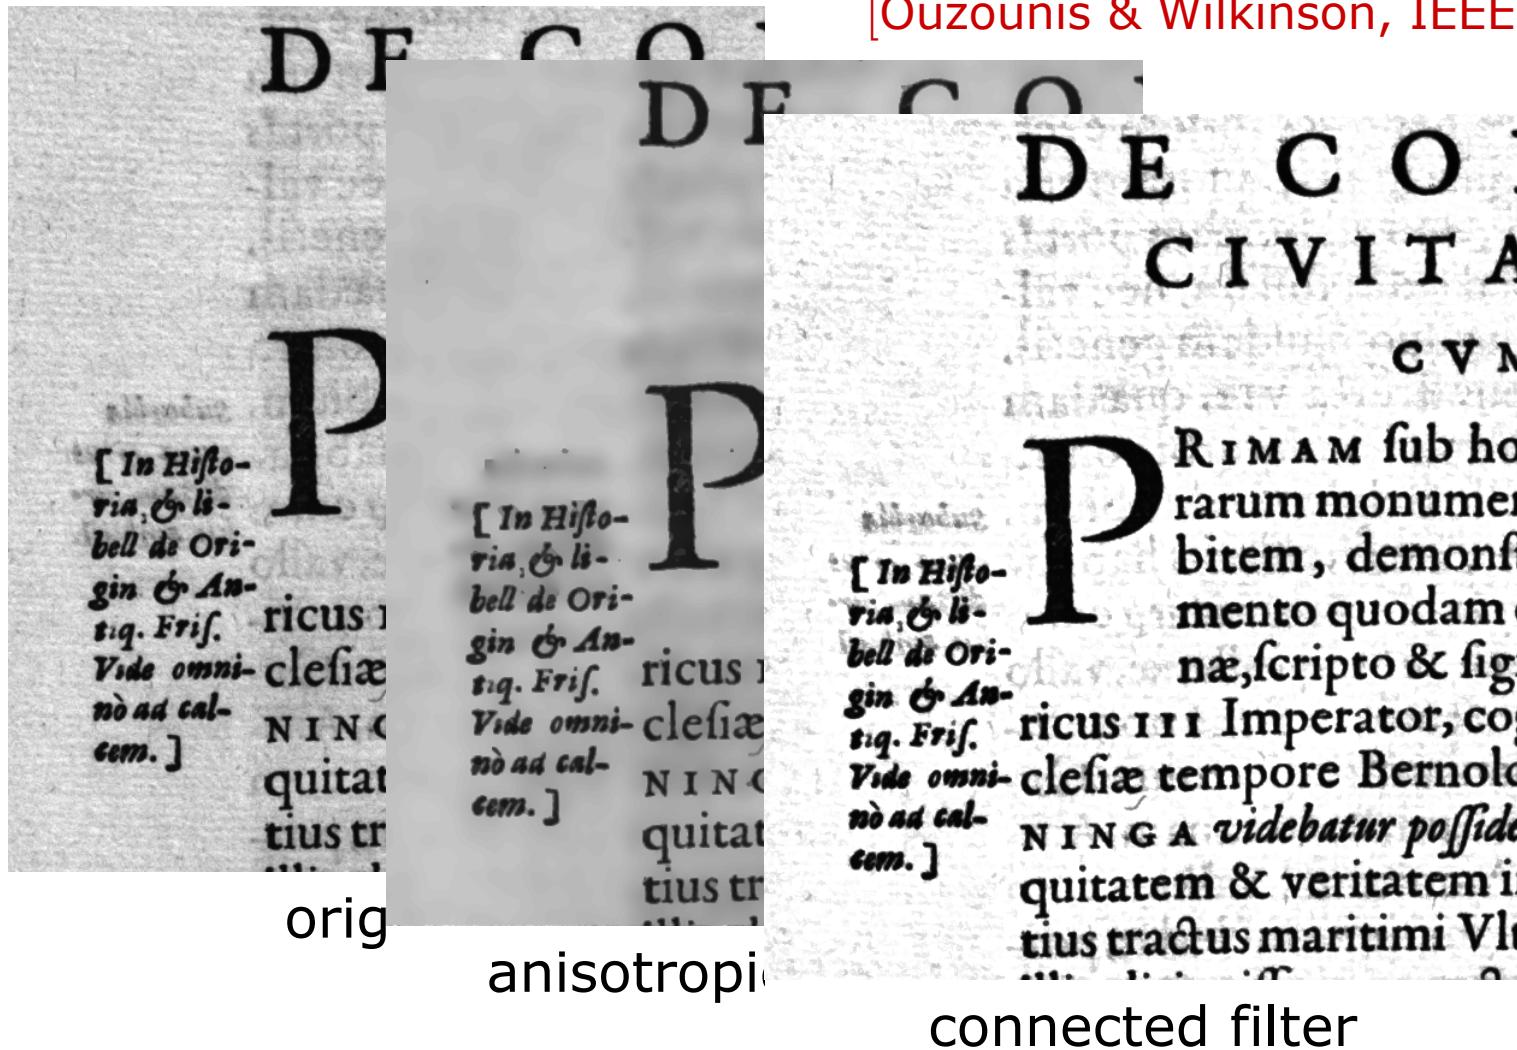
[Ouzounis & Wilkinson, IEEE Trans. PAMI 2011]



anisotropic diffusion

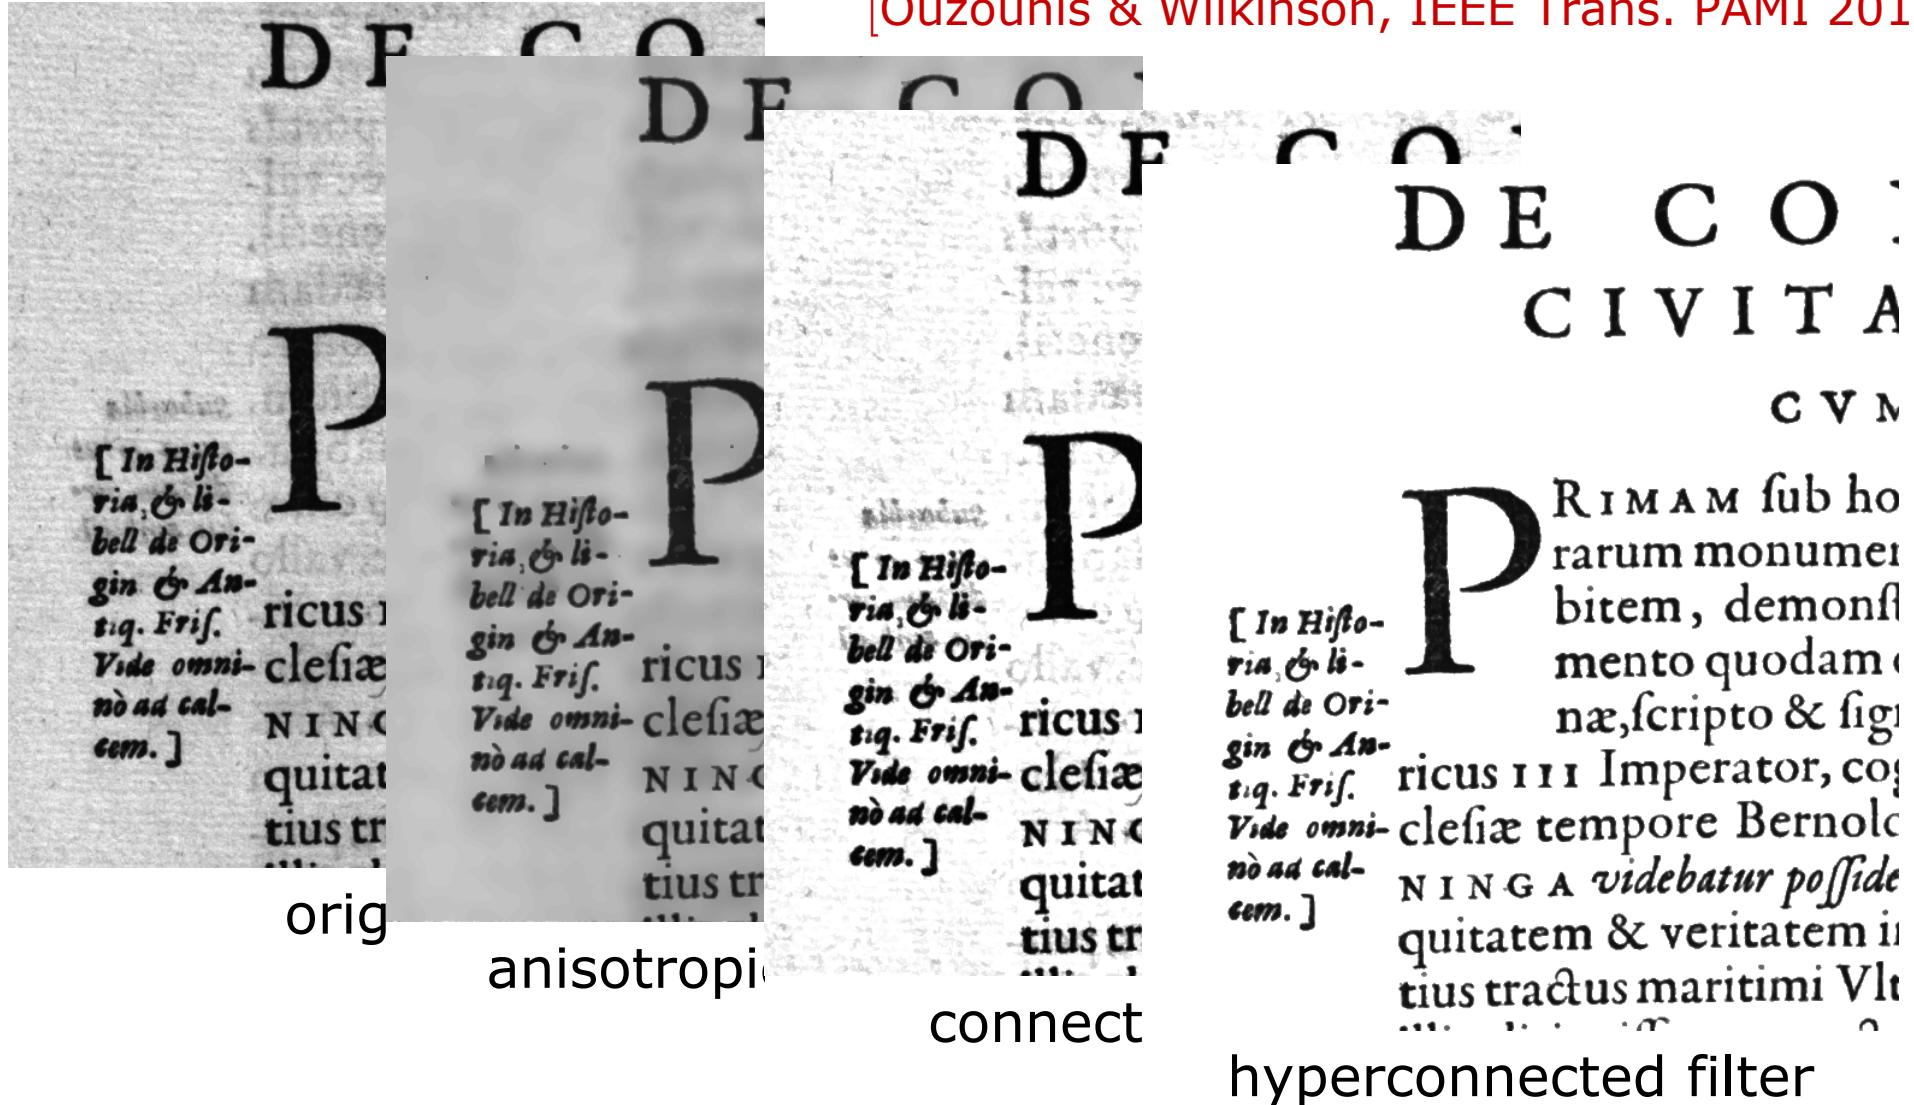


[Ouzounis & Wilkinson, IEEE Trans. PAMI 2011]



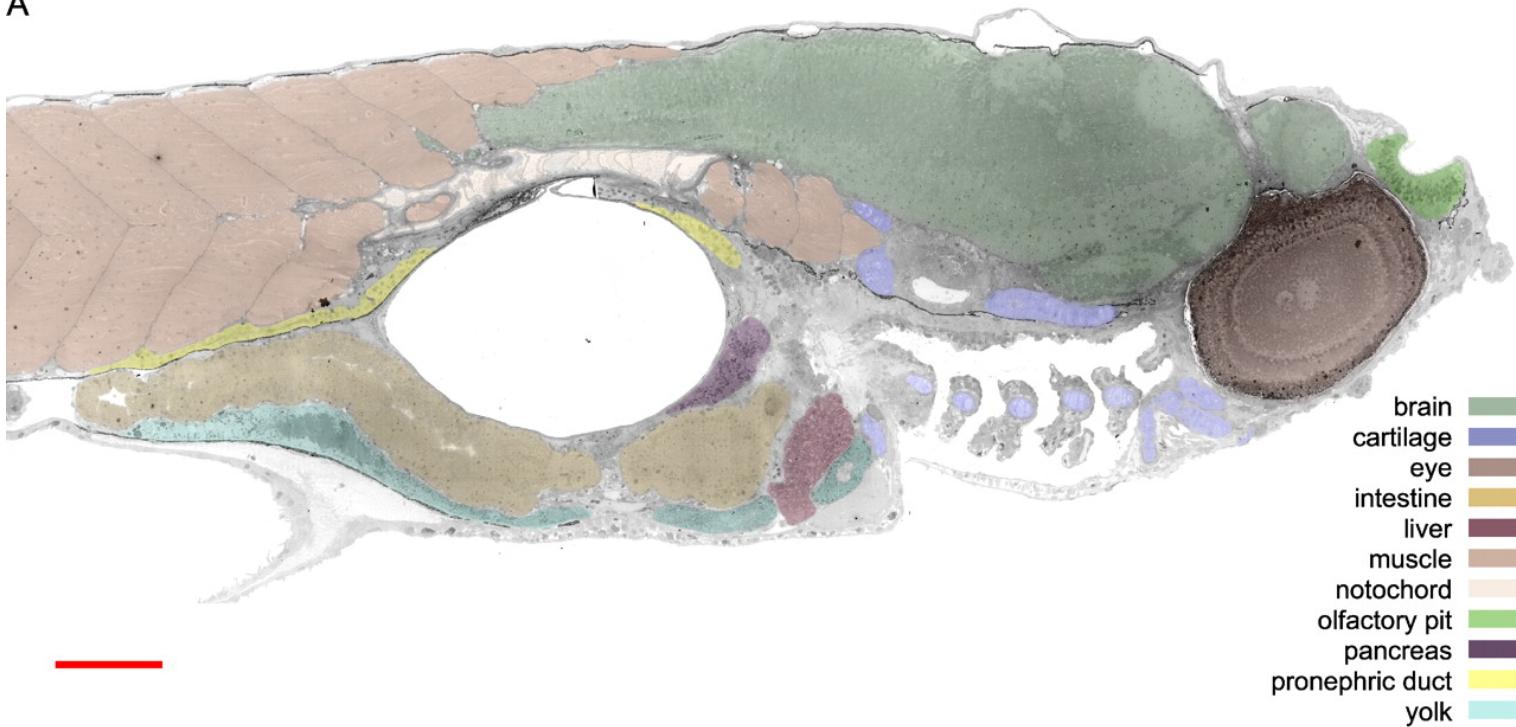


[Ouzounis & Wilkinson, IEEE Trans. PAMI 2011]

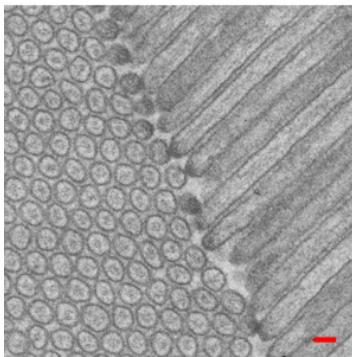




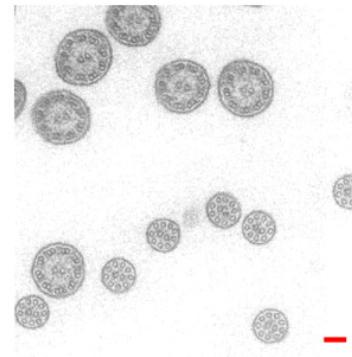
A



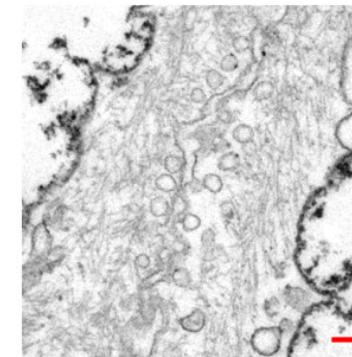
B



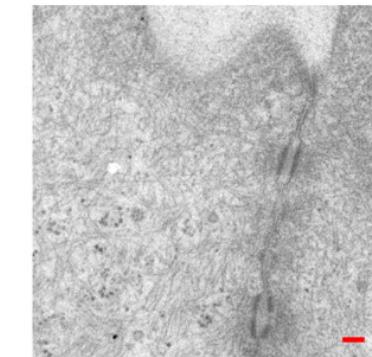
C



D



E

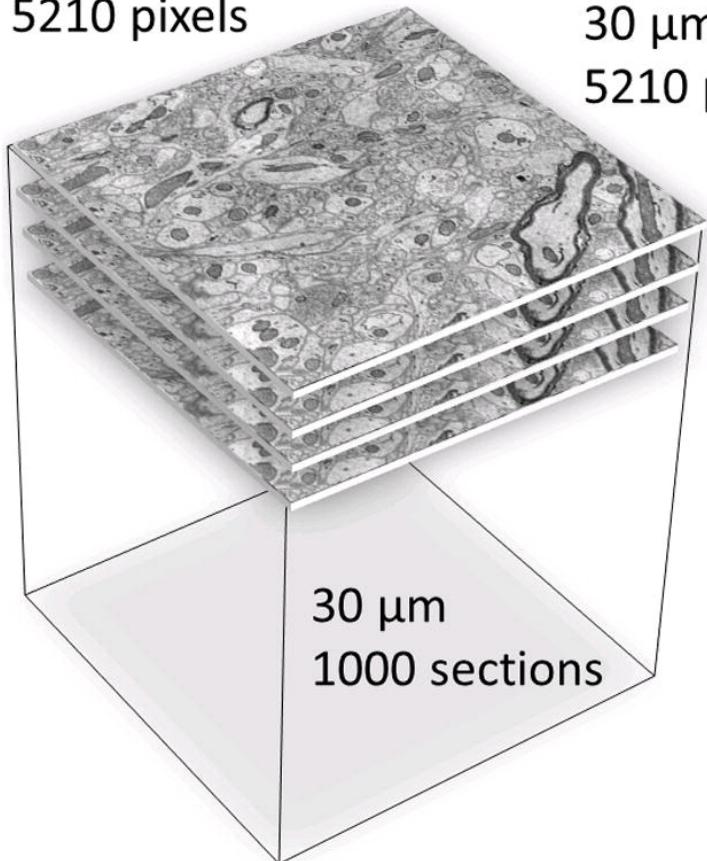


Faas et al (2012)



30 µm

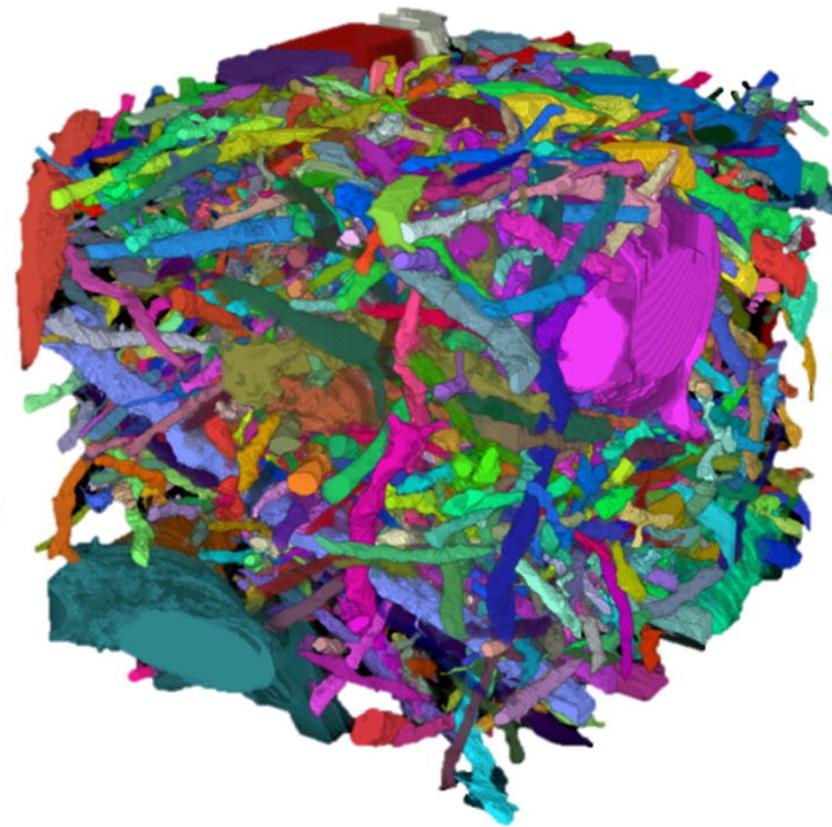
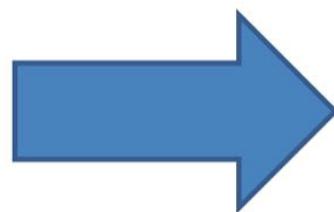
5210 pixels



EM images

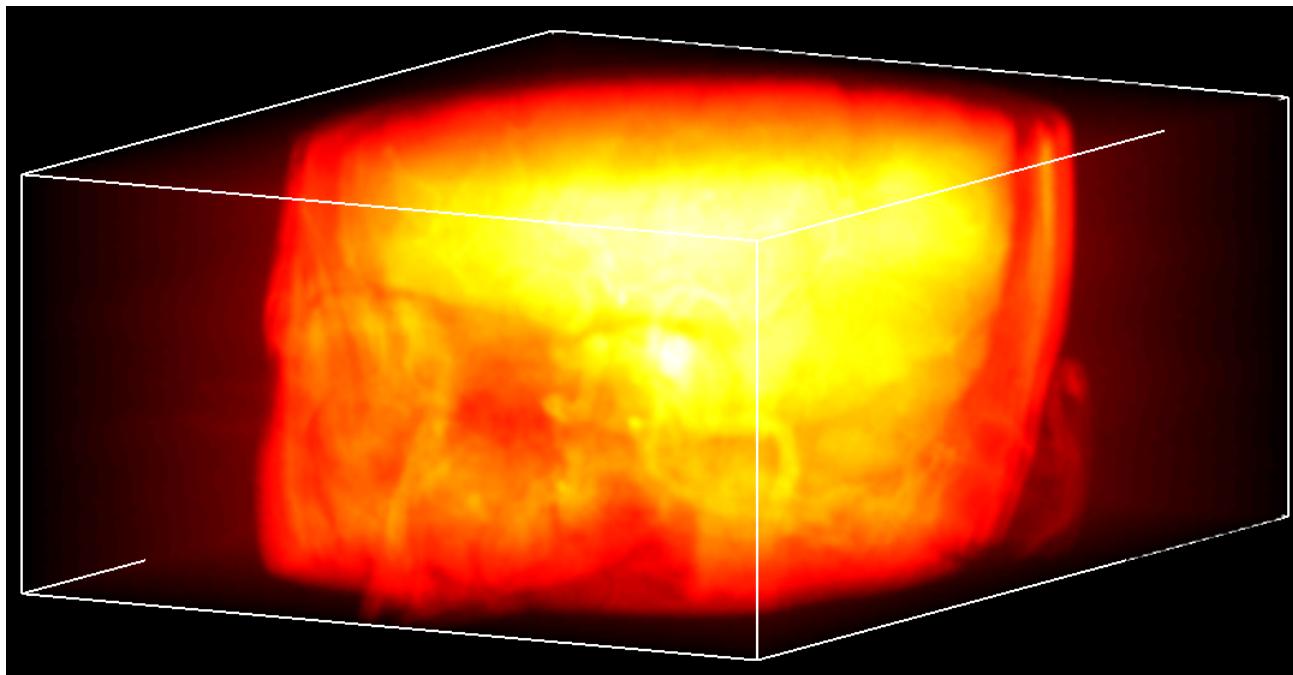
30 µm

5210 pixels

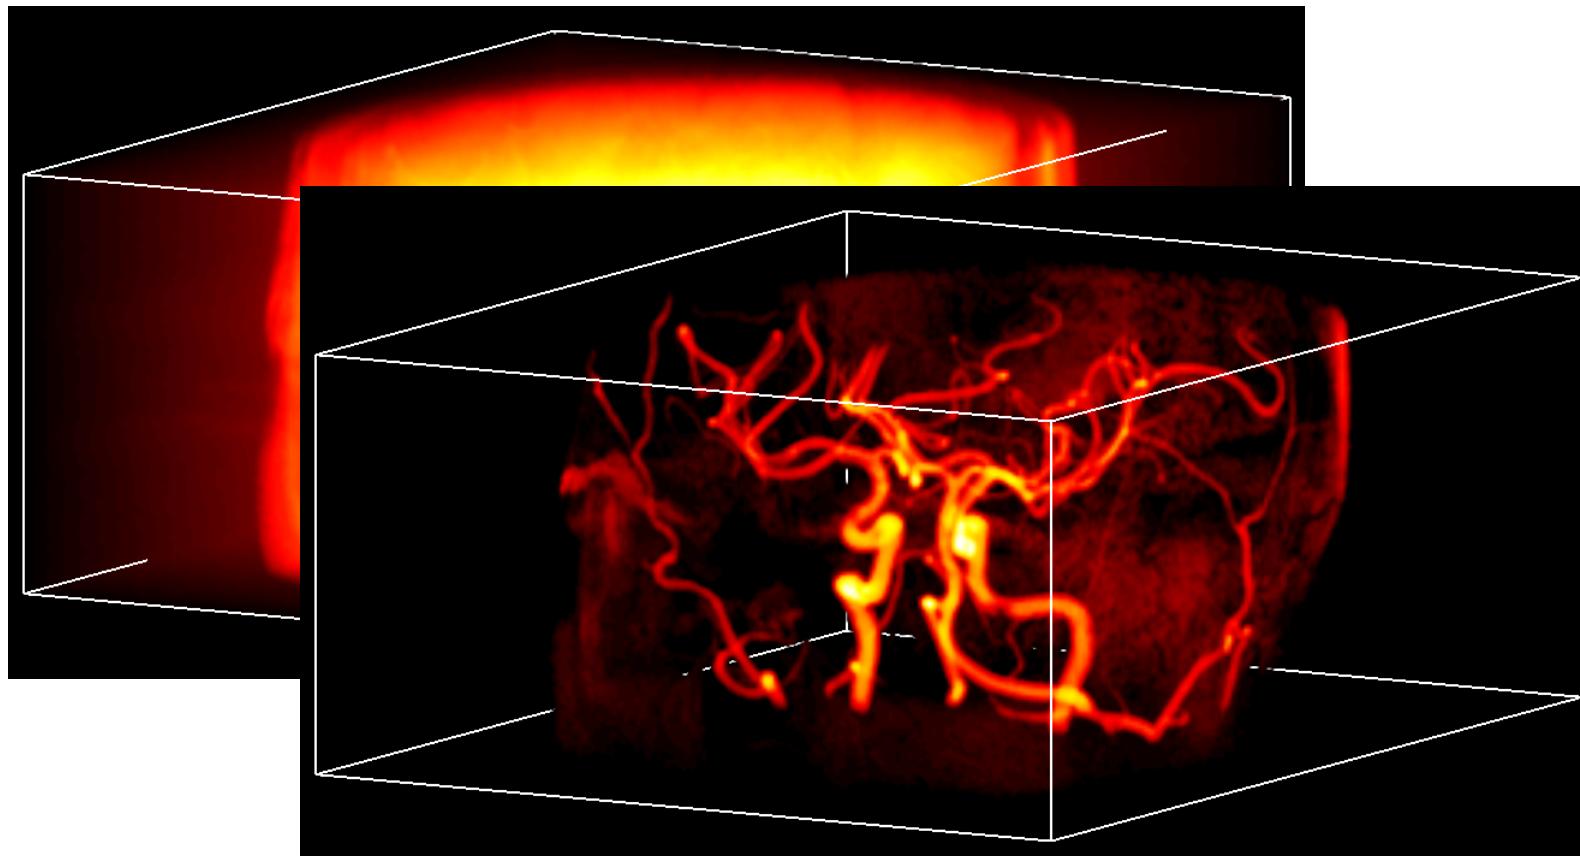


Reconstruction

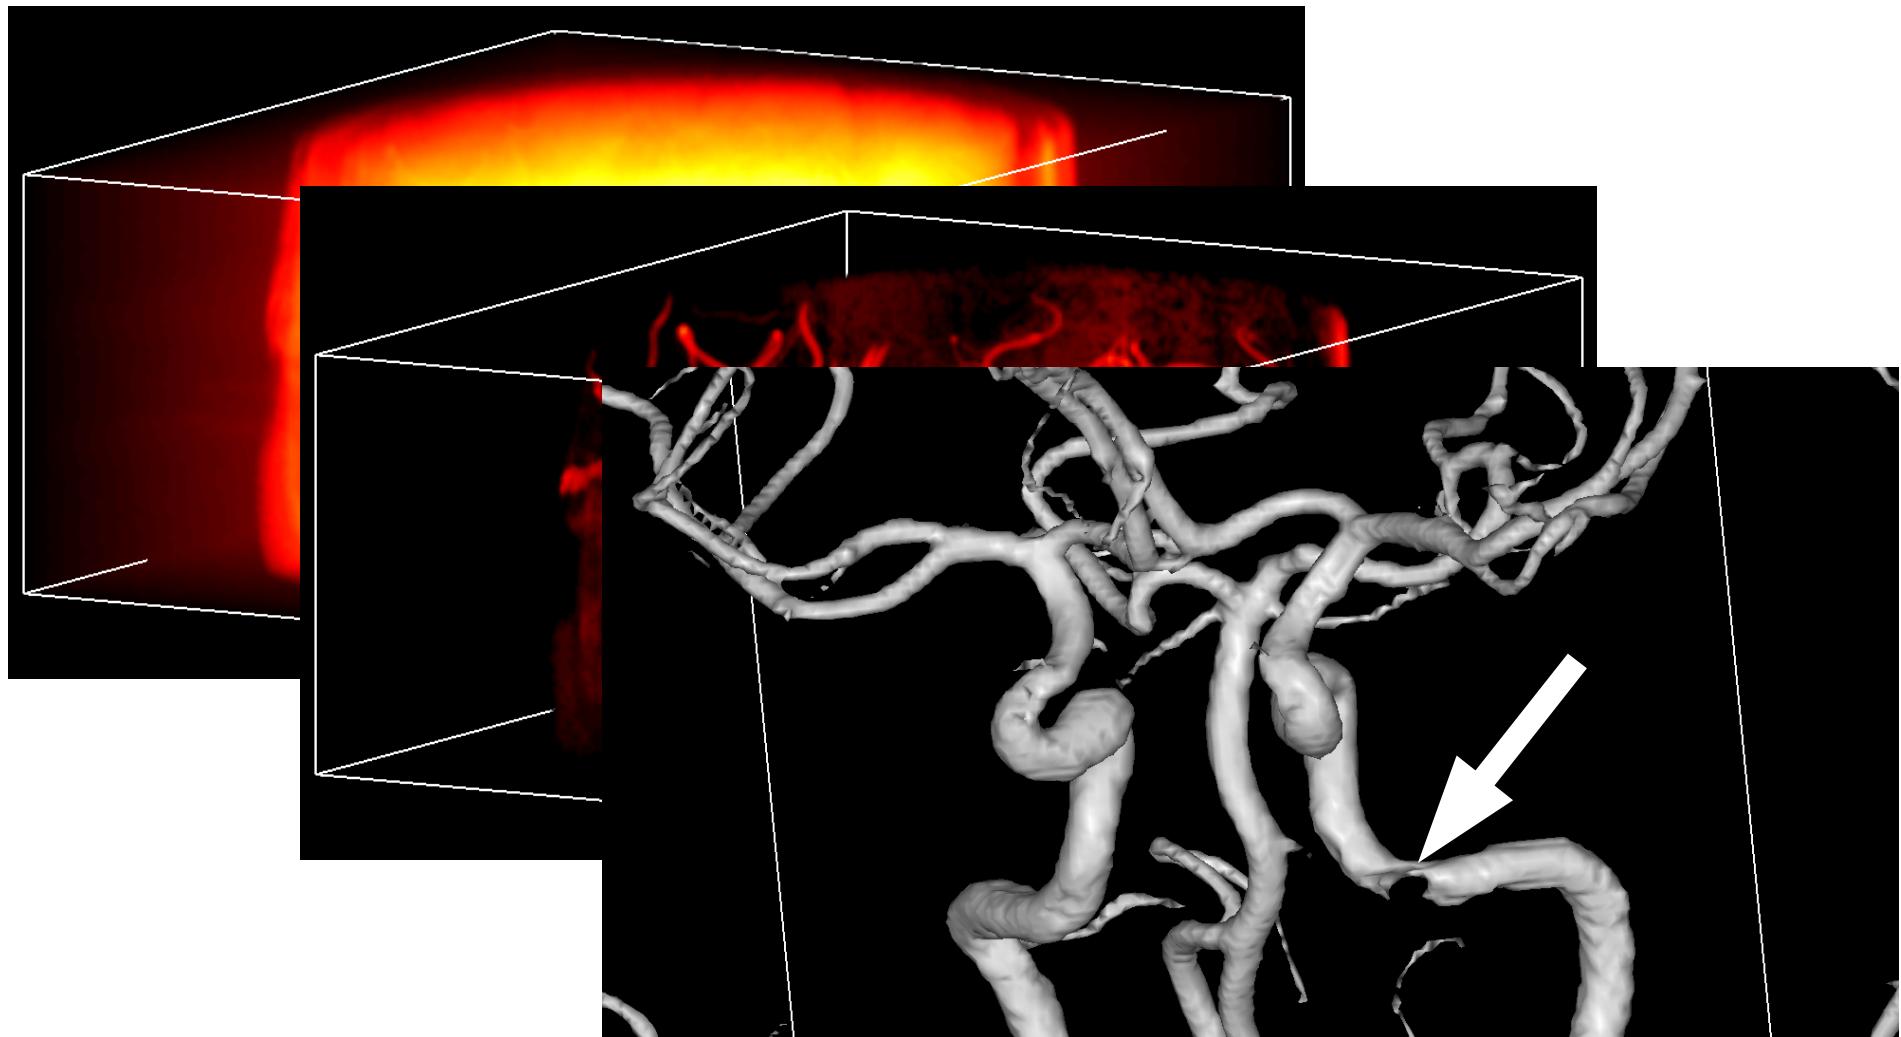
Kaynig et al (2015)



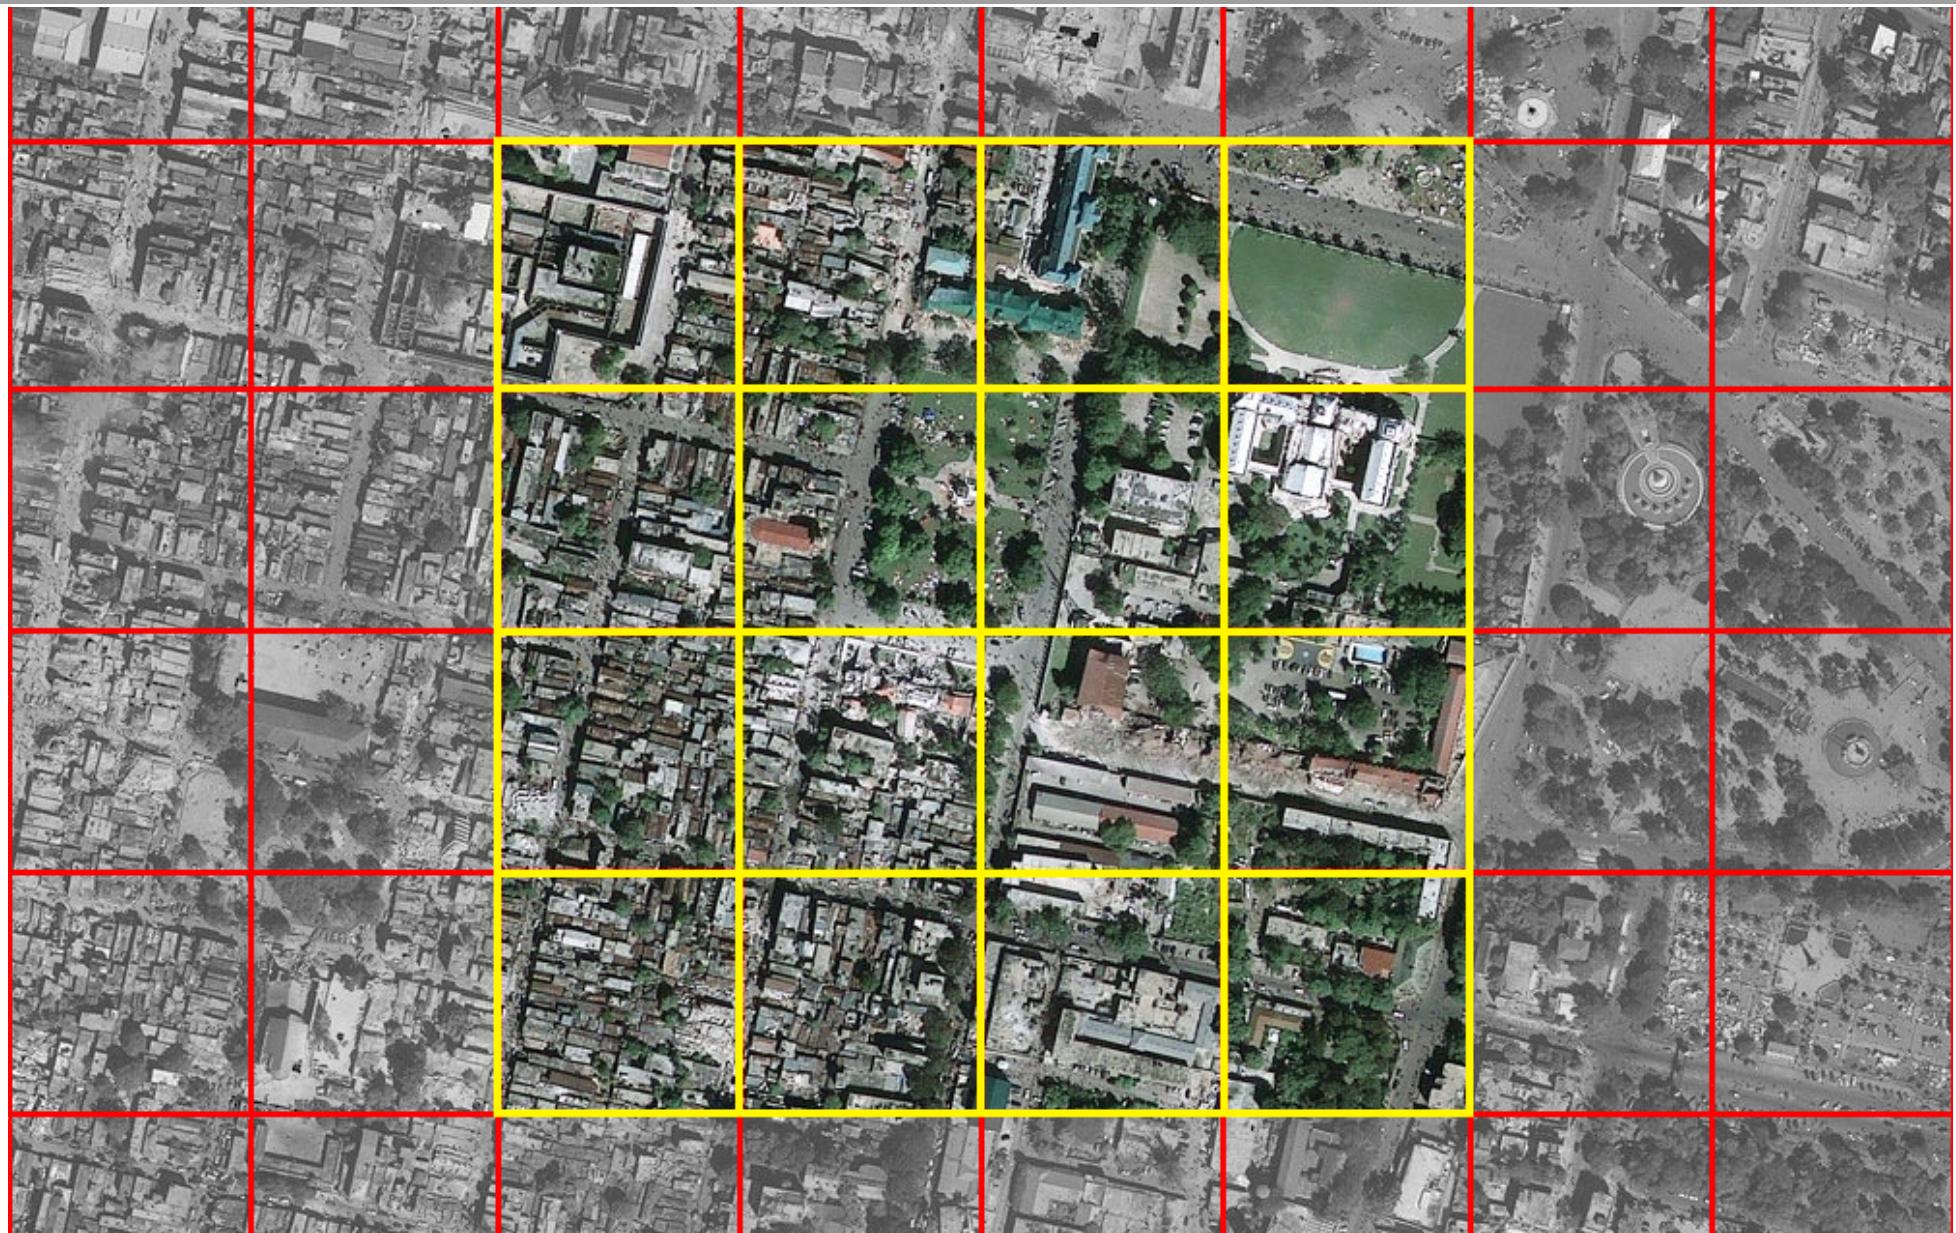
Raw data



Connected filter



Blood clot?

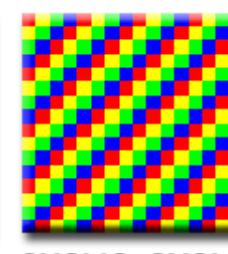
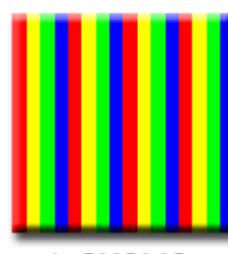
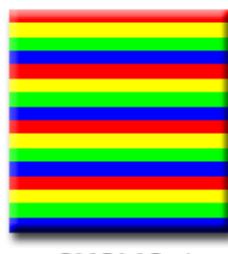
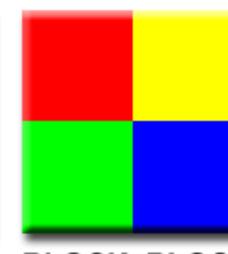




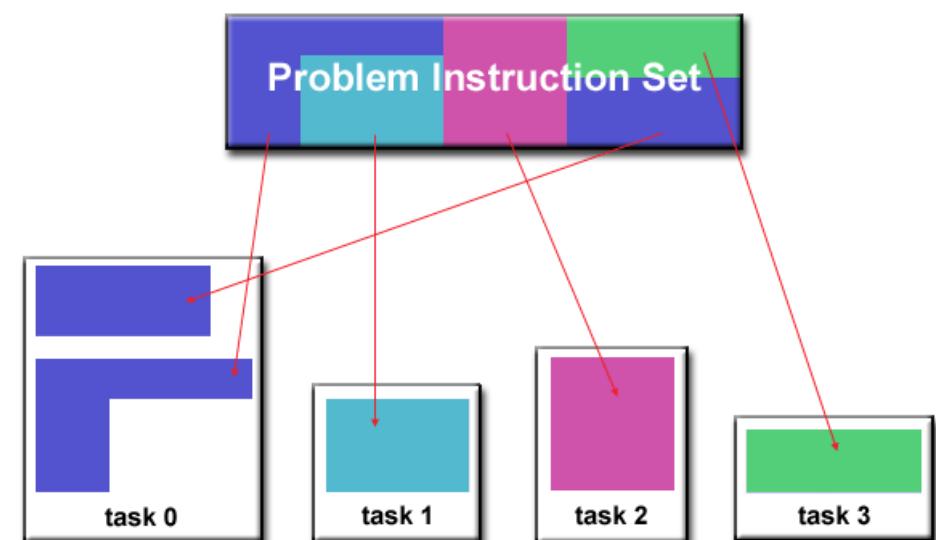
- Most computers are parallel machines
- Partitioning the task is essential



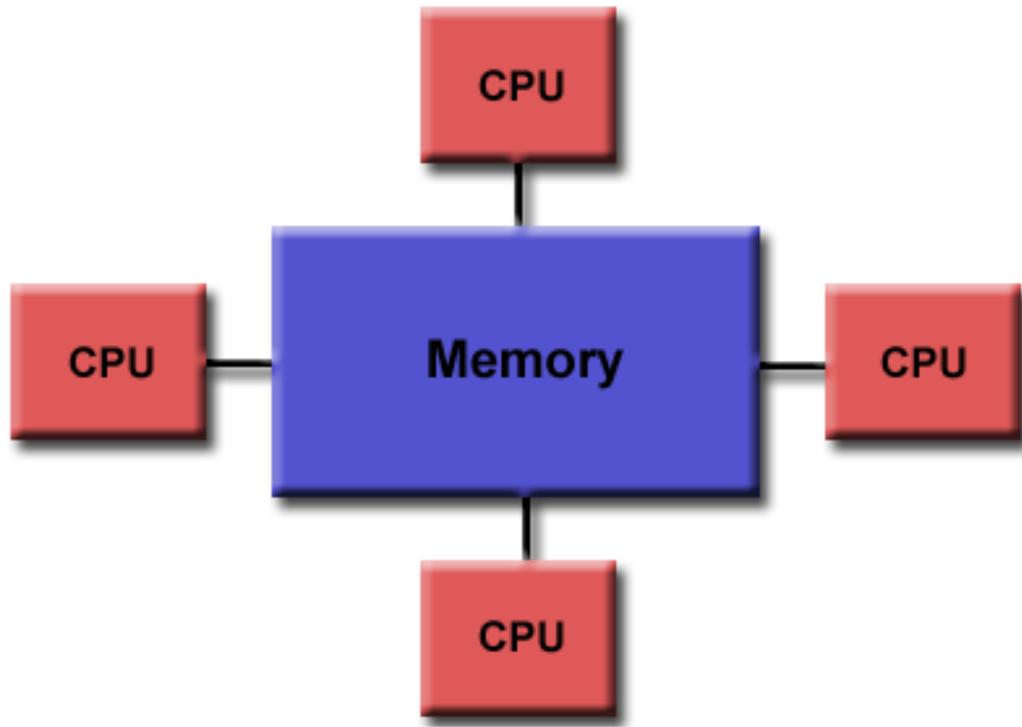
2D



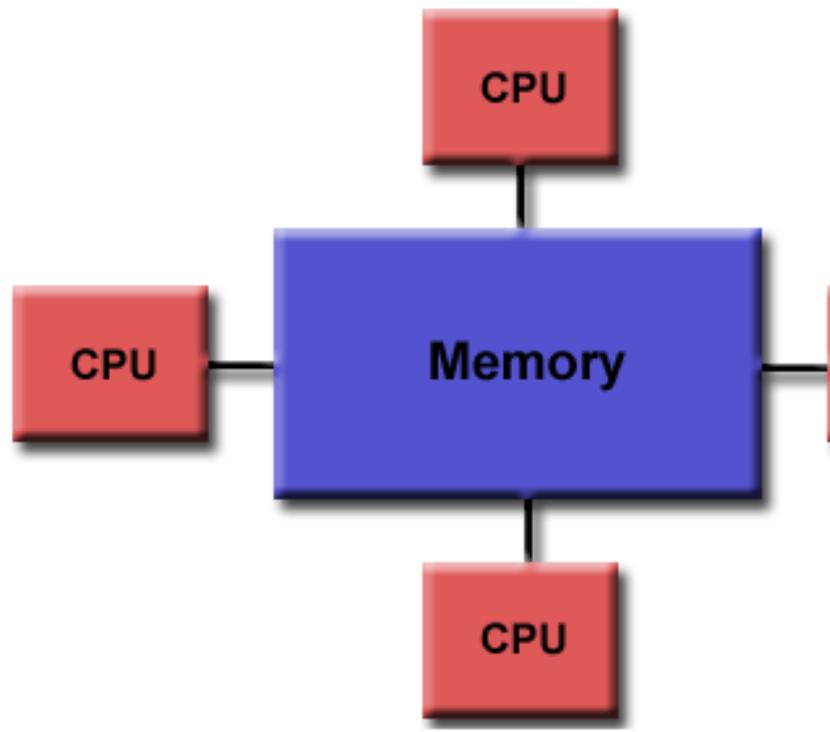
spatial



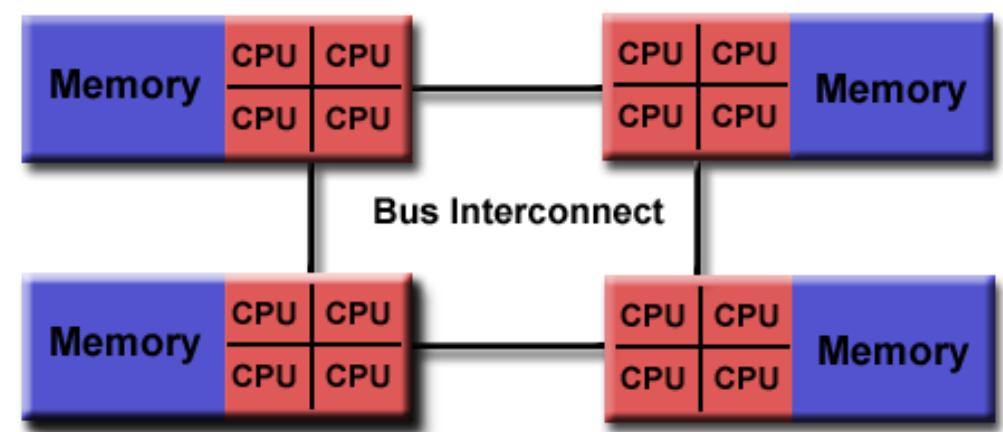
functional



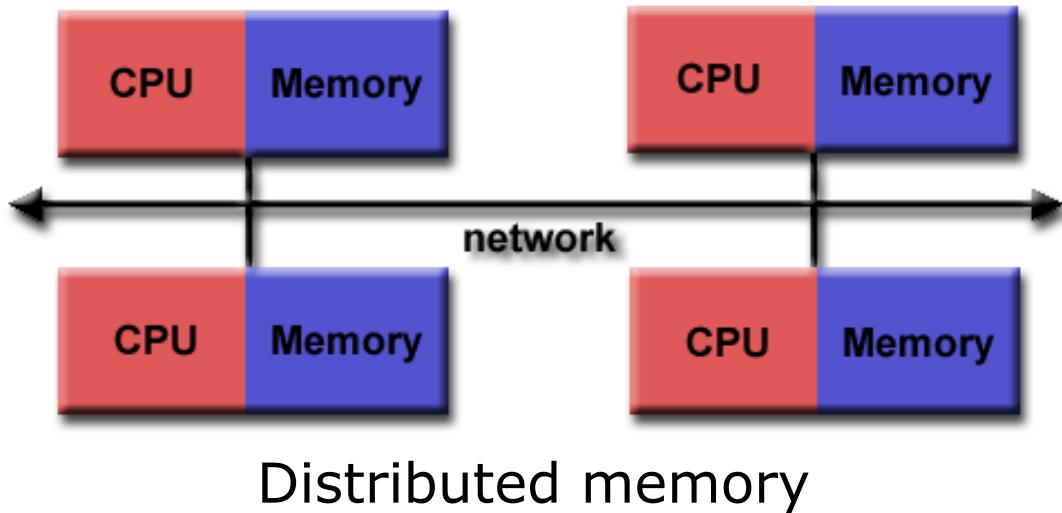
Simple shared memory

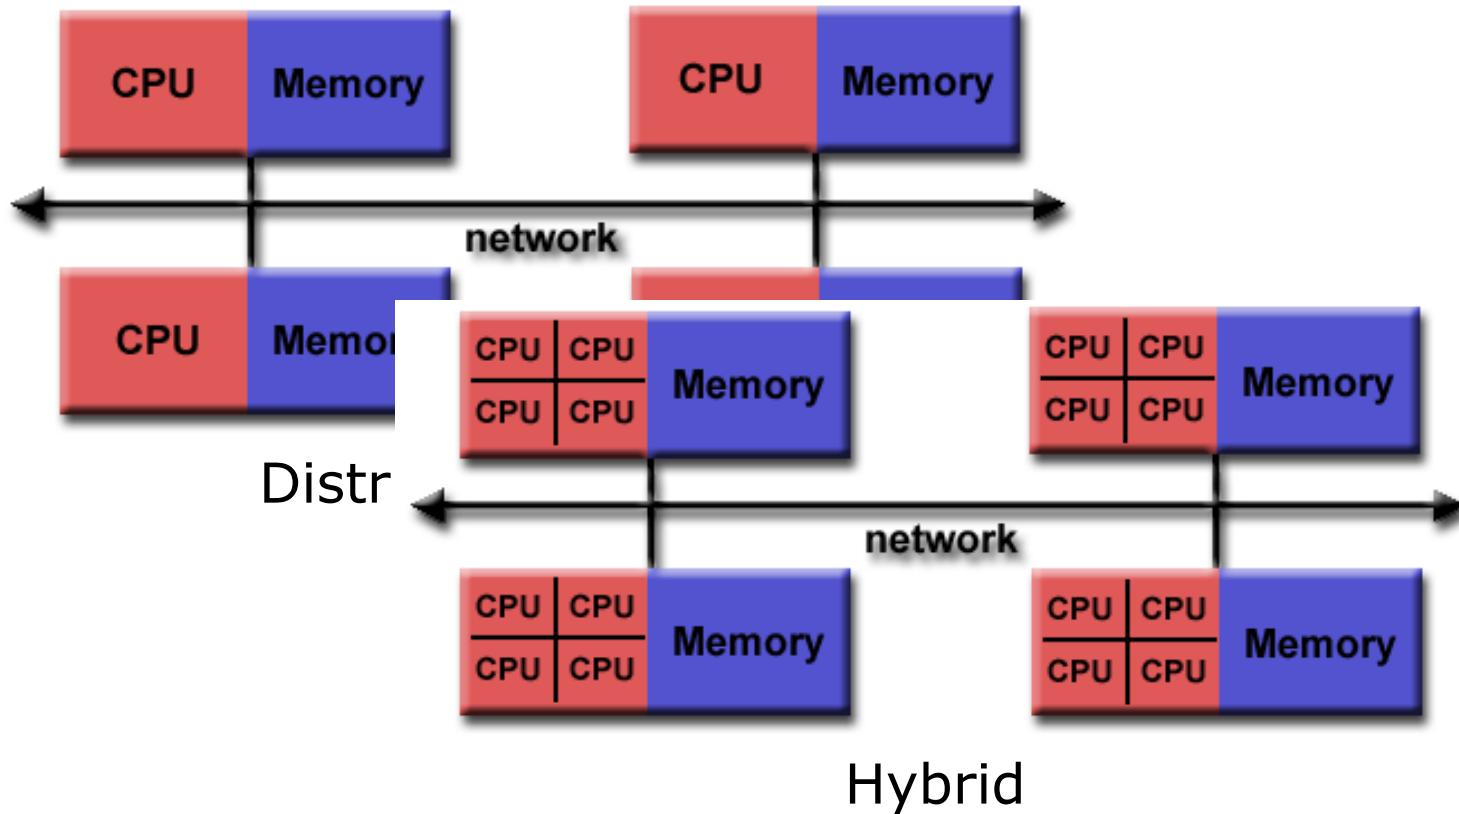


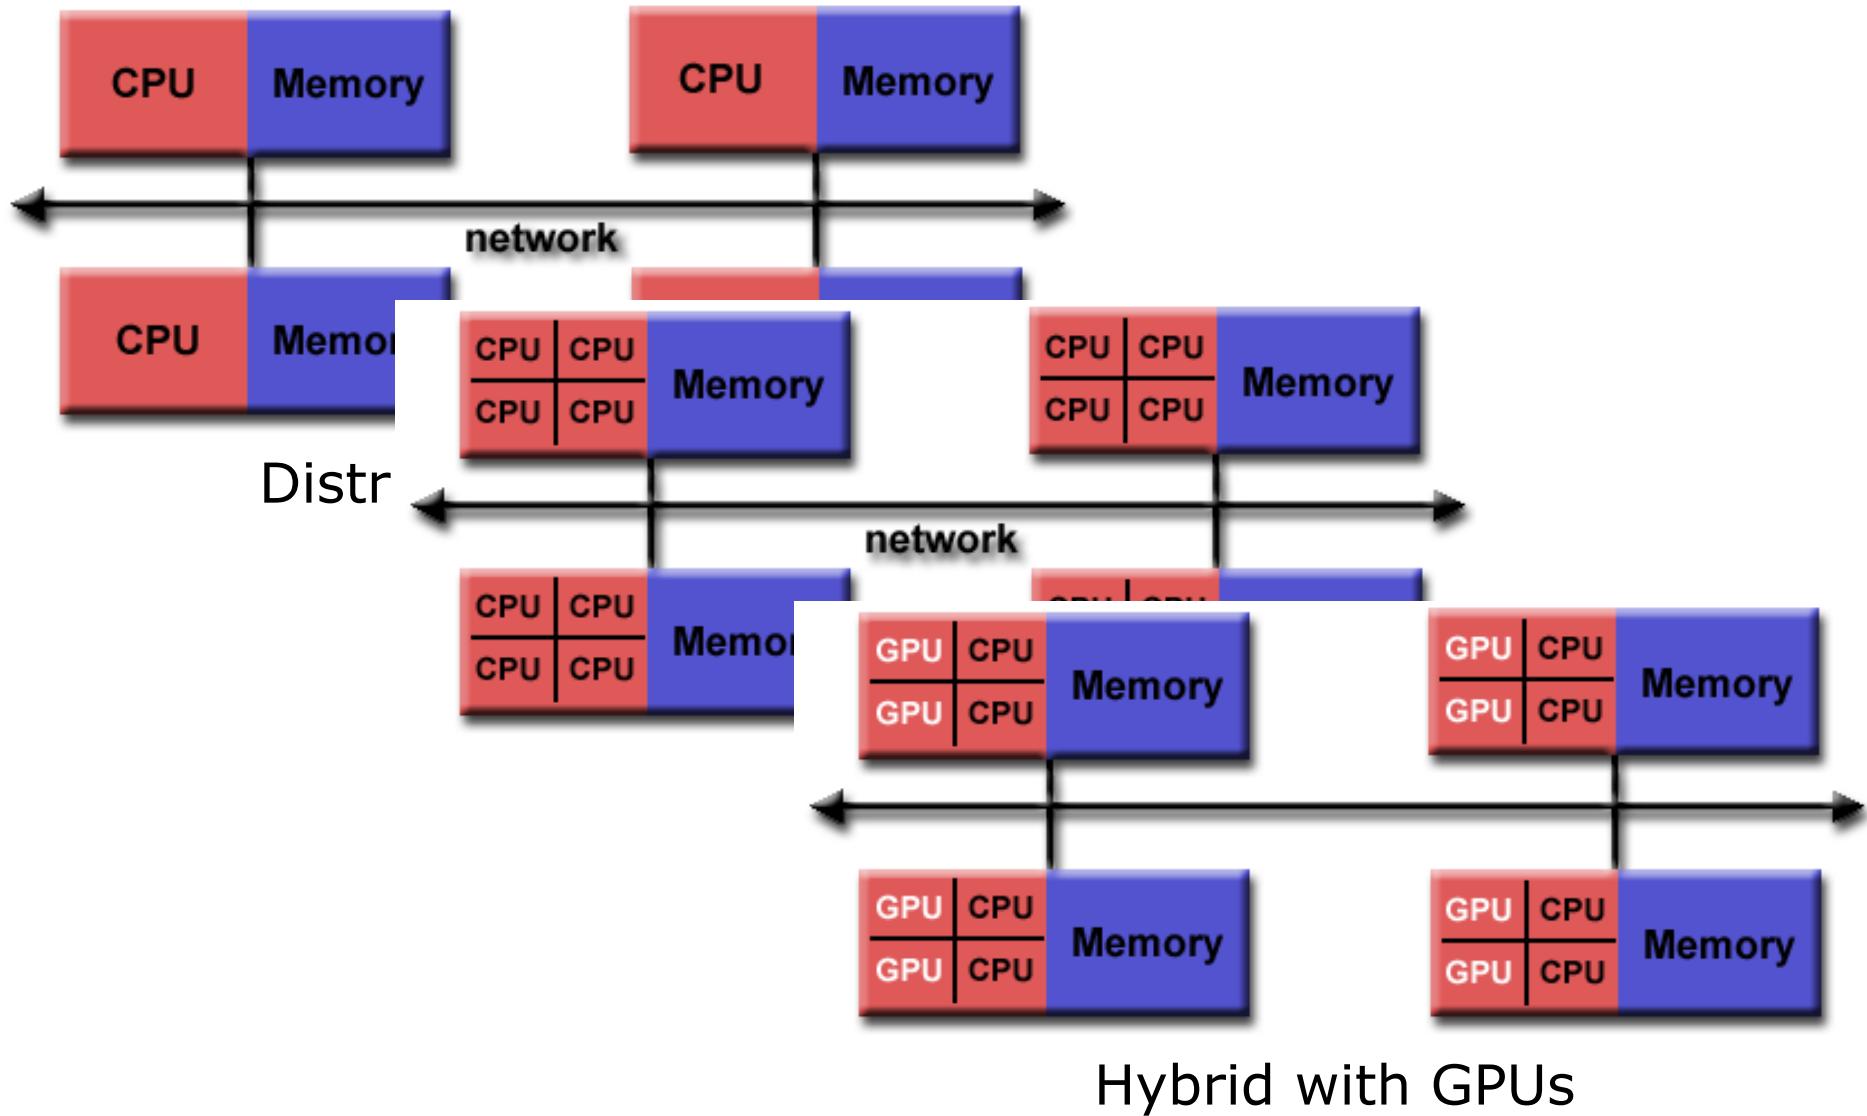
Simple shared memory



NUMA









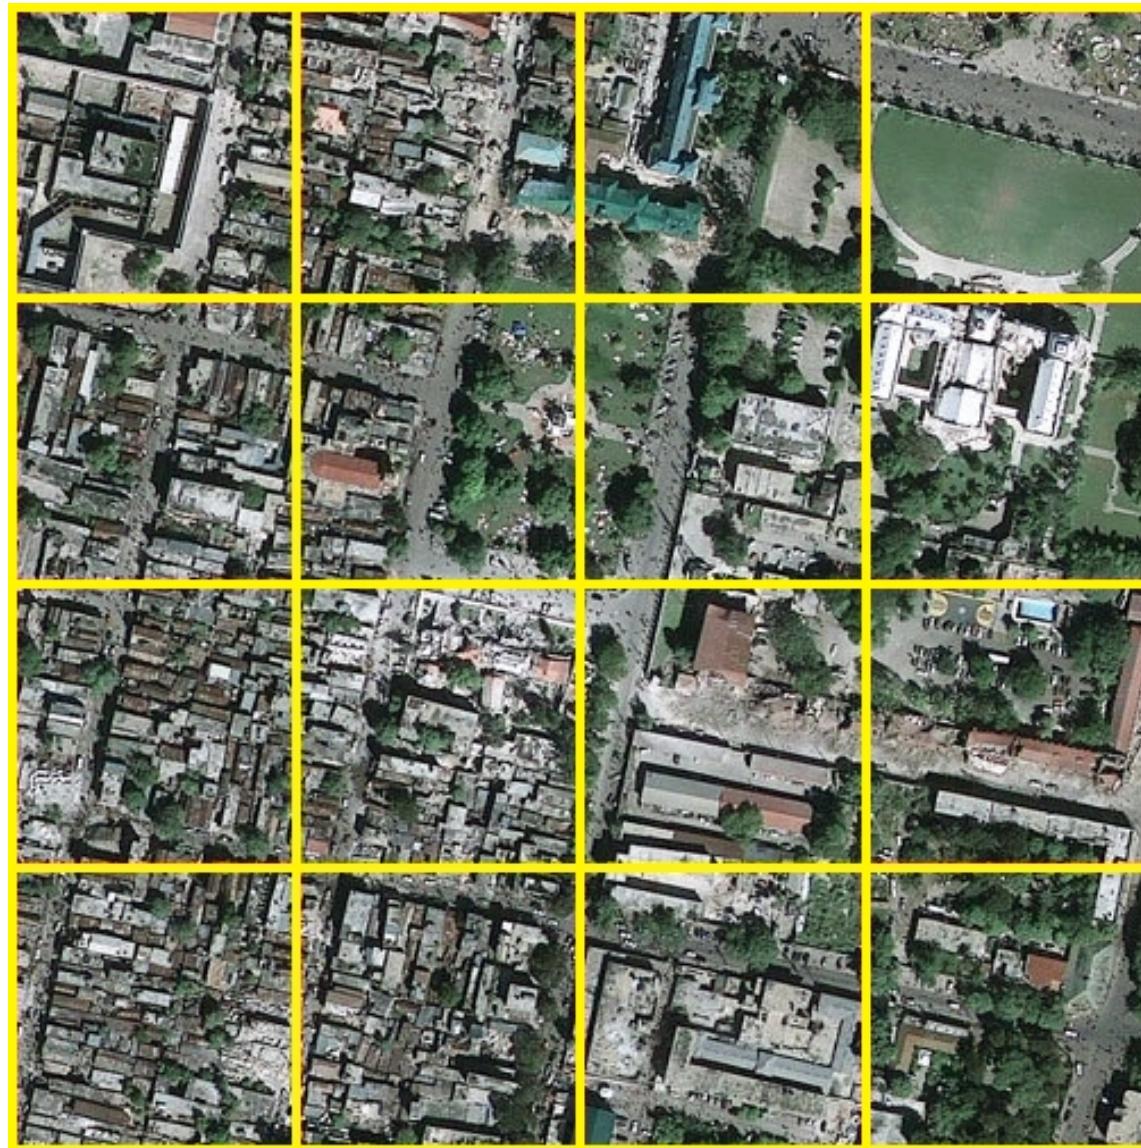
- The big challenge is to deal with many different architectures
- The approach most used is to identify algorithmic patterns
- Build a framework for generic programming based on these patterns
- Express new methods in terms of easily parallelized basic operations
- We then only need to develop specific parallel algorithms for these basic operators for each architecture
- A key problem can be domain decomposition



- Point operators
- Neighbourhood operators
- Geometric operators
- Global image → image operators
 - Separable
 - Non-separable
- Global reduction operators

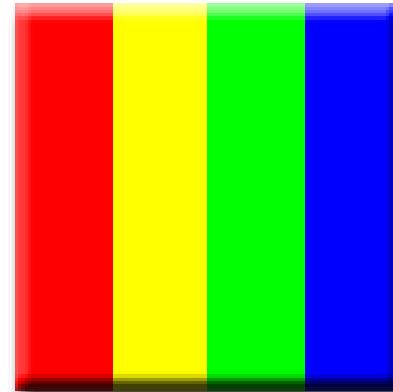


- Unary: Look-up-table (LUT) operators
- Image + constant operators
 - Thresholding
 - Arithmetic functions
- Binary: Arithmetic/logic operators
 - Addition/subtraction
 - Max/min
 - Image comparison
- Trivially parallelized

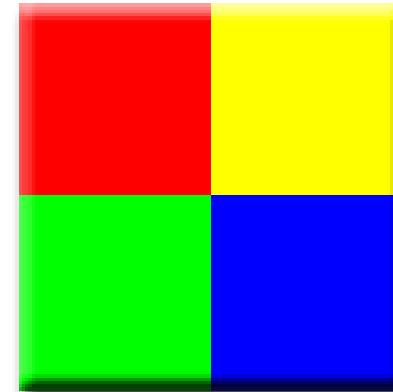




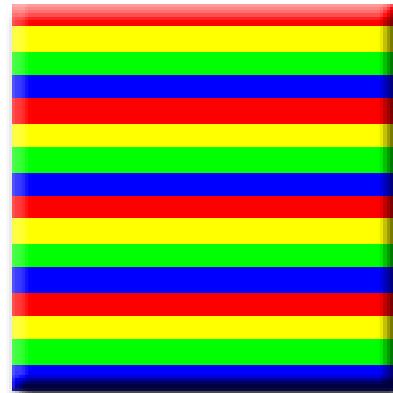
BLOCK, *



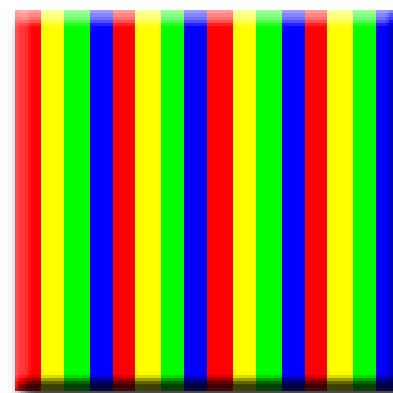
*, BLOCK



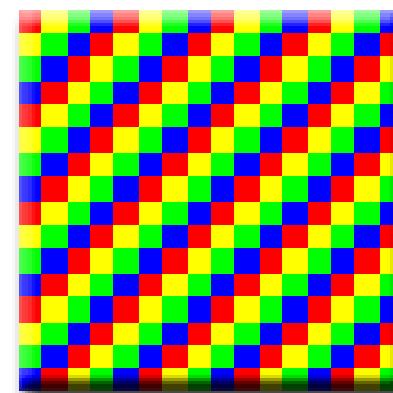
BLOCK, BLOCK



CYCLIC, *



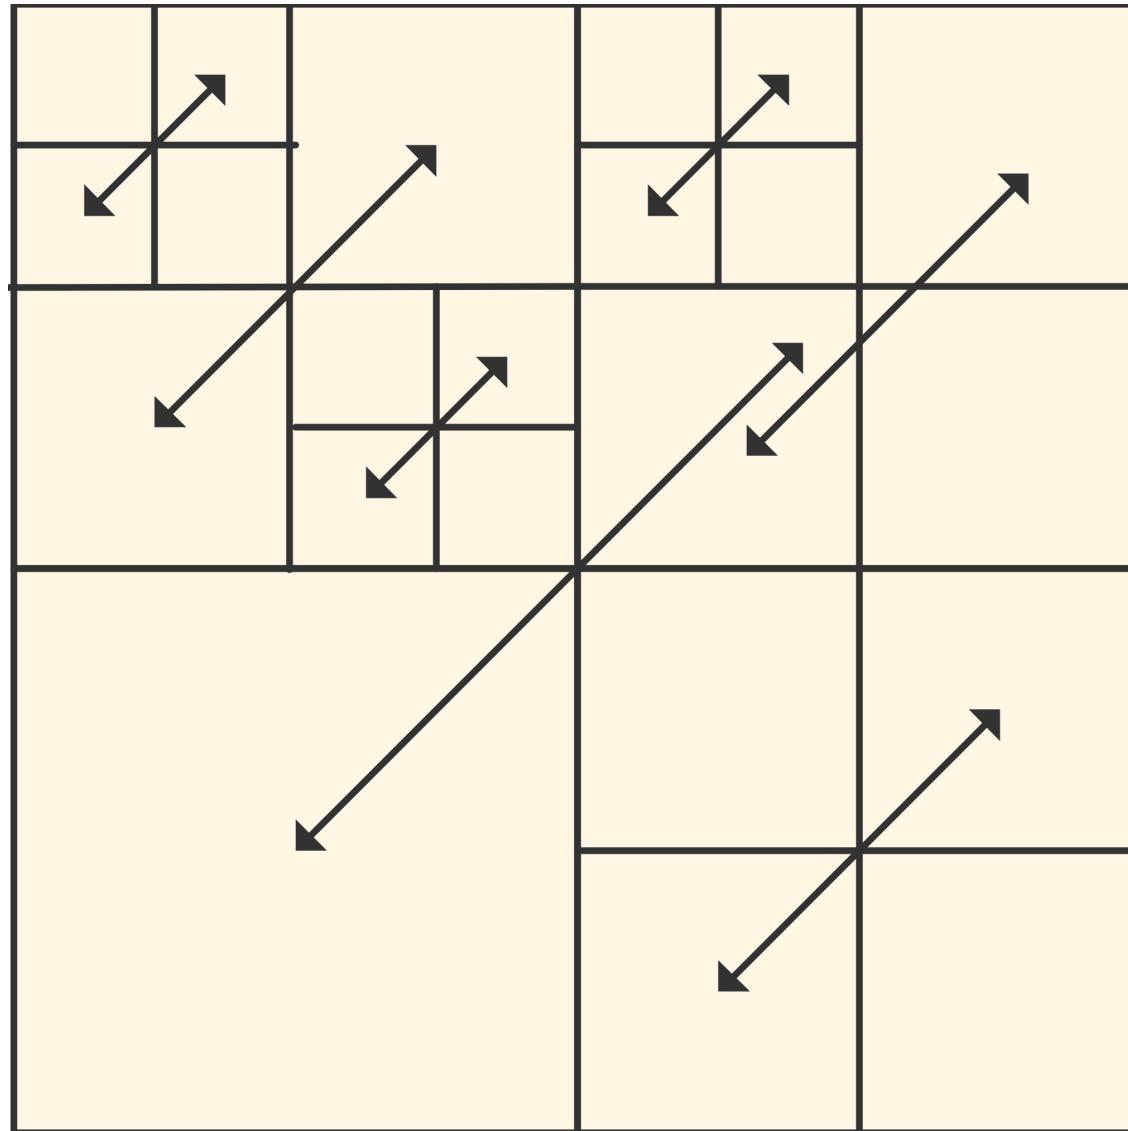
*, CYCLIC



CYCLIC, CYCLIC



- Translation: one-to-one
- Scaling: neighbourhood
- Transposition: one-to-one
- Rotation
 - 90° : one-to-one
 - arbitrary angle: neighbourhood
- Shearing: neighbourhood
- Generally memory-bound





- Convolutions with small kernels
- Morphological SE operators
 - dilations/erosions
 - openings/closings
 - hit-or-miss transforms
- Very suitable for SIMD machines



- Linear filters are implemented through convolutions

$$g(x, y) = (f * h)(x, y) = \sum_{(i,j) \in N} f(x - i, y - j)h(i, j) \quad (1)$$

- Morphological filters are based on Minkowski addition and subtraction

$$g(x, y) = (f \oplus h)(x, y) = \bigvee_{(i,j) \in N} f(x - i, y - j) + h(i, j) \quad (2)$$

and

$$g(x, y) = (f \ominus h)(x, y) = \bigwedge_{(i,j) \in N} f(x + i, y + j) - h(i, j) \quad (3)$$



- Convolutions are linear, shift-invariant, with simple distributivity

$$f * g * h = f * (g * h) = (f * g) * h \quad (4)$$

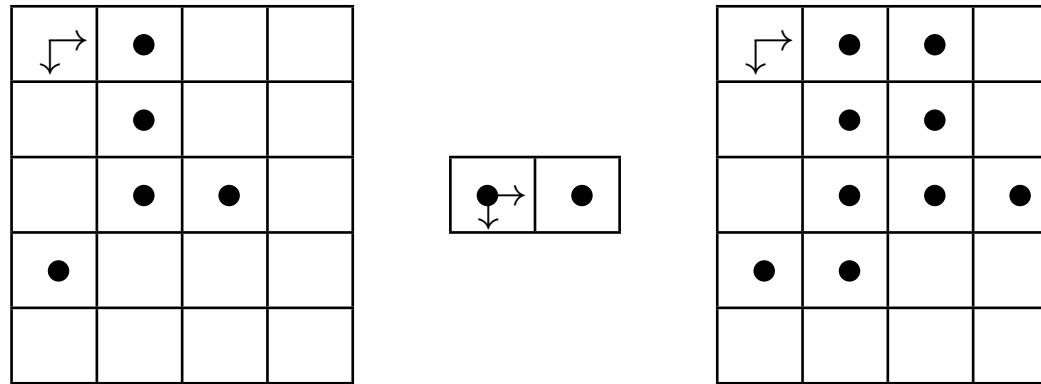
- For dilations we have

$$f \oplus g \oplus h = f \oplus (g \oplus h) = (f \oplus g) \oplus h \quad (5)$$

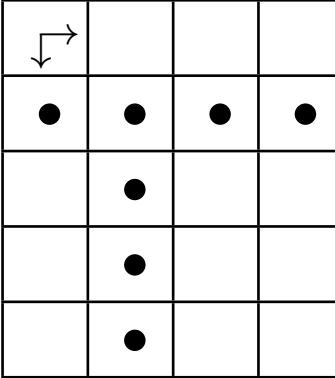
- For erosions things are slightly different

$$(f \ominus g) \ominus h = f \ominus (g \oplus h) \quad (6)$$

- All rules allow decomposition of kernels or structuring elements (SEs) into a number of smaller ones



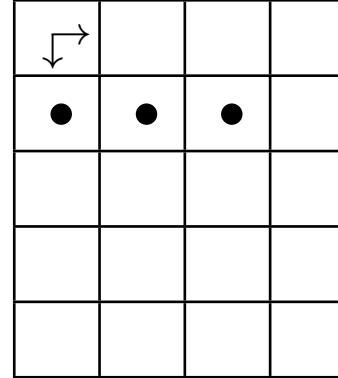
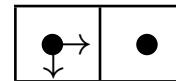
Left: binary image X . Middle: structuring element A . Right: dilation of X by A .



A 5x4 grid representing a binary image X . The grid contains the following pattern of black dots:

•	•	•	•
	•		
	•		
	•		
	•		

The first column has a small arrow pointing from top to bottom.

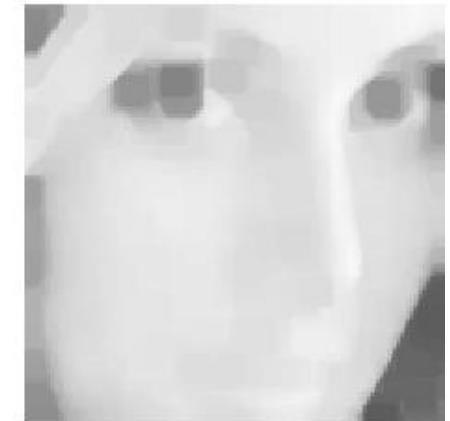


A 5x4 grid representing the erosion result of X by A . The grid contains the following pattern of black dots:

•	•	•	

The first column has a small arrow pointing from top to bottom.

Left: binary image X . Middle: structuring element A . Right: erosion of X by A .



Left: Lenna, the most famous woman in computer vision. Right panel:
upper left: dilation; upper right: erosion; lower left: opening; lower
right: closing.



- At every pixel we must find the minimum of maximum of all pixels within the S.E..
- Finding the maximum of N unsorted numbers kan be done as follows:

```
i = 1;  
maximum = a[i];  
  
while (i <= N) {  
    if (a[i] > maximum) {  
        maximum = a[i];  
    }  
    i++;  
}
```

- How many comparisons must I perform?



- The entire algorithm for an image of $H \times W$ pixels is:

```
for ( y = 0; y < H - 1; y++ ){
    for ( x = 0; x < W; x++ ){
        move S.E. to (x,y)
        result[x][y] = maximum(S.E.);
    }
}
```

- The number of comparisons for an $L \times L$ square S.E. is

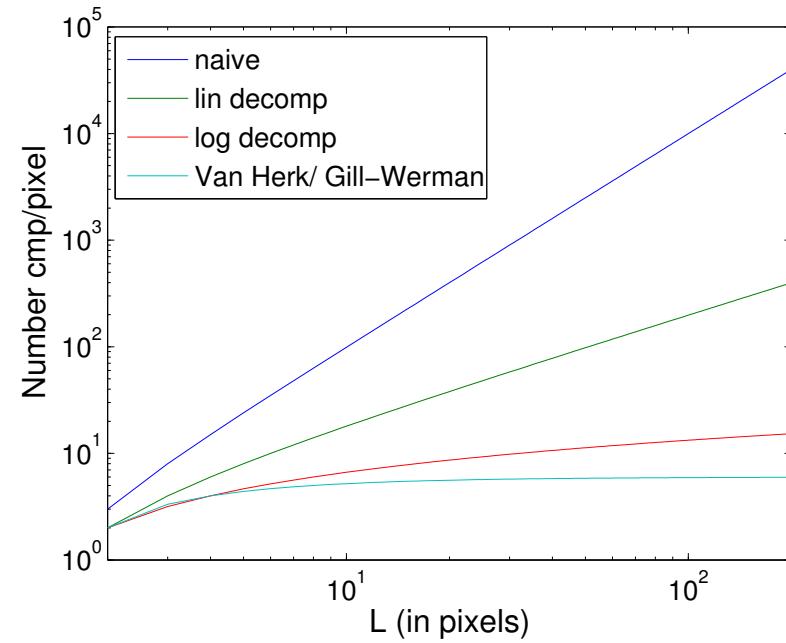
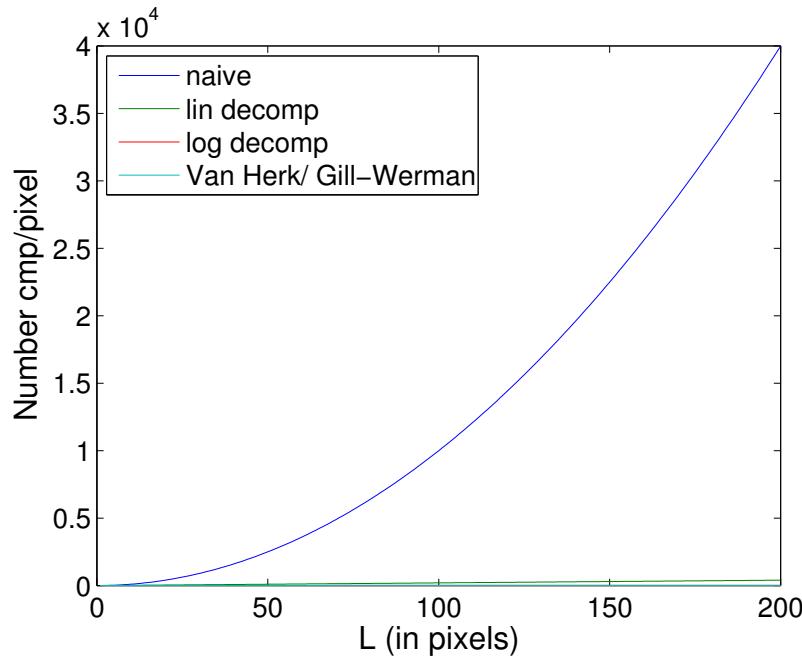
$$W \times H \times (L^2 - 1) \quad (7)$$

- Can we do better?

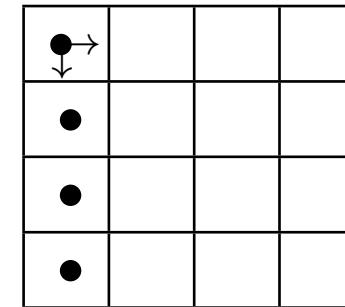
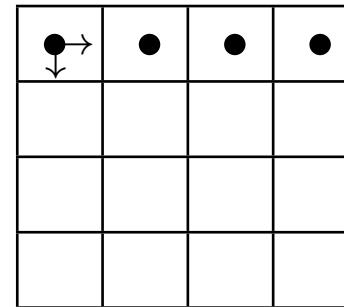
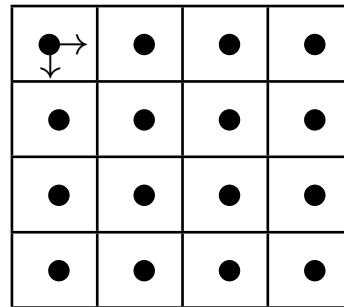
- The basic algorithm has complexity

$$O(W \times H \times L^2) \quad (8)$$

- Thus it is *linear* in W en H , and *quadratic* in L .



- A max-filter with square S.E. can be computed from two max-filters with linear S.E. :



- The number of comparisons is reduced to

$$2 \times W \times H \times (L - 1) \quad (9)$$

- The complexity is

$$O(W \times H \times L) \quad (10)$$



- A linear S.E. can be decomposed into $k = \lceil \log_2(L) \rceil$ S.E.'s with only points in each S.E.

$$L^i = L_0^i \oplus L_1^i \oplus \dots L_{k-1}^i, \quad i = r, c$$

with

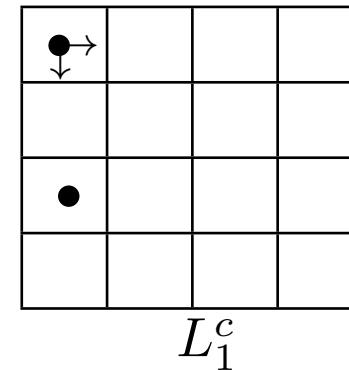
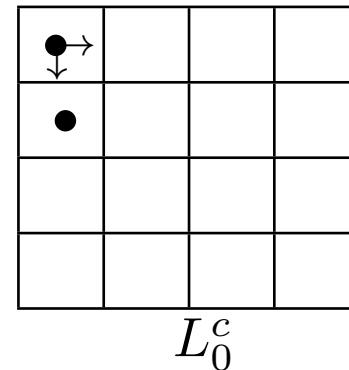
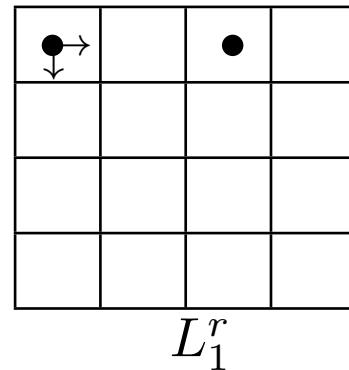
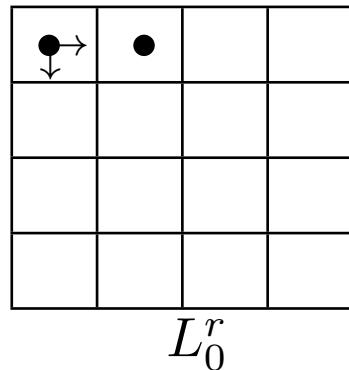
$$L_k^r = \left\{ \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} 2^k \\ 0 \end{pmatrix} \right\}, \quad L_k^c = \left\{ \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} 0 \\ 2^k \end{pmatrix} \right\}.$$

A square is decomposed as:

$$B_k = (L_0^r \oplus L_1^r \oplus \dots L_{k-1}^r) \oplus (L_0^c \oplus L_1^c \oplus \dots L_{k-1}^c)$$

with a total of $2k$ max-filters.

- Example for $L = 4$:



- The number of comparisons reduces to

$$2 \times W \times H \times \log_2(L) \quad (11)$$

- The complexity is

$$O(W \times H \times \log L) \quad (12)$$



- Instead of different decompositions, use different algorithm
- Divide the image row (column) into overlapping segments of $2L - 1$ pixels, centred on:

$$j = L - 1, 2L - 1, 3L - 1, \dots, kL - 1 \quad (13)$$

- Every pixel j is in L distinct shifted S.E.
- For every j compute arrays R_j and S_j :

$$R_j(k) = \max(p(j), p(j - 1), \dots, p(j - k)) \quad (14)$$

$$S_j(k) = \max(p(j), p(j + 1), \dots, p(j + k)) \quad (15)$$

with $k \in \{0, 1, \dots, L - 1\}$, and $p(i)$ the grey value at point i in the row (column).



5	4	6	2	1	4	11	10	8	2	0
---	---	---	---	---	---	----	----	---	---	---

j

6	6	6	2	1	R_j
---	---	---	---	---	-------

S_j	1	4	11	11	11
-------	---	---	----	----	----

- Note that we can compute $R_j(k)$ en $S_j(k)$ as

$$R_j(k) = \begin{cases} \max(R_j(k-1), p(j-k)) & \text{if } k > 0 \\ p(j) & \text{if } k = 0 \end{cases} \quad (16)$$

$$S_j(k) = \begin{cases} \max(S_j(k-1), p(j+k)) & \text{if } k > 0 \\ p(j) & \text{if } k = 0 \end{cases} \quad (17)$$

- Thus, we need $2(L - 1)$ comparisons.



- Another $L - 2$ comparisons to compute L points in the final result

$$\max(p(j-k), \dots, p(j+L-k-1)) = \max(R_j(k), S_j(L-k-1)) \quad (18)$$

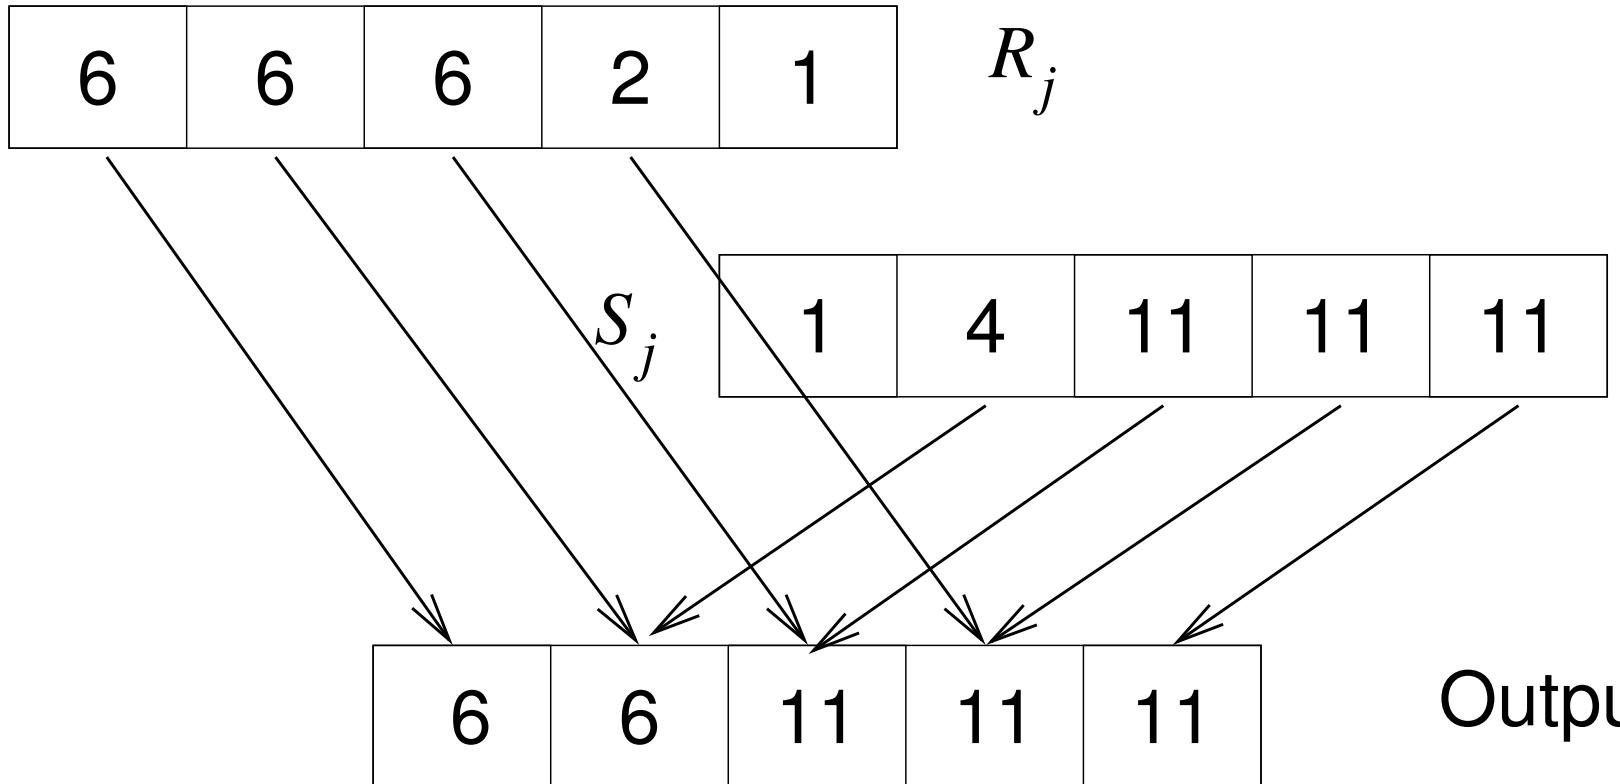
voor $k = 1, 2, \dots, p-2$

$$\max(p(j-p-1), \dots, p(j)) = R_j(p-1) \quad (19)$$

$$\max(p(j), \dots, p(j+p-1)) = S_j(p-1) \quad (20)$$



5	4	6	2	1	4	11	10	8	2	0
---	---	---	---	---	---	----	----	---	---	---

 j 



- To compute L results, we need a number of comparisons equal to

$$2(L - 1) + L - 1 \quad (21)$$

or, per pixel

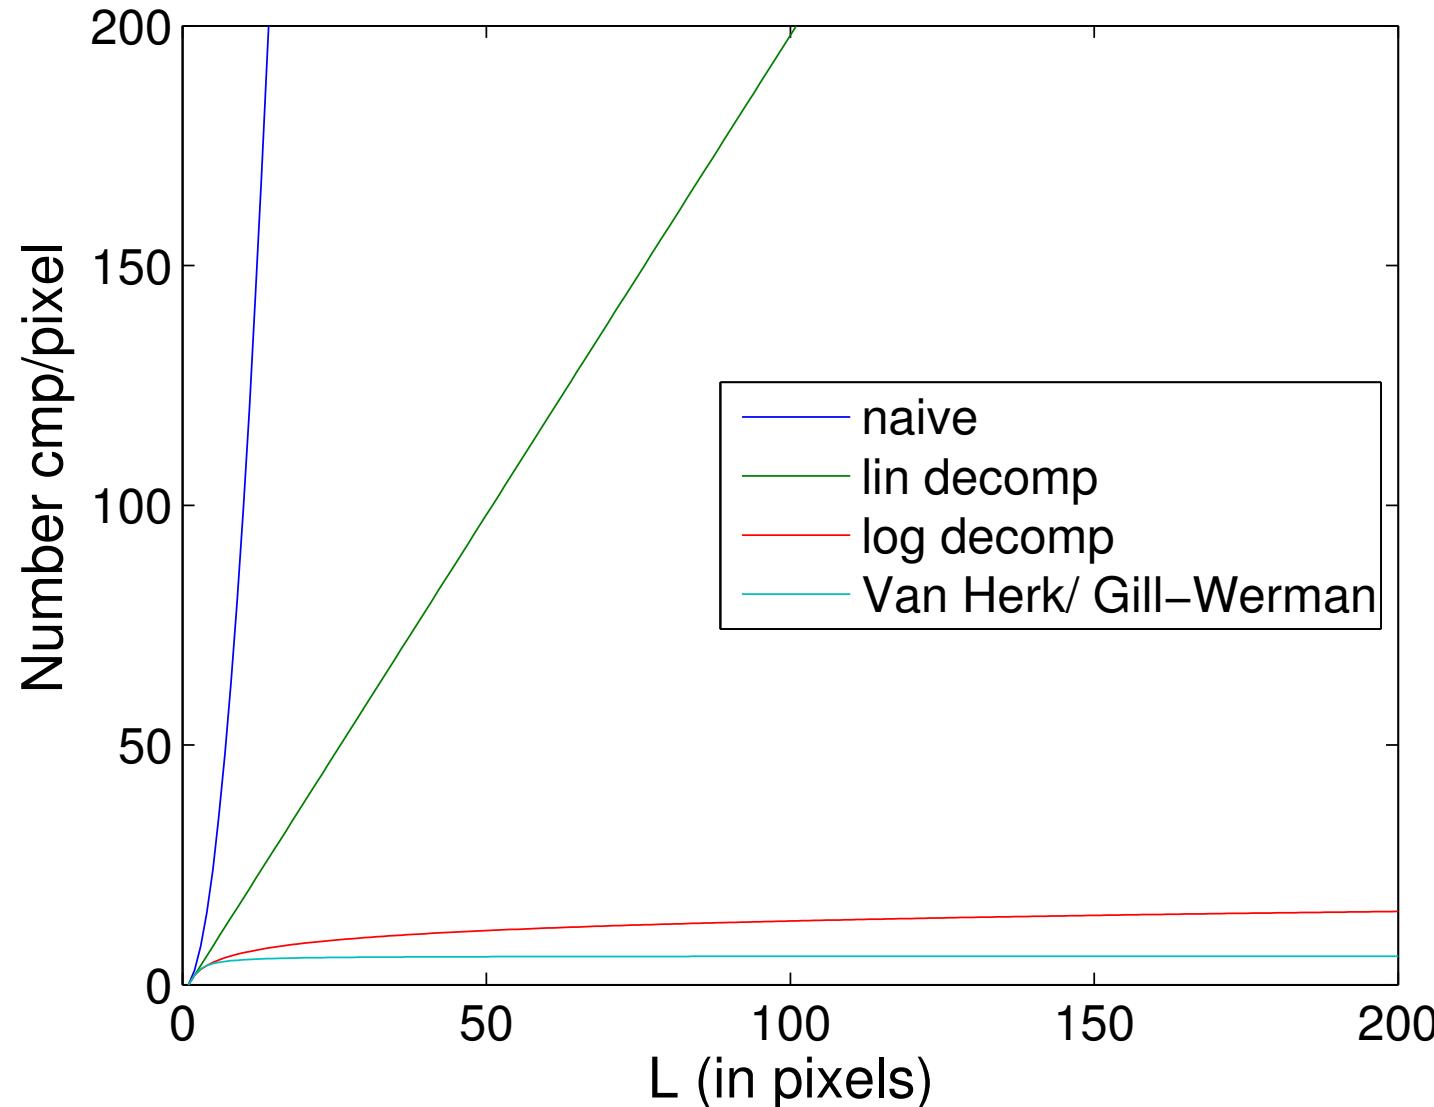
$$\frac{2(L - 1) + L - 1}{L} = 3 - \frac{4}{L} \quad (22)$$

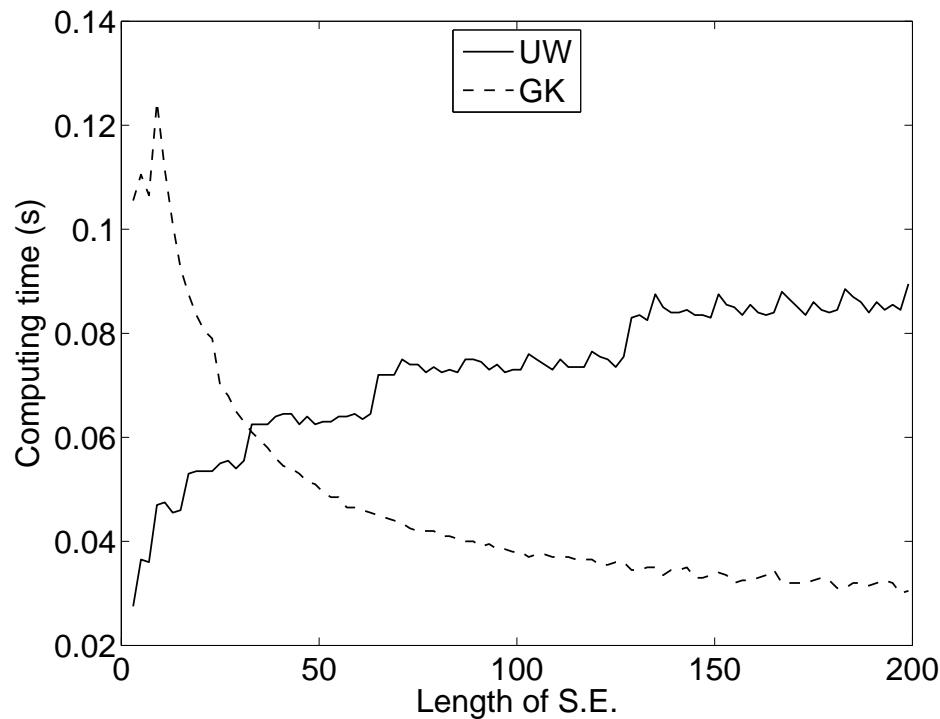
- The total number of comparisons becomes

$$W \times H \times \left(6 - \frac{8}{L}\right) \quad (23)$$

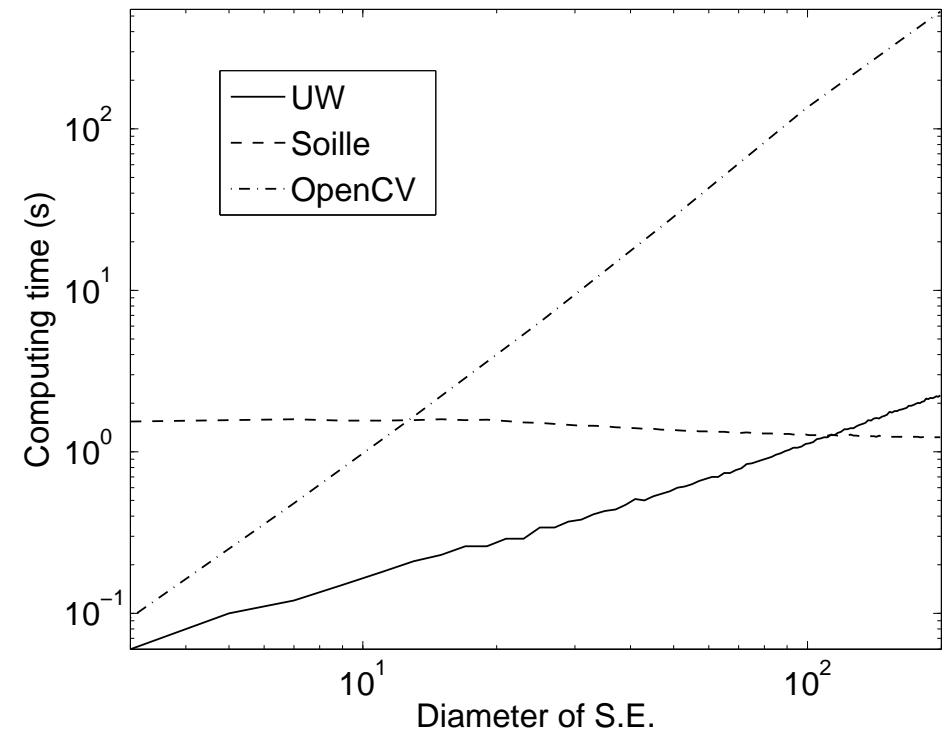
- And the complexity

$$O(W \times H). \quad (24)$$



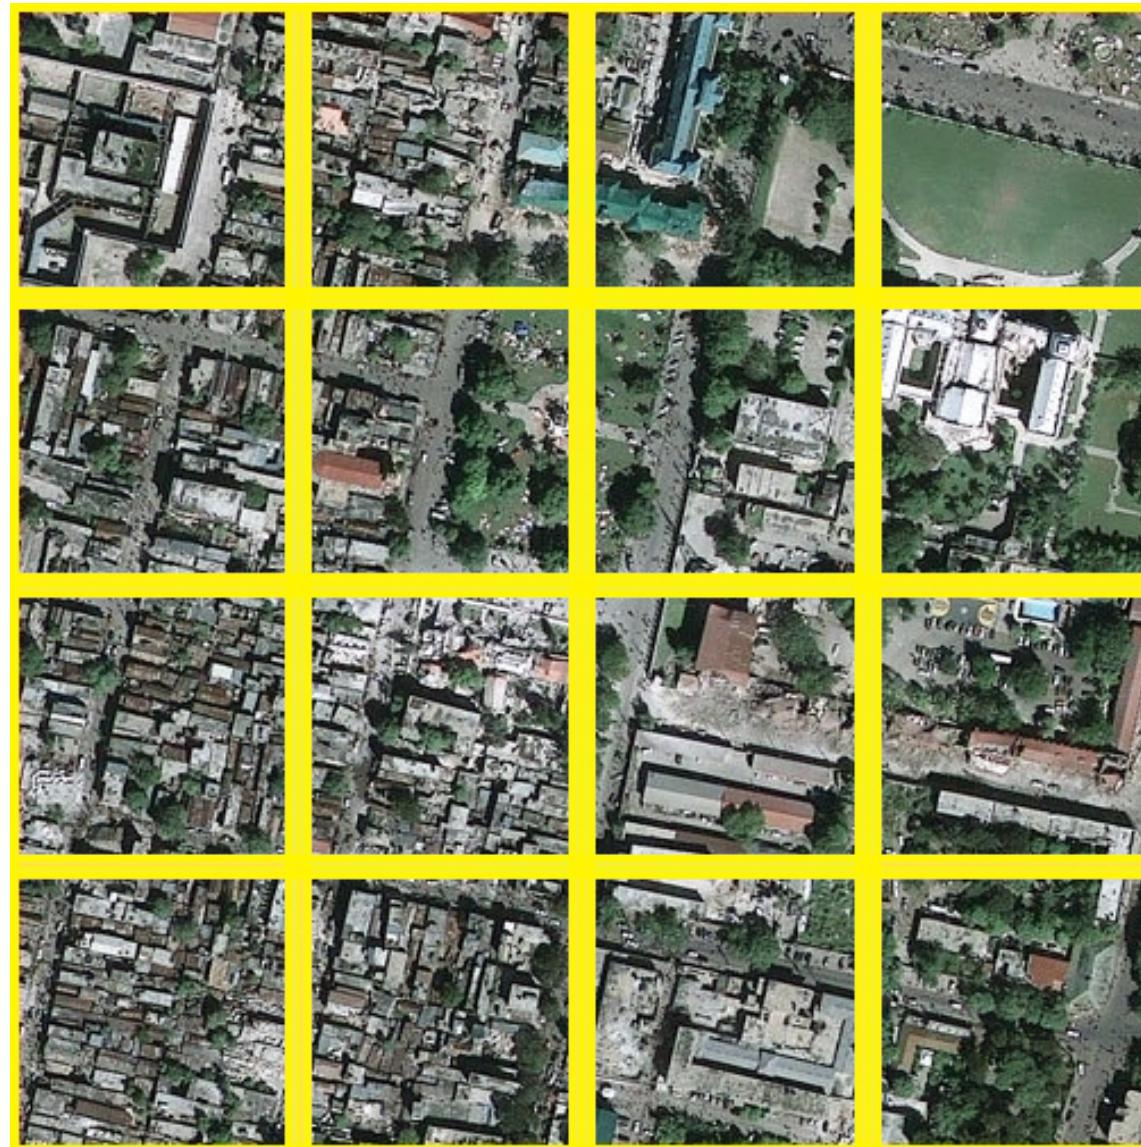


Lines



Circles

- Distribution with overlap
 - Overlap increases with kernel/SE size
 - Can be countered by kernel/SE decomposition
- Shift + combine strategies
 - Makes good use of easily parallelized point operations
 - Requires more barriers





- On shared-memory machines:
 - No overlap costs incurred
 - Data from adjacent block can be read directly
- On distributed-memory machines:
 - Data needs to be copied
 - Memory or communication cost $O(\sqrt{N_p}(W \times k\text{height} + H \times k\text{width}))$ in 2D
 - Rises steeply for higher-dimensional data



A simple sequential algorithm for convolution can be used per block for small kernels:

```
for ( y = 0; y < H - 1; y++ ){
    for ( x = 0; x < W; x++ ){

        result[x][y] = 0;

        for (j = -kheight/2; j <= kheight/2; j++){
            for (i = -kwidth/2; i <= kwidth/2; i++){
                result[x][y] += f[x-i][y-j] * h[i][j];
            }
        }
    }
}
```

Complexity: $O(W \times H \times kwidth \times kheight)$, or $O(N^2)$ if the kernel is as large as the image.



By moving the inner loops over the kernel outwards we can reuse parallel point operators and shifts:

```
ParImageSet(result,0);           // sets all pixels of result to zero

for (j = -kheight/2; j <= kheight/2; j++){
    for (i = -kwidth/2; i <= kwidth/2; i++){
        ParShiftImage(f,temp,i,j);      //creates copy temp of f
                                         // shifted by (i,j)
        ParMulImageConst(temp,h[i][j]);   // multiply temp by h[i][j]
        ParAddImages(result,temp,result); // adds to result
    }
}
```

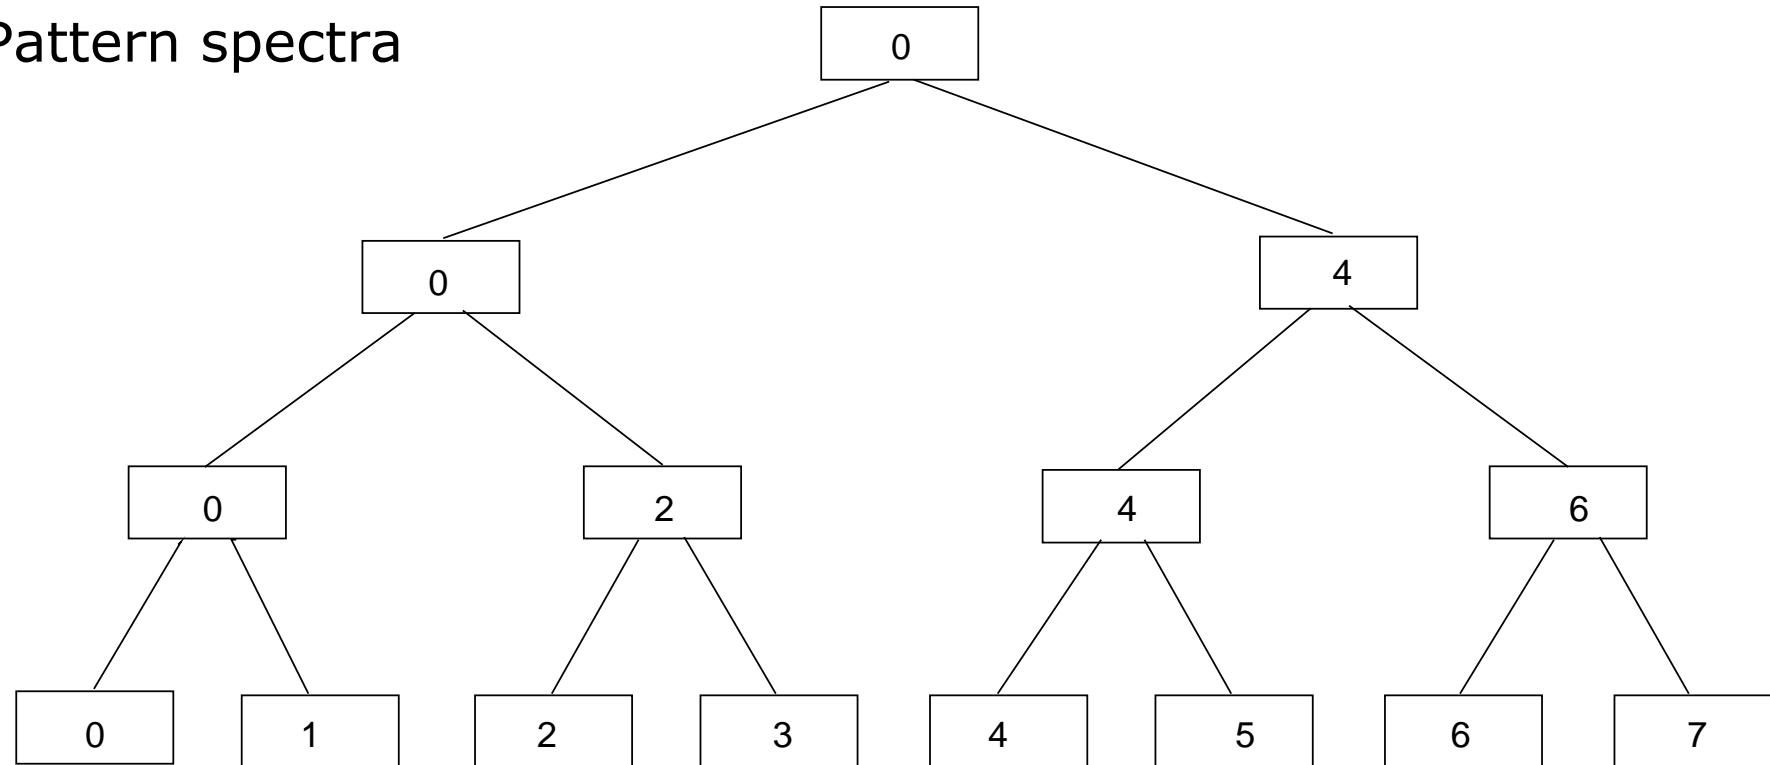


For dilation we obtain:

```
ParImageSet(result,-MAXINT); // sets all pixels of result to smallest value

for (j = -kheight/2; j <= kheight/2; j++){
    for (i = -kwidth/2; i <= kwidht/2; i++){
        ParShiftImage(f,temp,i,j);           //creates copy temp of f
                                                // shifted by (i,j)
        ParAddImageConst(temp,h[i][j]);      // add h[i][j] to temp
        ParMaxImages(result,temp,result);   // compute pixelwise maximum
    }
}
```

- Finding global minimum/maxmimum
- Histograms
- Pattern spectra



- Large kernel convolutions
- Large, non-flat SE-based morphological operators
- Fourier Transforms
- Distance transforms
- Watershed transforms
- Connected component labelling
- Connected filters



- For large kernels we use the following property:

$$f * h = \mathcal{F}^{-1}(\mathcal{F}(f)\mathcal{F}(h)) \quad (25)$$

with \mathcal{F} the Fourier transform, and \mathcal{F}^{-1} its inverse

- Using the Fast Fourier Transform we obtain a complexity of $O(N \log N)$.
- Parallel computation of the FFT can simply be obtained through separability
- Many parallel FFT algorithms for specific architectures exist



The convolution algorithm becomes

```
ParFFT(f, ft);           //Fourier transform f, store in ft
ParFFT(h, ht);           //Fourier transform h, store in ht

ParMulComplexImages(ht,ft,ft); // multiply and store in ft
ParInvFFT(ft,g);          // compute inverse FFT, store in g
```

For added speed, precompute FFTs of frequently used kernels.



- All methods which allow treatment per dimension
 - Fourier transform
 - Gaussian convolution
 - Parabolic dilation/erosion
- Distance Maps



- The 2D Discrete Fourier transform is given by

$$F[k, l] = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} f[m, n] e^{i2\pi(\frac{ml}{M} + \frac{nk}{N})} \quad (26)$$

- This can readily be rewritten as

$$F[k, l] = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} f[m, n] e^{i2\pi \frac{ml}{M}} e^{i2\pi \frac{nk}{N}}$$

or

$$F[k, l] = \sum_{n=0}^{N-1} e^{i2\pi \frac{nk}{N}} \left(\sum_{m=0}^{M-1} f[m, n] e^{i2\pi \frac{ml}{M}} \right) \quad (27)$$



- The 2D Gaussian kernel can be separated into an x and y kernel

$$k_{x,\sigma}(x, y) = \frac{\delta_y}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}} \quad k_{y,\sigma}(x, y) = \frac{\delta_x}{\sigma\sqrt{2\pi}} e^{-\frac{y^2}{2\sigma^2}} \quad (28)$$

- Using the fact that

$$(k_{x,\sigma} * k_{y,\sigma})(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} = k_\sigma(x, y) \quad (29)$$

we have

$$f * k_\sigma = f * (k_{x,\sigma} * k_{y,\sigma}) = f * k_{x,\sigma} * k_{y,\sigma} \quad (30)$$

- 1-D Gaussian convolution is very rapidly done using recursive filters (6 multiplies and adds per pixel) [Van Vliet and Young]



- The 2D parabolic SE kan be separated into an x and y SEs

$$P_x(x, y) = \bar{\delta}_y + ax^2 \quad P_y(x, y) = \bar{\delta}_x + ay^2 \quad (31)$$

with

$$\bar{\delta}_x = \begin{cases} 0 & \text{if } x = 0, \\ -\infty & \text{otherwise.} \end{cases}$$

- We then have

$$P_x \oplus P_y = a(x^2 + y^2) = ar^2 = P_r \quad (32)$$

and thus

$$f \oplus P_x \oplus P_y = f \oplus P_r \quad (33)$$

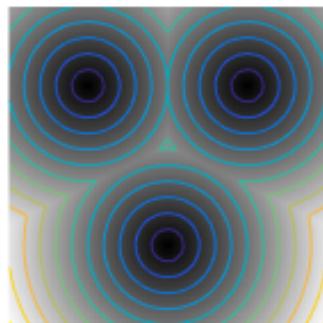


1. Distribute rows over processors
2. Perform transform per row
3. Rotate or transpose image
4. Perform transform per row
5. Rotate back or transpose image

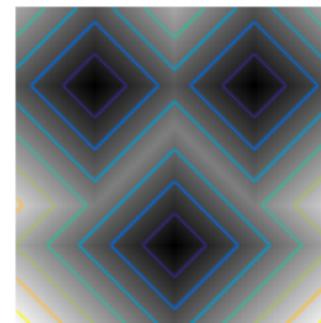


- Turn a binary image into a grey-scale image
- Each grey value indicates the distance to the nearest background pixel

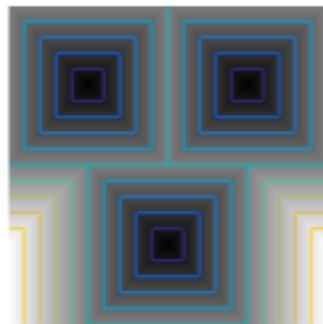
Euclidean



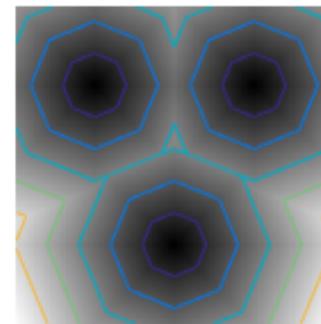
City block



Chessboard



Quasi-Euclidean





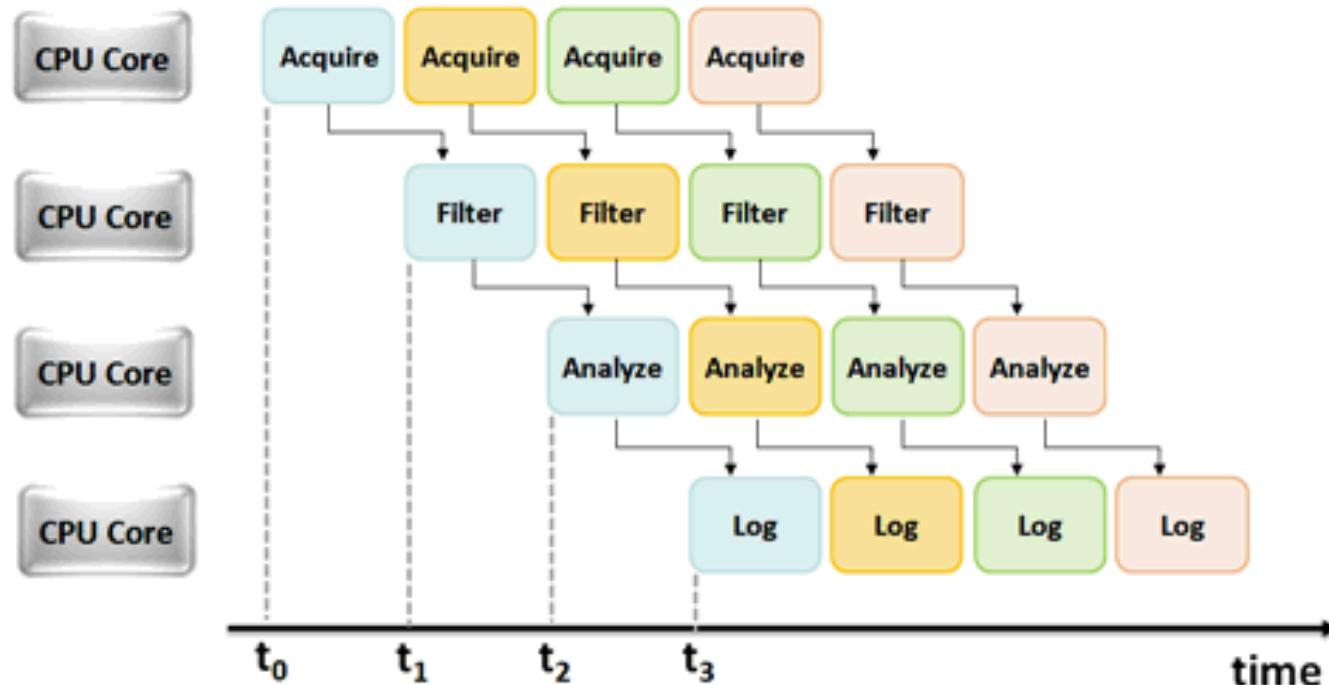
- An efficient, general algorithm has been developed
- It can handle Euclidean, Manhattan, and Chessboard distances
- Can be parallelized effectively by a three-pass method
 - One forward row-wise sweep
 - One reverse row-wise sweep
 - One column-wise sweep



- On shared memory machines:
 - No transposition strictly needed
 - Look out for cache-hostile scan orientations
- On distributed+hybrid machines:
 - Transposition can be costly
 - Tune transposition algorithm carefully to architecture



- Sequence processing
 - Pipe-lines: functional decomposition in time
 - Interleaving/multiplexing: temporal decomposition
- PDE-based filters:
 - PDE solvers readily expressed in matrix-vector operation
 - Use libraries like LINPACK, HPL, etc.
- Convolutional Neural Nets
 - Parallel by nature
 - Very heavy compute and memory load



- Watershed transforms
- Connected component labelling
- Connected filters
- Stuff for next time

