



Supercomputing for Giga- and Tera-Pixel Image Scales, Part III: Multi-Scale and Distributed Connected Filtering

Michael H. F. Wilkinson

*Johann Bernoulli Institute
University of Groningen*



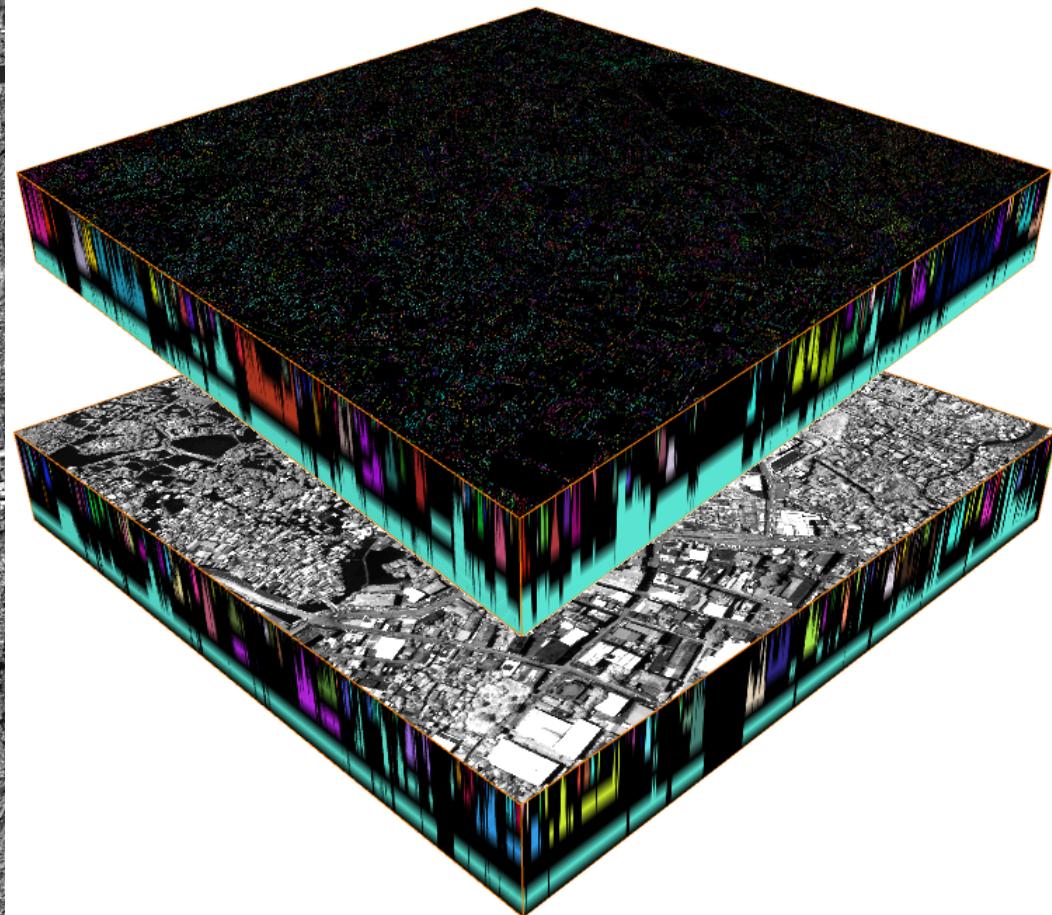
- Multi-Scale Analysis by Component Trees
 - Differential Morphological Profiles
 - Differential Area Profiles
 - Algorithms
- Distributed Computation of Connected Filters
 - Modified Max-tree Forests
 - Boundary Trees
 - Algorithms



Quickbird image



Quickbird image



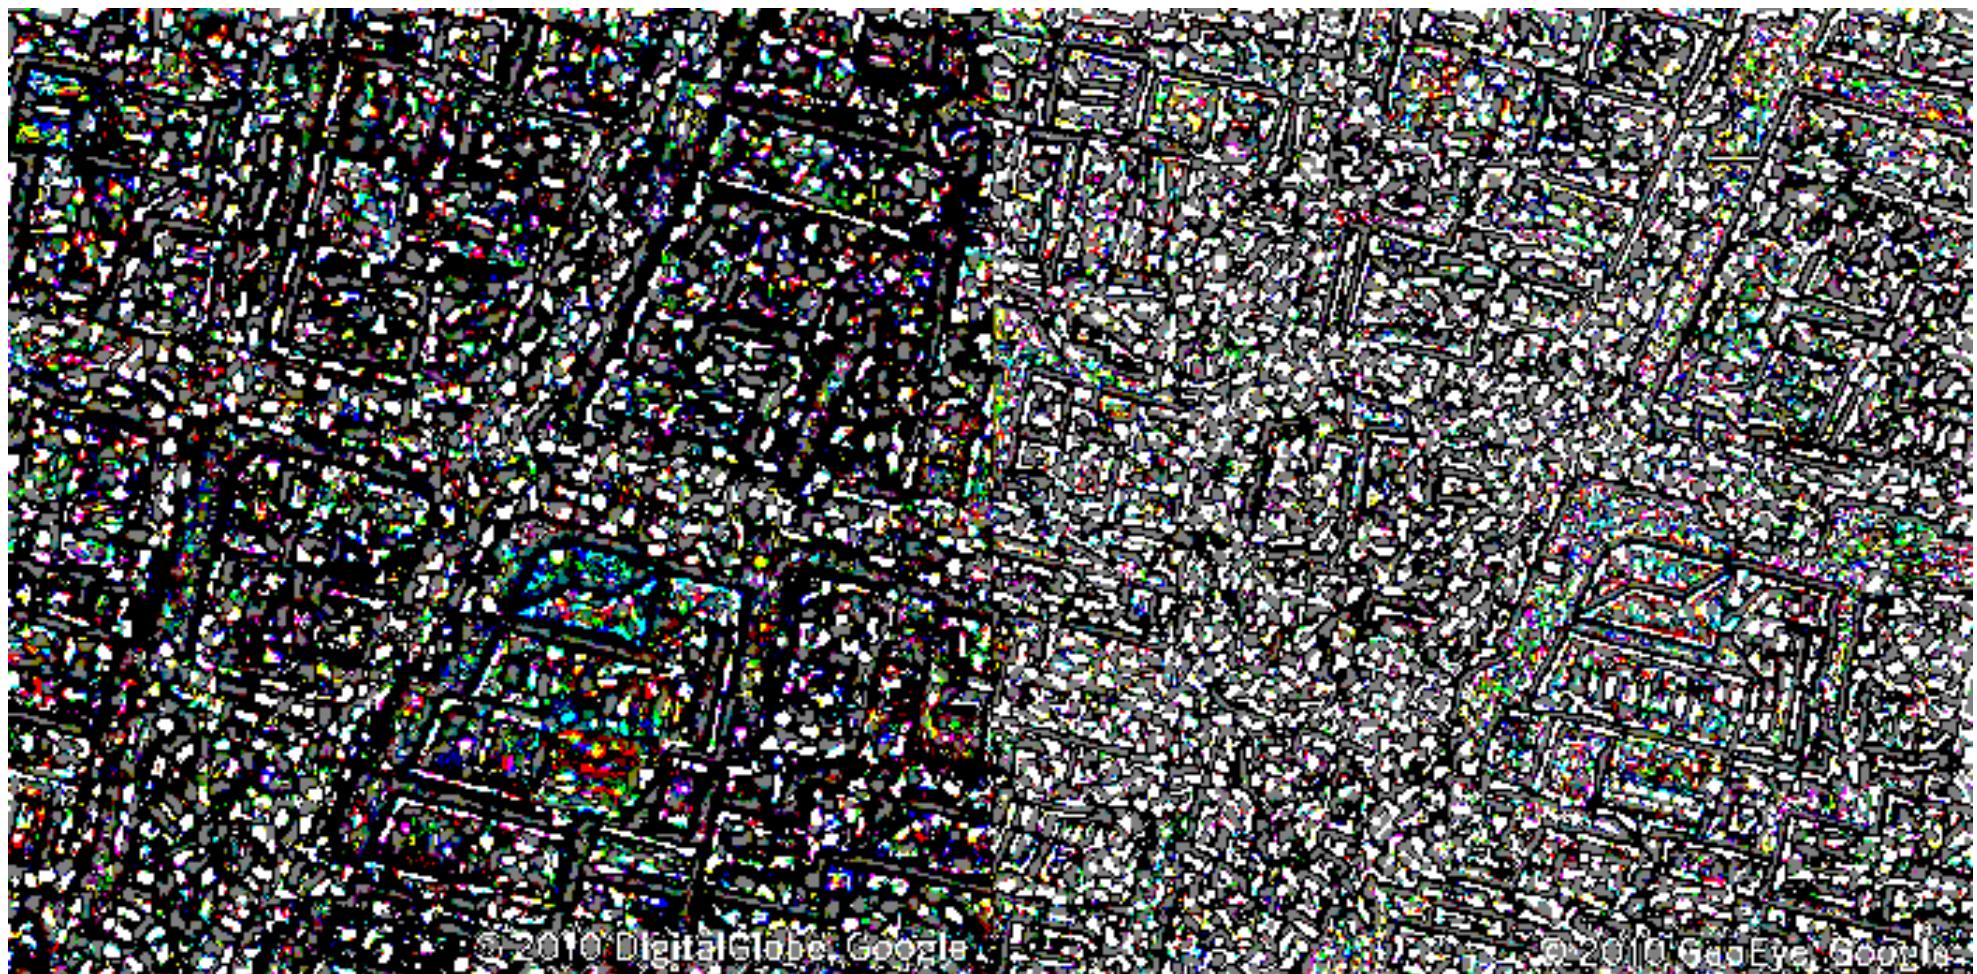
Differential Morphological Profile



High-performance **I**mage **Processing **A**lgorithms for **R**emote **S**sensing**

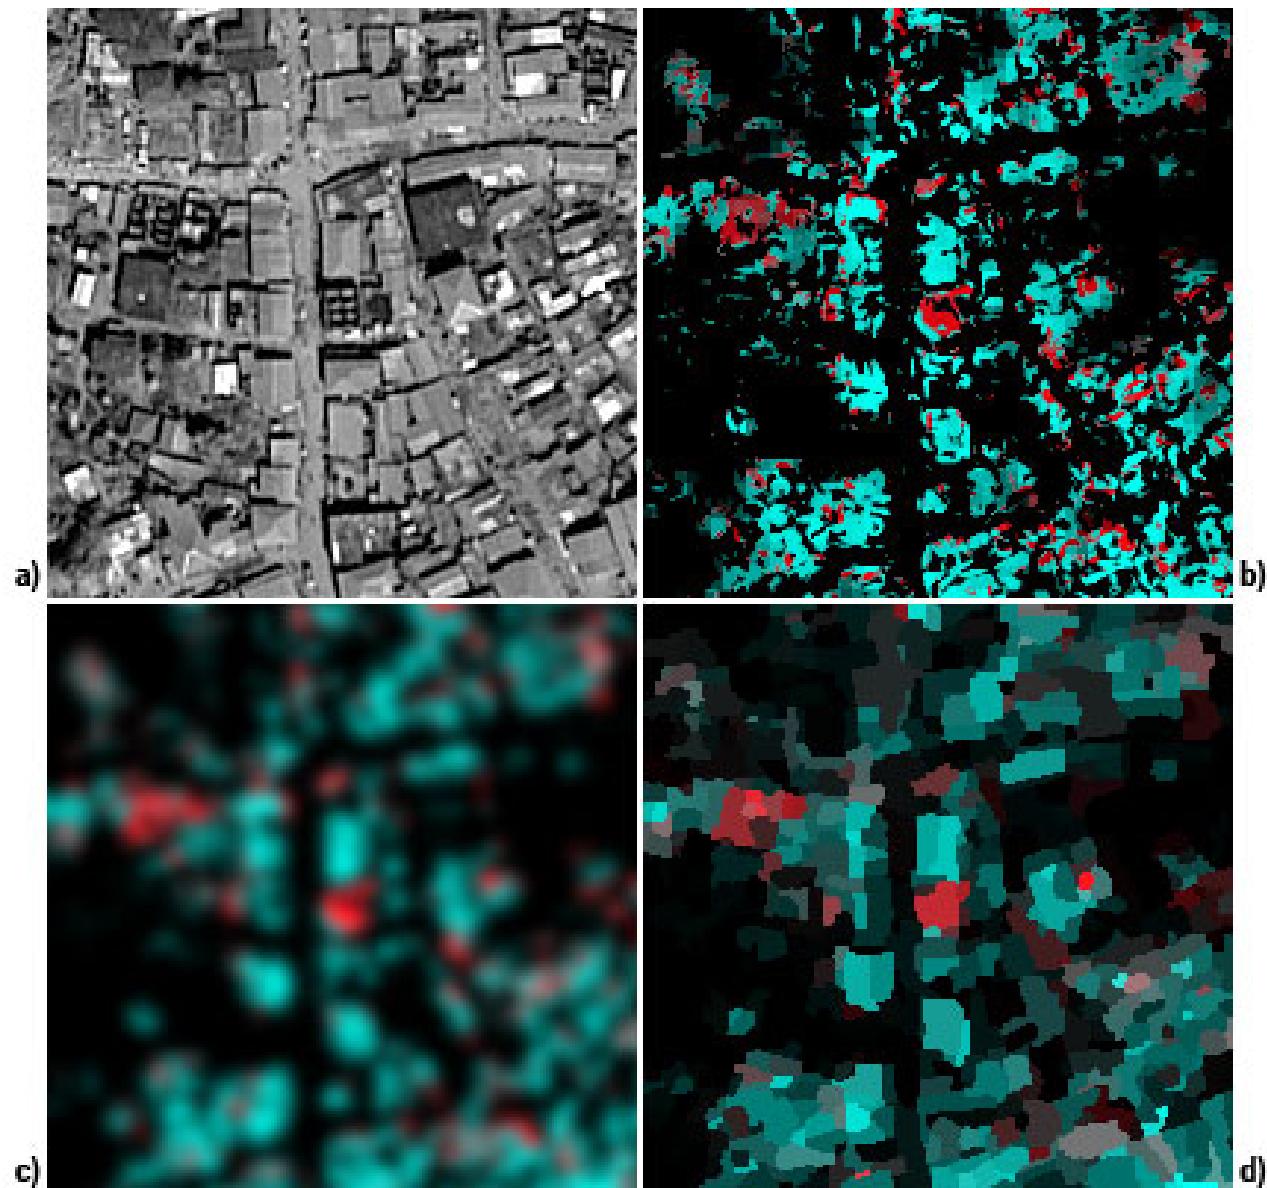
- Georgios Ouzounis
- Martino Pesaresi
- Pierre Soille
- Michael Wilkinson





before

after





Data set	Classical Algorithm	Improved Algorithm
Haiti 2010	34,000 years	104 days
Tsunami 2004	500,000 years	2.5 years



Data set	Classical Algorithm	Improved Algorithm
Haiti 2010	34,000 years	104 days
Tsunami 2004	500,000 years	2.5 years

This is useless!



Data set	Classical Algorithm	Improved Algorithm
Haiti 2010	34,000 years	104 days
Tsunami 2004	500,000 years	2.5 years

This is useless!

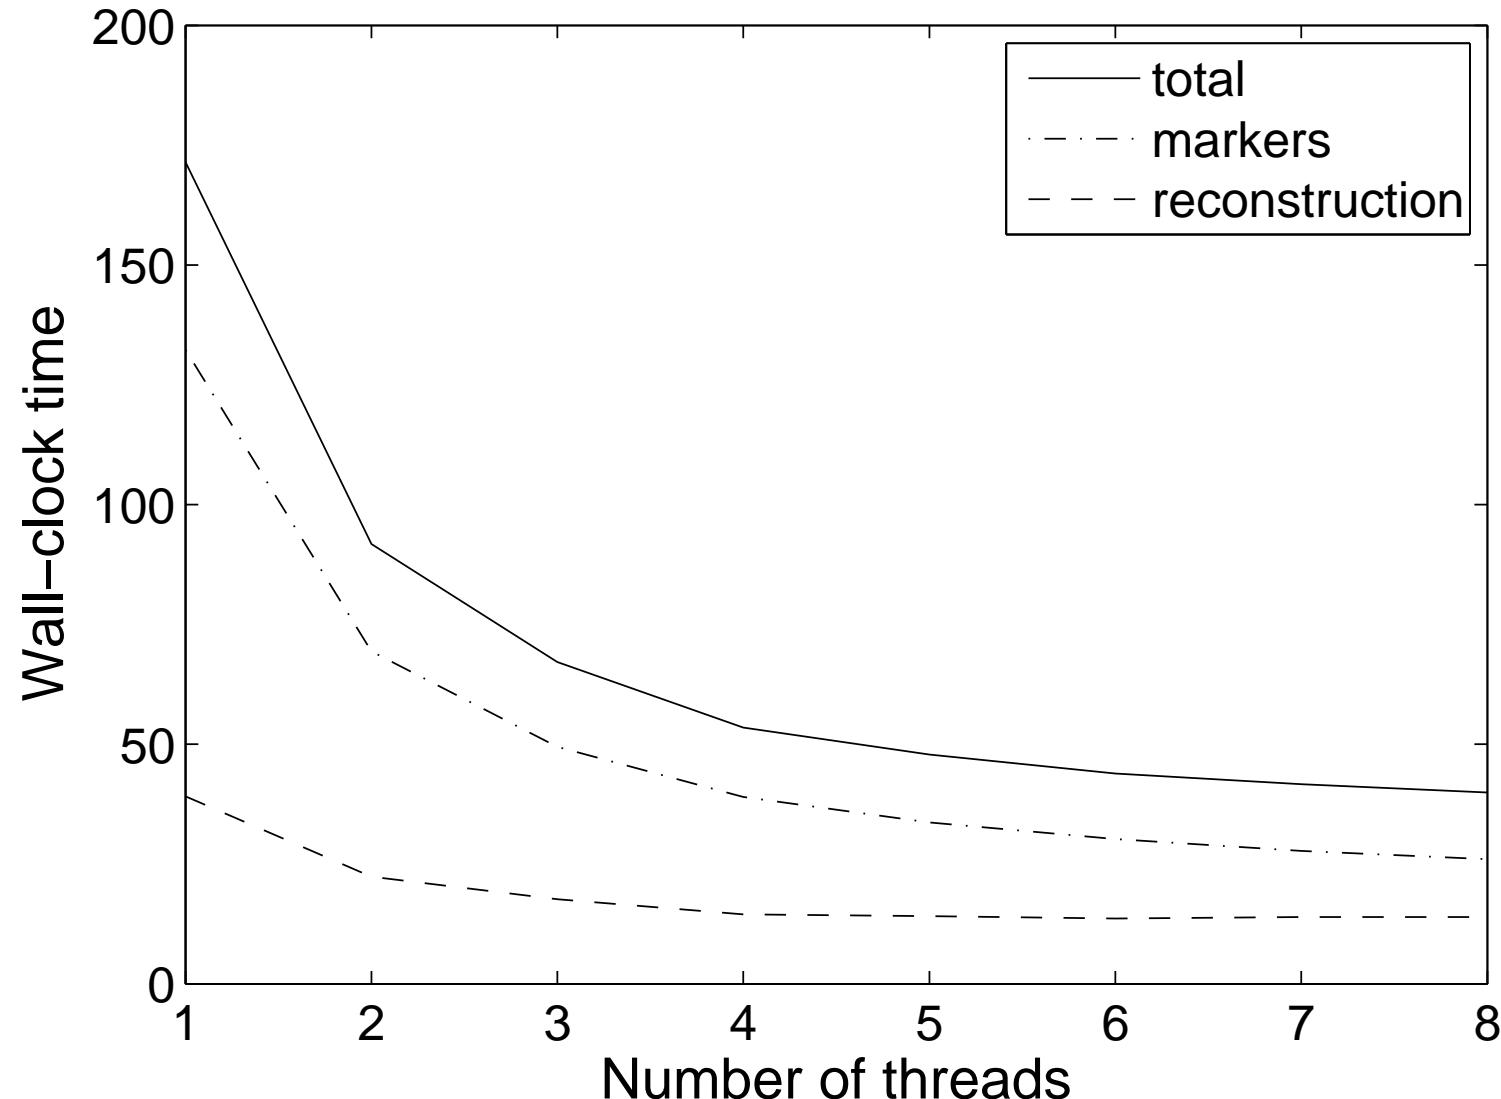
Haiti can now be done in
1 h 22 min

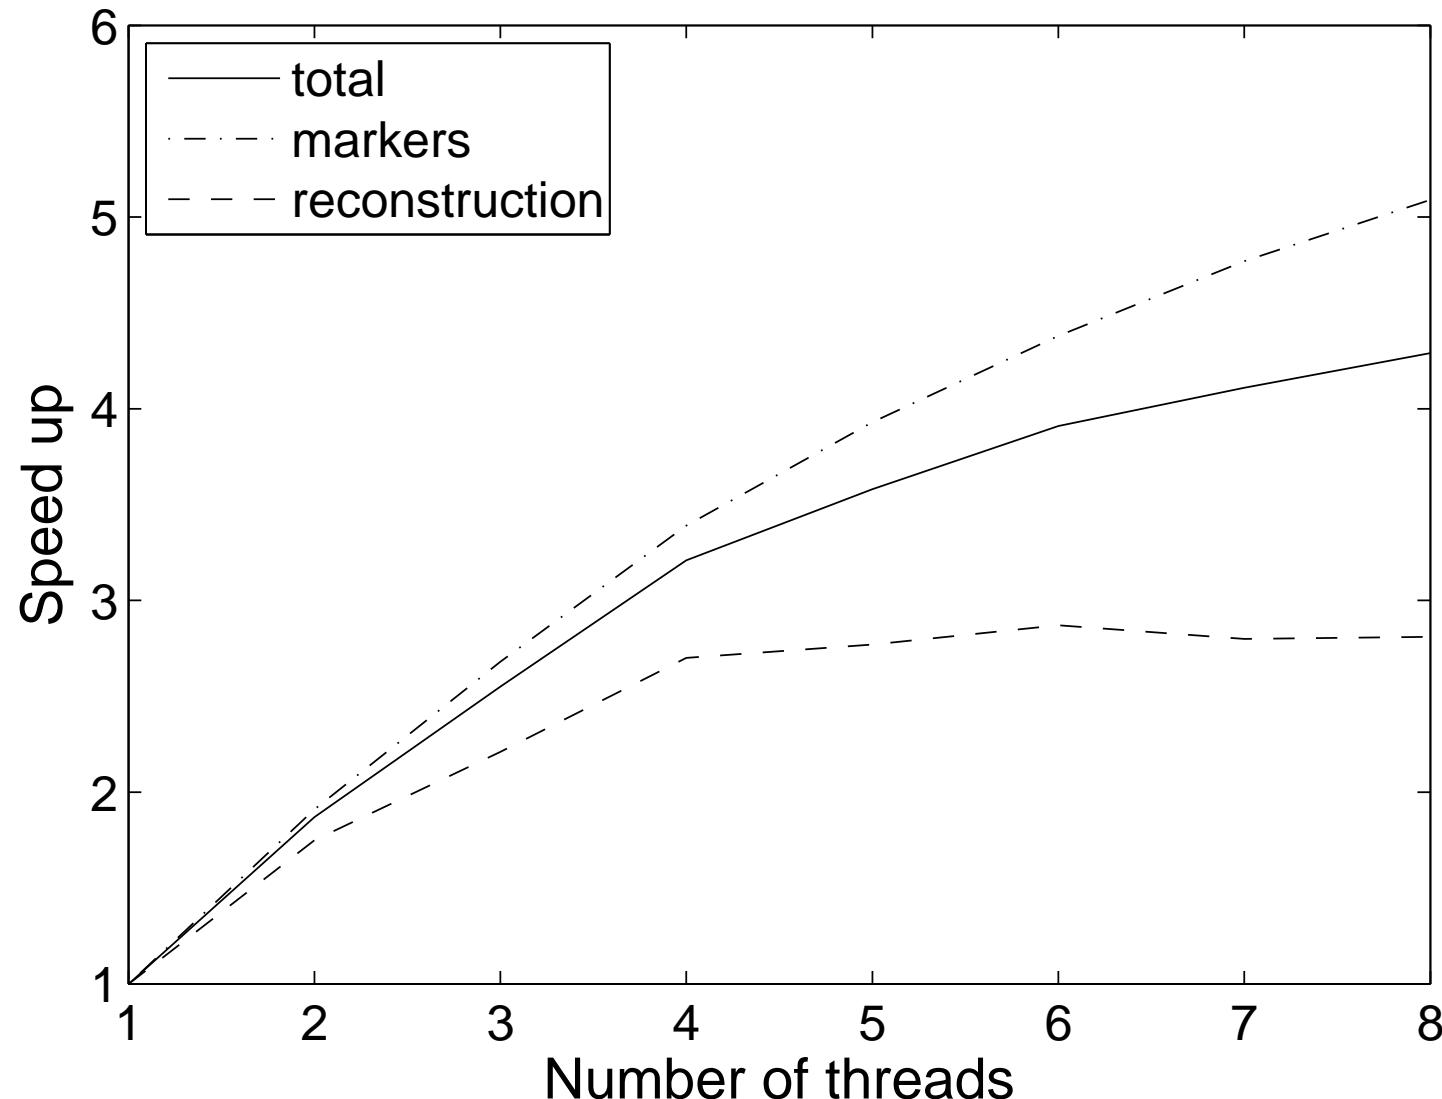


- Compute all markers using parallel erosions and store each scale i in an image $out[i]$;
- Compute a single Max-Tree in parallel, maintaining the maximum marker for all scales i for each node n in $out[i][n]$;
- For all scales i
 - Set all nodes to invalid indicating that they have not been filtered yet (in parallel);
 - Compute reconstruction (in parallel) for scale i , storing the result in $out[i]$;
- Compute differences between scales (in parallel);
- Output respective instance of DMP vector field.



```
procedure MaxTreeReconstruct ( $Vp$  : Section; var  $node$  : Max-Tree; var  $out$  : Image)
for all  $v \in Vp$  do
  if not  $node[v].valid$  then
     $w := v;$ 
    while  $Par(w) \neq \perp \wedge \text{not } node[w].valid \wedge f[Par(w)] > out[Par(w)]$  do
       $w := Par(w);$ 
    end;
    if  $node[w].valid$  then
       $val := out[w];$  (* Reconstructed node found *)
    else if  $Par(w) \neq \perp$  then
       $val := (Par(w) \neq \perp) ? out[w] \vee f[Par(w)] : 0;$ 
    end;
  end;
  for all  $u$  in root path from  $v$  to  $w$  inclusive do
    if  $u \in Vp$  then
       $out[u] := val; node[u].valid := \text{true};$ 
    end;
  end;
end;
end;
end.
```







- Compute all markers using parallel erosions and store each scale i in an image $out[i]$;
- Compute a single Max-Tree in parallel, maintaining the maximum marker for all scales i for each node n in $out[i][n]$;
- **Compute reconstruction (in parallel) for scale all scales i , storing the result in $out[i]$;**
- Compute differences between scales (in parallel);
- Output respective instance of DMP vector field.



```
procedure MaxTreeReconstruct ( $Vp$  : Section; var  $node$  : Max-Tree; var  $out$  : DMP)
  for all  $v \in Vp$  do
    if not  $node[v].valid$  then
       $w := v;$ 
      for all scales  $i$  increasing order do
        while  $Par(w) \neq \perp \wedge \text{not } node[w].valid \wedge f[Par(w)] > out[Par(w)]$  do
           $w := Par(w);$ 
        end;
         $ws[i] := w$ ; (* temporary storage of for each scale *)
        if  $node[w].valid$  then
          for all scales  $j \geq i$  do
             $val[j] := out[j][w]; ws[j] := w$  (* Reconstructed node found *)
          end;
        else if  $Par(w) \neq \perp$  then
           $val[i] := out[i][w] \vee f[Par(w)];$  (*  $w$  is reconstructed *)
        else (* marker was empty *)
          for all scales  $j \geq i$  do
             $val[j] := 0; ws[j] := w$ ; end;
          end;
        end;
      end;
    end;
```



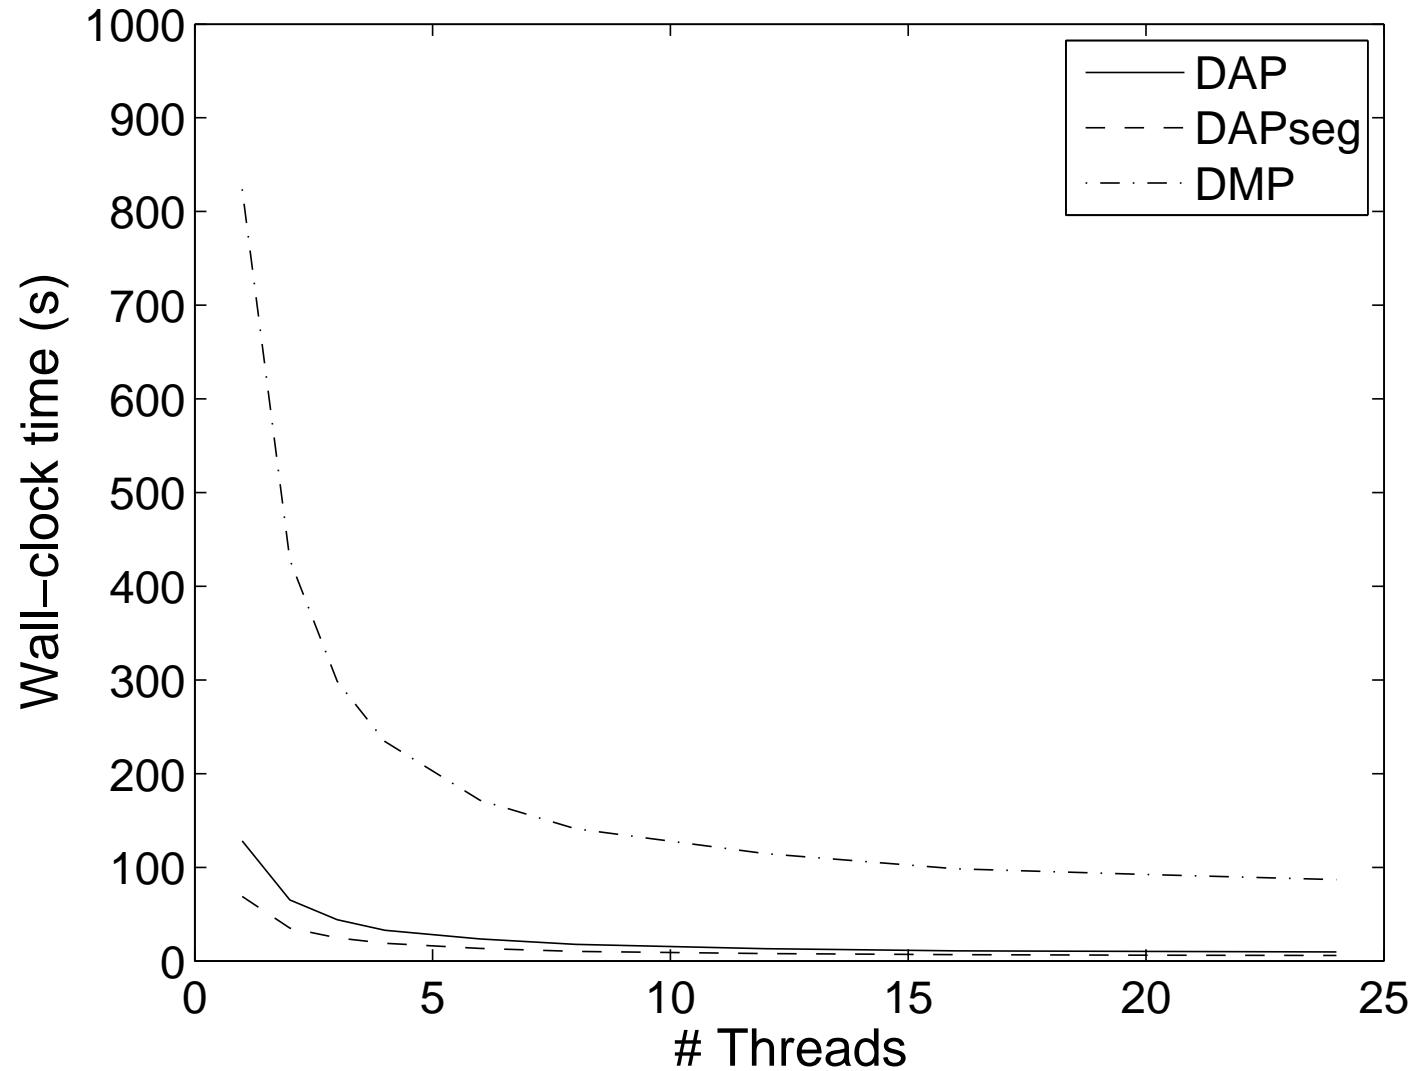
```
u := v;  
for all scales i increasing order do  
    repeat  
        if u ∈ Vp then  
            for all scales j < i do out[j][u] := f[u];end;  
            for all scales j ≥ i do out[j][u] := val[j]; end;  
            node[u].valid := true;  
        end;  
        u := node[u].parent;  
    until u = ws[i];  
 end;  
 if u ∈ Vp then (* Process ws[numscales - 1] *)  
    for all scales j < i do out[j][u] := f[u];end;  
    for all scales j ≥ i do out[j][u] := val[j]; end;  
    node[u].valid := true;  
 end;  
 end;  
 end;  
end.
```

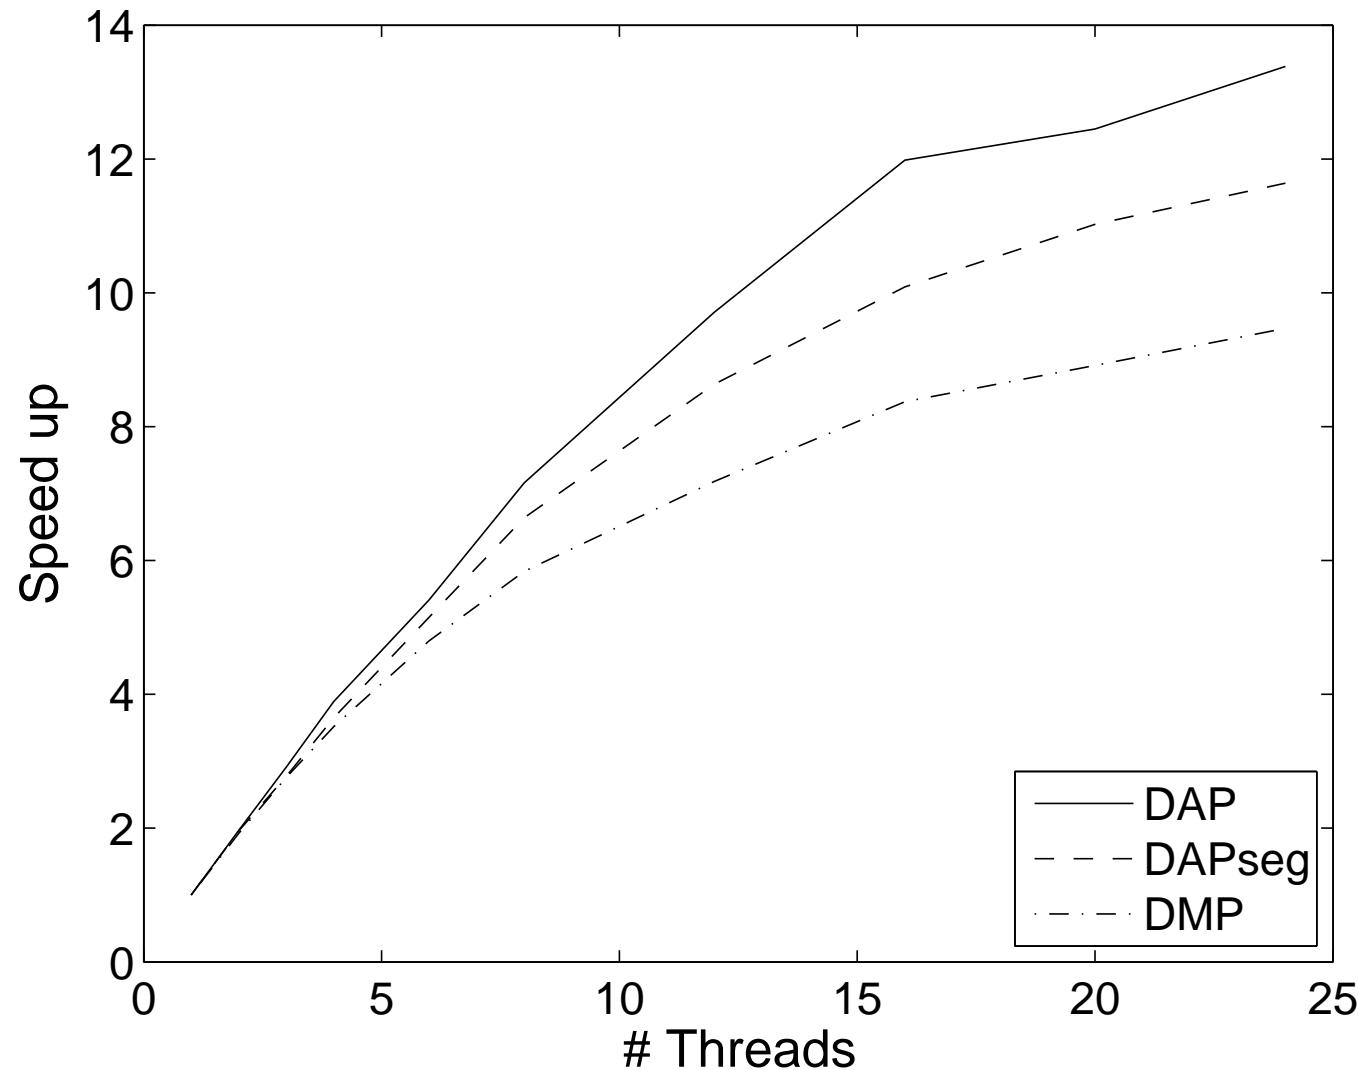


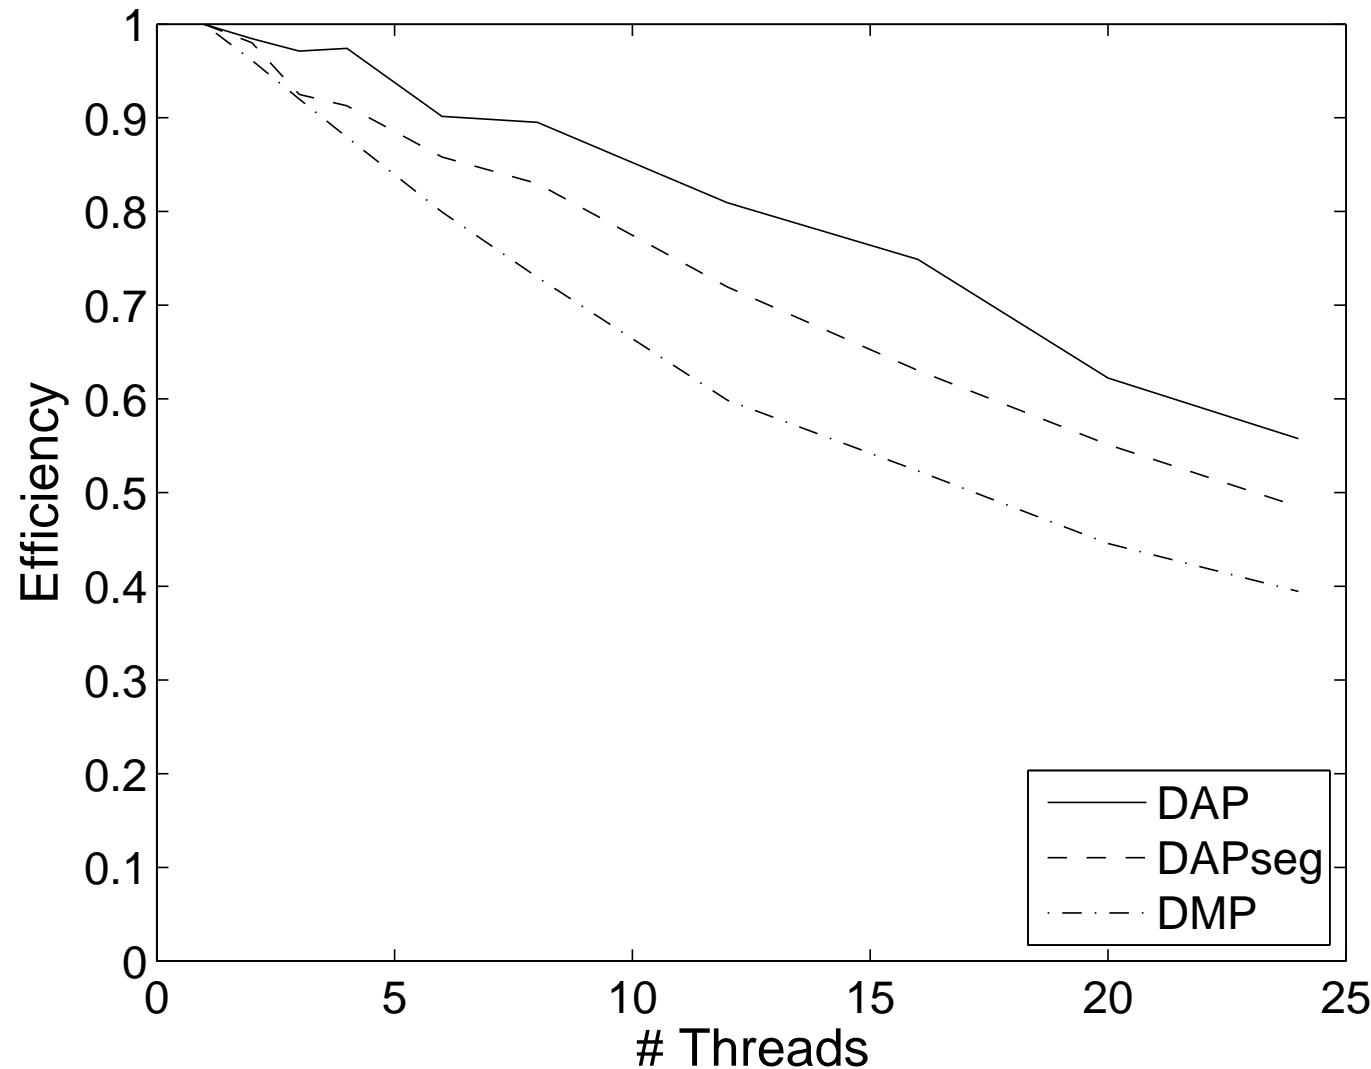
```
process parDMP( $p : ThreadID$  ;  $f : Image$  ; var  $node : Max\text{-Tree}$  ; var  $out : DMP$ )  
  find xm such that  $f(xm) \leq f(x) \forall x \in V^p$ ;  
   $hm := f(xm)$  ;  $set[hm] := \{xm\}$  ;  $lero[hm] := xm$  ;  
  for all scales s do  $maximum[s] := 0$ ; end;  
   $flood(V_p, hm, maximum, f, node, out)$ ;  
  var  $i := 1$  ,  $q := p$  ;  
  while  $p + i < K \wedge q \bmod 2 = 0$  do  
     $P(sa[p + i])$  (* wait to glue with right-hand neighbor *) ;  
    for all edges  $(x, y)$  between  $V(p)$  and  $V(p + i)$  do  
       $connect(x, y, f, node)$  ;  
    end ;  
     $i := 2 * i$  ;  $q := q/2$  ;  
  end ;  
  if  $p = 0$  then (* release the waiting threads *)  
    for  $i := 1$  to  $K - 1$  do  $V(sb[i])$  end  
  else  
     $V(sa[p]); P(sb[p]);$  (* signal left-hand neighbor; wait for thread 0 *)  
  end ;  
   $MaxTreeReconstruct(V_p, node, out)$  ;  
end.
```

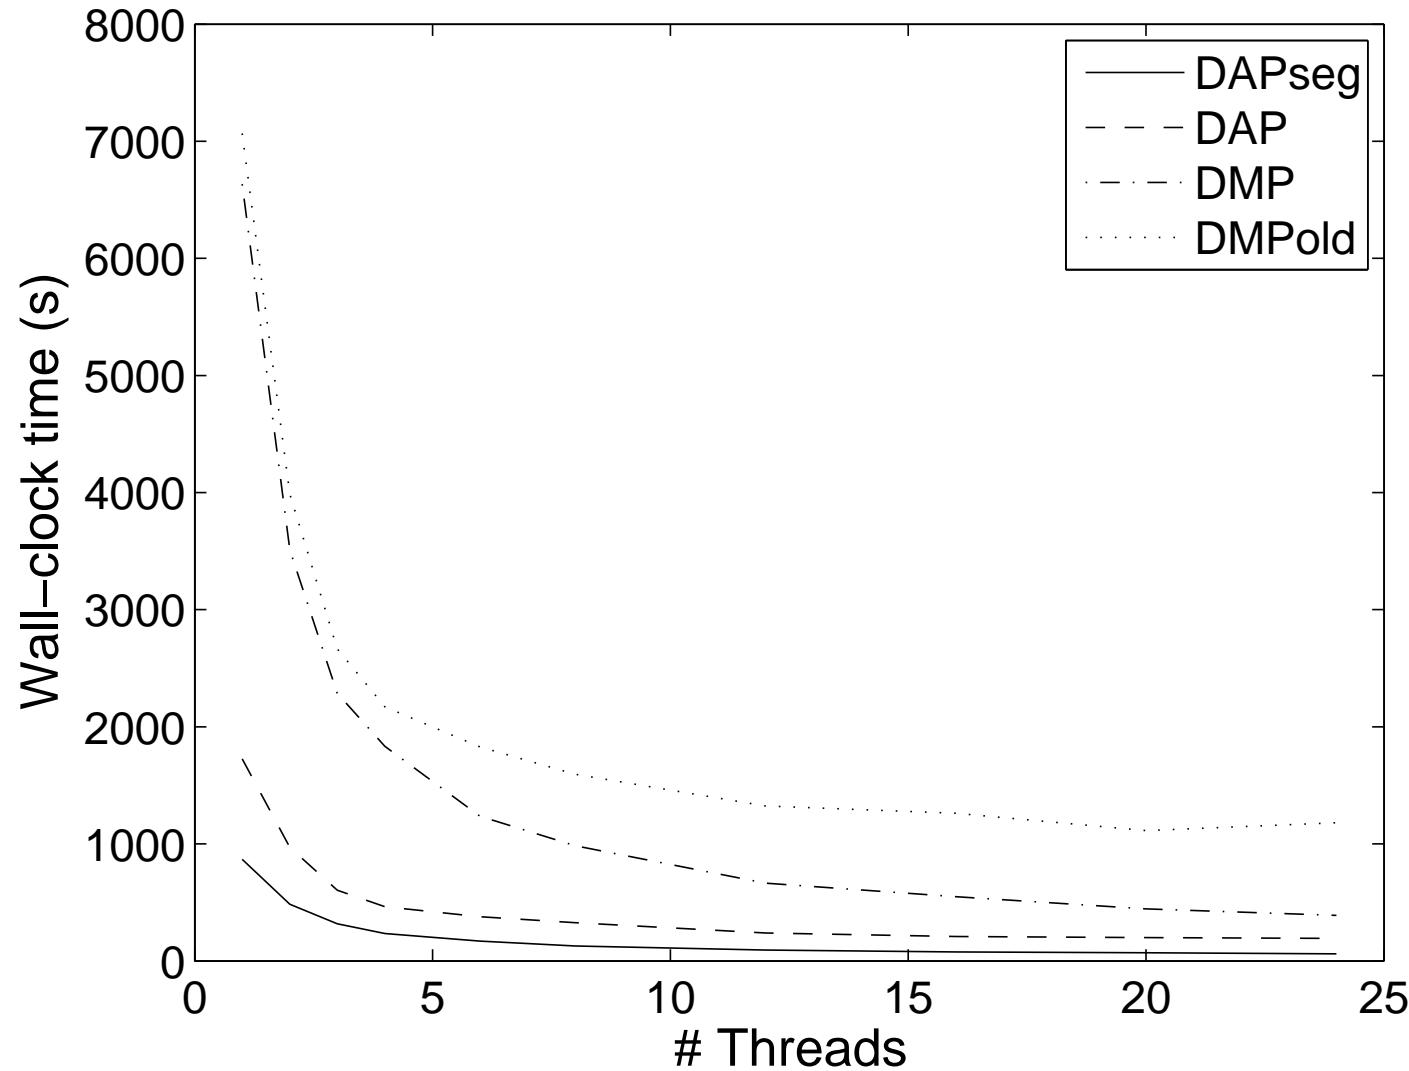


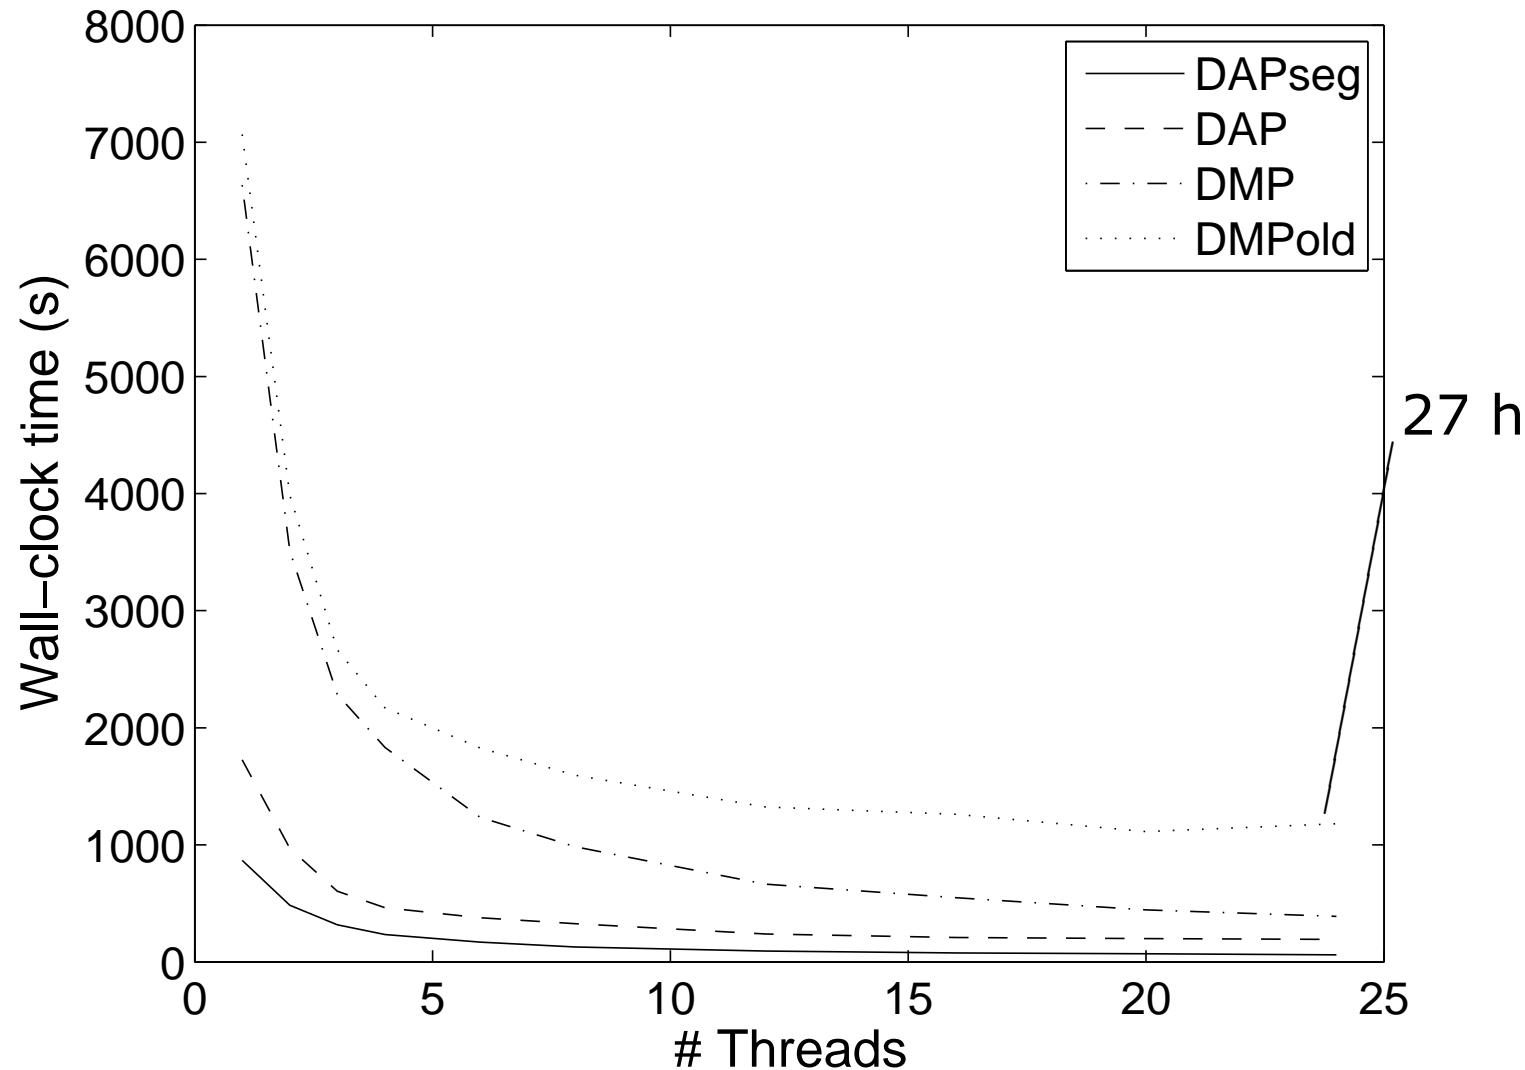
- Images: 120 MB and 1.2 GB
- Parallel Algorithms:
 - Differential Morphological Profile (DMP)
 - Differential Area Profile (DAP)
 - DAP-segmentation
- The Machine
 - Part of Millipede cluster
 - quad socket, opteron machine, 6 cores per socket
 - 128 GB memory

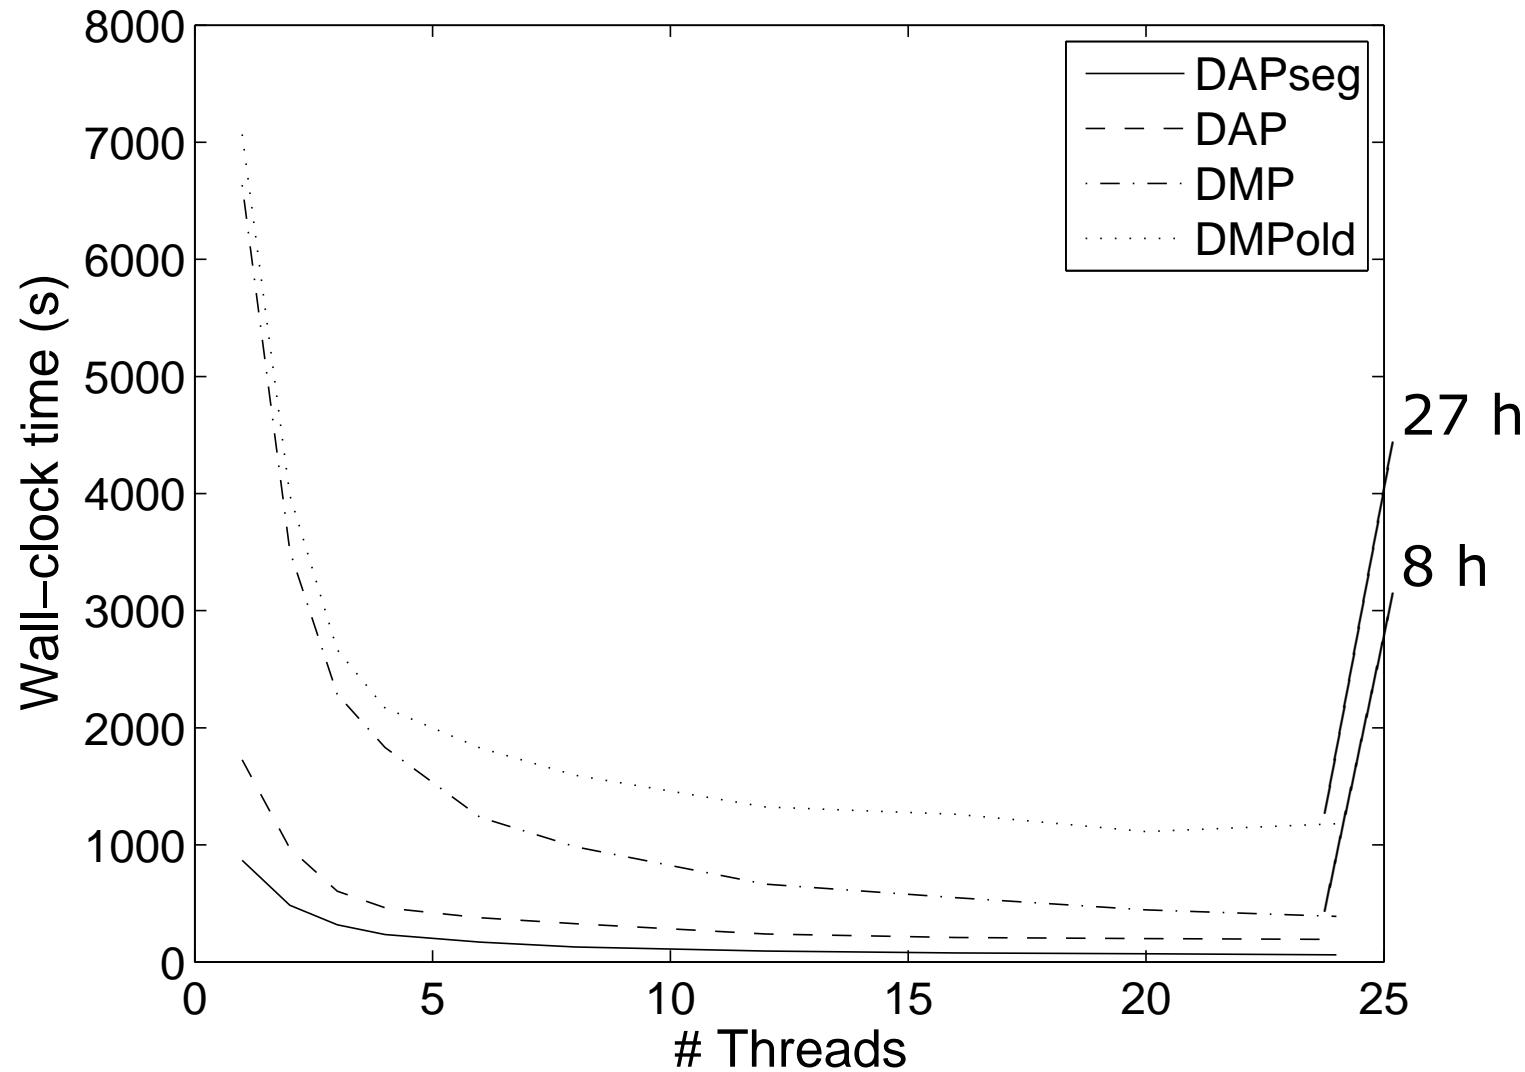


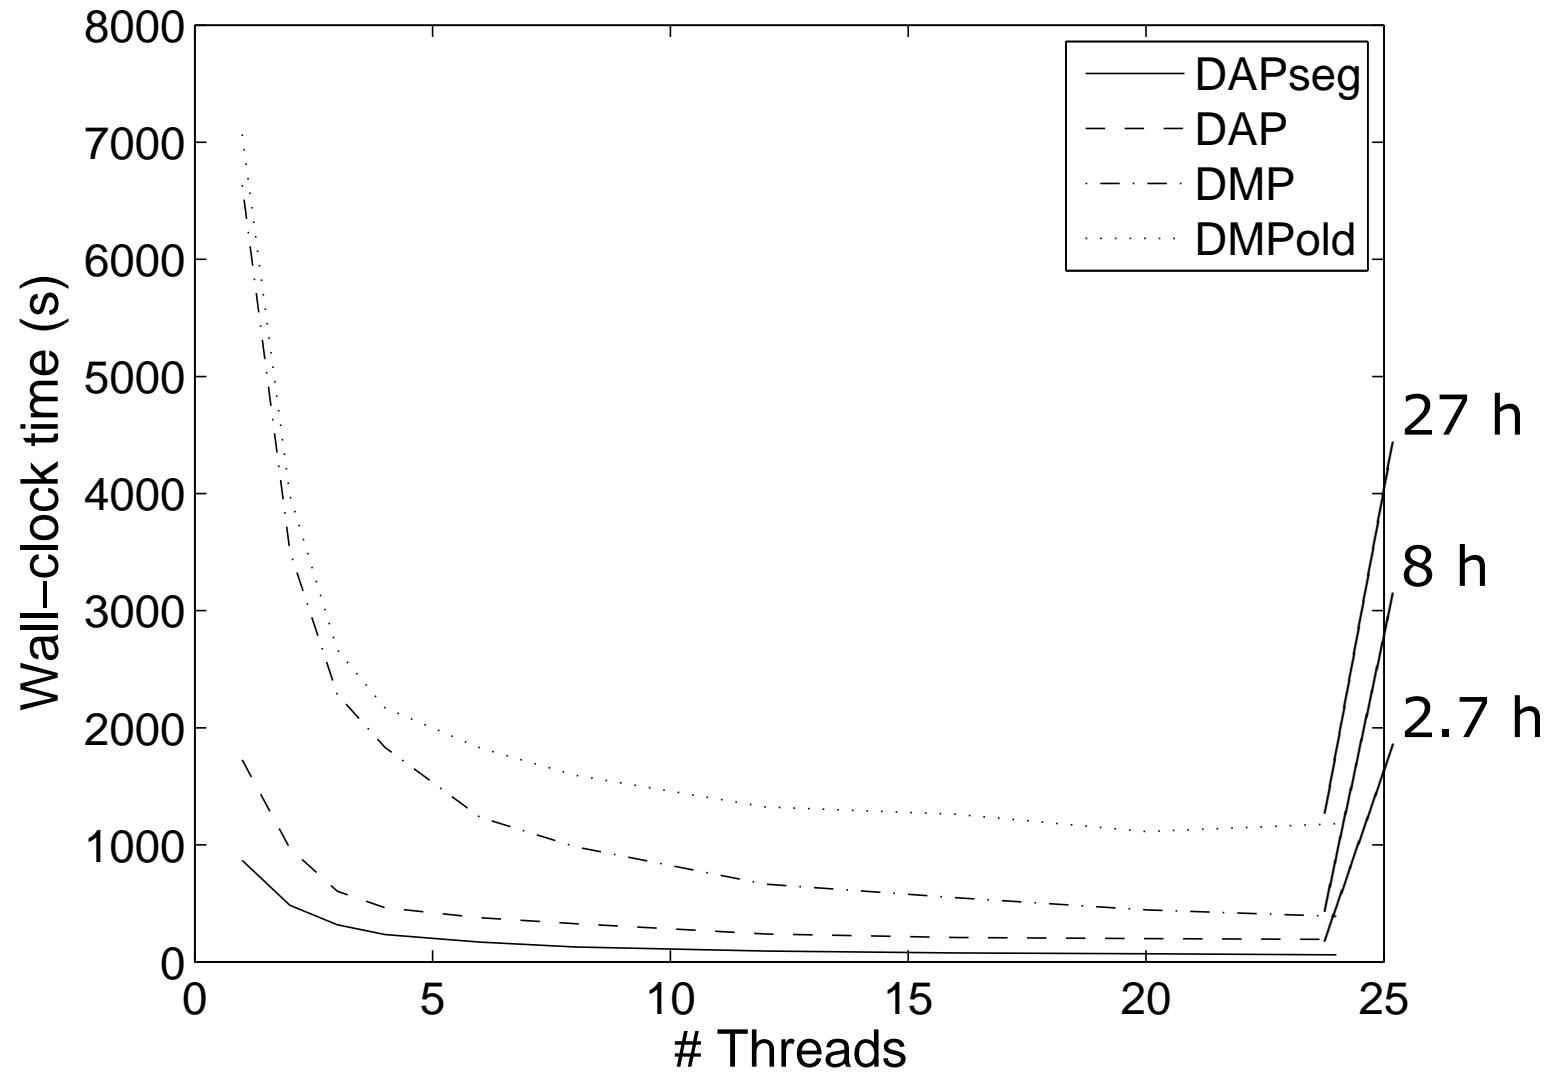


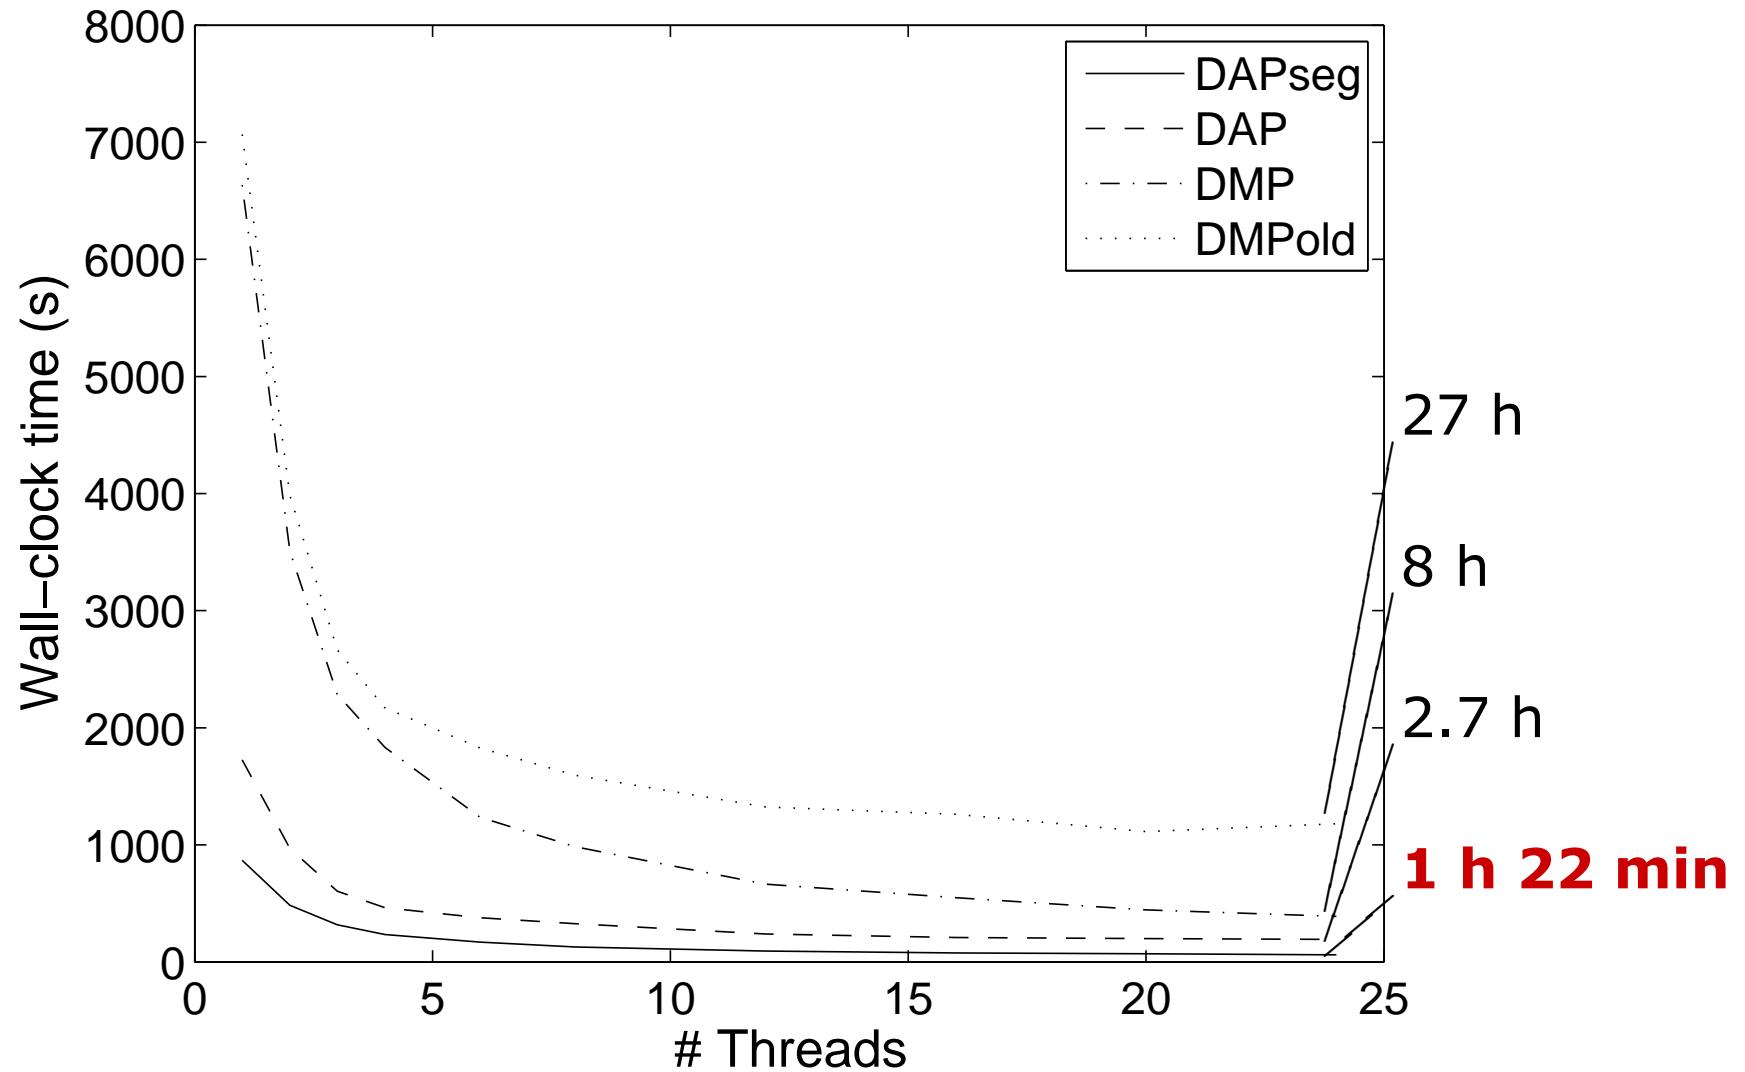


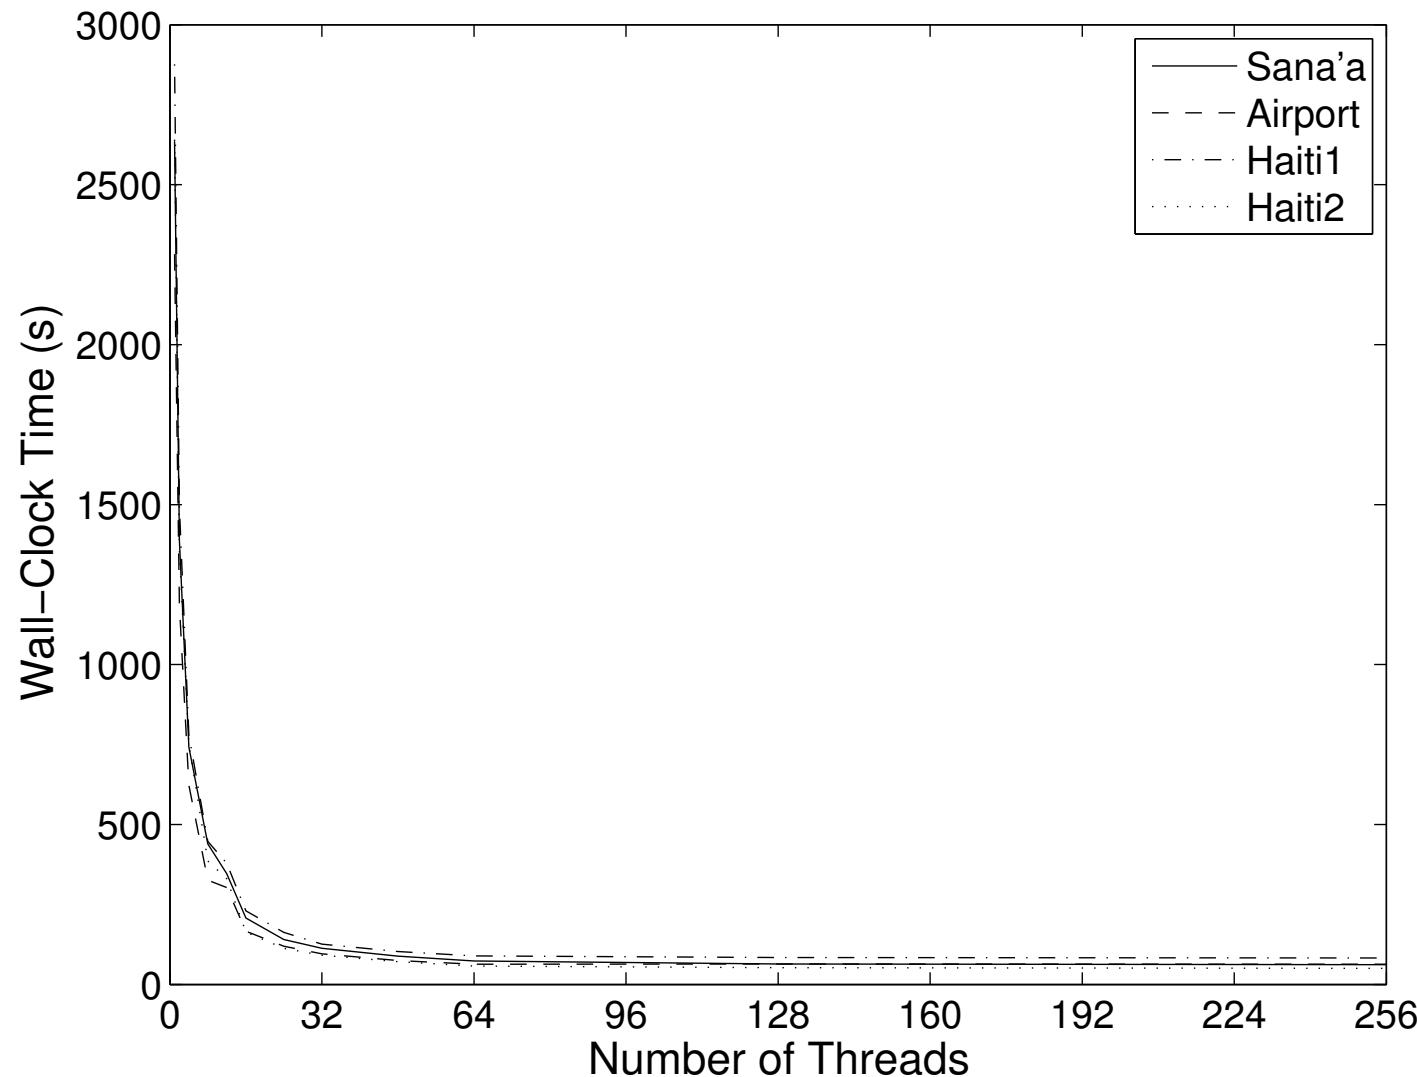






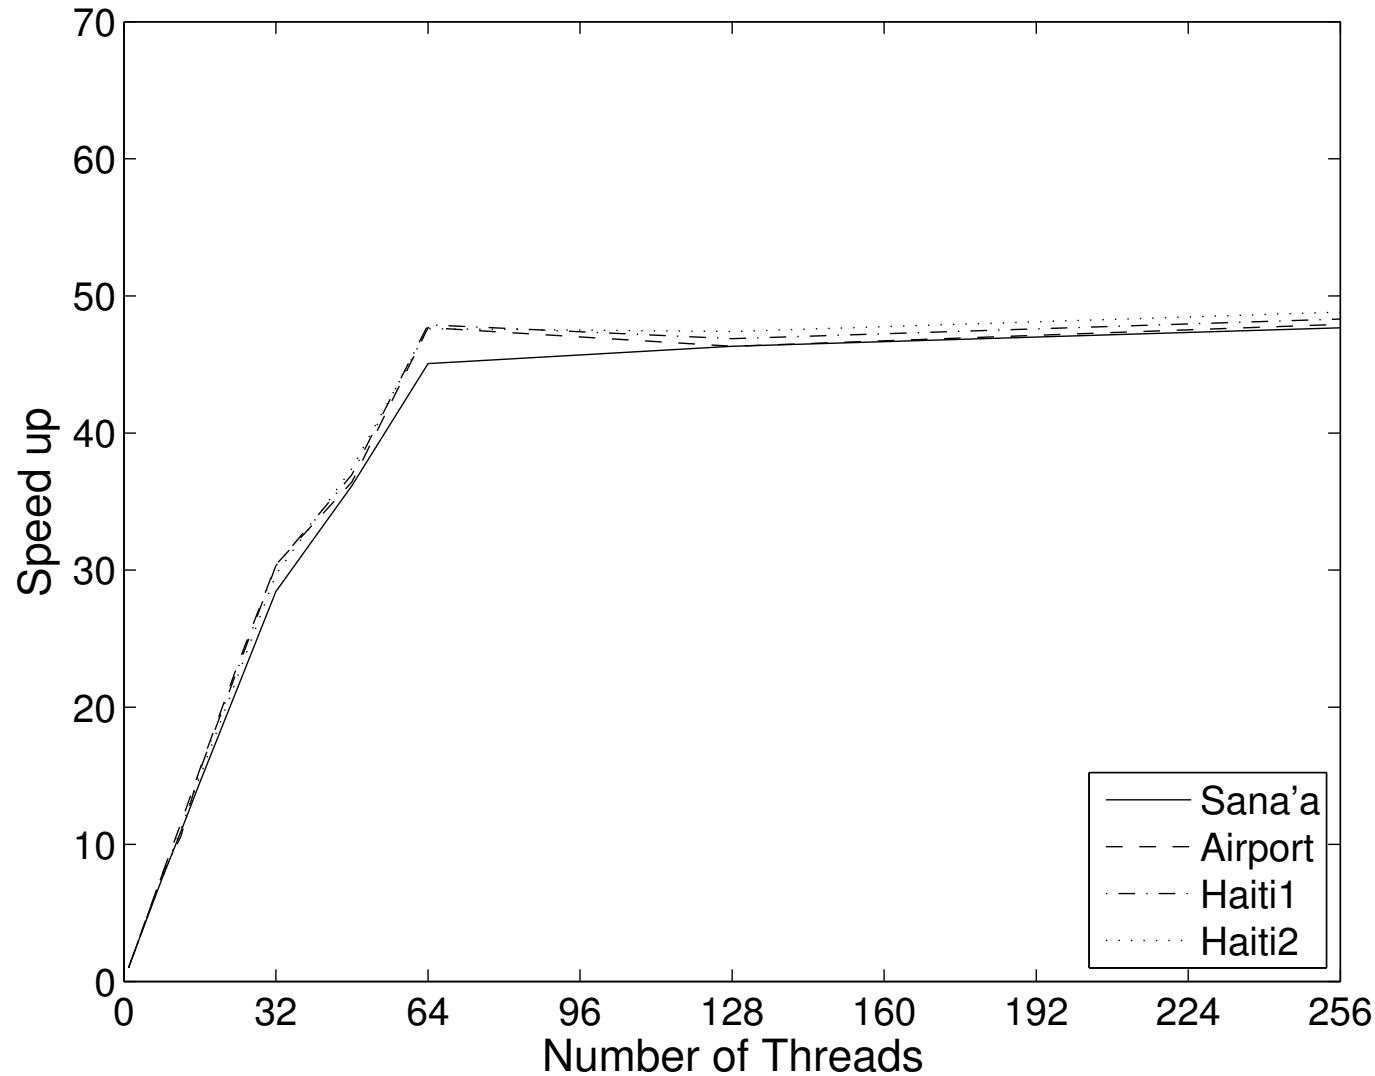






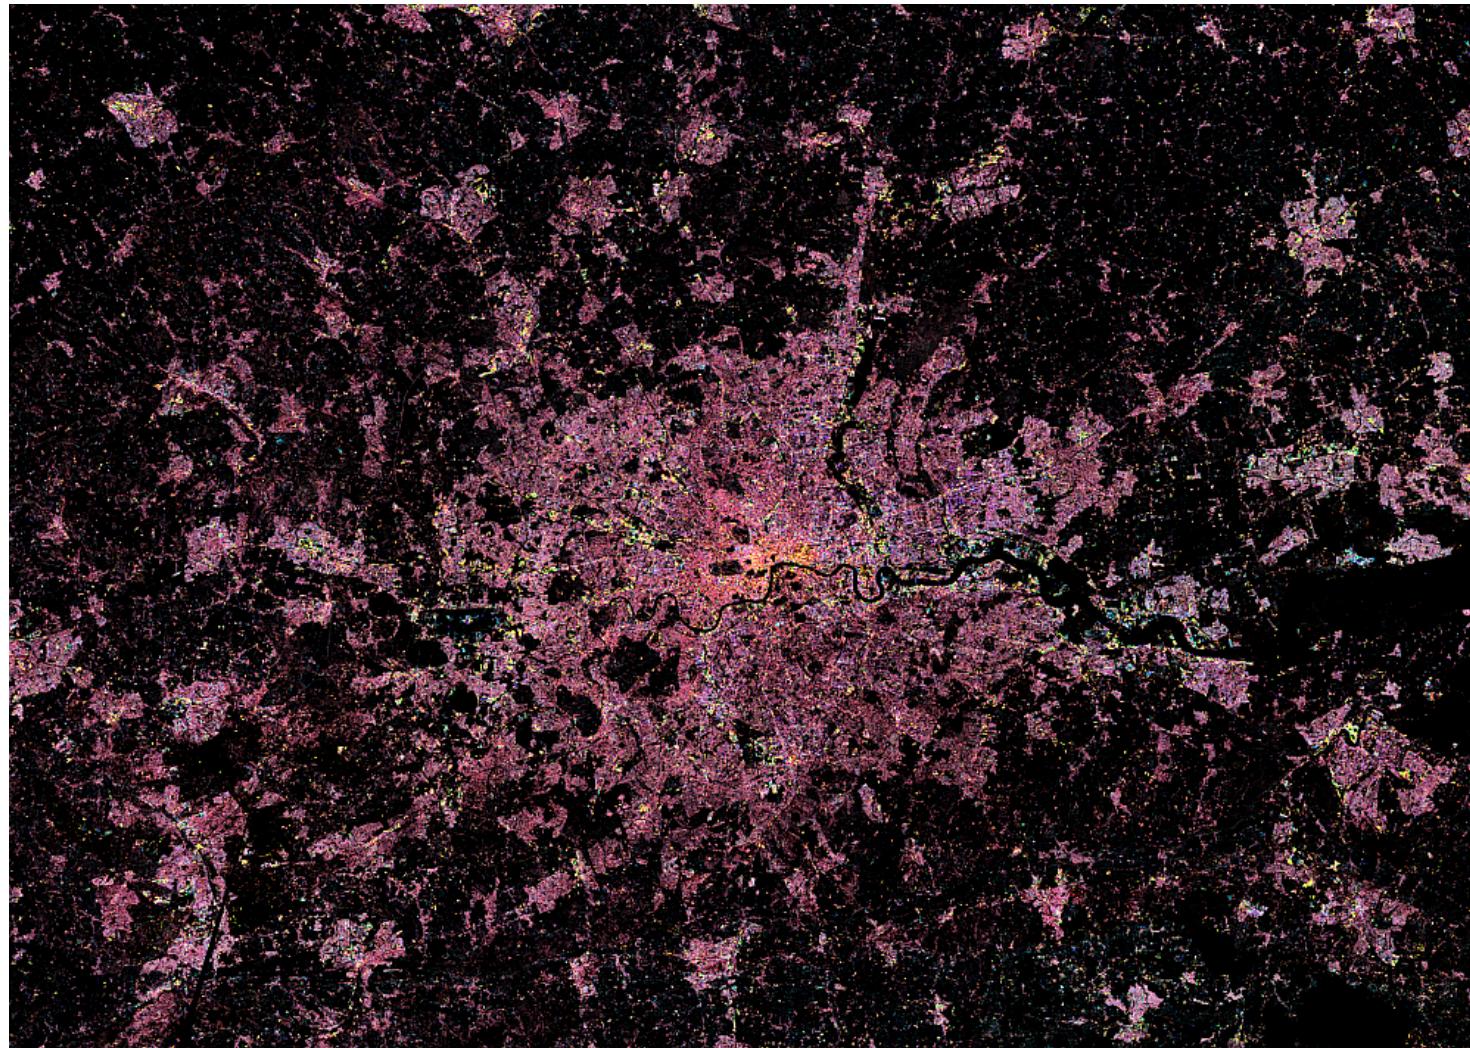
Bigger machines: 64 cores, part II

21 of 36





London, UK. © Landsat; 2012.



London, UK. © European Commission; Joint Research Centre; 2012.



London, UK. Night-lights by Andre Kuiper. Credit ESA/NASA. Posted at Kuipers' photostream in Flickr.com.
Taken from the ISS on April 5, 2012 using a Nikon D3S camera.



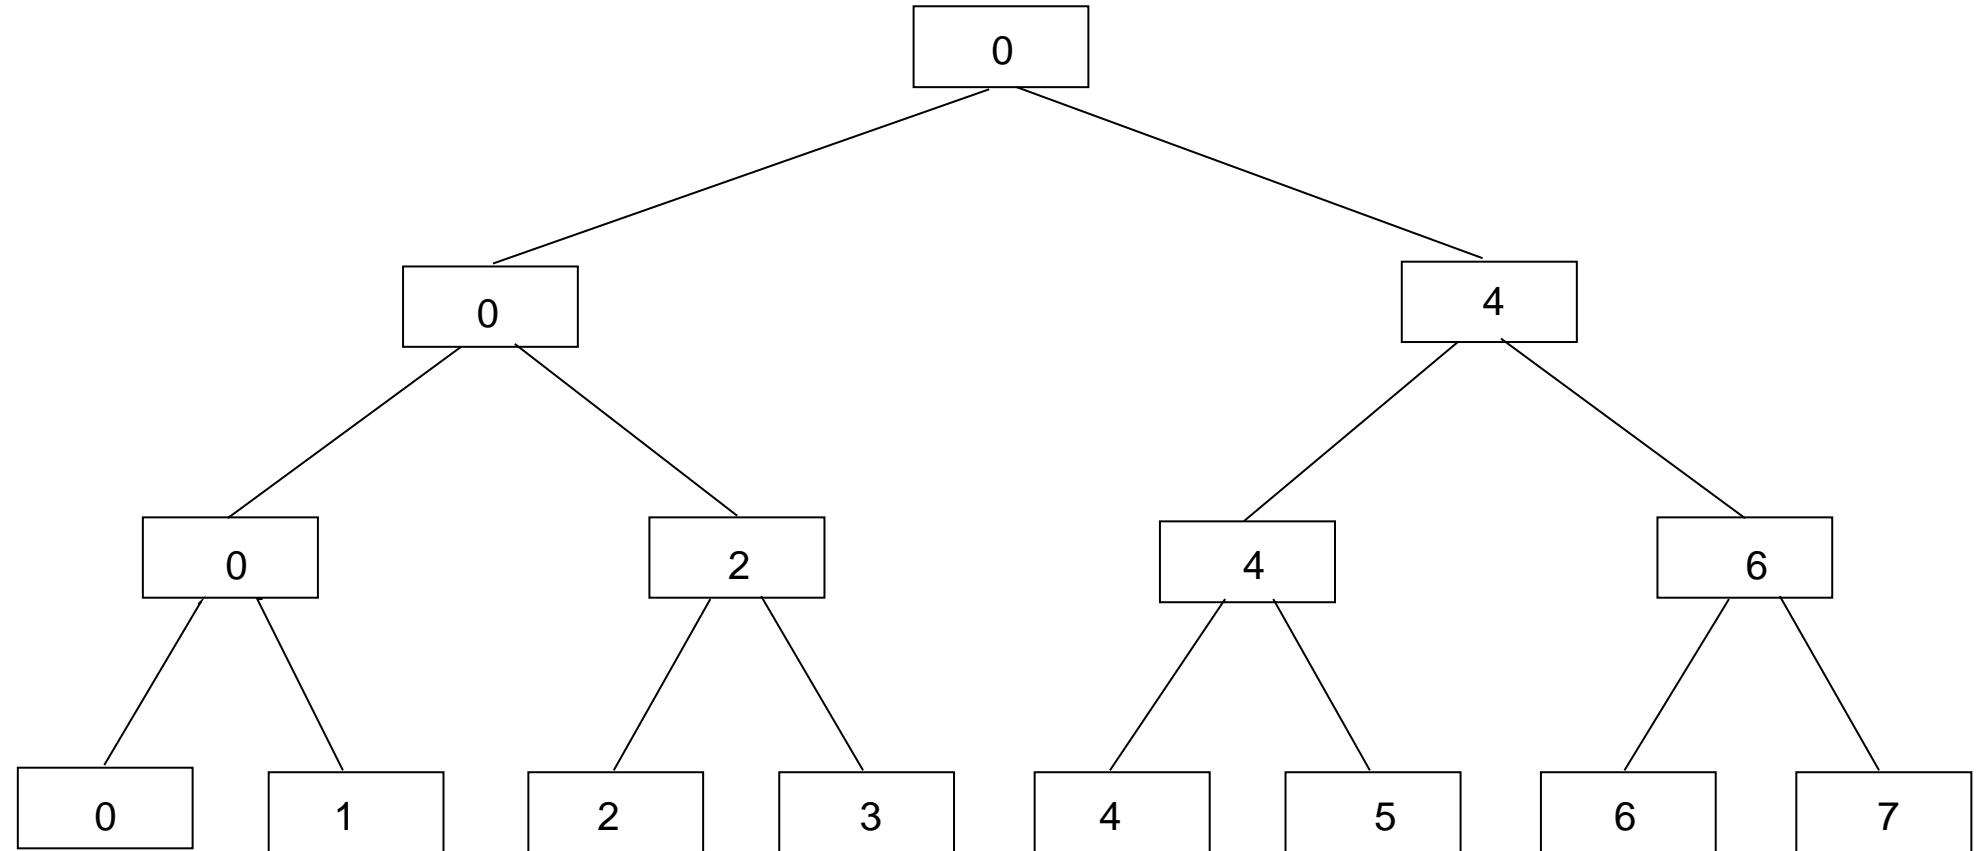
More data of course!

- Modern stitching methods generate image $\gg 10$ Gpixel
- Shared memory sizes insufficient
- Astronomical data sets astronomical in size
- Haiti aerial image data set 1.5 Tpixel

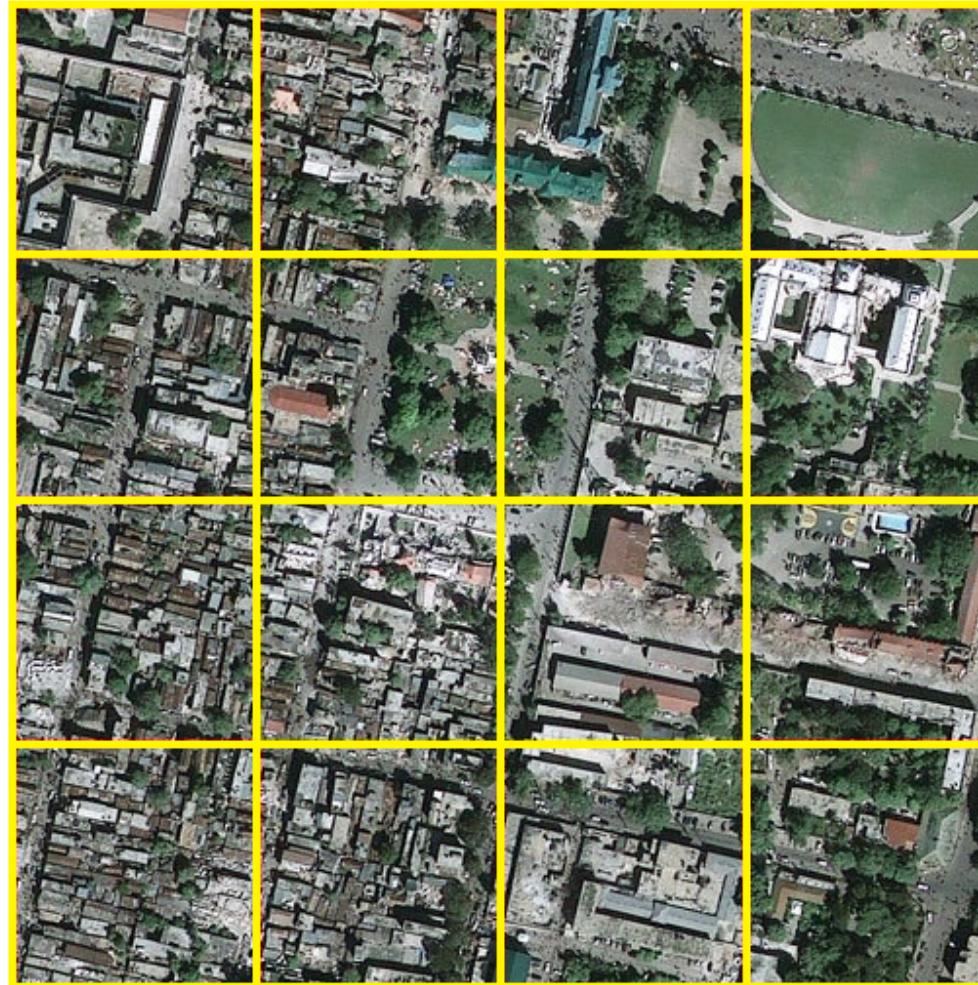
How can we handle these?

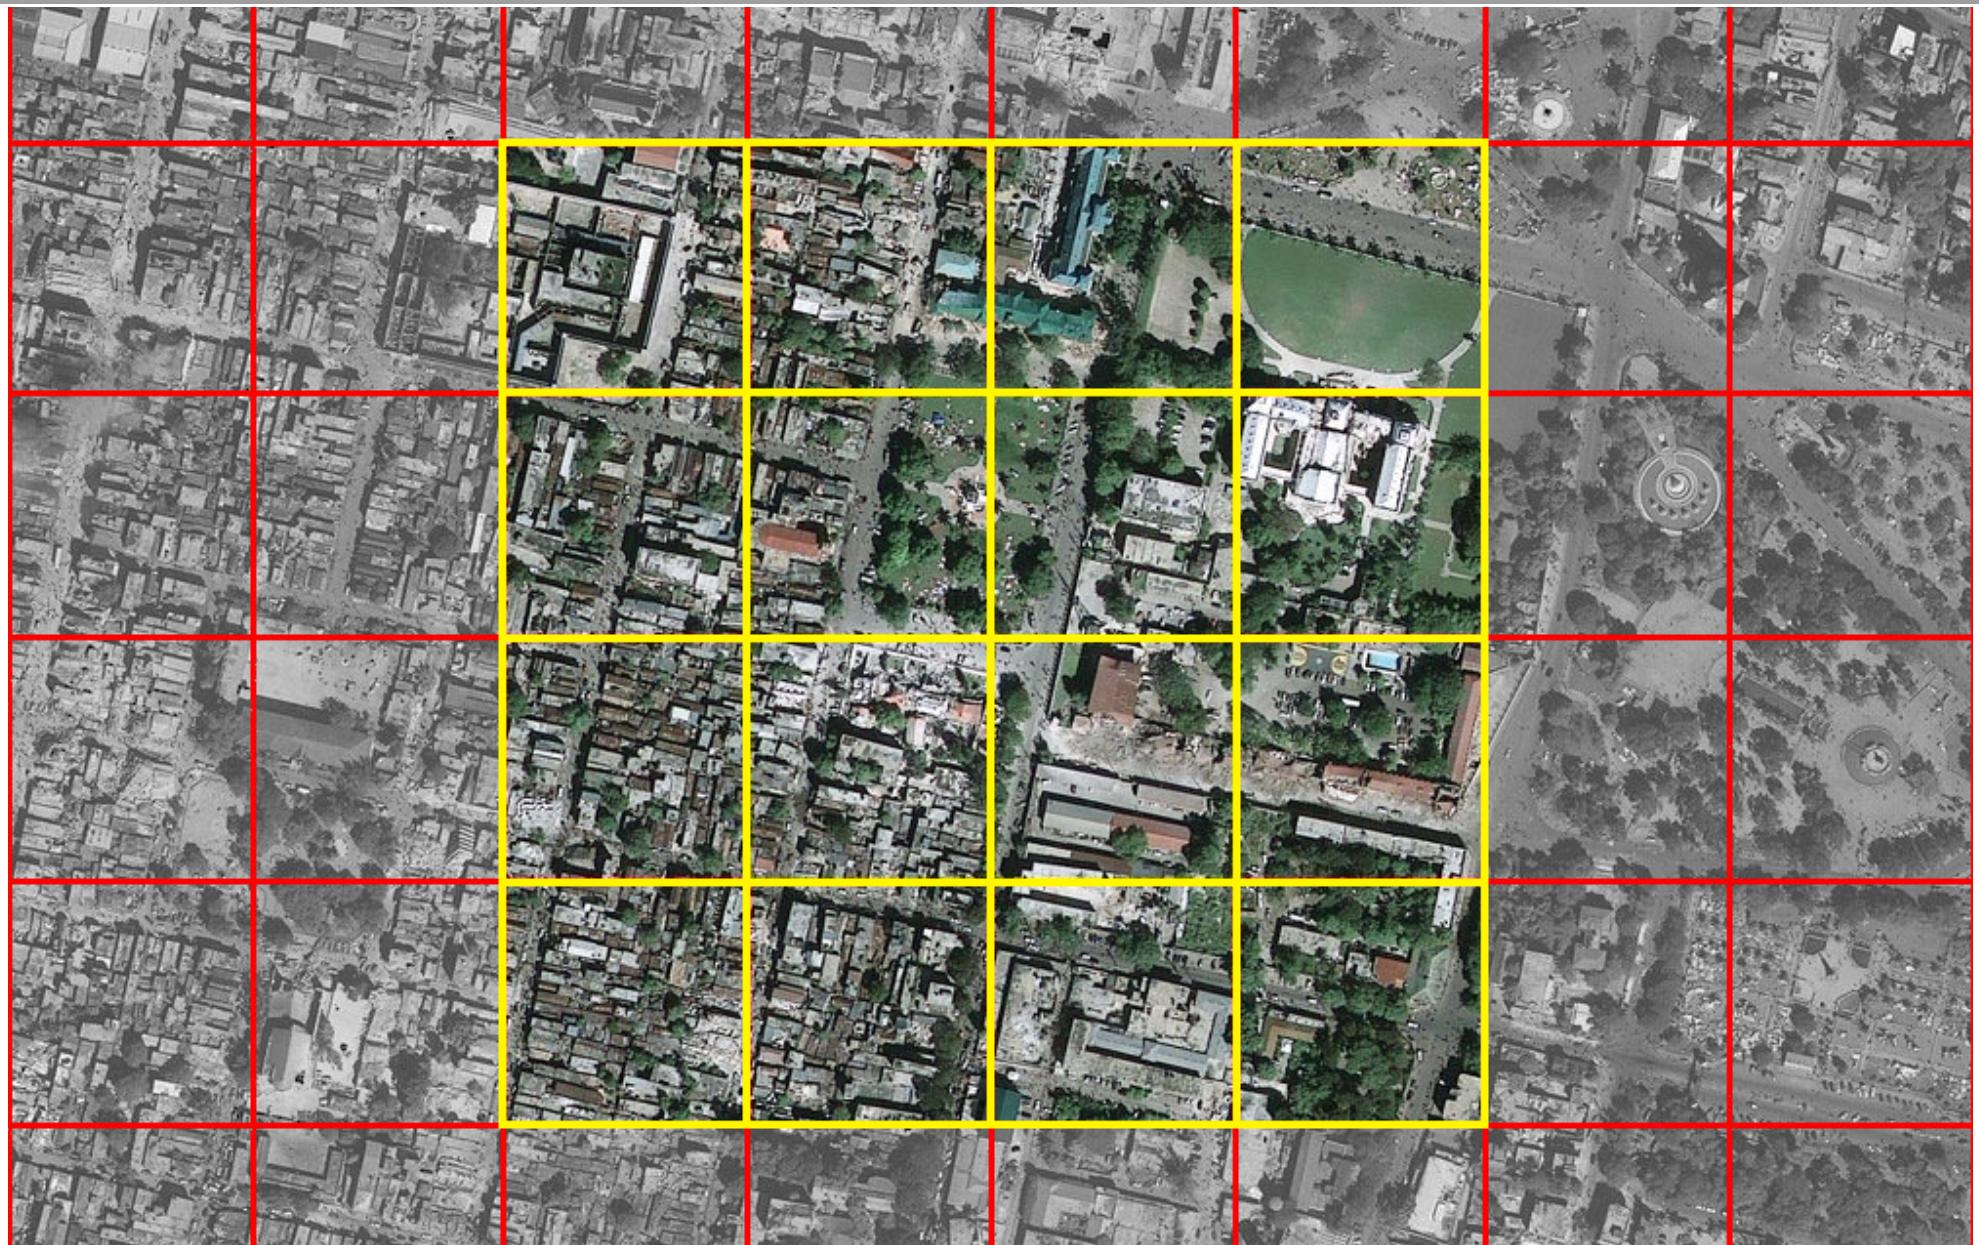


Folding@Home?



MapReduce?







A boundary tree of an image block $B_{i,j}$ is

- a subset of the Max-Tree of $B_{i,j}$
- containing only nodes C_h^k for which P_h^k touches the boundary $\delta B_{i,j}$ of $B_{i,j}$

Observations:

- During merging of max-trees of neighbouring blocks *only the boundary tree nodes are affected*
- Therefore, only boundary trees need be merged



- Compute Max-Trees of all blocks $B_{i,j}$
- Extract boundary trees
- Hierarchically merge boundary trees up to root
- Hierarchically transmit updated boundary trees down to leaves
- Update Max-Tree of each block $B_{i,j}$



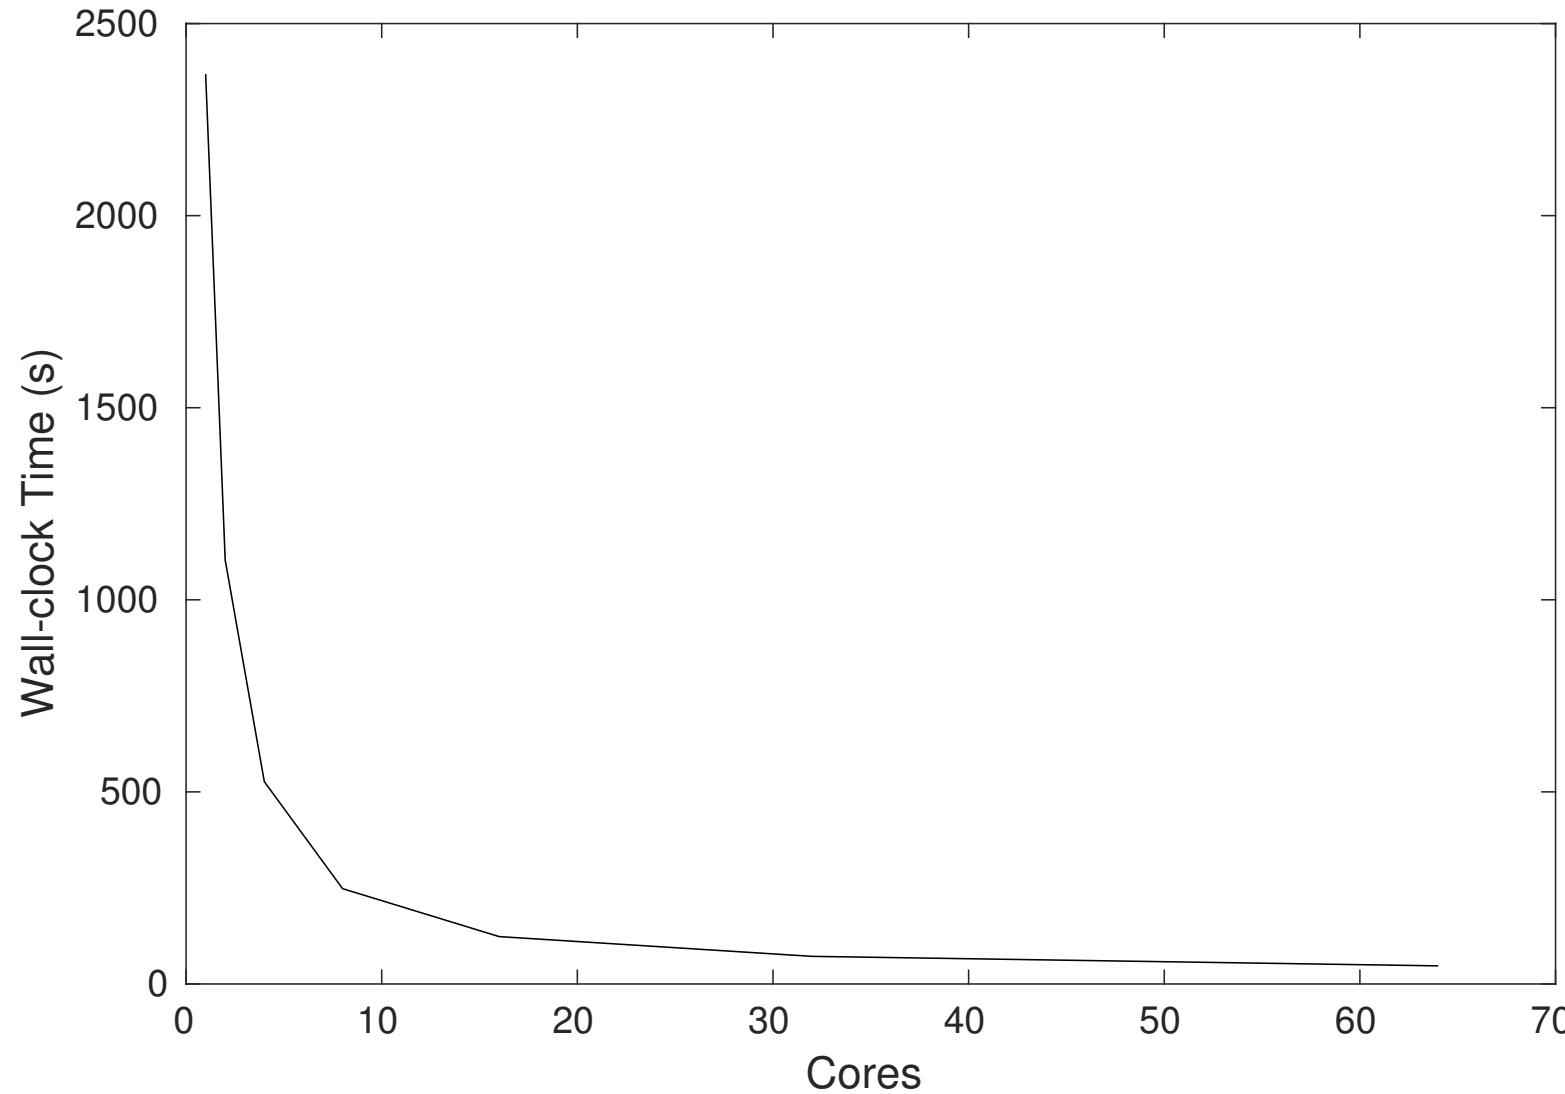
- The boundary $\delta B_{i,j}$ contains at most $\frac{|\delta B_{i,j}|}{2}$ local maxima
- The number of nodes is at most $\min(\frac{|\delta B_{i,j}|}{2}G, |B_{i,j}|)$
- Provided $\frac{|\delta B_{i,j}|}{2}G < |B_{i,j}|$ we have *reduction*
- For $G = 256$ and $|B_{i,j}| = 40960 \times 40960 = 1.6$ Gpixel we have $\frac{|\delta B_{i,j}|}{2}G = 512 \times 40960$ or $80\times$ reduction

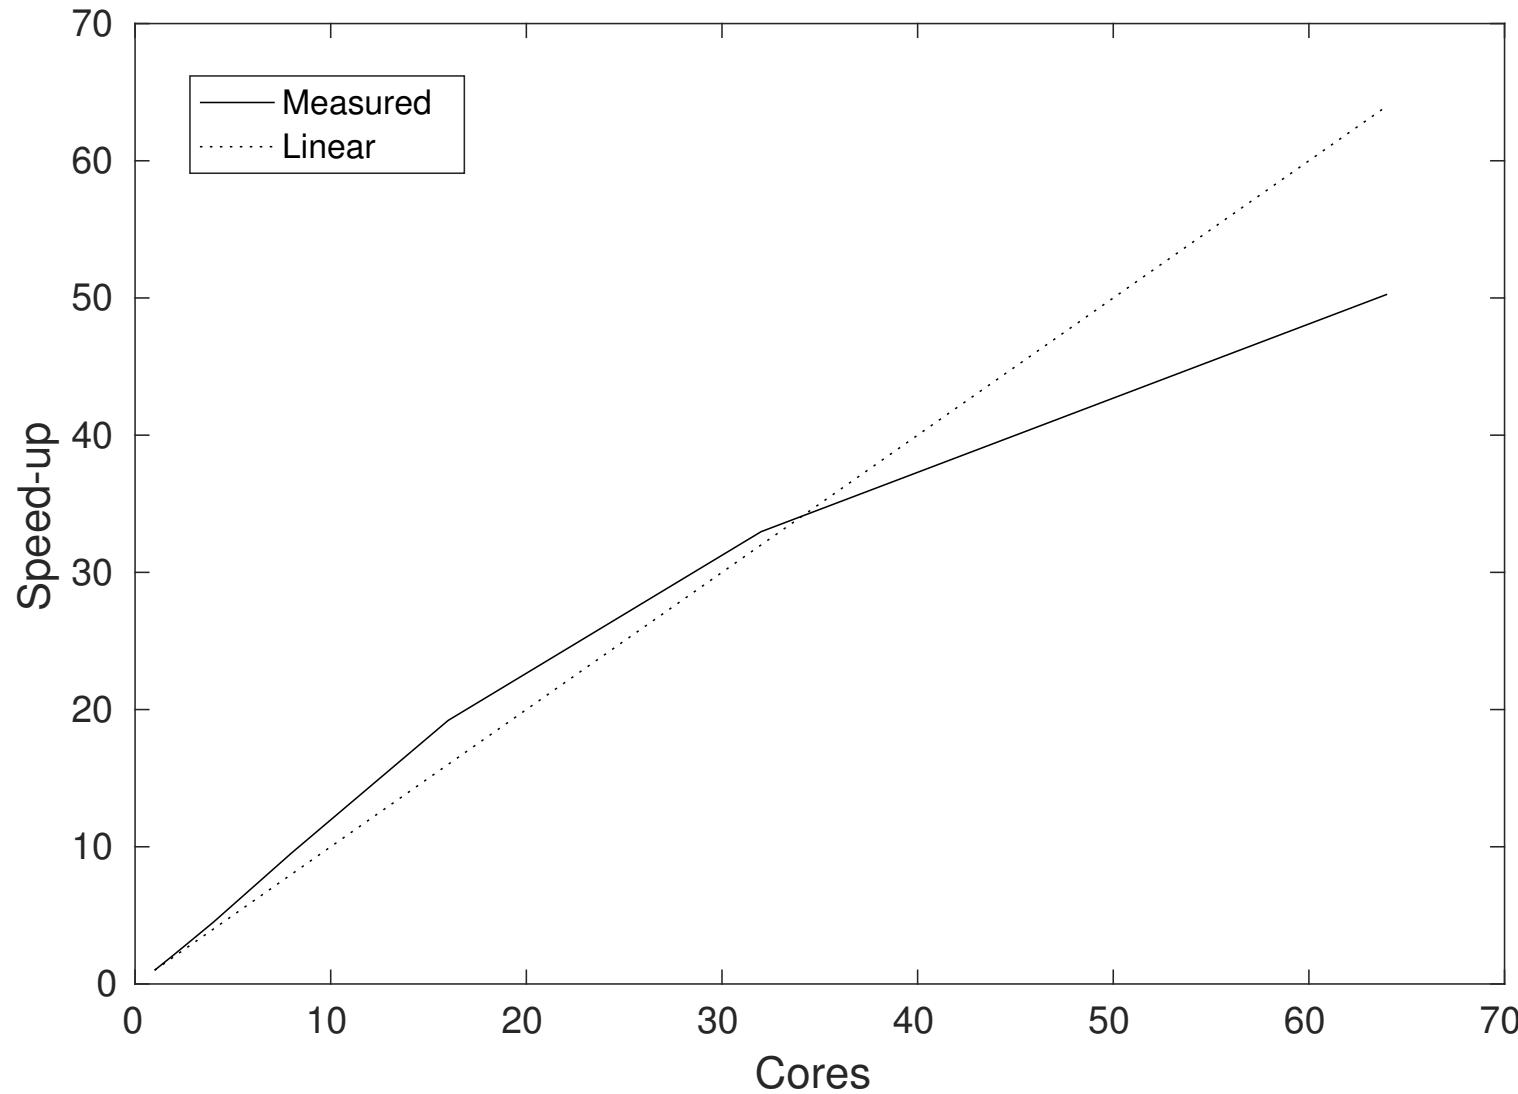


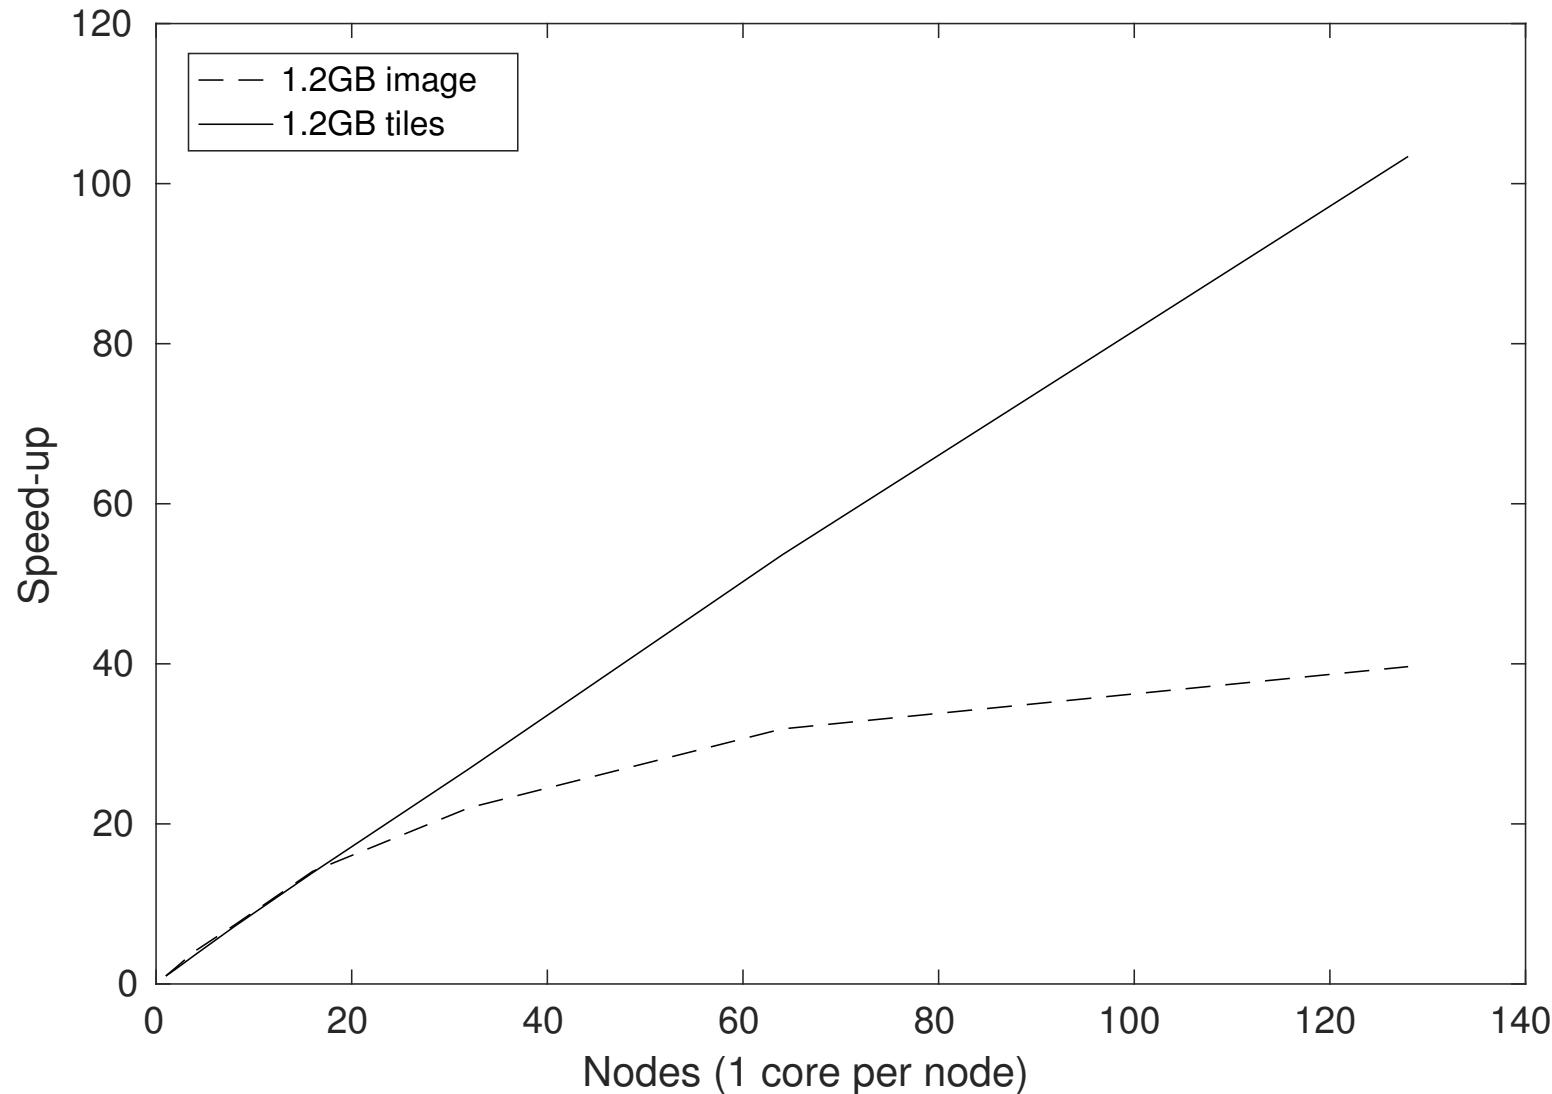
iteration	Imsize	$ B_{i,j} $	$\frac{ \delta B_{i,j} }{2}G$	Mem (orig)	Mem (new)
1	6.4 G	1.6 G	20 M	128.0 GB	33.6 GB
2	25.6 G	6.4 G	40 M	512.0 GB	35.2 GB
3	102.4 G	25.6 G	80 M	2048.0 GB	38.4 GB
4	409.6 G	102.4 G	160 M	8192.0 GB	44.8 GB
5	1.6 T	409.6 G	320 M	32.0 T	57.6 GB
6	6.4 T	1.6 T	640 M	128.0 T	83.2 GB
7	25.6 T	6.4 T	1280 M	512.0 T	134.4 GB

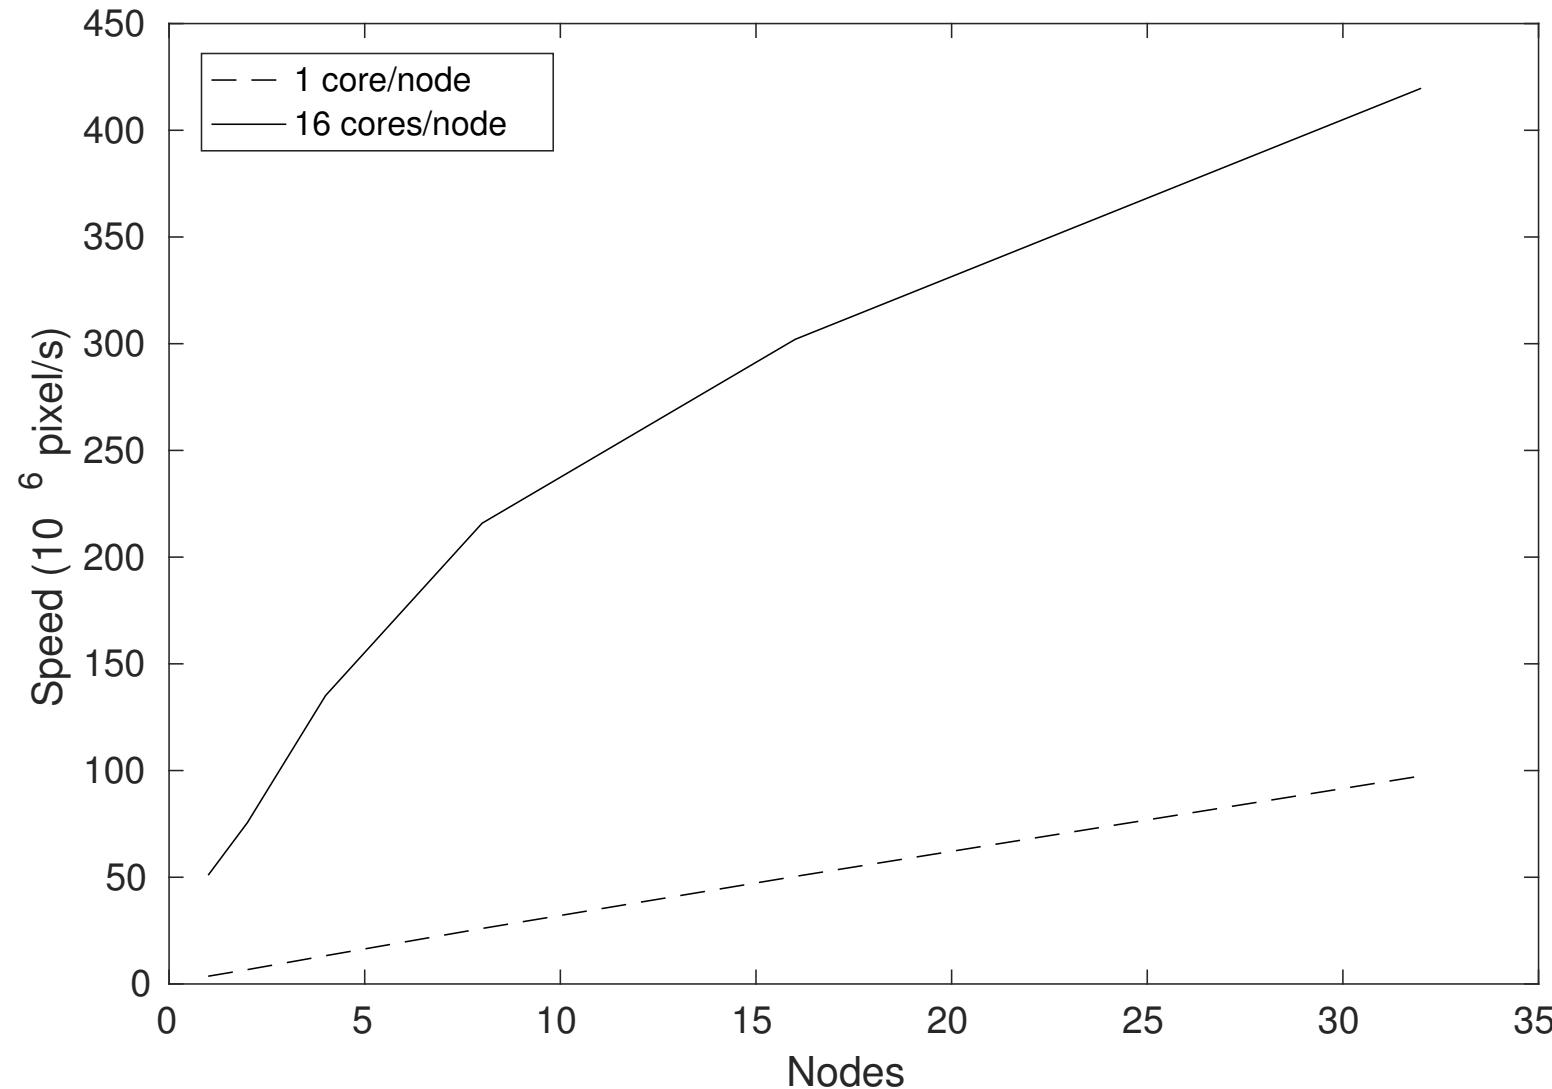


- We implemented the algorithm in MPI
- First test were done on a 3.8 Gpixel image on a 64-core compute server
- Three experiments were done on the Peregrine cluster
 - Hierarchically dividing a 1.2 Gpixel image into 1, 2, ..., 128 tiles, using one core per node
 - Using 1, 2, ..., 128 tiles of 1.2 Gpixel to create increasing image sizes, again one core per node
 - Same as above, but splitting each 1.2 Gpixel tile into 16, using 1, 2, ..., 32 nodes, using 16 cores per node











iteration	Imsize	$ B_{i,j} $	$\frac{ \delta B_{i,j} }{2}G$	Mem (orig)	Mem (new)
1	6.4 G	1.6 G	5.12 G	128.0 GB	37.12 GB
2	25.6 G	6.4 G	10.24 G	512.0 GB	42.24 GB
3	102.4 G	25.6 G	20.48 G	2048.0 GB	52.48 GB
4	409.6 G	102.4 G	40.96 G	8192.0 GB	72.96 GB
5	1.6 T	409.6 G	81.92 G	32.0 T	103.92 GB
6	6.4 T	1.6 T	163.84 G	128.0 T	195.84 GB
7	25.6 T	6.4 T	327.68 G	512.0 T	359.68 GB

Looks like 16 bits per pixel is possible!



- An 8-bits per pixel MPI implementation has been built
- It works on a shared memory machine
- Parity in performance and speed-up with shared memory algorithm (Wilkinson et al 2008)
- Tests on cluster show huge performance up to 128 nodes
- Same approach should work on α -partition tree
- 16 bits per pixel should be doable
- Out-of-memory computation possible
- What remains is high-dynamic-range distributed-memory algorithm

