

# Maximum likelihood estimation (of the parameters of class conditional probability)

# Problem formulation

- To design an optimal Bayesian classifier we need priors  $P(\omega_j)$  and class-conditional probabilities  $p(\mathbf{x}|\omega_j)$
- In practice, they are usually not available
- Available is some (hopefully representative) data
- Problem: how to design a classifier using this training data?
- Priors are easier to estimate  $P(\omega_j)$
- Specific problem: estimation of class-conditional probabilities  $p(\mathbf{x}|\omega_j)$
- Simplification: estimation of the parameters of a function of known type, e.g.  $\mu_i$  and  $\Sigma_i$  of normal density

# Purpose of MLE

Assume a data set  $\mathcal{D} = \{X_1, X_2, \dots, X_n\}$  of  $n$  feature vectors from class  $\omega$

Assume that  $p(x|\omega)$  has a known parametric form, such as  $p(x|\omega) \sim N(\mu, \Sigma)$


Denote by  $\theta$  the parameters of the distribution, e.g.  $\theta = [\mu, \Sigma]$

Hence, we know the form of  $p(x|\theta)$

but we do not know the values of the parameters  $\theta$

The purpose of MLE is to estimate the values of the parameters  $\theta$

using the observed data  $\mathcal{D} = \{X_1, X_2, \dots, X_n\}$



A hyper-ellipsoidal cluster formed by points drawn from a population which has normal density. What are the parameters of this normal density?

from Duda, Hart, Stork (2001) Pattern classification

The **maximum likelihood estimate**  $\hat{\theta}$  is the value of  $\theta$  that maximizes  $p(x_1, x_2, \dots, x_n \mid \theta)$

For analytical purposes, we use the logarithm of the likelihood, called **log-likelihood**

$$l(\theta) \equiv \ln p(x_1, x_2, \dots, x_n \mid \theta)$$

The solution is the value  $\hat{\theta}$  of the argument that maximizes the log-likelihood:

$$\hat{\theta} = \arg \max_{\theta} l(\theta)$$

# What to do?

The events  $x_i$ ,  $i = 1 \dots n$ , are statistically independent, i.e.

$$p(x_1, x_2, \dots, x_n | \theta) = \prod_{k=1}^n p(x_k | \theta)$$

or

$$l(\theta) = \sum_{k=1}^n \ln p(x_k | \theta)$$

Denote by  $\nabla_{\theta} = \begin{bmatrix} \frac{\partial}{\partial \theta_1} \\ \vdots \\ \frac{\partial}{\partial \theta_p} \end{bmatrix}$  the gradient operator.

# What to do?

The gradient of the log-likelihood is:

$$\nabla_{\theta} l = \sum_{k=1}^n \nabla_{\theta} \ln p(\mathbf{x}_k | \theta)$$

A set of necessary conditions for  $\hat{\theta}$  can be formulated as:

$$\nabla_{\theta} l = 0$$




Solve this (system of) equation(s)!

(The solutions of this equation can be *global*, *local* maxima or minima, or saddle points. Don't forget to check that it is a maximum!)

# Example

**Model: Gaussian  
with fixed variance**

**Estimated parameter:  
mean**



from Duda, Hart, Stork (2001) Pattern classification



# ML Estimation for normal distribution – unknown mean $\mu$

**Recall that:**  $\ln p(\mathbf{x}_k | \mu) = -\frac{1}{2} \ln[(2\pi)^d |\Sigma|] - \frac{1}{2} (\mathbf{x}_k - \mu)^t \Sigma^{-1} (\mathbf{x}_k - \mu)$

**It follows:**  $\nabla_{\mu} \ln p(\mathbf{x}_k | \mu) = \Sigma^{-1} (\mathbf{x}_k - \mu)$

**Solve:**  $\sum_{k=1}^n \Sigma^{-1} (\mathbf{x}_k - \mu) = 0$

$$\Rightarrow \hat{\mu} = \frac{1}{n} \sum_{k=1}^n \mathbf{x}_k$$

In plain text: the best estimate of the mean of a distribution is the mean of the sample!

*This kind of pedantic, algebra-filled and ultimately unsurprising fact is exactly the reason people throw down their “Statistics” book and pick up their “Agent Based Evolutionary Data Mining Using The Neuro-Fuzzy Transform” book. (from slides on MLE by Andrew W. Moore)*

# ML Estimation for normal distribution – unknown $\mu$ and $\Sigma$

In the univariate case,  $\theta_1 = \mu; \theta_2 = \sigma^2$

$$\ln p(\mathbf{x}_k | \theta) = -\frac{1}{2} \ln 2\pi\theta_2 - \frac{1}{2\theta_2} (\mathbf{x}_k - \theta_1)^2$$

The gradient is:

$$\nabla_{\theta} l = \nabla_{\theta} \ln p(\mathbf{x}_k | \theta) = \begin{bmatrix} \frac{1}{\theta_2} (\mathbf{x}_k - \theta_1) \\ -\frac{1}{2\theta_2} + \frac{(\mathbf{x}_k - \theta_1)^2}{2\theta_2^2} \end{bmatrix}$$

Imposing  $\nabla_{\theta} l = 0$ :

$$\sum_{k=1}^n \frac{1}{\hat{\theta}_2} (x_k - \hat{\theta}_1) = 0$$

$$-\sum_{k=1}^n \frac{1}{\hat{\theta}_2} + \sum_{i=1}^n \frac{(x_k - \hat{\theta}_1)^2}{\hat{\theta}_2^2} = 0$$

Substituting  $\hat{\mu} = \hat{\theta}_1$ ,  $\hat{\sigma}^2 = \hat{\theta}_2$  and rearranging:

$$\hat{\mu} = \frac{1}{n} \sum_{k=1}^n x_k$$

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{k=1}^n (x_k - \hat{\mu})^2$$

# MLE for normal distribution - multidimensional case

$$\hat{\mu} = \frac{1}{n} \sum_{k=1}^n \mathbf{x}^{(k)}$$

$$\hat{\Sigma} = \frac{1}{n} \sum_{k=1}^n (\mathbf{x}^{(k)} - \hat{\mu})(\mathbf{x}^{(k)} - \hat{\mu})^t$$

$$\hat{\Sigma}_{ij} = \frac{1}{n} \sum_{k=1}^n (x_i^{(k)} - \hat{\mu}_i)(x_j^{(k)} - \hat{\mu}_j)$$

where  $x_i^{(k)}$  is the  $i$ -th feature of the  $k$ -th feature vector  $\mathbf{x}^{(k)}$   
and  $\hat{\mu}_i$  is the  $i$ -th feature of the mean  $\hat{\mu}$  of all  $n$  feature vectors

# ML Estimation

The ML estimate for the variance  $\sigma^2$  is *biased*, i.e. the expected value over all possible (random!) data sets of size  $n$  is not equal to the true variance:

$$\mathcal{E} \left[ \frac{1}{n} \sum_{k=1}^n (x_k - \bar{x})^2 \right] = \frac{n-1}{n} \sigma^2 \neq \sigma^2$$

... but it is *asymptotically unbiased* (for large  $n$ )

# ML Estimation

An unbiased estimator for  $\Sigma$  is the *sample covariance matrix*:


$$C = \frac{1}{n-1} \sum_{k=1}^n (\mathbf{x}_k - \hat{\mu})(\mathbf{x}_k - \hat{\mu})^t$$

An *asymptotically* (i.e. large  $n$ ) unbiased estimator for  $\Sigma$  is:


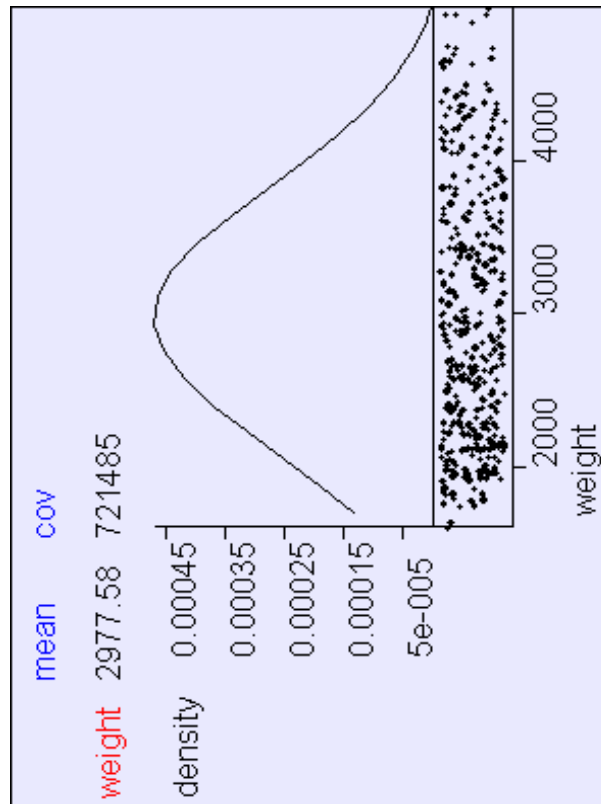
$$C = \frac{1}{n} \sum_{k=1}^n (\mathbf{x}_k - \hat{\mu})(\mathbf{x}_k - \hat{\mu})^t$$

# Gaussian MLE in action

Using  $n=392$  cars from the “MPG” UCI dataset supplied by Ross Quinlan



# Bivariate MLE in action





# Multivariate MLE

	mean	cov							
mpg	23.4459	60.9181	-10.3529	-657.585	-233.858	-5517.44	9.11551	16.6915	
cylinders	5.47194	-10.3529	2.9097	169.722	55.3482	1300.42	-2.37505	-2.17193	
displacement	194.412	-657.585	169.722	10950.4	3614.03	82929.1	-156.994	-142.572	
horsepower	104.469	-233.858	55.3482	3614.03	1481.57	28265.6	-73.187	-59.0364	
weight	2977.58	-5517.44	1300.42	82929.1	28265.6	721485	-976.815	-967.228	
acceleration	15.5413	9.11551	-2.37505	-156.994	-73.187	-976.815	7.61133	2.95046	
modelyear	75.9796	16.6915	-2.17193	-142.572	-59.0364	-967.228	2.95046	13.5699	

Covariance matrices are not exciting to look at